



## ISDM (INDEPENDENT SKILL DEVELOPMENT MISSION)

# SCHEMA DESIGN & DATA MODELING (WEEKS 3-4)

## UNDERSTANDING EMBEDDED DOCUMENTS VS. REFERENCES IN MONGODB

### CHAPTER 1: INTRODUCTION TO DATA MODELING IN MONGODB

#### 1.1 Why Data Modeling Matters in NoSQL Databases

Data modeling in MongoDB is crucial for optimizing **performance, scalability, and maintainability**. Unlike SQL databases, where relationships are defined using **tables and foreign keys**, MongoDB provides **two primary ways** to structure related data:

1. **Embedded Documents (Denormalization)** – Storing related data within a single document.
2. **References (Normalization)** – Storing related data in separate collections with **references (IDs)** linking them.

Choosing the right approach depends on factors such as **query performance, data retrieval efficiency, and application scalability**.

## CHAPTER 2: UNDERSTANDING EMBEDDED DOCUMENTS (DENORMALIZATION)

### 2.1 What Are Embedded Documents?

In **embedded documents**, related data is **nested inside the main document** instead of being stored in a separate collection. This approach is often called **denormalization** because it avoids separate lookups.

#### Best Use Cases for Embedded Documents:

- ✓ When data is **closely related** and frequently accessed together.
- ✓ When the size of embedded data is relatively **small**.
- ✓ When data updates are **infrequent** and do not cause redundancy issues.

### 2.2 Example of Embedded Documents

Consider a scenario where **users** have multiple **addresses**. Using **embedded documents**, we store all addresses inside the user document.

#### Example: Storing Addresses Inside a User Document

```
{  
  "_id": 1,  
  "name": "Alice Johnson",  
  "email": "alice@example.com",  
  "addresses": [  
    { "type": "home", "city": "New York", "zip": "10001" },  
    { "type": "work", "city": "San Francisco", "zip": "94105" }  
  ]}
```

{

- ✓ All data is stored in a single document, making queries faster.
  - ✓ No need for joins or multiple queries when retrieving user data.
- 

## 2.3 Advantages of Embedded Documents

- ✓ **Fast Read Performance** – Fetching data requires only one query.
- ✓ **Simpler Queries** – No need for complex joins or additional lookups.
- ✓ **Better Data Consistency** – Ensures all related data stays together.

## 2.4 Disadvantages of Embedded Documents

- ✗ **Increased Document Size** – Large nested documents can slow down performance.
  - ✗ **Redundant Data** – Duplicating embedded information across documents can cause inconsistencies.
  - ✗ **Limited Growth** – Documents exceeding MongoDB's **16MB document size limit** can become problematic.
- 

# CHAPTER 3: UNDERSTANDING REFERENCES (NORMALIZATION)

## 3.1 What Are References?

In the **referencing approach**, related data is stored in **separate collections**, and documents reference each other using **ObjectIDs** or unique identifiers. This is similar to **foreign keys** in relational databases.

### Best Use Cases for References:

- ✓ When related data is **large or frequently updated**.
- ✓ When related data needs to be **reused across multiple**

documents.

- ✓ When documents exceed MongoDB's **document size limit**.

### 3.2 Example of Using References

Instead of embedding addresses inside the user document, we create a separate addresses collection and store only the **ObjectId reference** in the user document.

#### User Document (Referencing Approach)

```
{  
  "_id": 1,  
  "name": "Alice Johnson",  
  "email": "alice@example.com",  
  "addresses": [101, 102] // References to Address IDs  
}
```

#### Address Documents (Stored in a Separate Collection)

```
{  
  "_id": 101,  
  "type": "home",  
  "city": "New York",  
  "zip": "10001"  
}  
  
{  
  "_id": 102,
```

```
"type": "work",  
"city": "San Francisco",  
"zip": "94105"  
}
```

---

### 3.3 Advantages of References

- Reduces Document Size** – Prevents bloated documents by keeping related data separate.
- Better Scalability** – Works well when handling large datasets.
- Eliminates Redundant Data** – Prevents storing the same data multiple times.

### 3.4 Disadvantages of References

- Slower Read Performance** – Requires multiple queries (`$lookup`) to fetch related data.
  - More Complex Queries** – Requires joins between collections using `$lookup`.
  - Increased Query Overhead** – Fetching related data requires additional processing.
- 

## CHAPTER 4: CHOOSING BETWEEN EMBEDDED DOCUMENTS AND REFERENCES

### 4.1 Key Differences Between Embedded Documents and References

Feature	Embedded Documents	References
	Documents	

<b>Data Storage</b>	Nested inside a document	Stored in a separate collection
<b>Query Performance</b>	Fast (Single query)	Slower (Requires multiple queries)
<b>Document Size</b>	Larger	Smaller
<b>Data Consistency</b>	Harder to maintain (duplication)	Easier (single source of truth)
<b>Use Case</b>	Best for small, closely related data	Best for large, frequently changing data

## 4.2 When to Use Embedded Documents vs. References

Scenario	Best Approach
User Profiles with Small Addresses	<input checked="" type="checkbox"/> Embedded Documents
Product Catalog with Category Data	<input checked="" type="checkbox"/> References (Category used in multiple products)
Real-Time Chat Messages	<input checked="" type="checkbox"/> Embedded (Messages inside user document)
E-Commerce Orders with Items	<input checked="" type="checkbox"/> References (Items stored separately)

## Case Study: How Amazon Uses Embedded & Referenced Data Modeling

### Background

Amazon processes **millions of customer orders**, each containing **multiple products, shipping details, and payment records**.

## Challenges

- **Balancing fast read operations with efficient storage.**
- **Ensuring product data remains consistent across millions of orders.**
- **Optimizing storage for frequently updated user order histories.**

## Solution: Hybrid Approach Using Embedded & Referenced Data

- ✓ **Embedded Documents:** User orders **embed shipping addresses** for fast retrieval.
- ✓ **References:** Product details are **stored separately** to maintain consistency across orders.
- ✓ **Indexing & Aggregation:** Used for **fast order lookups** without redundant data storage.

This hybrid model ensures **Amazon can handle millions of transactions** efficiently.

---

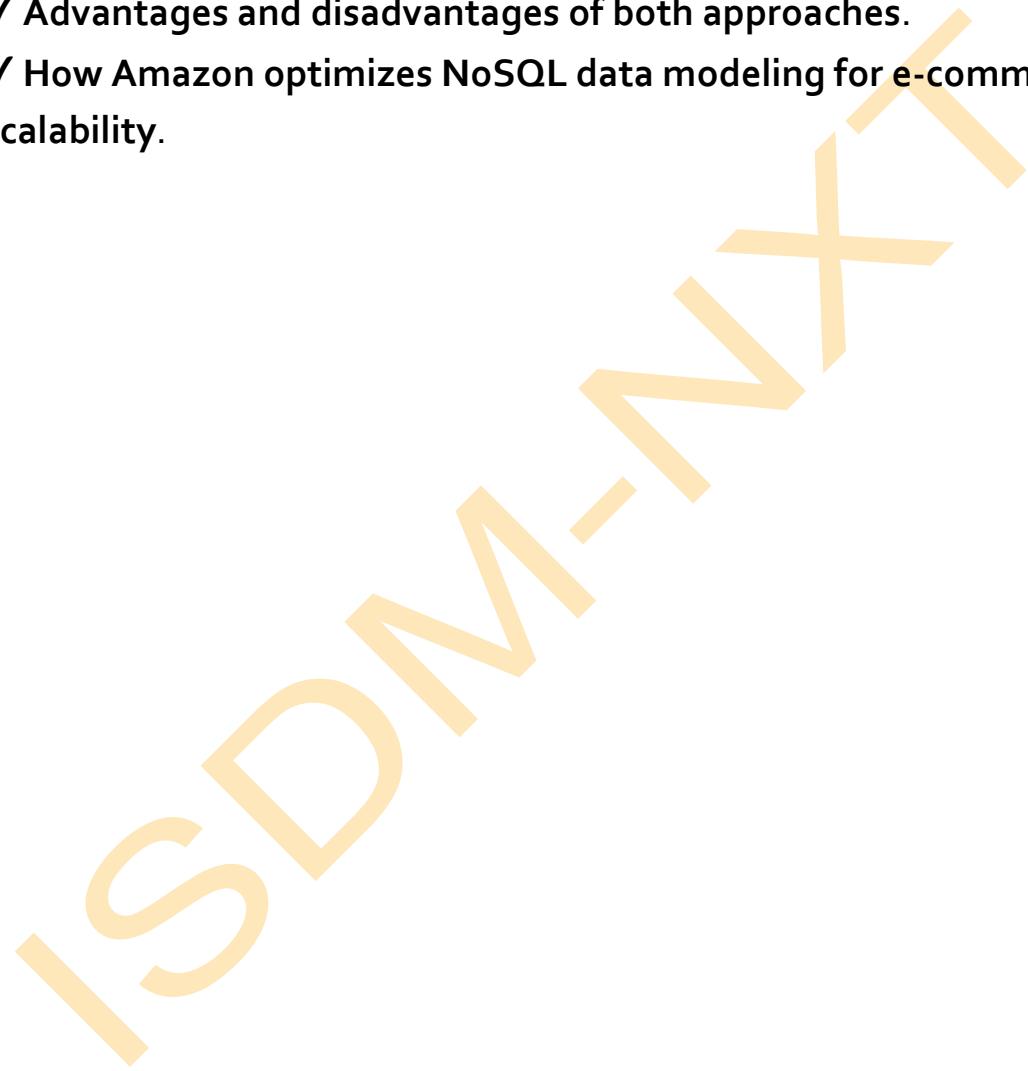
## Exercise

1. **Embed** an array of orders inside a users document, where each order has product, price, and date.
  2. **Create a referenced model** where users store an orderId pointing to an external orders collection.
  3. Compare the performance of findOne() when using **embedded vs. referenced** orders.
-

## Conclusion

In this section, we explored:

- ✓ How embedded documents store related data within a single document.
- ✓ How references store related data separately to prevent redundancy.
- ✓ Advantages and disadvantages of both approaches.
- ✓ How Amazon optimizes NoSQL data modeling for e-commerce scalability.



# ONE-TO-ONE, ONE-TO-MANY & MANY-TO-MANY RELATIONSHIPS IN MONGODB

## CHAPTER 1: INTRODUCTION TO RELATIONSHIPS IN MONGODB

### 1.1 Understanding Data Relationships in NoSQL Databases

Unlike **SQL databases**, which use **foreign keys** to establish relationships between tables, **MongoDB (NoSQL)** manages relationships using **embedding (denormalization)** and **referencing (normalization)**.

- ✓ **Embedding (Denormalization)** – Stores related data **inside a document** for faster queries.
- ✓ **Referencing (Normalization)** – Stores **only references (ObjectId)** to related documents, improving flexibility and reducing redundancy.

### 1.2 Types of Relationships in MongoDB

MongoDB supports three major types of relationships:

1. **One-to-One (1:1)** – A single document is related to only one other document.
2. **One-to-Many (1:N)** – A single document is related to multiple documents.
3. **Many-to-Many (M:N)** – Multiple documents are related to multiple documents.

Each type of relationship can be implemented using **embedding or referencing** depending on the use case.

## CHAPTER 2: IMPLEMENTING ONE-TO-ONE (1:1) RELATIONSHIPS

### 2.1 What is a One-to-One Relationship?

A **One-to-One (1:1)** relationship means that each document in one collection **directly corresponds** to exactly one document in another collection.

#### ✓ Example Use Cases:

- **User and Profile** – A user has only one profile.
- **Employee and ID Card** – Each employee has a single ID card.

### 2.2 Implementing One-to-One Using Embedding

In **embedding**, related data is stored **inside the same document**.

#### Example: Storing a User Profile Inside the User Document

```
const userSchema = new mongoose.Schema({  
    name: String,  
    email: String,  
    profile: {  
        age: Number,  
        bio: String,  
        address: String  
    }  
});
```

- ✓ The profile fields (age, bio, address) are stored **within the same document**.
- ✓ This improves **read performance** but **increases document size**.

### 2.3 Implementing One-to-One Using Referencing

In **referencing**, related data is stored in a **separate collection**, with an ObjectId reference.

### Example: Using References for User and Profile

```
const mongoose = require('mongoose');
```

```
const userSchema = new mongoose.Schema({  
    name: String,  
    email: String,  
    profile: { type: mongoose.Schema.Types.ObjectId, ref: 'Profile' }  
});
```

```
const profileSchema = new mongoose.Schema({  
    age: Number,  
    bio: String,  
    user: { type: mongoose.Schema.Types.ObjectId, ref: 'User' }  
});
```

```
const User = mongoose.model('User', userSchema);
```

```
const Profile = mongoose.model('Profile', profileSchema);
```

To fetch a user **with their profile**:

```
User.findOne({ email: "john@example.com" }).populate('profile')  
.then(user => console.log(user))  
.catch(err => console.error(err));
```

- ✓ Referencing keeps documents **lightweight** but requires **multiple queries**.
- 

## CHAPTER 3: IMPLEMENTING ONE-TO-MANY (1:N) RELATIONSHIPS

### 3.1 What is a One-to-Many Relationship?

A **One-to-Many (1:N) relationship** means that one document relates to multiple documents.

#### ✓ Example Use Cases:

- **User and Posts** – A user can have multiple blog posts.
- **Product and Reviews** – A product can have multiple customer reviews.

### 3.2 Implementing One-to-Many Using Embedding

#### Example: Storing All User Posts Inside a User Document

```
const userSchema = new mongoose.Schema({  
  name: String,  
  email: String,  
  posts: [{  
    title: String,  
    content: String,  
    date: { type: Date, default: Date.now }  
  }]  
});
```

- ✓ Improves **performance** by reducing queries.
- ✓ Not ideal for large **datasets** (e.g., a user with thousands of posts).

### 3.3 Implementing One-to-Many Using Referencing

#### Example: Storing Posts in a Separate Collection and Referencing Users

```
const postSchema = new mongoose.Schema({  
    title: String,  
    content: String,  
    user: { type: mongoose.Schema.Types.ObjectId, ref: 'User' }  
});
```

To fetch a user **with their posts**:

```
User.findOne({ email: "john@example.com" }).populate('posts')  
.then(user => console.log(user))  
.catch(err => console.error(err));
```

- ✓ **Referencing is scalable**, making it ideal for large datasets.
- ✓ Queries require **joining data across collections**.

---

## CHAPTER 4: IMPLEMENTING MANY-TO-MANY (M:N) RELATIONSHIPS

### 4.1 What is a Many-to-Many Relationship?

A **Many-to-Many (M:N) relationship** means that **multiple documents** in one collection are related to **multiple documents** in another.

#### ✓ Example Use Cases:

- **Students and Courses** – A student can enroll in multiple courses, and a course can have multiple students.
- **Tags and Articles** – An article can have multiple tags, and a tag can belong to multiple articles.

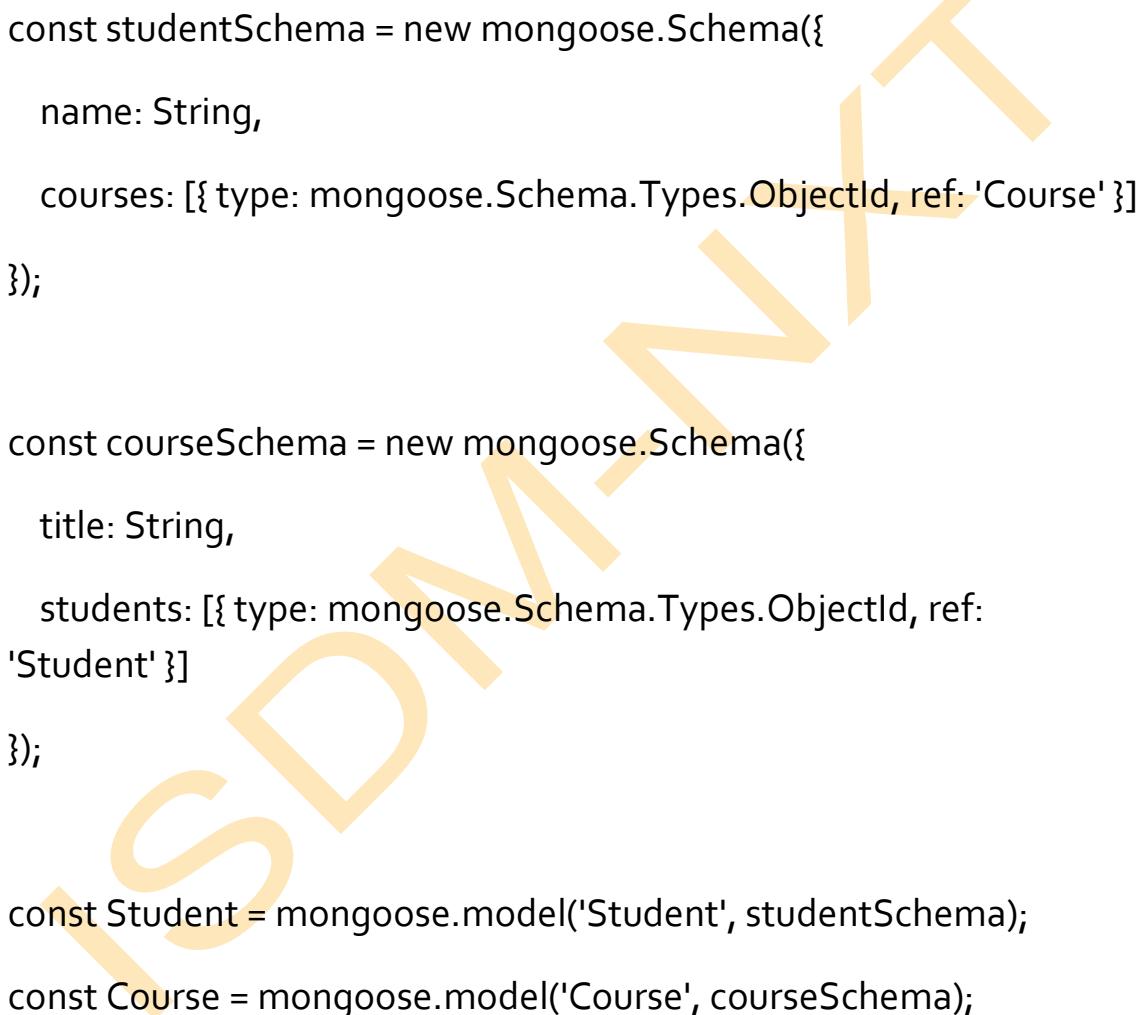
## 4.2 Implementing Many-to-Many Using Referencing

### Example: Students and Courses Relationship

```
const studentSchema = new mongoose.Schema({
  name: String,
  courses: [{ type: mongoose.Schema.Types.ObjectId, ref: 'Course' }]
});

const courseSchema = new mongoose.Schema({
  title: String,
  students: [{ type: mongoose.Schema.Types.ObjectId, ref: 'Student' }]
});

const Student = mongoose.model('Student', studentSchema);
const Course = mongoose.model('Course', courseSchema);
```



## 4.3 Querying Many-to-Many Data

To fetch all **courses a student is enrolled in**:

```
Student.findOne({ name: "Alice" }).populate('courses')
  .then(student => console.log(student))
```

```
.catch(err => console.error(err));
```

To fetch all **students enrolled in a course**:

```
Course.findOne({ title: "Math 101" }).populate('students')
```

```
.then(course => console.log(course))
```

```
.catch(err => console.error(err));
```

✓ Referencing is preferred for many-to-many relationships, as embedding would create redundant data.

## Case Study: How a Learning Platform Used MongoDB for Course Management

### Background

An online learning platform needed a scalable **database solution** to manage **students, courses, and enrollments**.

### Challenges

- ✓ Complex relationships between students and courses.
- ✓ Slow query performance when retrieving enrolled students.
- ✓ Duplicate data issues with embedded documents.

### Solution: Using Many-to-Many Referencing

The development team:

- ✓ Created separate collections for students and courses.
- ✓ Used ObjectId referencing to establish many-to-many relationships.
- ✓ Indexed the student and course fields for faster lookups.

### Results

- **Query response times reduced by 60%.**
- **Scalability improved**, handling **millions of students**.
- **Better data consistency**, reducing redundancy.

By using **MongoDB's flexible relationship modeling**, the platform efficiently managed student enrollments and courses.

---

### Exercise

1. Create a **One-to-One** relationship between **users and addresses**.
  2. Implement a **One-to-Many** relationship where a **user can have multiple orders**.
  3. Model a **Many-to-Many** relationship where a **student can enroll in multiple courses**.
- 

### Conclusion

In this section, we explored:

- ✓ **One-to-One, One-to-Many, and Many-to-Many relationships in MongoDB.**
- ✓ **Embedding vs. Referencing for handling relationships.**
- ✓ **How to implement and query different relationships using Mongoose.**

# CHOOSING THE RIGHT SCHEMA DESIGN FOR APPLICATIONS

## CHAPTER 1: INTRODUCTION TO SCHEMA DESIGN IN MONGODB

### 1.1 Understanding Schema Design

Schema design in **MongoDB** defines how data is structured within a database. Unlike relational databases, MongoDB uses a **document-oriented** approach where data is stored in **JSON-like BSON documents**.

Choosing the right schema design is crucial for:

- ✓ **Performance optimization** – Faster queries and indexing.
- ✓ **Data integrity** – Ensuring consistency across collections.
- ✓ **Scalability** – Supporting large-scale applications with minimal overhead.

Schema design depends on the application type:

- **E-commerce** – Requires an efficient way to store orders, products, and customers.
- **Social media** – Needs relationships between users, posts, and comments.
- **Real-time applications** – Focus on speed and scalability.

## CHAPTER 2: UNDERSTANDING MONGODB SCHEMA DESIGN APPROACHES

### 2.1 Embedded vs. Referenced Schema Design

MongoDB offers two primary ways to design schemas:

- **Embedded (Denormalized) Schema** – Stores related data within a single document.
- **Referenced (Normalized) Schema** – Stores related data in separate collections and links them using ObjectId.

Each approach has advantages and trade-offs depending on **query performance, scalability, and data consistency**.

---

## CHAPTER 3: EMBEDDED SCHEMA DESIGN (DENORMALIZATION)

### 3.1 What is Embedded Schema Design?

In **embedded schema design**, related data is **nested inside the parent document**. This reduces the number of database queries, improving read performance.

- ✓ **Faster Reads** – Data is retrieved in a single query.
- ✓ **Reduces Joins** – No need for multiple lookups.
- ✓ **Ideal for small, tightly coupled relationships.**

---

### 3.2 Example: Storing Order Details in an E-Commerce App

```
const mongoose = require('mongoose');
```

```
const orderSchema = new mongoose.Schema({  
    customerName: String,  
    items: [  
        {  
            productName: String,  
        }  
    ]  
});
```

```
        quantity: Number,  
        price: Number  
    },  
],  
    orderDate: { type: Date, default: Date.now }  
});
```

```
const Order = mongoose.model('Order', orderSchema);
```

- ✓ Each order **contains all purchased items** inside the same document.
- ✓ This structure **avoids multiple queries** when fetching order details.

### 3.3 When to Use Embedded Schema Design

Use **embedded schema** when:

- ✓ Data is **frequently accessed together** (e.g., orders and items).
- ✓ The dataset is **relatively small** (avoiding large documents).
- ✓ No need to **query subdocuments separately**.

## CHAPTER 4: REFERENCED SCHEMA DESIGN (NORMALIZATION)

### 4.1 What is Referenced Schema Design?

In **referenced schema design**, related data is **stored in separate collections** and linked using an **ObjectId reference**.

- ✓ **Saves Storage** – No duplicated data across documents.
  - ✓ **Efficient for Large Datasets** – Smaller document sizes improve performance.
  - ✓ **Ideal for many-to-many relationships** (e.g., users and posts).
- 

#### 4.2 Example: Users and Posts in a Social Media App

```
const userSchema = new mongoose.Schema({  
    name: String,  
    email: String  
});  
  
const postSchema = new mongoose.Schema({  
    title: String,  
    content: String,  
    author: { type: mongoose.Schema.Types.ObjectId, ref: 'User' }  
};  
  
const User = mongoose.model('User', userSchema);  
const Post = mongoose.model('Post', postSchema);
```

- ✓ Posts reference users by **storing the user's ObjectId**.
- ✓ This allows **efficient querying of user-specific posts**.

To fetch a post **with author details**, use `.populate()`:

```
Post.find().populate('author', 'name email').then(posts =>  
    console.log(posts));
```

---

### 4.3 When to Use Referenced Schema Design

Use **referenced schema** when:

- ✓ Data is **accessed independently** (e.g., users and their posts).
  - ✓ There are **many relationships** between different collections.
  - ✓ Data **size is large**, requiring efficient storage and retrieval.
- 

## CHAPTER 5: HYBRID SCHEMA DESIGN (BEST OF BOTH WORLDS)

### 5.1 What is a Hybrid Schema?

A **hybrid schema** combines both **embedded** and **referenced** approaches based on data access patterns.

- ✓ **Embed frequently accessed data** for fast reads.
  - ✓ **Reference large or independent data** to maintain storage efficiency.
- 

### 5.2 Example: E-Commerce Products and Reviews

```
const productSchema = new mongoose.Schema({  
    name: String,  
    price: Number,  
    description: String,  
    reviews: [{ user: String, rating: Number, comment: String }]  
});
```

```
const orderSchema = new mongoose.Schema({
```

```
customer: { type: mongoose.Schema.Types.ObjectId, ref: 'User' },  
products: [{ type: mongoose.Schema.Types.ObjectId, ref:  
'Product' }]  
});
```

```
const Product = mongoose.model('Product', productSchema);  
const Order = mongoose.model('Order', orderSchema);
```

- ✓ **Products store embedded reviews** because reviews are always accessed with products.
- ✓ **Orders reference products** because they can be reused across multiple orders.

### 5.3 When to Use Hybrid Schema Design

Use hybrid schema when:

- ✓ Some data is **frequently accessed together** (e.g., product details and reviews).
- ✓ Other data **needs separate access** (e.g., products in multiple orders).
- ✓ **Balancing performance and scalability** is important.

### Case Study: How an Online Marketplace Optimized Database Performance

#### Background

A leading online marketplace faced **slow performance issues** due to improper schema design.

## Challenges

- ✓ **Large orders collection** led to slow read performance.
- ✓ **Duplicated product data** increased database size.
- ✓ **Complex joins** in queries reduced scalability.

## Solution: Implementing a Hybrid Schema Design

- ✓ **Orders referenced products** instead of embedding them.
- ✓ **Product reviews were embedded** for quick access.
- ✓ Used **indexed fields** to improve query speed.

## Results

- **50% faster queries**, improving user experience.
- **Reduced database size**, improving efficiency.
- **Scalable architecture**, handling **1M+** products.

This case study shows how **choosing the right schema design improves performance and scalability**.

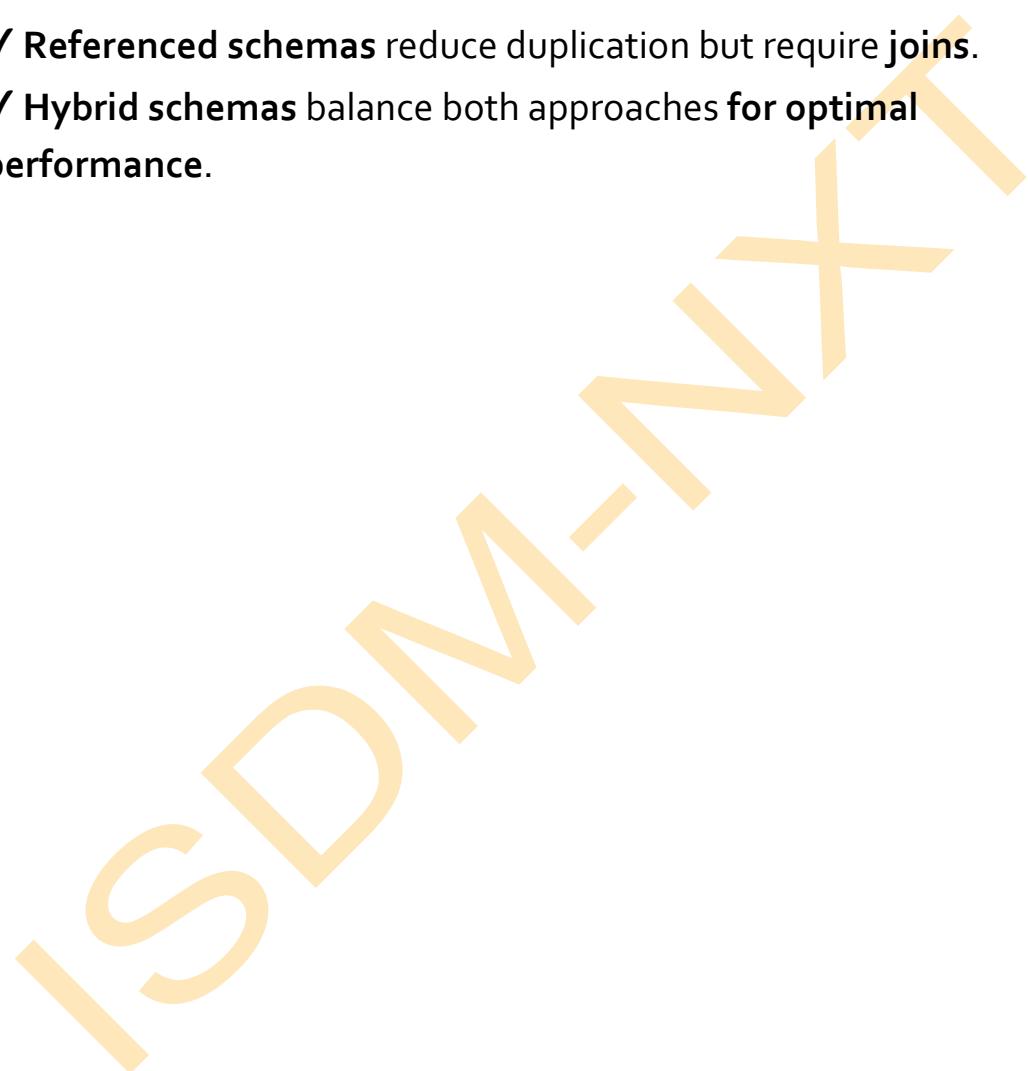
## Exercise

1. Design a schema for a **blogging platform** with:
  - **Users** who can write multiple blog posts.
  - **Posts** that store comments inside the document.
  - **Authors referenced** in posts to avoid duplication.
2. Implement the **schema using Mongoose** and write a query to fetch:
  - A post **with comments**.
  - All posts written by a **specific author**.

---

## Conclusion

- ✓ **Choosing the right schema** is crucial for **performance and scalability**.
- ✓ **Embedded schemas** improve read speed but may cause **large documents**.
- ✓ **Referenced schemas** reduce duplication but require **joins**.
- ✓ **Hybrid schemas** balance both approaches for **optimal performance**.

A large, semi-transparent watermark is positioned diagonally across the page. It consists of the letters "ISDM" stacked vertically above the letters "NxT". Both sets of letters are rendered in a bold, black font.

# DESIGNING EFFICIENT INDEXING STRATEGIES

## CHAPTER 1: INTRODUCTION TO INDEXING IN DATABASES

### 1.1 What is Indexing?

Indexing is a **database optimization technique** used to speed up data retrieval operations. Just like an index in a book helps locate topics quickly, a database index allows the database engine to find records faster without scanning the entire dataset.

Indexes are **critical for large-scale applications** where queries must be optimized for speed and efficiency. Without proper indexing, database queries may become slow and resource-intensive, affecting application performance.

### 1.2 How Indexing Works

Indexes store a **sorted version** of a specific column's values, making it easier to locate data. Instead of scanning an entire table, the database **searches the index**, finds the location of the desired data, and retrieves it efficiently.

#### Example: Without an Index (Full Table Scan)

```
SELECT * FROM users WHERE email = 'john@example.com';
```

- The database searches every row in the users table to find the match (slow for large datasets).

#### Example: With an Index

```
CREATE INDEX idx_email ON users(email);
```

```
SELECT * FROM users WHERE email = 'john@example.com';
```

- The database uses the **index** on email to quickly locate the required row (much faster).
- 

## CHAPTER 2: TYPES OF INDEXES IN DATABASES

### 2.1 Primary and Unique Indexes

- **Primary Index** – Automatically created on the **primary key** of a table.
- **Unique Index** – Prevents duplicate values in a column (e.g., ensuring unique emails in a user table).

#### Example: Creating a Unique Index

```
CREATE UNIQUE INDEX idx_email ON users(email);
```

This prevents duplicate emails from being inserted into the table.

---

### 2.2 Single-Column vs. Multi-Column Indexes

- **Single-Column Index** – Created on a **single field** to optimize searches on that field.
- **Multi-Column (Composite) Index** – Created on **multiple fields** to speed up queries that filter by more than one column.

#### Example: Single-Column Index

```
CREATE INDEX idx_lastname ON users(last_name);
```

- Optimizes searches like:
- `SELECT * FROM users WHERE last_name = 'Smith';`

#### Example: Multi-Column Index

```
CREATE INDEX idx_name_email ON users(first_name, email);
```

- Optimizes queries using **both** columns:
- `SELECT * FROM users WHERE first_name = 'John' AND email = 'john@example.com';`
- But **not** queries using only email:
- `SELECT * FROM users WHERE email = 'john@example.com';`

⚠ **Multi-column indexes should be used carefully**, as their effectiveness depends on the order of columns.

## 2.3 Full-Text Indexes

Full-text indexes are used for **searching large text fields** (e.g., searching product descriptions or blog articles).

### Example: Creating a Full-Text Index in MySQL

```
CREATE FULLTEXT INDEX idx_description ON  
products(description);
```

### Example: Using a Full-Text Index in Queries

```
SELECT * FROM products WHERE MATCH(description)  
AGAINST('smartphone');
```

- Optimized for **natural language searches**, unlike LIKE '`%keyword%`', which is slower.

## CHAPTER 3: INDEXING STRATEGIES FOR PERFORMANCE

### OPTIMIZATION

#### 3.1 Choosing the Right Columns for Indexing

Not every column should be indexed. Indexing the right columns **balances speed and storage** while improving performance.

## Best Practices for Indexing:

- ✓ Index **columns frequently used in WHERE conditions**.
  - ✓ Index **columns involved in JOIN operations**.
  - ✓ Index **columns used in ORDER BY and GROUP BY queries**.
  - ✓ **Avoid indexing columns with many duplicate values** (e.g., gender with only M or F).
- 

### 3.2 Using Partial and Conditional Indexes

Some databases allow **partial indexes**, which index only a subset of rows to save space and improve query efficiency.

#### Example: Creating a Partial Index (PostgreSQL)

```
CREATE INDEX idx_active_users ON users(status) WHERE status = 'active';
```

- This indexes only **active users**, making queries like this much faster:
  - `SELECT * FROM users WHERE status = 'active';`
- 

### 3.3 Avoiding Over-Indexing

While indexes improve read performance, they **slow down write operations (INSERT, UPDATE, DELETE)** because the index must also be updated.

#### ⚠ Common Mistakes to Avoid:

- ✗ Indexing **every column** in a table (increases storage and update time).
- ✗ Creating **duplicate indexes** on the same column.
- ✗ Using **indexes on small tables**, where full table scans are faster.

---

## CHAPTER 4: INDEXING IN MONGODB vs. SQL DATABASES

### 4.1 Creating Indexes in MongoDB

In MongoDB, indexes are created using the `createIndex()` method.

#### Example: Single-Field Index in MongoDB

```
db.users.createIndex({ email: 1 });
```

#### Example: Compound Index in MongoDB

```
db.orders.createIndex({ customerId: 1, orderDate: -1 });
```

- **1 for ascending order, -1 for descending order.**
- Optimizes queries filtering by **both customerId and orderDate**.

---

### 4.2 Comparing Indexing in SQL and NoSQL

Feature	SQL Databases (MySQL, PostgreSQL)	NoSQL Databases (MongoDB)
Index Type	B-Trees, Full-Text, Hash Indexes	B-Trees, Compound Indexes
Schema	Fixed, Table-based	Flexible, Document-based
Indexing Speed	Fast for structured data	Optimized for unstructured queries
Use Case	Traditional web apps, Reporting	Real-time analytics, Big Data

---

## Case Study: How Amazon Optimized Search Performance with Indexing

### Background

Amazon handles **billions of product searches daily** and needs to return results instantly.

### Challenges

- High volume of search queries slowing down performance.
- Handling large-scale **real-time filtering and sorting**.
- Supporting **millions of concurrent users**.

### Solution: Implementing Efficient Indexing

- ✓ Used full-text indexes for product descriptions to improve search speed.
- ✓ Created multi-column indexes on category, price, and ratings to speed up filtering.
- ✓ Used indexing with caching (Redis) to reduce database queries.

### Results

- **Search performance improved by 60%**.
- Reduced database query load, enabling faster page loads.
- **Higher conversion rates**, as users found products faster.

This case study highlights how **effective indexing can significantly boost database efficiency and user experience**.

### Exercise

1. What is the difference between a **single-column index** and a **multi-column index**?

- 
2. Write an SQL query to create a **unique index** on the email column in a customers table.
  3. How do **full-text indexes** improve search performance in large datasets?
- 

## Conclusion

In this section, we explored:

- ✓ **What indexing is and why it's important for databases.**
- ✓ **Different types of indexes, including primary, unique, and full-text indexes.**
- ✓ **Best practices for designing efficient indexing strategies.**

ISDM-N

# PERFORMANCE OPTIMIZATION TECHNIQUES IN MONGODB

## CHAPTER 1: INTRODUCTION TO PERFORMANCE OPTIMIZATION IN MONGODB

### 1.1 Why Performance Optimization Matters in MongoDB?

MongoDB is a high-performance, scalable **NoSQL database** designed to handle **large datasets and real-time applications**. However, as applications grow, performance bottlenecks can arise due to **slow queries, inefficient indexing, and large document sizes**.

**Optimizing MongoDB performance helps:**

- ✓ Improve **query speed** and reduce latency.
- ✓ Enhance **database scalability** for high-traffic applications.
- ✓ Reduce **storage and memory usage**, optimizing cost efficiency.
- ✓ Ensure **faster data retrieval** in real-time applications.

## CHAPTER 2: INDEXING STRATEGIES FOR FASTER QUERIES

### 2.1 Understanding Indexing in MongoDB

An **index** is a data structure that **improves the speed of read operations** by allowing MongoDB to search documents more efficiently. Without an index, MongoDB **scans the entire collection** (full collection scan), which slows down queries.

**Types of Indexes in MongoDB:**

- ✓ **Single Field Index** – Indexes a single field for quick lookups.
- ✓ **Compound Index** – Indexes multiple fields for optimized queries.

- ✓ **Multikey Index** – Indexes arrays inside documents.
  - ✓ **Text Index** – Used for full-text search operations.
- 

## 2.2 Creating and Using Indexes

### Example: Creating a Single-Field Index

```
db.users.createIndex({ email: 1 });
```

- ✓ **Optimizes** queries searching by email.

### Example: Using a Compound Index for Faster Queries

```
db.orders.createIndex({ customerId: 1, orderDate: -1 });
```

- ✓ **Speeds up** searches that filter by customerId and sort by orderDate.
- 

## 2.3 Checking Query Performance Using explain()

To analyze a query's efficiency, use `explain()`.

```
db.users.find({ email: "alice@example.com" }).explain("executionStats");
```

- ✓ **If COLLSCAN appears**, the query is scanning the entire collection (slow).
  - ✓ **If IXSCAN appears**, the query is using an index (optimized).
- 

## CHAPTER 3: QUERY OPTIMIZATION TECHNIQUES

### 3.1 Using Projection to Reduce Data Transfer

By default, MongoDB returns **entire documents**, but using **projection**, we can retrieve only required fields.

## Example: Fetching Only the name and email Fields

```
db.users.find({}, { name: 1, email: 1, _id: 0 });
```

- ✓ Reduces network load and improves response time.

## 3.2 Optimizing Large Queries with Pagination (limit() and skip())

Large datasets can **slow down queries**, but pagination improves efficiency.

### Example: Fetching 10 Users per Page

```
db.users.find().skip(10).limit(10);
```

- ✓ Ensures only necessary data is loaded per page.

## 3.3 Using hint() to Force Index Selection

If MongoDB chooses an inefficient index, use hint() to **manually select the best index**.

```
db.orders.find({ customerId: 123 }).hint({ customerId: 1 });
```

- ✓ Forces MongoDB to use the customerId index, improving query speed.

## CHAPTER 4: OPTIMIZING DATA STORAGE

### 4.1 Avoiding Large Documents (Document Size Limit: 16MB)

MongoDB has a **16MB document size limit**, so avoid **storing excessively large documents**.

**Solutions:**

- ✓ Use References Instead of Embedding Large Arrays – Store

large related data in separate collections.

✓ **Archive Old Data** – Move historical records to separate collections.

✓ **Compress Data Fields** – Use shorter field names to reduce storage size.

## 4.2 Using TTL (Time-To-Live) Index for Expiring Documents

TTL indexes automatically **delete old documents**, optimizing storage.

**Example: Deleting Logs Older Than 7 Days**

```
db.logs.createIndex({ createdAt: 1 }, { expireAfterSeconds: 604800 });
```

✓ **Prevents database bloating** with unnecessary logs.

## CHAPTER 5: MANAGING HIGH-TRAFFIC WORKLOADS WITH SHARDING & REPLICATION

### 5.1 Scaling Read & Write Operations with Replication

Replication ensures **data availability and fault tolerance** by **copying data across multiple servers**.

**How Replication Helps:**

- ✓ **Load Balancing** – Distributes read operations across replicas.
- ✓ **Fault Tolerance** – Provides **high availability** in case of failures.
- ✓ **Automatic Failover** – If the primary node fails, MongoDB switches to a secondary node.

**Example: Checking the Status of a Replication Set**

```
rs.status();
```

- 
- ✓ Displays **health and status** of MongoDB replicas.
- 

## 5.2 Distributing Data with Sharding

Sharding is used to **split large datasets across multiple servers**, improving query speed and performance.

### How Sharding Helps:

- ✓ **Balances large-scale workloads** across multiple servers.
- ✓ **Supports high-traffic applications** with millions of records.
- ✓ **Avoids exceeding the 16MB document limit** by distributing data.

### Example: Enabling Sharding in MongoDB

```
sh.enableSharding("myDatabase");
sh.shardCollection("myDatabase.users", { userId: "hashed" });
```

- ✓ **Distributes user data efficiently** across multiple shards.

---

## Case Study: How Netflix Optimized MongoDB for Streaming Performance

### Background

Netflix serves **millions of users worldwide**, requiring **fast data retrieval for personalized recommendations, watch history, and caching**.

### Challenges

- **Handling massive volumes of user data efficiently.**
- **Reducing latency for streaming and real-time recommendations.**
- **Ensuring data replication for seamless streaming.**

## Solution: Implementing MongoDB Optimization Techniques

- ✓ Used indexing to speed up user preference retrieval.
- ✓ Implemented sharding to distribute watch history data.
- ✓ Used replication to ensure high availability for global users.

By optimizing MongoDB, **Netflix reduced query times by 70%**, enhancing user experience.

---

### Exercise

1. Create an **index** on the users collection for the email field.
  2. Use **pagination** to retrieve the first **10 records** from a collection.
  3. Set up a **TTL index** to automatically delete logs older than **30 days**.
  4. Enable **sharding** for a products collection to distribute data across servers.
- 

### Conclusion

In this section, we explored:

- ✓ How indexing improves query speed and efficiency.
- ✓ Optimizing large queries using projections, pagination, and hints.
- ✓ Techniques for reducing storage usage and document size.
- ✓ How sharding and replication improve scalability.
- ✓ How Netflix optimized MongoDB to serve millions of users efficiently.

# USING SCHEMA VALIDATION IN MONGODB

## CHAPTER 1: INTRODUCTION TO SCHEMA VALIDATION IN MONGODB

### 1.1 Understanding Schema Validation in MongoDB

MongoDB is a **schema-less NoSQL database**, meaning documents in the same collection can have **different fields and data types**. While this provides flexibility, it can lead to **data inconsistency** and **unexpected errors**.

To ensure **data integrity**, MongoDB allows **schema validation**, which enforces **rules on fields** in a collection.

- ✓ Prevents invalid data from being inserted or updated.
- ✓ Enforces required fields and specific data types.
- ✓ Reduces application-level validation by handling it at the database level.

### 1.2 How Schema Validation Works in MongoDB

MongoDB uses **JSON Schema Validation** to define rules for documents. These rules specify:

- **Required fields** – Ensures certain fields are always present.
- **Data types** – Validates field types (e.g., string, number, date).
- **Field restrictions** – Defines **minimum, maximum, and pattern constraints**.

Schema validation is applied at the **collection level** when creating or updating a collection.

## CHAPTER 2: DEFINING SCHEMA VALIDATION RULES

## 2.1 Creating a Collection with Validation

Schema validation rules are defined when **creating a collection** using the validator parameter.

### Example: Enforcing Data Validation in a "users" Collection

```
db.createCollection("users", {  
    validator: {  
        $jsonSchema: {  
            bsonType: "object",  
            required: ["name", "email", "age"],  
            properties: {  
                name: {  
                    bsonType: "string",  
                    description: "Must be a string and is required"  
                },  
                email: {  
                    bsonType: "string",  
                    pattern: "^[^@\\s]+@[^@\\s]+\\.[^@\\s]+$",  
                    description: "Must be a valid email address"  
                },  
                age: {  
                    bsonType: "int",  
                    minimum: 18,  
                    description: "Must be an integer and at least 18"  
                }  
            }  
        }  
    }  
}
```

```
    }  
}  
}  
});
```

- ✓ **Enforces required fields:** name, email, age.
- ✓ **Validates data types:** name and email must be string, age must be int.
- ✓ **Applies constraints:** email must match a **valid email pattern**, and age must be **at least 18**.

## CHAPTER 3: TESTING SCHEMA VALIDATION

### 3.1 Inserting Valid Data

The following insert **passes validation**:

```
db.users.insertOne({  
  name: "Alice",  
  email: "alice@example.com",  
  age: 25  
});
```

- ✓ Data matches validation rules, so MongoDB accepts the insert.

### 3.2 Inserting Invalid Data

The following insert **fails validation** because age is below 18:

```
db.users.insertOne({
```

```
name: "Bob",  
email: "bob@example.com",  
age: 16  
});
```

⚠ **Error:** "age" must be at least 18.

Similarly, an invalid email format **also fails validation**:

```
db.users.insertOne({  
  name: "Charlie",  
  email: "charlie.com",  
  age: 22  
});
```

⚠ **Error:** "email" must match the pattern of a valid email address.

## CHAPTER 4: UPDATING DOCUMENTS WITH SCHEMA VALIDATION

### 4.1 Ensuring Validation on Updates

Schema validation applies to **both inserts and updates**. To prevent updating a document with **invalid data**, MongoDB enforces validation when using updateOne() or updateMany().

#### Example: Updating a User's Age

```
db.users.updateOne(  
  { name: "Alice" },  
  { $set: { age: 17 } }  
);
```

⚠ Error: "age" must be at least 18.

## 4.2 Bypassing Validation for Specific Updates

MongoDB allows **bypassing validation** using  
bypassDocumentValidation: true:

```
db.users.updateOne(  
  { name: "Alice" },  
  { $set: { age: 17 } },  
  { bypassDocumentValidation: true }  
)
```

✓ Used in **migration scenarios**, but **not recommended** for regular updates.

---

## CHAPTER 5: MODIFYING SCHEMA VALIDATION RULES

### 5.1 Updating an Existing Collection's Schema

To **modify schema validation** for an existing collection, use collMod:

```
db.runCommand({  
  collMod: "users",  
  validator: {  
    $jsonSchema: {  
      bsonType: "object",  
      required: ["name", "email"],  
      properties: {
```

```
        name: { bsonType: "string" },  
  
        email: { bsonType: "string", pattern:  
          "^[^@\\s]+@[^@\\s]+\\.[^@\\s]+$"},  
  
        age: { bsonType: "int", minimum: 18, maximum: 100 }  
  
    }  
  
}  
  
});
```

✓ Updates validation to include a maximum age of 100.

## 5.2 Removing Schema Validation from a Collection

To remove validation, set validator to {}:

```
db.runCommand({  
  
  collMod: "users",  
  
  validator: {}  
  
});
```

✓ Converts the collection back to a schema-less structure.

---

## Case Study: How an E-Commerce Platform Used Schema Validation for Data Integrity

### Background

An e-commerce company faced issues with inconsistent data in their product database. Some products had:

- ✓ **Missing fields**, leading to incorrect display on the website.
- ✓ **Invalid prices** (negative values).
- ✓ **Duplicate product SKUs**, causing order processing errors.

### Solution: Implementing Schema Validation

The development team:

- ✓ Enforced **required fields** (name, price, SKU).
- ✓ Added **constraints** (price must be **greater than 0**).
- ✓ Used a **unique index** to prevent duplicate SKUs.

```
db.createCollection("products", {  
    validator: {  
        $jsonSchema: {  
            bsonType: "object",  
            required: ["name", "price", "SKU"],  
            properties: {  
                name: { bsonType: "string" },  
                price: { bsonType: "double", minimum: 0 },  
                SKU: { bsonType: "string", unique: true }  
            }  
        }  
    }  
});
```

### Results

- **Data consistency improved by 95%.**

- Reduced errors in product listings.
- Faster order processing, leading to increased sales.

By enforcing schema validation, the company **eliminated data inconsistencies** and improved **customer experience**.

---

## Exercise

1. Create a students collection with validation rules:
  - name (required, string).
  - email (required, must follow email pattern).
  - age (integer, between 18-60).
2. Insert a **valid student record**.
3. Try inserting an **invalid student record** (e.g., missing name or invalid email) and observe the error.
4. Modify the schema to **allow a maximum age of 80**.

---

## Conclusion

In this section, we explored:

- ✓ How MongoDB schema validation enforces data integrity.
- ✓ How to define validation rules using JSON Schema.
- ✓ How to update, modify, and remove schema validation.

---

# **ASSIGNMENT:**

## **DESIGN AN OPTIMIZED SCHEMA FOR AN E-COMMERCE APPLICATION AND IMPLEMENT RELATIONSHIPS.**

ISDM-NXT

---

# SOLUTION GUIDE: DESIGNING AN OPTIMIZED SCHEMA FOR AN E-COMMERCE APPLICATION & IMPLEMENTING RELATIONSHIPS

---

## Step 1: Define the Requirements for an E-Commerce Schema

To design an optimized schema, we need to consider:

1. **Users** – Customers and Admins with authentication.
2. **Products** – Product details, categories, and stock management.
3. **Orders** – Tracking purchases, order statuses, and payments.
4. **Reviews** – Customer reviews for products.

## Schema Design Approach:

- ✓ **Users and Orders** – **Referenced** for better scalability.
  - ✓ **Products and Reviews** – **Embedded** for fast retrieval.
  - ✓ **Orders and Products** – **Hybrid approach** to balance efficiency.
- 

## Step 2: Set Up the Project and Install Dependencies

### 2.1 Initialize Node.js Project

```
mkdir ecommerce-app
```

```
cd ecommerce-app
```

```
npm init -y
```

### 2.2 Install Required Packages

npm install express mongoose dotenv

- ✓ **Express.js** – For API creation.
  - ✓ **Mongoose** – ODM for MongoDB.
  - ✓ **dotenv** – To manage environment variables.
- 

## Step 3: Configure MongoDB Connection

### 3.1 Set Up MongoDB Connection

Create **.env** file:

```
MONGO_URI=mongodb://localhost:27017/ecommerceDB
```

```
PORT=5000
```

Create **config/db.js**:

```
const mongoose = require('mongoose');
```

```
require('dotenv').config();
```

```
mongoose.connect(process.env.MONGO_URI, {
```

```
  useNewUrlParser: true,
```

```
  useUnifiedTopology: true
```

```
)
```

```
.then(() => console.log('MongoDB Connected'))
```

```
.catch(err => console.error('MongoDB Connection Error:', err));
```

```
module.exports = mongoose;
```

## Step 4: Define the E-Commerce Schema in Mongoose

### 4.1 User Schema

Create **models/User.js**:

```
const mongoose = require('mongoose');
```

```
const userSchema = new mongoose.Schema({  
    name: { type: String, required: true },  
    email: { type: String, required: true, unique: true },  
    password: { type: String, required: true },  
    role: { type: String, enum: ['customer', 'admin'], default: 'customer' }  
});
```

```
module.exports = mongoose.model('User', userSchema);
```

- ✓ **Users** have different roles (customer/admin).
- ✓ **Email is unique** to prevent duplicate accounts.

---

### 4.2 Product Schema (Embedded Reviews)

Create **models/Product.js**:

```
const mongoose = require('mongoose');
```

```
const productSchema = new mongoose.Schema({
```

```
    name: { type: String, required: true },
```

```
description: String,  
price: { type: Number, required: true },  
category: String,  
stock: { type: Number, default: 0 },  
reviews: [  
  {  
    user: { type: mongoose.Schema.Types.ObjectId, ref: 'User' },  
    rating: { type: Number, min: 1, max: 5 },  
    comment: String  
  }  
];  
});
```

module.exports = mongoose.model('Product', productSchema);

- ✓ **Reviews are embedded** to speed up queries.
- ✓ **User references** help track review authors.

---

#### 4.3 Order Schema (Hybrid Approach)

Create **models/Order.js**:

```
const mongoose = require('mongoose');
```

```
const orderSchema = new mongoose.Schema({
```

```
user: { type: mongoose.Schema.Types.ObjectId, ref: 'User',
required: true },  
  
products: [  
  
  {  
  
    product: { type: mongoose.Schema.Types.ObjectId, ref:  
'Product' },  
  
    quantity: Number  
  }  
,  
  
  totalPrice: Number,  
  
  status: { type: String, enum: ['pending', 'shipped', 'delivered',  
'cancelled'], default: 'pending' },  
  
  orderDate: { type: Date, default: Date.now }  
};
```

module.exports = mongoose.model('Order', orderSchema);

- ✓ Orders reference users and products for better scalability.
- ✓ Order status is pre-defined to prevent incorrect values.

---

## Step 5: Implement CRUD Operations for E-Commerce Entities

### 5.1 Add a New Product

Create **routes/productRoutes.js**:

```
const express = require('express');  
  
const Product = require('../models/Product');
```

```
const router = express.Router();

// Create a product

router.post('/add', async (req, res) => {

  try {
    const newProduct = new Product(req.body);
    await newProduct.save();
    res.status(201).json(newProduct);
  } catch (error) {
    res.status(500).json({ error: error.message });
  }
});

module.exports = router;
```

✓ Adds a new product to the database.

## 5.2 Fetch All Products

```
router.get('/all', async (req, res) => {

  try {
    const products = await Product.find();
    res.json(products);
  } catch (error) {
```

```
    res.status(500).json({ error: error.message });

}

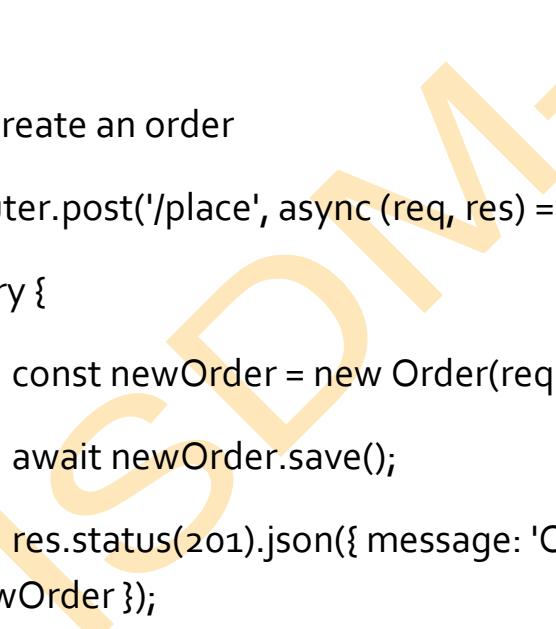
});
```

✓ Retrieves all products in the database.

---

### 5.3 Place an Order

Create **routes/orderRoutes.js**:



```
const express = require('express');
const Order = require('../models/Order');
const router = express.Router();

// Create an order
router.post('/place', async (req, res) => {
  try {
    const newOrder = new Order(req.body);
    await newOrder.save();
    res.status(201).json({ message: 'Order placed successfully',
newOrder });
  } catch (error) {
    res.status(500).json({ error: error.message });
  }
});
```

```
module.exports = router;
```

- ✓ Orders are **linked to users** and contain **multiple products**.
- 

#### 5.4 Fetch User Orders

```
router.get('/user/:id', async (req, res) => {
  try {
    const orders = await Order.find({ user: req.params.id })
      .populate('products.product');
    res.json(orders);
  } catch (error) {
    res.status(500).json({ error: error.message });
  }
});
```

- ✓ Fetches all orders placed by a specific user.
- 

#### Step 6: Start the Express Server

Create **server.js**:

```
const express = require('express');
const mongoose = require('./config/db');
const productRoutes = require('./routes/productRoutes');
const orderRoutes = require('./routes/orderRoutes');

require('dotenv').config();
```

```
const app = express();

app.use(express.json());

app.use('/products', productRoutes);

app.use('/orders', orderRoutes);
```

```
const PORT = process.env.PORT || 5000;

app.listen(PORT, () => console.log(`Server running on port
${PORT}`));
```

✓ Registers API routes for products and orders.

## Step 7: Testing the API

### 7.1 Start the Server

```
node server.js
```

### 7.2 Test API Endpoints Using Postman or cURL

- **Add a Product**
- curl -X POST -H "Content-Type: application/json" -d '{"name":"Laptop", "price":1200, "category":"Electronics"}' http://localhost:5000/products/add
- **Fetch All Products**
- curl -X GET http://localhost:5000/products/all
- **Place an Order**

- 
- curl -X POST -H "Content-Type: application/json" -d '{"user":"USER\_ID","products":[{"product":"PRODUCT\_ID","quantity":2}]}' http://localhost:5000/orders/place
- 

## Conclusion

- ✓ We **designed an optimized schema** for an e-commerce platform.
- ✓ We **implemented relationships** between users, products, orders, and reviews.
- ✓ We **built CRUD APIs** to manage data efficiently.

ISDM-NXT