



#### ISDM (INDEPENDENT SKILL DEVELOPMENT MISSION

# AWS DATABASE SERVICES

# AMAZON RDS (RELATIONAL DATABASE SERVICE) – STUDY MATERIAL

#### Introduction to Amazon RDS

#### What is Amazon RDS?

Amazon RDS (Relational Database Service) is a fully managed relational database service that simplifies the setup, operation, and scaling of databases in the AWS cloud. It automates database management tasks, including provisioning, backups, patching, and scaling, reducing the burden on database administrators.

#### Key Features of Amazon RDS

- ✓ Fully Managed AWS handles database maintenance, backups, and patching.
- ✓ **High Availability** Supports **Multi-AZ deployments** for failover protection.
- ✓ Automatic Backups Enables point-in-time recovery for data protection.
- ✓ **Scalability** Easily scale storage and compute resources as needed.
- ✓ Security Supports VPC, IAM, KMS encryption, and network

#### isolation.

- ✓ Monitoring & Performance Optimization Integrated with CloudWatch, Performance Insights, and Enhanced Monitoring.
- ✓ Multi-Region Replication Supports read replicas for performance and disaster recovery.

CHAPTER 1: SUPPORTED DATABASE ENGINES IN AMAZON RDS
Amazon RDS supports multiple database engines to accommodate different application needs:

Database	Best For	Features
Engine		
Amazon	High-performance	5x faster than MySQL,
Aurora	cloud-native database	automatic failover, auto-
		scaling
MySQL	Web applications,	Replication, high
	startups, open-source	availability, cost-
	projects	effective
PostgreSQL	Advanced analytics,	Supports JSON, GIS, and
	enterprise workloads	complex queries
MariaDB	MySQL-compatible,	Replication, clustering,
	open-source	high performance
	applications	
Microsoft SQL	Enterprise	Windows
Server	applications, .NET	authentication, SQL
	workloads	Server features

Oracle	Enterprise applications	Supports Oracle Data
	requiring Oracle	Guard, licensing
	compatibility	flexibility

#### \* Example:

A startup uses Amazon Aurora for its high-performance SaaS application, benefiting from automatic scaling and fault tolerance.

#### CHAPTER 2: HOW AMAZON RDS WORKS

#### 1. RDS Architecture Overview

- Primary Instance Main database handling read/write operations.
- Multi-AZ Replica Standby instance for failover protection.
- Read Replicas Additional replicas to handle high-read traffic.
- Storage Uses GP3, IO1, or Magnetic storage for different performance needs.

# 2. Amazon RDS Deployment Models

- **Single-AZ Deployment** Runs in one Availability Zone (cheaper, but no failover).
- Multi-AZ Deployment Creates standby instance in another
   AZ for high availability.
- Read Replicas Supports read scaling and cross-region replication.

# Example:

A financial company uses Multi-AZ RDS for PostgreSQL to ensure database reliability and automatic failover.

#### CHAPTER 3: SETTING UP AN AMAZON RDS DATABASE

#### Step 1: Create an Amazon RDS Instance

- Go to AWS Console → Open Amazon RDS.
- 2. Click "Create Database".
- 3. Select **Standard Create**.

#### Step 2: Choose Database Engine

- Choose a database engine (e.g., MySQL, PostgreSQL, or Aurora).
- 2. Select Engine Version (latest stable version recommended).

#### **Step 3: Configure Database Instance**

- DB Instance Identifier: my-rds-instance.
- 2. Username: admin.
- 3. Password: Set a strong password.
- 4. Choose **DB Instance Class** (e.g., db.t<sub>3</sub>.micro for small workloads).

# Step 4: Storage and Connectivity

- Storage Type: Choose GP3 for general workloads.
- Allocate Storage: Start with 20GB (auto-scale enabled).
- 3. **VPC & Subnet:** Choose an existing **VPC and Public/Private**Subnet.
- Enable Multi-AZ Deployment (Recommended for production).

# Step 5: Configure Security and Monitoring

1. **Enable Encryption** (for data security).

- 2. **Set up IAM authentication** for secure access.
- 3. Enable CloudWatch Monitoring.

#### **Step 6: Create and Connect to RDS Instance**

- 1. Click "Create Database".
- 2. Once the instance is available, copy the endpoint URL.
- 3. Connect to RDS using MySQL CLI:
- 4. mysql -h my-rds-instance.cndpjzabc123.us-east-1.rds.amazonaws.com -u admin -p
- \* Expected Outcome: Successfully created and connected to an RDS database.

CHAPTER 4: MANAGING AMAZON RDS (BACKUPS, SCALING, SECURITY)

- 1. Automated Backups & Snapshots
  - Automated Backups: AWS automatically backs up data and transaction logs for point-in-time recovery.
  - Manual Snapshots: Users can create on-demand backups and restore from snapshots.
  - Restore from Backup:
  - aws rds restore-db-instance-from-db-snapshot --db-instanceidentifier my-restored-db --db-snapshot-identifier my-dbsnapshot

#### 2. Scaling Amazon RDS

 Vertical Scaling: Increase CPU, RAM, and storage by modifying the instance size.  Horizontal Scaling: Use Read Replicas for read-heavy workloads.

#### 3. Security Best Practices for Amazon RDS

- ✓ Enable IAM authentication instead of traditional credentials.
- ✓ Restrict access with Security Groups and VPC settings.
- ✓ Encrypt data at rest using AWS KMS.
- ✓ Monitor database performance with Amazon CloudWatch.

# **\*** Example:

A healthcare application stores patient records in Amazon RDS with KMS encryption to meet HIPAA compliance.

CHAPTER 5: MONITORING & TROUBLESHOOTING AMAZON RDS

1. Monitoring with Amazon CloudWatch

CloudWatch provides key RDS metrics, including:

- ✓ CPU Utilization Tracks database performance.
- ✓ Free Storage Space Prevents storage exhaustion.
- ✓ **Read/Write IOPS** Monitors database query performance.
- ✓ **Database Connections** Tracks concurrent user sessions.

# \* Example:

An e-commerce platform sets up CloudWatch Alarms to receive alerts if CPU usage exceeds 80%.

# 2. Troubleshooting Common Issues

Issue	Solution
High CPU Utilization	Scale up the instance, optimize queries,
	or use read replicas

Slow Query	Enable <b>Performance Insights</b> , use
Performance	indexing
Database Connectivity	Check security groups, VPC settings, and
Issues	IAM policies
Storage Full	Enable auto-scaling or manually increase
	storage

#### \* Example:

A data analytics company uses Amazon RDS Performance Insights to detect slow SQL queries and optimize database performance.

#### CHAPTER 6: AMAZON RDS USE CASES

✓ E-Commerce Websites – Stores product catalogs, customer orders, and transaction data.

✓ SaaS Applications – Uses multi-tenant RDS databases for scalable web apps.

✓ Financial Systems – Implements Multi-AZ RDS for high availability.

✓ IoT & Big Data Analytics – Stores real-time sensor data in RDS.

★ Case Study: Netflix Uses Amazon RDS

Problem: Needed scalable, fault-tolerant databases.

**Solution:** Implemented **Amazon Aurora** with read replicas.

Outcome: 99.99% uptime and cost efficiency.

CONCLUSION: MASTERING AMAZON RDS

By using **Amazon RDS**, businesses can:

Deploy fully managed relational databases with ease.

- Ensure high availability with Multi-AZ.
- Scale seamlessly based on demand.
- Secure and monitor databases with AWS tools.

#### FINAL EXERCISE:

- 1. Create an RDS MySQL database, connect to it, and create a test table.
- 2. **Set up a read replica** and measure query performance improvements.
- 3. **Implement a CloudWatch alarm** for high CPU utilization on an RDS instance.

# AMAZON DYNAMODB (NOSQL DATABASE) – STUDY MATERIAL

#### INTRODUCTION TO AMAZON DYNAMODB

#### What is Amazon DynamoDB?

Amazon **DynamoDB** is a **fully managed NoSQL database service** designed for high availability, scalability, and low-latency performance. It provides **key-value and document-based storage** with **automatic scaling**, making it ideal for modern applications requiring real-time performance.

#### **Key Features of Amazon DynamoDB**

- ✓ Fully Managed No need to manage servers, backups, or scaling.
- ✓ High Availability & Durability Data is automatically replicated across multiple Availability Zones.
- ✓ Performance at Scale Handles millions of requests per second with single-digit millisecond latency.
- ✓ Flexible Data Model Supports key-value and document-based storage.
- ✓ On-Demand & Provisioned Capacity Choose between pay-perrequest or fixed throughput.
- ✓ Security & Compliance Supports encryption, IAM authentication, and VPC integration.
- ✓ Event-Driven Architecture Integrates with AWS Lambda, Kinesis, and S<sub>3</sub> for real-time data processing.

#### CHAPTER 1: DYNAMODB ARCHITECTURE & CORE CONCEPTS

#### 1. How DynamoDB Works

Tables: Collections of data with no fixed schema.

- Items: Similar to rows in SQL databases.
- Attributes: Individual fields within an item.
- Primary Keys: Unique identifiers for items (either Partition Key or Partition + Sort Key).

#### 2. DynamoDB Key Design

- Partition Key (Hash Key): A unique identifier that determines which partition the item is stored in.
- Composite Key (Partition + Sort Key): Enables querying on multiple attributes.

# Example:

#### A customer database uses:

- CustomerID (Partition Key) to uniquely identify customers.
- OrderID (Sort Key) to store multiple orders under one customer.

# 3. Data Consistency in DynamoDB

- Eventually Consistent Reads: Faster reads with possible slight delays in data propagation.
- Strongly Consistent Reads: Ensures the most up-to-date data but may increase latency.

#### CHAPTER 2: CREATING A DYNAMODB TABLE

# Step 1: Navigate to DynamoDB Console

- 1. Open AWS Management Console.
- 2. Search for **DynamoDB** and open the service.
- 3. Click "Create Table".

#### Step 2: Define Table Schema

- 1. Table Name: CustomerOrders.
- 2. Primary Key:
  - Partition Key: CustomerID (String).
  - Sort Key (Optional): OrderID (String).
- 3. Choose "On-Demand Capacity Mode" (auto-scaling).
- 4. Click "Create Table".

#### Step 3: Insert Sample Data

- Click "Explore Items" → "Create Item".
- 2. Add sample data:
- 3. {
- 4. "CustomerID": "CUST1001",
- "OrderID": "ORD5001",
- "ProductName": "Laptop",
- 7. "OrderDate": "2025-02-20"
- 8. }
- Click "Save Item".
- **Expected Outcome:** The item is successfully stored in **DynamoDB**.

# CHAPTER 3: QUERYING & SCANNING DATA IN DYNAMODB

1. Querying Items Using Primary Key

- Use Partition Key and optionally Sort Key to retrieve items efficiently.
- Example query: Find all orders for CustomerID = "CUST1001".

#### 2. Scanning Items (Full Table Search)

- Scans every item in the table (less efficient than queries).
- Use **filters to refine results** (e.g., ProductName = "Laptop").

# **\*** Example:

Retrieve all orders for CUST1001 using AWS CLI:

aws dynamodb query --table-name CustomerOrders \

- --key-condition-expression "CustomerID = :cust" \
- --expression-attribute-values '{":cust":{"S":"CUST1001"}}'

#### CHAPTER 4: INDEXING & PERFORMANCE OPTIMIZATION

# Global Secondary Index (GSI)

- Allows querying on attributes other than the primary key.
- Example: Create a GSI on ProductName to find all orders for a specific product.

#### 2. Local Secondary Index (LSI)

- Uses the same Partition Key but a different Sort Key for faster lookups.
- Example: Store OrderDate as an LSI to retrieve recent orders for a customer.

#### 3. Choosing Between GSIs & LSIs

Feature	Global Secondary Index (GSI)	Local Secondary Index (LSI)
Partition	Different from the	Same as the base table
Key	base table	
Sort Key	Can be any attribute	Must be a different attribute
		from the base table
Use Case	Querying data by	Sorting within a partition
	alternative fields	

# **\*** Example:

A retail store creates a GSI on "ProductName" to quickly find all orders for a specific product.

CHAPTER 5: DYNAMODB SECURITY & BEST PRACTICES

- 1. Security Best Practices
- ✓ Enable Encryption at Rest using AWS KMS.
- ✓ Use IAM Roles & Policies to restrict access.
- ✓ Enable VPC Endpoints for private connectivity.
- 2. Cost Optimization Best Practices
- ✓ Use On-Demand Mode for unpredictable workloads.
- ✓ Use Provisioned Mode for steady traffic with auto-scaling enabled.
- ✓ Delete unused Global Secondary Indexes (GSIs) to reduce costs.

# \* Example:

A **fintech company** encrypts DynamoDB **with AWS KMS** and limits access via **IAM roles** to protect sensitive data.

# CHAPTER 6: DYNAMODB STREAMS & EVENT-DRIVEN ARCHITECTURES

#### 1. What is DynamoDB Streams?

- Captures data changes in real-time.
- Integrates with AWS Lambda for event-driven applications.

#### 2. Use Case: Serverless Order Processing

- 1. New order is inserted into DynamoDB.
- 2. DynamoDB Streams triggers AWS Lambda.
- Lambda processes the order & updates inventory.

# \* Example:

An e-commerce store uses DynamoDB Streams + Lambda to automatically process orders in real-time.

# CHAPTER 7: REAL-WORLD USE CASES OF DYNAMODB

- ✓ E-Commerce Applications Store product catalogs, orders, and user sessions.
- √ Gaming Leaderboards Track real-time player scores.
- ✓ IoT Data Storage Store sensor data for real-time analytics.
- ✓ Social Media Apps Handle user posts, likes, and comments at scale.
- 📌 Case Study: Amazon Uses DynamoDB

**Problem:** Amazon needed a **high-speed NoSQL database** to handle **Black Friday sales traffic**.

Solution: Used DynamoDB with auto-scaling & caching.

Outcome: Reduced latency & handled millions of transactions per second.

CONCLUSION: MASTERING AMAZON DYNAMODB

By using Amazon DynamoDB, businesses can:

- Store & retrieve data at scale with high availability.
- Ensure low-latency performance with GSIs & caching.
- Implement event-driven architectures with DynamoDB Streams.

#### FINAL EXERCISE:

- Create a DynamoDB table, add sample data, and query using AWS CLI.
- 2. Implement a Global Secondary Index (GSI) to optimize queries.
- 3. Use **DynamoDB Streams with AWS Lambda** to **automate** order processing.

# AMAZON AURORA – STUDY MATERIAL

#### INTRODUCTION TO AMAZON AURORA

#### What is Amazon Aurora?

Amazon Aurora is a fully managed, high-performance relational database service designed for enterprise applications, SaaS solutions, and mission-critical workloads. Aurora is compatible with MySQL and PostgreSQL while providing up to 5x the performance of standard MySQL and 3x the performance of standard PostgreSQL.

#### **Key Features of Amazon Aurora**

- ✓ High Performance Delivers low-latency, high-throughput performance.
- ✓ Fault Tolerance & High Availability Data is automatically replicated across multiple Availability Zones (Multi-AZ).
- ✓ Auto-Scaling Storage Expands from 10GB to 128TB automatically.
- ✓ Fully Managed by AWS Automates backups, patching, and monitoring.
- ✓ Security & Compliance Supports encryption, IAM authentication, and VPC isolation.
- ✓ Global Database Support Enables multi-region replication for disaster recovery.
- ✓ Serverless Option (Aurora Serverless) Scales automatically based on demand.

#### CHAPTER 1: AMAZON AURORA VS. OTHER DATABASES

1. Aurora vs. RDS MySQL & PostgreSQL

Feature	Amazon Aurora	RDS MySQL	RDS
			PostgreSQL
Performance	5x MySQL, 3x	Standard	Standard
	PostgreSQL	MySQL	PostgreSQL
Replication	Multi-AZ, Multi-	Read	Read Replicas
	Region	Replicas	
Scalability	Auto-scales up to	Manual	Manual scaling
	128TB	scaling	
Failover	<30 seconds	Minutes	Minutes
Time			
Backup	Continuous,	Scheduled	Scheduled
	automated		

- 2. Aurora vs. Traditional Databases (Oracle, SQL Server)
- ✓ No licensing costs Uses open-source MySQL/PostgreSQL.
- ✓ Fully managed No need for manual backups or clustering.
- ✓ Better scalability Grows dynamically without downtime.

# \* Example:

A large e-commerce website migrates from Oracle to Aurora PostgreSQL for cost savings and better performance.

#### CHAPTER 2: AMAZON AURORA ARCHITECTURE & COMPONENTS

# 1. Aurora Cluster Components

- Primary Instance Handles read and write operations.
- Read Replicas Up to 15 read replicas for load balancing.
- Aurora Storage Layer Automatically scales from 10GB to 128TB.

#### 2. High Availability & Replication

- Multi-AZ Replication Replicates six copies of data across three Availability Zones.
- Aurora Global Database Allows low-latency reads across multiple regions.
- Failover Support Automatically switches to a read replica in case of primary failure.

# \* Example:

A fintech company uses Aurora Global Database to replicate customer transactions across the US & Europe.

#### CHAPTER 3: SETTING UP AMAZON AURORA

#### Step 1: Create an Aurora Database Cluster

- Open AWS Console → Go to RDS Dashboard.
- 2. Click "Create Database".
- 3. Choose **Amazon Aurora** as the database engine.
- 4. Select "Aurora MySQL" or "Aurora PostgreSQL".
- 5. Choose "Production" or "Dev/Test" environment.

#### Step 2: Configure the Cluster

- 1. DB Cluster Identifier: my-aurora-cluster.
- 2. **Username:** admin, **Password:** mypassword123.
- 3. Choose Instance Class:
  - db.t3.medium for development.
  - o db.r5.large for production.

#### Step 3: Enable High Availability & Security

- Select Multi-AZ Deployment.
- 2. Enable Encryption (AWS KMS).
- Set Public Access = No (for security).
- 4. Configure **Security Groups** to allow **only specific IPs**.

#### Step 4: Launch & Connect to Aurora

- 1. Click "Create Database".
- 2. Once available, copy the **Endpoint URL**.
- 3. Connect using MySQL CLI:
- 4. mysql -h my-aurora-cluster.cluster-abc123.us-east-1.rds.amazonaws.com -u admin -p
- **Expected Outcome:** Successfully created and connected to an Aurora database cluster.

# CHAPTER 4: SCALING, PERFORMANCE, AND OPTIMIZATION

- 1. Auto-Scaling Aurora Storage & Compute
- ✓ Storage Auto-Scaling Expands from 10GB to 128TB automatically.
- ✓ Aurora Auto-Scaling Read Replicas Automatically adds or removes read replicas based on traffic.
- 2. Performance Insights & Optimization
  - Enable Performance Insights to analyze slow queries.
  - Use Aurora Query Plan Caching to speed up repetitive queries.

• Implement Connection Pooling for high-traffic workloads.

# **\*** Example:

A gaming company uses Aurora Auto-Scaling to handle traffic spikes during peak hours.

CHAPTER 5: AURORA SERVERLESS – ON-DEMAND SCALING

#### 1. What is Aurora Serverless?

- Auto-scales compute capacity based on demand.
- No need for provisioning fixed instances.
- Cost-efficient for intermittent workloads.

#### 2. Setting Up Aurora Serverless

- 1. In the RDS Console, click "Create Database".
- 2. Choose Amazon Aurora → Select Aurora Serverless.
- 3. Configure Auto-Scaling Minimum & Maximum Capacity (e.g., 1–10 ACUs).
- 4. Click "Create Database".

# **Example:**

A start-up with unpredictable traffic uses Aurora Serverless to reduce costs by automatically scaling.

#### CHAPTER 6: SECURITY & BACKUP STRATEGIES IN AURORA

# 1. Security Best Practices

- ✓ Enable IAM Authentication for passwordless access.
- ✓ Use Security Groups & VPC Peering for private access.
- ✓ Enable Multi-Factor Authentication (MFA) for database access.

#### 2. Backup & Disaster Recovery

- Automated Backups Retains up to 35 days of backups.
- Manual Snapshots Take snapshots before updates
- Aurora Global Database Replicates data across regions for DR.

# **\*** Example:

A healthcare company uses Aurora Global Database to store patient records across continents securely.

#### CHAPTER 7: AMAZON AURORA USE CASES & CASE STUDY

- ✓ E-Commerce Applications Handles large-scale transactions and analytics.
- ✓ Financial & Banking Services Supports secure and high-speed processing.
- ✓ SaaS Applications Powers multi-tenant applications efficiently.
- ✓ AI/ML & Analytics Works with AWS Glue, Redshift, and S3.
- CASE STUDY: NETFLIX USES AURORA

**Problem:** Needed a **highly scalable relational database** for content personalization.

Solution: Migrated from traditional MySQL to Aurora MySQL. Outcome: Improved performance, reduced costs, and better scalability.

**CONCLUSION: MASTERING AMAZON AURORA** 

By using Amazon Aurora, businesses can:

- Improve performance compared to traditional RDS databases.
- Scale dynamically with Aurora Auto-Scaling & Serverless.
- Ensure high availability with Multi-AZ & Global Replication.
- Enhance security with IAM authentication and KMS encryption.

#### FINAL EXERCISE:

- Create an Amazon Aurora MySQL database, connect to it, and insert test data.
- 2. **Enable Auto-Scaling Read Replicas** and measure performance improvements.
- 3. **Set up a global Aurora cluster** and test cross-region replication.

# **AWS SECURITY ESSENTIALS**

# AWS SHARED RESPONSIBILITY MODEL – STUDY MATERIAL

INTRODUCTION TO THE AWS SHARED RESPONSIBILITY MODEL What is the AWS Shared Responsibility Model?

The **AWS Shared Responsibility Model** defines who is responsible for securing different aspects of the cloud environment:

- AWS is responsible for the security of the cloud (infrastructure, hardware, networking).
- Customers are responsible for security in the cloud (data, applications, IAM, encryption).

This model ensures that both AWS and customers play a role in maintaining security and compliance.

Key Features of the AWS Shared Responsibility Model

- ✓ Clear separation of responsibilities between AWS and customers.
- ✓ AWS secures the global infrastructure (data centers, networking, hypervisors).
- ✓ Customers control and manage their cloud environment (data, IAM, OS security).
- √ Compliance with industry standards (ISO, SOC, PCI DSS, HIPAA).

CHAPTER 1: UNDERSTANDING THE AWS SHARED RESPONSIBILITY

MODEL

#### 1. AWS Responsibility: Security OF the Cloud

AWS ensures the security of:

- Physical Security Data centers, hardware, access controls.
- Networking Security Firewalls, DDoS protection, AWS Shield.
- Storage & Compute Infrastructure EC2, S3, RDS, Lambda security.
- ✓ Compliance & Certifications GDPR, ISO 27001, SOC 2, HIPAA.

#### 2. Customer Responsibility: Security IN the Cloud

Customers must secure:

- **✓ Data & Application Security** Encrypt data at rest and in transit.
- ✓ IAM & Access Control Manage permissions with IAM roles & policies.
- ✓ Operating System & Patching Keep OS and applications updated.
- ✓ **Network Configuration** Secure VPC, Security Groups, NACLs.
- **\*** Example:

If a company hosts a website on AWS EC2, AWS is responsible for server infrastructure, but the company must secure the OS, data, and user permissions.

#### CHAPTER 2: SHARED RESPONSIBILITY ACROSS AWS SERVICES

1. Infrastructure Services (EC2, S3, VPC, RDS, Lambda)

AWS Responsibility	Customer Responsibility

Physical security of data centers	Configure IAM roles & permissions
Securing hardware, storage, and networking	Secure OS, applications, and databases
Ensuring availability & redundancy	Encrypt data and manage backup policies

# Example:

- EC2 Instance Security → AWS secures the hardware, but users must install security patches and configure firewalls.
- S3 Bucket Security → AWS manages the storage infrastructure, but users must enable encryption and restrict public access.

# 2. AWS Managed Services (Lambda, DynamoDB, RDS, ECS, CloudFront)

AWS Responsibility	Customer Responsibility
Securing the runtime environment	Secure application code & API permissions
Automated patching & maintenance	Manage user access & authentication
Ensuring availability & scaling	Monitor logs for anomalies (CloudWatch, GuardDuty)

# \* Example:

In **AWS Lambda**, AWS manages the **underlying servers**, but the customer must ensure **secure API permissions & data encryption**.

#### **CHAPTER 3: SECURITY BEST PRACTICES FOR CUSTOMERS**

- 1. Identity & Access Management (IAM)
- ✓ Use IAM roles instead of root credentials.
- ✓ Enable Multi-Factor Authentication (MFA) for users.
- ✓ Apply least privilege access (grant only necessary permissions).
- 2. Data Protection
- ✓ Encrypt S<sub>3</sub> buckets, databases, and EC<sub>2</sub> storage (AWS KMS).
- ✓ Enable automatic backups and versioning for data recovery.
- ✓ Use **AWS Secrets Manager** to store API keys & passwords.
- 3. Network Security
- ✓ Restrict access using VPC Security Groups and NACLs.
- ✓ Enable AWS Shield & WAF to protect against DDoS attacks.
- ✓ Implement CloudTrail & GuardDuty to monitor security events.

# **\*** Example:

A finance company using AWS RDS should enable encryption, restrict public access, and use IAM authentication.

#### CHAPTER 4: AWS COMPLIANCE AND SHARED RESPONSIBILITY

#### 1. AWS Compliance Certifications

AWS maintains **global security and compliance** certifications, including:

- **ISO 27001, SOC 1/2/3** − Cloud security and operational excellence.
- ✓ HIPAA Protects healthcare data.
- PCI DSS Secures payment processing systems.
- GDPR Ensures data privacy for European users.

- 2. Customer Responsibilities for Compliance
- ✓ Enable AWS Config & Audit Logs to track changes.
- ✓ Use AWS Organizations to enforce security policies across accounts.
- ✓ Follow AWS Well-Architected Security Framework for compliance.

# Example:

A hospital storing patient records in AWS must encrypt data, control IAM access, and follow HIPAA guidelines.

CHAPTER 5: REAL-WORLD USE CASES OF AWS SHARED RESPONSIBILITY MODEL

1. E-Commerce Website Security (Using EC2, S3, RDS, CloudFront)

#### √ AWS Responsibilities:

- Secures data center infrastructure & networking.
- Ensures high availability of EC2, S3, and CloudFront.

# ✓ Customer Responsibilities:

- Encrypt S<sub>3</sub> bucket storing user orders.
- **Secure EC2 instance** with a firewall (Security Groups).
- Implement IAM roles for developers and support teams.
- Banking & Financial Application Security (Using RDS, Lambda, S3, VPC)

#### √ AWS Responsibilities:

- Secures RDS database infrastructure with automated failover.
- Provides DDoS protection via AWS Shield.

#### ✓ Customer Responsibilities:

- Restrict database access using IAM roles & Security Groups.
- Enable data encryption (KMS) for S<sub>3</sub> backups.
- Monitor transactions with AWS CloudTrail.

#### **\*** Example:

A bank uses Amazon RDS for financial transactions but must enforce strong IAM policies and encryption.

#### 3. Healthcare Application (HIPAA-Compliant AWS Setup)

#### **✓** AWS Responsibilities:

- Provides HIPAA-compliant infrastructure.
- Ensures data encryption and compliance certifications.

#### ✓ Customer Responsibilities:

- Encrypt patient data in Amazon RDS using AWS KMS.
- Restrict public access to sensitive information using IAM.
- Enable logging & monitoring (AWS CloudWatch, GuardDuty).

# 📌 Example:

A **telemedicine company** storing patient records in AWS must **encrypt data and follow HIPAA security policies**.

CONCLUSION: MASTERING THE AWS SHARED RESPONSIBILITY MODEL

By understanding the **AWS Shared Responsibility Model**, businesses can:

- Ensure a secure cloud environment by managing responsibilities effectively.
- Encrypt and protect their own data while AWS secures the infrastructure.
- Follow compliance standards like HIPAA, PCI DSS, and GDPR.
- Implement best security practices using IAM, encryption, and monitoring tools.

#### FINAL EXERCISE:

- List five AWS responsibilities and five customer responsibilities for securing an EC2 instance.
- Configure an S3 bucket with encryption and IAM permissions to ensure security.
- Enable AWS CloudTrail and CloudWatch to monitor API activity in your AWS account.

# IAM POLICIES AND PERMISSIONS – STUDY MATERIAL

INTRODUCTION TO AWS IAM (IDENTITY AND ACCESS MANAGEMENT)

#### What is AWS IAM?

AWS Identity and Access Management (IAM) is a security service that enables you to manage users, groups, and roles, and control their permissions to AWS resources. It follows the principle of least privilege, ensuring that users only have the permissions necessary for their tasks.

#### **Key Features of AWS IAM**

- ✓ Fine-Grained Access Control Assign permissions to specific AWS resources.
- ✓ User Authentication & Authorization Securely manage users and
  roles.
- ✓ Multi-Factor Authentication (MFA) Adds an extra layer of security.
- ✓ Federated Access Allows SSO integration with external identity providers.
- ✓ **Policy-Based Permissions** Uses JSON-based policies to define permissions.
- ✓ **AWS** Organizations Integration Enforce security policies across multiple AWS accounts.

#### CHAPTER 1: IAM POLICIES AND PERMISSIONS – CORE CONCEPTS

# 1. IAM Users, Groups, and Roles

• IAM User – Represents an individual user (e.g., developero1).

- IAM Group Collection of users with shared permissions (e.g., Developers).
- IAM Role A temporary access identity used by AWS services or external users.

#### 2. IAM Policies – Defining Permissions

IAM policies are JSON-based permission rules that define who can access what in AWS.



# Example:

An S3 read-only policy allows users to view (but not modify) files in an S3 bucket.

#### CHAPTER 2: TYPES OF IAM POLICIES

# AWS-Managed Policies (Predefined by AWS)

AWS provides **predefined policies** for common use cases.

Policy Name	Purpose	
AdministratorAccess	Grants full access to AWS	
PowerUserAccess	Full access except IAM & security settings	
AmazonS3ReadOnlyAccess	Allows users to read but not modify S3 objects	
AmazonEC2FullAccess	Grants full control over EC2 instances	

#### 2. Customer-Managed Policies (Custom Policies)

- ✓ Created by AWS customers for **specific business needs**.
- ✓ Provides more control over permissions than AWS-managed policies.
- ✓ Useful for restricting access to specific AWS resources.

\*

Example:

A company creates a **custom IAM policy** to allow **developers to start and stop EC2 instances** but not delete them.

```
"Resource": "arn:aws:ec2:region:account-id:instance/*"
},

{

"Effect": "Deny",

"Action": "ec2:TerminateInstances",

"Resource": "arn:aws:ec2:region:account-id:instance/*"
}

3. Inline Policies (Attached Directly to Users, Groups, or Roles)

✓ Inline policies are attached directly to an IAM entity (user, group, or role).
```

✓ Unlike managed policies, inline policies cannot be reused across

CHAPTER 3: IAM POLICY STRUCTURE AND ELEMENTS IAM policies consist of five key elements:

✓ Best for one-time or highly specific permissions.

#### 1. Version

multiple

Specifies the **policy language version** (always set to "2012-10-17").

#### 2. Statement

Defines permissions within a policy.

#### 3. Effect

• "Allow" → Grants permissions.

groups.

"Deny" → Explicitly denies access.

#### 4. Action

Defines the **specific AWS service actions** (e.g., s3:PutObject, ec2:StartInstances).

#### 5. Resource

Specifies **which AWS resources** are affected (e.g., "arn:aws:s3:::my-bucket/\*").

**Example Policy:** Full access to EC2 but **denies terminating** instances.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
        "Effect": "Allow",
        "Action": "ec2:*",
        "Resource": "*"
    },
    {
        "Effect": "Deny",
        "Action": "ec2:TerminateInstances",
        "Resource": "*"
    }
}
```

}

#### CHAPTER 4: APPLYING IAM POLICIES TO AWS RESOURCES

# 1. Assigning a Policy to a User

- Open AWS IAM Console → Click "Users".
- 2. Select a user → Click "Permissions".
- Click "Attach Policy" → Select a policy (e.g., AmazonS<sub>3</sub>ReadOnlyAccess).
- 4. Click "Apply Changes".

#### 2. Assigning a Policy to a Role

- 1. Go to IAM Console → Click "Roles".
- 2. Select an existing role or create a new one.
- 3. Attach a policy (e.g., AmazonEC2FullAccess).
- 4. Save changes and assign the role to an EC2 instance or Lambda function.

#### CHAPTER 5: IAM BEST PRACTICES

- ✓ Follow the Principle of Least Privilege Grant only necessary permissions.
- ✓ Enable Multi-Factor Authentication (MFA) for all users.
- ✓ Use IAM Roles Instead of IAM Users for EC2, Lambda, and other AWS
  services.
- ✓ Rotate IAM Credentials Regularly Avoid long-term access keys.
- ✓ Use AWS Organizations to apply policies at the account level.
- ✓ Monitor IAM Activity with AWS CloudTrail.



#### **Example:**

A finance company enables MFA for all IAM users and uses AWS CloudTrail to monitor login attempts.

#### CHAPTER 6: REAL-WORLD USE CASES OF IAM POLICIES

#### 1. Restricting Developers to Read-Only Access

A company allows **developers to read AWS resources** but prevents them from modifying or deleting them.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "ec2:Describe*",
      "Resource": "*"
   },
    {
      "Effect": "Deny",
      "Action": ["ec2:StopInstances", "ec2:TerminateInstances"],
      "Resource": "*"
    }
  ]
}
```

#### 2. Granting Read/Write Access to a Specific S3 Bucket

A marketing team needs access to only one S<sub>3</sub> bucket.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
        "Effect": "Allow",
        "Action": ["s3:GetObject", "s3:PutObject"],
        "Resource": "arn:aws:s3:::marketing-bucket/*"
    }
}
```

#### 3. Temporary Access for Contractors Using IAM Roles

A company allows **external contractors** to temporarily **access AWS Lambda**.

- Uses IAM Roles instead of IAM Users.
- Sets session duration limit for security.

\*

Example:

A contractor accesses AWS **for 12 hours using an IAM role** but cannot modify policies.

```
{
    "Version": "2012-10-17",
    "Statement": [
```

```
"Effect": "Allow",

"Action": "lambda:InvokeFunction",

"Resource": "*"
}
]
```

CONCLUSION: MASTERING IAM POLICIES AND PERMISSIONS

By using AWS IAM, businesses can:

- Secure AWS resources with fine-grained access control.
- Protect user identities and enforce security best practices.
- Monitor and audit IAM activity for compliance.
- Ensure least-privilege access across AWS accounts.

#### FINAL EXERCISE:

- Create an IAM user and attach AmazonS3ReadOnlyAccess.
- 2. Create a custom policy to allow only EC2 start/stop actions.
- 3. Monitor IAM access logs using AWS CloudTrail.

## AWS KMS (KEY MANAGEMENT SERVICE) – STUDY MATERIAL

INTRODUCTION TO AWS KMS (KEY MANAGEMENT SERVICE)
What is AWS KMS?

AWS **Key Management Service (KMS)** is a **managed encryption service** that enables users to **create, store, and manage cryptographic keys** securely. It helps **encrypt and decrypt data** across AWS services and applications while maintaining **security, compliance, and auditing** capabilities.

#### **Key Features of AWS KMS**

- ✓ Fully Managed Key Storage AWS securely stores and manages encryption keys.
- ✓ Seamless Integration Works with S<sub>3</sub>, RDS, DynamoDB, Lambda, EBS, and other AWS services.
- ✓ **Key Rotation & Lifecycle Management** Automatically rotates keys for security.
- ✓ Access Control & IAM Integration Restricts key usage via IAM policies.
- ✓ Audit Logging with CloudTrail Logs key usage and access history.
- ✓ FIPS 140-2 Compliance Meets stringent security standards.
- ✓ **Multi-Region Keys** Enables cross-region encryption and replication.

CHAPTER 1: UNDERSTANDING AWS KMS CONCEPTS

#### 1. AWS KMS Key Types

AWS KMS supports different types of encryption keys:

Key Type	Description
Customer Managed Keys (CMK)	User-created keys with full control over permissions.
AWS Managed Keys	Automatically created and managed by AWS for specific services.
AWS Owned Keys	Used internally by AWS for encryption across services.
KMS Multi-Region Keys	Enable cross-region encryption for disaster recovery.

## **\*** Example:

A company encrypts sensitive customer data in Amazon S<sub>3</sub> using a Customer Managed Key (CMK) to control access.

#### 2. Symmetric vs. Asymmetric Keys in AWS KMS

Key Type Use Case
 Symmetric Encrypt & decrypt with the same key (used in most AWS services).
 Asymmetric Uses a public-private key pair for digital signatures & encryption.

## **\*** Example:

- Symmetric keys encrypt S<sub>3</sub> files, while
- Asymmetric keys verify digital signatures in blockchain applications.

#### CHAPTER 2: SETTING UP AWS KMS FOR ENCRYPTION

#### Step 1: Create a Customer Managed Key (CMK)

- Open AWS Console → Navigate to KMS (Key Management Service).
- 2. Click "Create a Key".
- 3. Select "Symmetric Key" for general encryption use.
- 4. Choose "Customer Managed Key (CMK)".
- Define a Key Alias (e.g., MyAppEncryptionKey).
- 6. Set **Key Administrators** (IAM users who can manage the key).
- 7. Set **Key Users** (IAM roles that can use the key for encryption).
- 8. Enable Automatic Key Rotation (Recommended).
- 9. Click "Create Key".
- **Expected Outcome:** A **CMK** is **created** and ready to be used for encryption.

#### Step 2: Encrypt and Decrypt Data Using AWS KMS

1. Encrypting Data with AWS KMS (CLI Command)

aws kms encrypt --key-id alias/MyAppEncryptionKey \

- --plaintext "SensitiveData" --output text --query CiphertextBlob
- 2. Decrypting Data with AWS KMS (CLI Command)

aws kms decrypt --ciphertext-blob fileb://encryptedfile \

--output text --query Plaintext

## \* Example:

An IoT device encrypts sensor data using KMS before sending it to an AWS database.

#### Step 3: Enable S3 Bucket Encryption with AWS KMS

- 1. Open Amazon S3 Console.
- 2. Select a **Bucket** → Click **Properties**.
- 3. Under **Default Encryption**, select **AWS KMS**.
- 4. Choose the Customer Managed Key (CMK) created earlier.
- 5. Click Save.
- \* Expected Outcome: All new objects in the S<sub>3</sub> bucket are automatically encrypted using AWS KMS.

#### CHAPTER 3: USING AWS KMS WITH OTHER AWS SERVICES

- 1. AWS KMS with Amazon RDS (Encrypting Databases)
  - AWS KMS encrypts RDS databases and snapshots.
  - Can be enabled during RDS instance creation.

## **Example:**

A banking application encrypts its PostgreSQL RDS database with a KMS CMK to protect customer transactions.

- 2. AWS KMS with Amazon DynamoDB (Encrypting NoSQL Data)
  - AWS KMS encrypts **DynamoDB tables at rest**.
  - Uses either AWS Managed Keys or CMK.

## 📌 Example:

A social media platform encrypts user profile data in DynamoDB using AWS KMS for compliance.

#### 3. AWS KMS with Amazon Lambda (Securing API Keys & Secrets)

AWS Lambda can decrypt KMS-encrypted environment variables.

## **\*** Example:

A serverless application decrypts database credentials before connecting to an RDS instance.

#### CHAPTER 4: AWS KMS SECURITY BEST PRACTICES

- ✓ Use Customer Managed Keys (CMKs) for full control over encryption policies.
- ✓ Enable Key Rotation to automatically generate new keys.
- ✓ Use IAM Policies & Key Policies to restrict access.
- ✓ Monitor Key Usage with AWS CloudTrail.
- ✓ Use Multi-Region Keys for Global Applications.

## **\*** Example:

A healthcare company follows HIPAA compliance by restricting access to encryption keys using IAM policies.

CHAPTER 5: MONITORING AND AUDITING AWS KMS WITH CLOUDTRAIL

AWS CloudTrail records **all AWS KMS API requests**, allowing users to monitor:

√ Who accessed encryption keys.

- √ When decryption events occurred.
- √ Which services used AWS KMS.

## **\*** Example:

A cybersecurity team detects unauthorized decryption attempts in CloudTrail logs and revokes IAM permissions.

CHAPTER 6: AWS KMS MULTI-REGION KEYS (CROSS-REGION ENCRYPTION)

- Multi-Region Keys enable disaster recovery and global applications.
- Allows encrypting data in one AWS region and decrypting in another.
- Useful for multi-region SaaS applications.

## **\*** Example:

A global financial app encrypts data in US-East-1 and decrypts in Europe for compliance.

CHAPTER 7: REAL-WORLD USE CASES OF AWS KMS

- ✓ E-Commerce & Payment Processing Encrypts credit card details before storage.
- ✓ Healthcare & Compliance (HIPAA, GDPR) Encrypts medical records & patient data.
- ✓ SaaS & Enterprise Applications Protects customer credentials & API keys.
- ✓ Government & Defense Ensures secure encryption for classified information.
- ★ CASE STUDY: AMAZON USES AWS KMS

**Problem:** Needed **secure key management** for S<sub>3</sub> & DynamoDB encryption.

Solution: Implemented AWS KMS with automated key rotation.

Outcome: Improved security, compliance, and scalability.

CONCLUSION: MASTERING AWS KMS

By using **AWS KMS**, businesses can:

- Securely encrypt sensitive data across AWS services.
- Manage cryptographic keys easily with full control.
- Meet compliance standards like PCI DSS, HIPAA, and GDPR.
- Monitor and audit key usage for security best practices.

#### FINAL EXERCISE:

- 1. Create a Customer Managed Key (CMK) in AWS KMS.
- 2. Encrypt a file using AWS KMS and decrypt it using CLI.
- 3. **Enable S3 Bucket Encryption** using a KMS key and upload a test file.

# AWS SHIELD & AWS WAF (WEB APPLICATION FIREWALL) – STUDY MATERIAL

## INTRODUCTION TO AWS SHIELD & AWS WAF What is AWS Shield?

AWS **Shield** is a **managed Distributed Denial of Service (DDoS) protection service** that safeguards applications running on AWS. It provides **automatic attack mitigation**, ensuring application availability and security.

AWS Shield comes in two versions:

- AWS Shield Standard Free, always-on protection against common DDoS attacks.
- AWS Shield Advanced Premium protection with real-time monitoring, cost protection, and attack response support.

#### What is AWS WAF (Web Application Firewall)?

AWS WAF (Web Application Firewall) protects web applications from malicious traffic, SQL injection, cross-site scripting (XSS), and bot attacks. It monitors, filters, and blocks HTTP(S) requests based on predefined security rules.

#### Key Features of AWS Shield & AWS WAF

- ✓ Real-Time DDoS Protection Detects & mitigates volumetric, state-exhaustion, and application-layer attacks.
- ✓ Managed Rules for Web Security Predefined WAF rules to block threats like SQL injection & XSS.
- ✓ Integration with AWS Services Works with CloudFront, ALB, API Gateway, and Route 53.

- ✓ Custom Security Rules Define rules to block specific IPs, countries, or bot patterns.
- ✓ Bot Control Mitigates automated bot traffic & scraper attacks.
- ✓ Threat Intelligence from AWS Advanced security insights with AWS Shield Advanced.

CHAPTER 1: AWS SHIELD - DDOS PROTECTION

#### 1. How AWS Shield Works

AWS Shield provides **always-on DDoS protection to** mitigate attacks targeting:

- Network Layer (Layer 3 & 4) Protects against UDP flood,
   SYN flood, and reflection attacks.
- Application Layer (Layer 7) Prevents HTTP floods and botbased attacks.

AWS Shield Standard	AWS Shield Advanced		
Free tier for all AWS	Paid service for <b>enterprise-level</b>		
customers	DDoS protection		
Protects against common	Real-time attack visibility &		
volumetric DDoS attacks	dedicated AWS response team		
Automatic attack mitigation	Cost protection against		
	unexpected DDoS-related		
	expenses		
Works with CloudFront, ALB,	Works with AWS WAF for		
Route 53	advanced threat intelligence		

#### 📌 Example:

An **e-commerce website** uses **AWS Shield Standard** to prevent **DDoS attacks** from affecting checkout pages.

#### 2. Setting Up AWS Shield Advanced

- Open AWS Shield Console → Click "Enable AWS Shield Advanced".
- Select AWS resources to protect (ALB, CloudFront, API Gateway, Route 53).
- Enable AWS WAF integration for additional filtering.
- Configure AWS DDoS Response Team (DRT) support (for assistance during attacks).
- 5. Click "Enable".

#### Expected Outcome:

AWS Shield Advanced **monitors and mitigates DDoS threats** in real-time.

CHAPTER 2: AWS WAF (WEB APPLICATION FIREWALL) OVERVIEW

#### 1. How AWS WAF Works

AWS WAF filters HTTP(S) traffic based on:

- ✓ Rate-Based Rules Limits requests per IP.
- ✓ IP Blocking Rules Blocks traffic from specific IPs or countries.
- ✓ SQL Injection & XSS Protection Prevents web application vulnerabilities.
- ✓ Bot Control Blocks malicious bots and scrapers.

#### \* Example:

A **news website** uses AWS WAF **rate-based rules** to prevent excessive traffic from bots scraping their content.

#### 2. AWS WAF Rule Types

Rule Type	Purpose	
IP Match Rule	Blocks or allows traffic from specific IPs or IP ranges.	
Geo Match Rule	Restricts access from specific countries or regions.	
Rate-Based Rule	Limits the number of req	uests per IP
SQL Injection Rule	Detects and blocks SQL i attacks.	<mark>njectio</mark> n
Cross-Site Scripting (XSS) Rule	Prevents cross-site scripting attacks.	

## **\*** Example:

A financial application blocks all traffic except from the US using Geo Match Rules.

CHAPTER 3: SETTING UP AWS WAF FOR WEB PROTECTION

Step 1: Create a Web ACL (Access Control List)

- Open AWS WAF Console → Click "Create Web ACL".
- Select Resource Type (CloudFront, ALB, API Gateway, AppSync).
- 3. Enter **Web ACL Name** (MyWebAppWAF).

#### Step 2: Add Security Rules

Click "Add Rules" → Choose AWS Managed Rules.

- Select "SQL Injection Rule Set" & "XSS Rule Set".
- 3. Add Rate-Based Rule (Limit 1000 requests per 5 minutes).
- 4. Add **Geo Restriction Rule** (Allow only US, Canada).

#### Step 3: Associate WAF with AWS Services

- Select CloudFront Distribution or Application Load Balancer (ALB).
- 2. Click "Associate Web ACL".
- 3. Save and **Deploy WAF Rules**.

#### **\*** Expected Outcome:

AWS WAF **blocks malicious requests** while allowing legitimate traffic.

#### CHAPTER 4: ADVANCED AWS SHIELD & WAF CONFIGURATIONS

- AWS WAF Logging & Monitoring
- ✓ Enable AWS CloudWatch Metrics to track blocked and allowed requests.
- ✓ Use AWS WAF Logs to analyze malicious request patterns.
- ✓ Integrate with AWS Security Hub & GuardDuty for advanced threat detection.

#### **Example:**

A healthcare website logs WAF events to CloudWatch and triggers alerts for unusual activity.

#### 2. AWS Shield Advanced Cost Protection

- ✓ AWS reimburses DDoS-related AWS charges under Shield Advanced.
- ✓ Ensures unexpected scaling costs from DDoS attacks are covered.
- ✓ Monitors attack traffic in real-time via AWS Security Hub.

## \* Example:

A stock trading platform prevents DDoS-related auto-scaling costs using Shield Advanced.

#### CHAPTER 5: BEST PRACTICES FOR AWS SHIELD & WAF

- ✓ Use AWS Shield Standard for basic DDoS protection (free tier).
- ✓ Enable AWS Shield Advanced for mission-critical applications.
- ✓ **Use AWS WAF rate-limiting rules** to prevent excessive bot traffic.
- ✓ Enable multi-layered security (AWS WAF + Shield + CloudFront).
- ✓ Regularly monitor WAF logs for suspicious activity.

#### \* Example:

A government website implements Shield Advanced + AWS WAF + CloudFront to protect against DDoS and bot traffic.

#### CHAPTER 6: REAL-WORLD USE CASES OF AWS SHIELD & WAF

- 1. Protecting an E-Commerce Website from DDoS Attacks
- ✓ Shield Standard prevents volumetric DDoS attacks.
- ✓ AWS WAF blocks malicious bots scraping product prices.
- ✓ Rate-Based Rules prevent excessive checkout requests from fraudulent users.

#### 2. Securing a Banking API from Injection Attacks

- ✓ AWS WAF SQL Injection Rule blocks attacks on login & payment APIs.
- ✓ Geo Restriction Rule allows access only from trusted regions.
- ✓ Shield Advanced protects against DDoS attacks targeting online banking.
- 3. Preventing Fake Reviews on a Social Media Platform
- ✓ AWS WAF Bot Control blocks fake users from posting spam comments.
- ✓ Rate-Based Rules limit the number of reviews per user per hour.
- ✓ CloudWatch Logs + WAF Monitoring detects unusual login patterns.

CONCLUSION: MASTERING AWS SHIELD & WAF

By using **AWS Shield & AWS WAF**, businesses can:

- Prevent DDoS attacks and ensure uptime.
- Block SQL injection, XSS, and bot attacks.
- Enforce rate-based limits to prevent API abuse.
- Monitor and audit security threats in real-time.

#### FINAL EXERCISE:

- 1. **Enable AWS Shield Standard** for an ALB or CloudFront distribution.
- 2. Create a Web ACL in AWS WAF with rate-limiting and SQL injection protection.

3. **Analyze WAF logs** in CloudWatch and identify **potential** threats.



## **IDENTITY & COMPLIANCE**

## AWS COGNITO FOR USER AUTHENTICATION – STUDY MATERIAL

#### INTRODUCTION TO AWS COGNITO

#### What is AWS Cognito?

AWS Cognito is a fully managed authentication and user management service that enables developers to add sign-up, signin, and access control to web and mobile applications. It allows users to authenticate using social, enterprise, or federated identity providers.

#### **Key Features of AWS Cognito**

- ✓ **User Pool & Identity Pool** Supports user authentication and federated access.
- ✓ Multi-Factor Authentication (MFA) Enhances security with OTP-based login.
- ✓ OAuth 2.0, SAML & OpenID Connect Enables single sign-on (SSO).
- ✓ Federated Identity Supports Google, Facebook, Apple, and Amazon login.
- ✓ Adaptive Authentication Uses AI to detect suspicious login attempts.
- ✓ Secure Access to AWS Services Grants temporary AWS credentials via Identity Pools.

#### CHAPTER 1: AWS COGNITO COMPONENTS

#### AWS Cognito consists of two main components:

#### 1. User Pools (Authentication & User Management)

- Manages user registration, sign-in, and account recovery.
- Supports email/phone-based authentication.
- Integrates with social login providers (Google, Facebook, Apple).
- Provides OAuth 2.0 authentication tokens (ID, Access, and Refresh tokens).

#### 📌 Example:

A mobile banking app uses Cognito User Pools for secure user authentication and password resets.

#### 2. Identity Pools (Authorization & AWS Resource Access)

- Provides temporary AWS credentials to authenticated users.
- Grants users access to S3, DynamoDB, Lambda, and other AWS services.
- Works with User Pools or third-party identity providers.

#### **Example:**

An e-learning platform uses Cognito Identity Pools to grant temporary access to S<sub>3</sub>-hosted video content.

#### CHAPTER 2: SETTING UP AWS COGNITO USER POOL

#### Step 1: Create a Cognito User Pool

- Open AWS Cognito Console → Click "Create a User Pool".
- 2. Enter a **Pool Name** (MyAppUserPool).

- 3. Choose **Sign-in options**:
  - Email, Phone, or Username-based authentication.
  - Enable Multi-Factor Authentication (MFA) (Optional).

#### **Step 2: Configure Security Policies**

- 1. Set **Password Policy** (minimum length, special characters).
- Enable Multi-Factor Authentication (MFA) via SMS or Authenticator App.
- Define User Attributes (e.g., email, phone\_number).

#### Step 3: Create an App Client (OAuth2 Token Issuer)

- Click "Add App Client".
- 2. Enter an **App Name** (e.g., MyWebAppClient).
- 3. Enable **OAuth 2.0** scopes (openid, profile, email).
- 4. Disable Client Secret (for mobile apps).
- 5. Click Create App Client.
- Expected Outcome:

A **User Pool is created**, ready for authentication.

CHAPTER 3: INTEGRATING AWS COGNITO WITH WEB/MOBILE APPS

1. Using Cognito Hosted UI for Authentication

AWS Cognito provides a **pre-built Hosted UI** for authentication.

#### Steps to Use Hosted UI

 Open Cognito Console → Select User Pool → Click "App Integration".

- 2. Copy the **Hosted UI URL** (e.g., https://myapp.auth.us-east-1.amazoncognito.com).
- Configure Allowed Callback URLs (e.g., https://myapp.com/login).
- 4. Redirect users to the Hosted UI to authenticate via Google, Facebook, or AWS Cognito.

#### \* Example:

A **React web app** redirects users to the Cognito Hosted UI for **sign-in and registration**.

#### 2. Authenticating via AWS SDK (Using Amplify)

AWS **Amplify SDK** simplifies Cognito integration in web and mobile apps.

#### Install Amplify in a React App

npm install aws-amplify @aws-amplify/auth

#### Configure Cognito in React App

import Amplify from 'aws-amplify';

import awsconfig from './aws-exports';

Amplify.configure(awsconfig);

## Sign Up a New User

import { Auth } from 'aws-amplify';

Auth.signUp({

```
username: 'john_doe',
password: 'StrongPass123!',
attributes: { email: 'johndoe@example.com' }
})
.then(data => console.log("User signed up!", data))
.catch(err => console.log("Error signing up", err));
Sign In an Existing User
Auth.signIn('john_doe', 'StrongPass123!')
.then(user => console.log("User signed in!", user))
.catch(err => console.log("Error signing in", err));
```

#### Expected Outcome:

Users can sign up and sign in using Cognito User Pool authentication.

CHAPTER 4: USING COGNITO IDENTITY POOLS FOR AWS ACCESS

1. Granting AWS Access to Authenticated Users

Cognito Identity Pools allow users to access AWS resources securely.

#### Steps to Create an Identity Pool

- Open AWS Cognito Console → Click "Create Identity Pool".
- 2. Enable Authenticated and Unauthenticated Users.
- Choose IAM Roles for user access (CognitoAuthenticatedRole).
- 4. Click "Create Pool" → Attach IAM policies for AWS access.

#### **Example IAM Policy for S3 Access**

```
{
  "Effect": "Allow",
  "Action": ["s3:GetObject"],
  "Resource": "arn:aws:s3:::my-secure-bucket/*"
}
```

#### **\*** Example:

A mobile photo-sharing app allows authenticated users to upload images to S3 using temporary credentials.

#### CHAPTER 5: AWS COGNITO SECURITY BEST PRACTICES

- ✓ Enforce Strong Password Policies Minimum length, special characters, expiration.
- ✓ Enable Multi-Factor Authentication (MFA) Protects against password theft.
- ✓ **Use Least Privilege Access** Restrict IAM roles with minimal permissions.
- ✓ Enable Adaptive Authentication Blocks suspicious login attempts.
- ✓ Monitor Cognito Logs with CloudTrail Detect unauthorized access.

#### 📌 Example:

A financial services app enforces MFA and logs all Cognito activity in AWS CloudTrail.

#### CHAPTER 6: AWS COGNITO REAL-WORLD USE CASES

#### 1. User Authentication for E-Commerce Applications

- ✓ Cognito User Pools manage customer logins.
- ✓ Social login (Google, Facebook, Apple) for seamless onboarding.
- ✓ Cognito Identity Pools grant temporary AWS credentials for accessing order history.

#### **\*** Example:

An online store lets users sign in via Google and access their purchase history in DynamoDB.

#### 2. Secure API Authentication with Cognito & API Gateway

- ✓ Cognito issues JWT tokens for API authentication.
- ✓ AWS API Gateway **validates Cognito tokens** before processing requests.
- ✓ API Gateway allows or denies access based on Cognito groups.

## **\*** Example:

A ride-sharing app uses Cognito to authenticate drivers and passengers before calling APIs.

#### 3. Mobile Banking Security with Cognito MFA

- ✓ SMS-based MFA for customer logins.
- ✓ Cognito Identity Pools for temporary AWS resource access.
- ✓ Adaptive Authentication blocks login attempts from unusual locations.

#### **\*** Example:

A mobile banking app requires MFA and adaptive authentication for fraud detection.

CONCLUSION: MASTERING AWS COGNITO

By using **AWS Cognito**, businesses can:

- Provide secure user authentication with social and enterprise logins.
- Implement role-based access control using IAM roles.
- Protect user data with MFA and adaptive authentication.
- ☑ Ensure scalability for millions of users with AWS's secure infrastructure.

#### FINAL EXERCISE:

- 1. Create a Cognito User Pool and authenticate a test user.
- 2. Integrate Cognito authentication into a React or mobile appusing Amplify.
- 3. **Set up Cognito Identity Pools** and grant AWS S<sub>3</sub> access for authenticated users.

## AWS ORGANIZATIONS AND ACCESS CONTROL – STUDY MATERIAL

INTRODUCTION TO AWS ORGANIZATIONS AND ACCESS CONTROL What is AWS Organizations?

AWS **Organizations** is a service that helps businesses **manage multiple AWS accounts centrally**. It allows administrators to apply **security policies, consolidated billing, access controls, and service permissions** across accounts in an **organization hierarchy**.

#### **Key Features of AWS Organizations**

- ✓ **Centralized Account Management** Create and manage multiple AWS accounts under a single organization.
- ✓ **Service Control Policies (SCPs)** Enforce fine-grained security controls across accounts.
- ✓ **Consolidated Billing** Simplifies cost management by grouping billing under one account.
- ✓ **Organizational Units (OUs)** Group AWS accounts for easy policy management.
- ✓ Integration with AWS IAM Apply policies to control IAM access across multiple accounts.
- ✓ Cross-Account Access Management Enables secure resource sharing between AWS accounts.

CHAPTER 1: UNDERSTANDING AWS ORGANIZATIONS STRUCTURE

#### 1. AWS Organizations Hierarchy

AWS Organizations consists of the following key components:

Component	Description

Management Account	The root AWS account that controls all
	member accounts.
Member Accounts	AWS accounts managed under AWS
	Organizations.
Organizational Units	Logical groups of AWS accounts that can
(OUs)	have common policies.
Service Control	Security policies that restrict permissions
Policies (SCPs)	across member accounts.

#### **\*** Example:

A company with multiple teams (Finance, Development, Security) creates **OUs for each team** to enforce different security and billing policies.

#### CHAPTER 2: SETTING UP AWS ORGANIZATIONS

#### Step 1: Enable AWS Organizations

- 1. Open AWS Organizations Console.
- Click "Create an Organization" → Select "Enable All Features".
- 3. Choose "Create Organization".

#### Step 2: Add Member Accounts to an Organization

- Click "Add Account" → Choose "Create an AWS Account" or "Invite an Existing AWS Account".
- 2. Enter the **Email ID** of the account to invite.
- 3. Click "Send Invitation".

#### \* Expected Outcome:

The invited account joins the AWS Organization, and policies can be applied centrally.

#### Step 3: Create Organizational Units (OUs)

- Open AWS Organizations Console → Click "Organizational Units".
- 2. Click "Create Organizational Unit".
- 3. Enter **OU Name** (e.g., Security-Team).
- 4. Drag and drop AWS accounts into the new OU.
- 5. Click "Save".

## Expected Outcome:

AWS accounts are grouped into OUs for better management.

CHAPTER 3: IMPLEMENTING SERVICE CONTROL POLICIES (SCPS) IN AWS ORGANIZATIONS

What are SCPs (Service Control Policies)?

SCPs are organization-wide permission policies that enforce security best practices.

Policy Name	Description
Deny All Except Billing	Restricts all actions except billing.
Prevent IAM User Creation	Blocks creation of new IAM users.

**Restrict S3 Public Access** Denies public access to S3 buckets.

Policy Name	Description		
Restrict Root Account	Prevents root user from performing		

operations.

Usage

#### \* Example:

A company applies an SCP to prevent unauthorized users from **creating IAM roles** in production accounts.

#### 2. Creating an SCP to Restrict Specific AWS Services

- Open AWS Organizations Console → Click "Policies".
- 2. Click "Create Policy" → Enter Policy Name (Deny-S3-Public-Access).
- 3. Enter the following SCP JSON Policy:

```
4. {
```

- 5. "Version": "2012-10-17",
- "Statement": [
- 7. {
- "Effect": "Deny", 8.
- "Action": "s3:PutBucketPublicAccessBlock", 9.
- "Resource": "\*" 10.
- 11.
- 1 12.
- 13.}
- Click "Attach Policy" → Select the OU or Account. 14.

#### Expected Outcome:

All AWS accounts under this OU are restricted from making S3 buckets public.

CHAPTER 4: AWS ACCESS CONTROL AND IAM INTEGRATION

#### 1. IAM vs. AWS Organizations

Feature	IAM	AWS Organizations
Scope	User and role permissions inside an	Organization-wide security and service
	AWS account	restrictions
	AVV3 account	restrictions
Policy Type	IAM Policies	Service Control Policies
		(SCPs)
Management	Managed wi <mark>thi</mark> n a	Centrally applied across
	single AWS account	multiple accounts

#### **\*** Example:

A company applies IAM policies to limit S3 access within an account, while AWS Organizations applies SCPs to all accounts in an OU.

## 2. Cross-Account Access with AWS Organizations

AWS Organizations enables secure access between AWS accounts using:

- ✓ IAM Roles Allows temporary access to another AWS account.
- √ Resource Sharing via AWS Resource Access Manager (RAM) Shares AWS resources securely.

Example: Setting Up Cross-Account IAM Roles

1. In Account A (Sender) – Create an IAM Role:

- 2. {
- 3. "Effect": "Allow",
- 4. "Principal": { "AWS": "arn:aws:iam::123456789012:root" },
- 5. "Action": "sts:AssumeRole"
- 6. }
- 7. In Account B (Receiver) Assume Role using AWS CLI:
- 8. aws sts assume-role --role-arn arn:aws:iam::123456789012:role/CrossAccountRole

#### Expected Outcome:

Users from Account B can securely access resources in Account A using a temporary role.

#### CHAPTER 5: COST MANAGEMENT WITH AWS ORGANIZATIONS

- 1. Consolidated Billing in AWS Organizations
- ✓ Single Payment Method All AWS accounts are billed under the Management Account.
- ✓ Cost Savings with Reserved Instances Discounts apply across multiple accounts.
- ✓ **Detailed Billing Reports** View costs per **OU, service, or** individual account.

#### \* Example:

A company centralizes billing for multiple AWS accounts, simplifying cost tracking and optimization.

#### CHAPTER 6: AWS ORGANIZATIONS BEST PRACTICES

- ✓ Use OUs to Separate Production, Development, and Security
  Accounts.
- ✓ Apply SCPs to Restrict Unauthorized Actions Across AWS Accounts.
- ✓ Enable MFA for All Accounts and IAM Users.
- ✓ Use IAM Roles for Secure Cross-Account Access.
- ✓ Monitor AWS Accounts Using AWS Security Hub and CloudTrail.

## 📌 Example:

A large enterprise sets up AWS Organizations with:

- Separate OUs for production and development.
- SCPs to restrict unauthorized changes.
- Consolidated billing for cost efficiency.

CHAPTER 7: REAL-WORLD USE CASES OF AWS ORGANIZATIONS & ACCESS CONTROL

- 1. Enterprise Security and Compliance Enforcement
- ✓ OUs group accounts by department (HR, Finance, IT).
- √ SCPs restrict IAM modifications in security accounts.
- √ CloudTrail logs all security events for auditing.

#### **\*** Example:

A healthcare company enforces HIPAA security policies across all AWS accounts using SCPs.

#### 2. Multi-Account Cost Optimization for SaaS Companies

- ✓ Consolidated billing across multiple AWS accounts.
- √ Reserved instance savings applied globally.
- √ OUs separate customer environments.

## 📌 Example:

A SaaS company optimizes costs by grouping accounts under AWS Organizations.

CONCLUSION: MASTERING AWS ORGANIZATIONS AND ACCESS CONTROL

By using AWS Organizations, businesses can:

- Centrally manage multiple AWS accounts with security policies.
- Apply Service Control Policies (SCPs) for organization-wide restrictions.
- Enable cross-account access securely using IAM roles.
- Optimize costs through consolidated billing.

#### FINAL EXERCISE:

- 1. Create an AWS Organization and add member accounts.
- 2. Set up an Organizational Unit (OU) and apply an SCP.
- 3. Configure cross-account IAM roles and test secure access.

## DATA ENCRYPTION & COMPLIANCE FRAMEWORKS – STUDY MATERIAL

Introduction to Data Encryption & Compliance Frameworks

#### What is Data Encryption?

Data encryption is the **process of converting data into an unreadable format** (ciphertext) to prevent unauthorized access.
Only users with the correct **decryption key** can access the original data. Encryption ensures **data confidentiality, integrity, and compliance** with security regulations.

#### **Key Features of Data Encryption**

- ✓ Protects sensitive data at rest and in transit.
- ✓ Uses cryptographic algorithms (AES, RSA, SHA).
- ✓ Prevents unauthorized access and data breaches.
- ✓ Ensures compliance with regulatory standards (HIPAA, GDPR, PCI DSS).

#### What are Compliance Frameworks?

Compliance frameworks are industry-specific regulations and standards that define how organizations should store, process, and protect data to ensure security and privacy.

#### Key Compliance Frameworks for Data Protection

- ✓ **GDPR (General Data Protection Regulation)** Protects personal data in the EU.
- ✓ HIPAA (Health Insurance Portability and Accountability Act) Secures healthcare information.
- ✓ PCI DSS (Payment Card Industry Data Security Standard) Safeguards credit card transactions.

✓ ISO 27001 (Information Security Standard) – Provides a global framework for data security.

CHAPTER 1: UNDERSTANDING DATA ENCRYPTION

#### 1. Types of Data Encryption

Encryption	Description	Use Cases
Туре		
Symmetric	Uses a single key for	Secure database
Encryption	encryption & decryption.	sto <mark>ra</mark> ge, file
		encryption.
Asymmetric	Uses a public key for	Digital signatures,
Encryption	encryption and a private	secure key
	key for decryption.	exchange.
Hashing	Converts data into a fixed-	Password storage,
	length hash value	data integrity
	(irreversible).	verification.

## **\*** Example:

A financial institution encrypts customer bank records using **AES-256 encryption** before storing them in AWS S<sub>3</sub>.

## 2. Encryption Algorithms & Standards

Algorithm	Туре	Key Length	Use Cases
AES (Advanced	Symmetric	128, 192,	Secure file
Encryption		256-bit	storage, cloud
Standard)			encryption.

RSA (Rivest-	Asymmetric	1024,	Digital signatures,
Shamir-Adleman)		2048,	SSL/TLS
		4096-bit	encryption.
SHA (Secure Hash	Hashing	256, 512-	Password hashing,
Algorithm)		bit	blockchain
			security.

#### \* Example:

A web application encrypts API traffic using RSA 2048-bit encryption with TLS 1.3.

#### 3. Data Encryption in AWS Cloud

AWS provides several services for **encrypting data at rest and in transit**:

- ✓ AWS Key Management Service (KMS) Manages encryption keys securely.
- ✓ AWS Secrets Manager Stores & rotates API keys and credentials.
- ✓ AWS Certificate Manager (ACM) Issues and manages SSL/TLS certificates.
- ✓ AWS CloudHSM (Hardware Security Module) Provides dedicated encryption hardware.

#### **Example:**

A SaaS company encrypts Amazon RDS databases using AWS KMS Customer Managed Keys (CMKs).

#### CHAPTER 2: COMPLIANCE FRAMEWORKS & REGULATIONS

## 1. GDPR (General Data Protection Regulation)

- ✓ Enforced by the **European Union** to protect user data privacy.
- ✓ Requires data encryption, user consent, and breach notifications.
- ✓ Grants users the "Right to be Forgotten" (data deletion upon request).

#### \* Example:

An e-commerce company encrypts customer payment details to comply with GDPR Article 32 on data security.

- 2. HIPAA (Health Insurance Portability and Accountability Act)
- ✓ Regulates healthcare data (ePHI Electronic Protected Health Information).
- ✓ Requires encryption of patient records, audit logging, and access controls.
- ✓ Enforces Business Associate Agreements (BAAs) for third-party service providers.

#### \* Example:

A hospital encrypts patient medical records in Amazon S<sub>3</sub> using AWS KMS to comply with HIPAA encryption guidelines.

- 3. PCI DSS (Payment Card Industry Data Security Standard)
- ✓ Governs credit card transactions and payment processing.
- ✓ Requires encryption of cardholder data (CHD) and secure network policies.
- ✓ Enforces regular security audits and penetration testing.

#### 📌 Example:

A payment gateway encrypts credit card transactions using AES-256 and stores keys securely in AWS CloudHSM.

- 4. ISO 27001 (International Information Security Standard)
- ✓ Global standard for managing information security risks.
- ✓ Requires encryption, risk assessment, and continuous monitoring.
- ✓ Enforces security policies across organizations.

#### **\*** Example:

A tech company implements ISO 27001 standards by encrypting all user passwords with SHA-256 hashing.

#### CHAPTER 3: IMPLEMENTING ENCRYPTION IN AWS

- 1. Enabling AWS KMS for Data Encryption
  - Open AWS KMS Console → Click "Create Key".
  - 2. Choose **Key Type** (Symmetric or Asymmetric).
  - Set Key Alias (e.g., MyAppEncryptionKey).
  - 4. Assign IAM Users/Roles for encryption permissions.
  - 5. Click "Create Key".

#### Expected Outcome:

A **Customer Managed Key (CMK) is created** and can be used to encrypt AWS services.

#### 2. Encrypting Data in Amazon S3

- Open S<sub>3</sub> Console → Select a Bucket → Click Properties.
- 2. Under **Default Encryption**, choose **AWS KMS**.
- 3. Select **an existing KMS key** or create a new one.
- 4. Click "Save Changes".
- Expected Outcome:

All new objects in the S<sub>3</sub> bucket are encrypted with AWS KMS.

- 3. Enabling Transparent Data Encryption (TDE) in RDS
  - Open AWS RDS Console → Select a Database Instance.
  - Click Modify → Enable Encryption under Additional Configurations.
  - 3. Select AWS KMS Key.
  - 4. Click "Apply Changes".
- Expected Outcome:

Amazon RDS encrypts all stored data using AWS KMS.

CHAPTER 4: AWS COMPLIANCE & AUDITING TOOLS

AWS provides tools to monitor and ensure compliance:

- ✓ **AWS CloudTrail** Logs all API calls for security audits.
- ✓ **AWS Config** Monitors resource configurations for compliance violations.
- ✓ AWS Security Hub Centralized security monitoring and compliance reporting.
- ✓ AWS Audit Manager Automates compliance assessments for frameworks like GDPR, HIPAA, PCI DSS.



#### \* Example:

A financial company enables AWS Security Hub to monitor PCI DSS compliance across AWS accounts.

#### Chapter 5: Best Practices for Secure Data Encryption

- ✓ Use AWS KMS Customer Managed Keys (CMKs) for full control over encryption.
- ✓ Enable S<sub>3</sub> Bucket Encryption and Block Public Access.
- ✓ Rotate Encryption Keys Regularly for better security.
- √ Use IAM Policies to Restrict Encryption Key Usage.
- √ Monitor Encryption Events with AWS CloudTrail.

#### \* Example:

A government agency follows encryption best practices by implementing AWS CloudHSM for secure key storage.

CHAPTER 6: REAL-WORLD USE CASES OF DATA ENCRYPTION & **COMPLIANCE** 

- 1. Financial Data Protection for Banks
- ✓ Encrypts customer transaction data using AWS KMS.
- √ Uses CloudTrail for auditing all data access attempts.
- ✓ Implements IAM least privilege access policies.

#### \* Example:

A bank ensures PCI DSS compliance by encrypting credit card transactions with AES-256.

#### 2. Healthcare Data Security for Hospitals

- ✓ Stores patient medical records in encrypted RDS databases.
- ✓ Enforces HIPAA-compliant access controls with IAM roles.
- √ Logs all API calls in AWS CloudTrail for security audits.

#### **\*** Example:

A hospital encrypts medical imaging files in Amazon S<sub>3</sub> to comply with HIPAA security standards.

CONCLUSION: MASTERING DATA ENCRYPTION & COMPLIANCE
By implementing data encryption & compliance frameworks,
businesses can:

- Secure sensitive data across AWS services.
- Ensure compliance with GDPR, HIPAA, PCI DSS, and ISO 27001.
- Use AWS KMS and CloudHSM for key management.
- Automate security audits using AWS Security Hub & Audit Manager.

#### FINAL EXERCISE:

- Encrypt an S<sub>3</sub> bucket using AWS KMS.
- 2. Enable PCI DSS compliance monitoring in AWS Security Hub.
- 3. Use AWS CloudTrail to audit data access events.

#### **ASSIGNMENT**

## DEPLOY A MYSQL DATABASE ON RDS AND CONNECT TO AN APPLICATION



# SOLUTION: DEPLOY A MYSQL DATABASE ON AWS RDS AND CONNECT TO AN APPLICATION

This step-by-step guide will walk you through deploying a **MySQL** database on Amazon RDS and connecting it to a **web application** (e.g., a PHP or Node.js app).

#### Step 1: Log in to AWS and Navigate to RDS

- 1. Open **AWS Console**  $\rightarrow$  Search for **RDS**.
- 2. Click "Create database".

#### Step 2: Create an RDS MySQL Database

#### 1. Choose Database Creation Method

- Select Standard create.
- Under Engine Options, choose MySQL.
- Select MySQL Version 8.0 (or the latest stable version).

#### 2. Configure Database Instance

- DB Instance Identifier: my-mysql-db.
- Master Username: admin.
- Master Password: YourSecurePassword123! (Store it securely).

#### 3. Choose Instance Type and Storage

• **DB Instance Class:** db.t3.micro (for free tier) or db.t3.medium (for production).

- Storage Type: General Purpose (SSD).
- Allocated Storage: 20 GB (Auto-scaling enabled).

#### 4. Configure Connectivity

- VPC: Choose default VPC or create a new one.
- Subnet Group: Select default subnet group.
- Public Access: Enable (for testing purposes; disable in production).
- Security Group: Choose "Create new security group".
- Modify Security Group to Allow Access:
  - Go to EC2 → Security Groups.
  - Select the RDS security group → Click Inbound Rules → Edit.
  - Add Rule:
    - Type: MySQL/Aurora
    - Port Range: 3306
    - Source: o.o.o.o/o (for testing, restrict to your IP in production).
    - Click Save Changes.

#### 5. Configure Database Settings

- Initial Database Name: myappdb.
- Enable Automated Backups: Yes (Backup retention: 7 days).
- Enable Performance Insights: Yes (for monitoring).

#### 6. Create the RDS Instance

Click "Create Database" and wait for it to be Available (may take a few minutes).

#### **Find RDS Endpoint:**

- Navigate to RDS Console → Databases.
- 2. Click "my-mysql-db".
- 3. Copy **"Endpoint"** (e.g., my-mysql-db.xyz123.us-east-1.rds.amazonaws.com).

#### Step 3: Connect to RDS MySQL Database

- 1. Connect Using MySQL CLI
  - 1. Install MySQL Client on your local machine:
    - o Linux/Mac:
    - sudo apt install mysql-client -y # Ubuntu/Debian
    - sudo yum install mysql -y # Amazon Linux
    - Windows: Install MySQL Workbench from MySQL Website.
  - 2. Connect to RDS:
  - mysql -h my-mysql-db.xyz123.us-east-1.rds.amazonaws.com u admin -p

Enter your **password** when prompted.

\*

#### Expected

Outcome:

You are now connected to your AWS RDS MySQL database.

#### 2. Create a Sample Table

```
Once connected, run:

CREATE TABLE users (

id INT AUTO_INCREMENT PRIMARY KEY,

name VARCHAR(100),

email VARCHAR(100) UNIQUE,

created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP

);

Expected Outcome:

The users table is created.
```

#### Step 4: Connect a Web Application to MySQL RDS

#### 1. Connecting a PHP Application

Install MySQL PDO extension and update your PHP app's database configuration:

```
$pdo->setAttribute(PDO::ATTR_ERRMODE,
PDO::ERRMODE_EXCEPTION);
echo "Connected successfully!";
} catch (PDOException $e) {
    die("Connection failed: " . $e->getMessage());
}
?>
```

Expected Outcome:

If connected successfully, the script prints "Connected successfully!".

#### 2. Connecting a Node.js Application

- Install MySQL package;
- 2. npm install mysql2
- 3. Create a **db.js** file in your Node.js project:
- const mysql = require('mysql2');
- 5.
- 6. const connection = mysql.createConnection({
- 7. host: "my-mysql-db.xyz123.us-east-1.rds.amazonaws.com",
- 8. user: "admin",
- 9. password: "YourSecurePassword123!",
- 10. database: "myappdb"
- 11.});
- 12.

```
13.connection.connect(err => {
14.
          if (err) {
      console.error("Database connection failed: " + err.stack);
15.
16.
            return;
17. }
          console.log("Connected to MySQL database.");
18.
        });
19.
20.
        module.exports = connection;
21.
        Run your Node.js app:
22.
        node db.js
23.
                       Expected
                                                        Outcome:
```

Step 5: Secure & Optimize Your RDS MySQL Database

Terminal prints "Connected to MySQL database."

- 1. Restrict Public Access (Production Only)
  - 1. Go to RDS Console → Select your Database.
  - 2. Click Modify → Set Public Access = No.
  - 3. Apply changes and **restart the instance**.
- 2. Enable Multi-AZ Deployment (For High Availability)
  - 1. Open **RDS Console** → Click **Modify**.
  - 2. Enable Multi-AZ Deployment for automatic failover.
  - 3. Click Save Changes.

- 3. Enable IAM Authentication for Secure Login
  - Open AWS RDS Console → Select MySQL Database.
  - 2. Click Modify → Enable IAM Database Authentication.
  - 3. Use **IAM roles** instead of passwords for authentication.

#### Step 6: Monitoring & Scaling MySQL Database

- 1. Monitor RDS Performance with Amazon CloudWatch
- ✓ Open AWS CloudWatch → Metrics → RDS to monitor:
  - CPU Utilization
  - Read/Write IOPS
  - Database Connections
- Example: Set up a CloudWatch Alarm for CPU Usage > 75%.
- Scale RDS MySQL Instance (Vertical Scaling)
  - Open AWS RDS Console → Click Modify.
  - 2. Change Instance Class (e.g., db.t3.medium  $\rightarrow$  db.m5.large).
  - 3. Click **Apply Changes** (Apply immediately or during a maintenance window).
- **\*** Expected Outcome:

RDS scales **vertically** to handle higher loads.

- 3. Scale RDS MySQL with Read Replicas (Horizontal Scaling)
  - Open RDS Console → Click Create Read Replica.

- 2. Choose a region and instance class.
- 3. Click **Create** → Use replica for **read-heavy workloads**.

**\*** Expected Outcome:

The **read replica** helps offload database read operations.

CONCLUSION: SUCCESSFULLY DEPLOYED MYSQL RDS & CONNECTED TO AN APPLICATION

By following this guide, you have:

- Created an AWS RDS MySQL Database.
- Configured security & networking settings.
- Connected the RDS MySQL database to a web application.
- Enabled monitoring & performance optimization.
- Secured the database for production use.

#### FINAL EXERCISE:

- 1. Create an IAM-based authentication for MySQL RDS.
- 2. Implement automatic backup retention for 7 days.
- 3. Deploy an AWS Lambda function to fetch data from RDS.

### IMPLEMENT IAM POLICIES AND TEST ROLE-BASED ACCESS



### SOLUTION: IMPLEMENT IAM POLICIES AND TEST ROLE-BASED ACCESS

This step-by-step guide will walk you through **creating an IAM policy, assigning it to a role, and testing role-based access** in AWS.

#### Step 1: Log in to AWS IAM Console

- Open the AWS Management Console.
- 2. Navigate to IAM (Identity and Access Management).

#### Step 2: Create a Custom IAM Policy

- 1. Navigate to Policies
  - 1. In the IAM Console, click "Policies" in the left menu.
  - 2. Click "Create Policy".

#### 2. Define the **Policy Permissions**

 Select "JSON" and paste the following policy that grants readonly access to Amazon S3:

```
{
  "Version": "2012-10-17",
  "Statement": [
     {
        "Effect": "Allow",
        "Action": [
```

```
"s3:ListBucket",

"s3:GetObject"

],

"Resource": [

"arn:aws:s3:::my-secure-bucket",

"arn:aws:s3:::my-secure-bucket/*"

]

}
```

#### **\*** Explanation:

- ✓ Allows listing the bucket and reading objects in "my-secure-bucket".
- ✓ Does NOT allow modifying or deleting objects.
- 3. Name and Create the Policy
  - 1. Click "Next: Tags"  $\rightarrow$  (Optional: Add tags).
  - 2. Click "Next: Review".
  - 3. Policy Name: S3ReadOnlyAccess.
  - 4. Click "Create Policy".
- **Expected Outcome:** A **custom IAM policy** is created successfully.

#### Step 3: Create an IAM Role and Attach the Policy

#### 1. Navigate to IAM Roles

- Click "Roles" → "Create Role".
- Choose Trusted Entity: Select AWS service.
- 3. **Select Use Case:** Choose **EC2** (if assigning to an instance).
- 4. Click "Next".

#### 2. Attach the Custom Policy

- Search for S<sub>3</sub>ReadOnlyAccess (the policy created earlier).
- 2. Select it  $\rightarrow$  Click "Next".

#### 3. Name and Create the Role

- 1. **Role Name:** S3ReadOnlyRole.
- 2. **Description:** "Grants read-only access to S<sub>3</sub> bucket".
- 3. Click "Create Role".

#### Expected Outcome:

The **IAM role is created and ready to be assigned** to users or services.

#### Step 4: Assign the IAM Role to an EC2 Instance (Optional)

- Navigate to EC2 Console → Click "Instances".
- Select an EC2 instance → Click "Actions" → "Security" →
   "Modify IAM Role".
- Select S<sub>3</sub>ReadOnlyRole → Click "Update IAM Role".

#### Expected Outcome:

The EC2 instance now has read-only access to the S3 bucket.

#### Step 5: Testing Role-Based Access (AWS CLI or Console)

#### 1. Testing S<sub>3</sub> Access via AWS CLI

- SSH into the EC2 instance:
- 2. ssh -i my-key.pem ec2-user@your-ec2-public-ip
- 3. Run the following AWS CLI command to list S3 buckets:
- 4. aws s<sub>3</sub> ls
- 5. Try to list objects in the bucket:
- 6. aws s3 ls s3://my-secure-bucket
- 7. Attempt to upload a file (which should fail):
- 8. echo "Test file" > test.txt
- aws s3 cp test.txt s3://my-secure-bucket/
- Expected Outcome:
- ✓ Listing objects should work ✓.
- X Uploading files should fail (due to read-only access) X.

#### Step 6: Modify IAM Policy to Grant Full Access (Optional)

If you need to grant **full access** instead of read-only, modify the policy:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
        "Effect": "Allow",
```

```
"Action": "s3:*",

"Resource": [

"arn:aws:s3:::my-secure-bucket",

"arn:aws:s3:::my-secure-bucket/*"

]

}
```

#### Expected Outcome:

Users with this role can now read, write, and delete objects in S3.

#### Step 7: Monitor IAM Role Activity with AWS CloudTrail

- 1. Navigate to AWS CloudTrail Console.
- 2. Click "Event History" → Filter by IAM Role Activity.
- 3. Check if unauthorized access attempts appear (e.g., trying to delete objects).

#### **Expected Outcome:**

IAM access logs show successful and failed API calls, helping in security audits.

CONCLUSION: SUCCESSFULLY IMPLEMENTED IAM POLICIES & ROLE-BASED ACCESS

By following this guide, you have:

Created a custom IAM policy for S3 access.

- Assigned the policy to an IAM role.
- Attached the IAM role to an EC2 instance.
- Tested role-based access using AWS CLI.
- Monitored IAM activity using AWS CloudTrail.

#### FINAL EXERCISE:

- 1. **Modify the IAM policy** to grant write access to a specific folder in S<sub>3</sub> (my-secure-bucket/logs/).
- 2. Create an IAM policy that restricts EC2 instance termination.
- 3. Enable IAM MFA (Multi-Factor Authentication) for IAM users.

## SECURE A WEB APPLICATION USING AWS WAF



# SOLUTION: SECURE A WEB APPLICATION USING AWS WAF (WEB APPLICATION FIREWALL)

This step-by-step guide will walk you through securing a web application using AWS WAF by creating a Web ACL, setting up WAF rules, and integrating it with Amazon CloudFront or an Application Load Balancer (ALB).

#### Step 1: Log in to AWS and Navigate to AWS WAF Console

- 1. Open the AWS Management Console.
- 2. Search for "AWS WAF & Shield" in the Services search bar.
- 3. Click "AWS WAF" to open the console.

#### Step 2: Create a Web ACL (Access Control List)

- Click "Create Web ACL".
- 2. Enter Web ACL Name (MyWebAppWAF).
- 3. Select Resource Type (Choose where you want to apply WAF):
  - CloudFront (Recommended for global applications).
  - Application Load Balancer (ALB) (Best for backend services).
  - API Gateway (Secures REST APIs).
- 4. Click "Next".
- **Expected Outcome:** A **Web ACL** is created and ready for rules.

#### Step 3: Configure Security Rules in AWS WAF

#### 1. Block SQL Injection Attacks

- Click "Add Rules" → Select "Add Managed Rule Groups".
- Search for "AWS WAF SQL Injection Rule Set".
- Select it and click "Add Rules".

#### 2. Block Cross-Site Scripting (XSS) Attacks

- Click "Add Rules" → Select "AWS WAF XSS Rule Set".
- 2. Click "Add Rules".

#### Restrict Traffic from Specific Countries (Geo-Blocking)

- Click "Add Rules" → Select "Add My Own Rules and Rule Groups".
- Click "Create Rule".
- 3. Rule Name: BlockCountryRule.
- Choose "Geographic Match" → Select countries to block (e.g., China, Russia).
- 5. Click "Save Rule" → "Add to Web ACL".

#### 4. Rate-Limit Requests to Prevent DDoS & Bots

- Click "Add Rules" → Choose "Rate-Based Rule".
- 2. Enter Rule Name: RateLimitRequests.
- 3. Set **Rate Limit** (e.g., 1000 requests per 5 minutes).
- 4. Click "Save Rule" → "Add to Web ACL".

#### 5. Allow Only Specific IPs (Whitelist)

- Click "Add Rules" → Select "IP Set".
- 2. Enter Rule Name: AllowedIPs.
- 3. Add **specific IP addresses** (e.g., 192.168.1.10/32).
- 4. Click "Save Rule"  $\rightarrow$  "Add to Web ACL".
- ★ Expected Outcome: The Web ACL now includes rules to prevent SQL Injection, XSS, Geo-blocking, Rate Limiting, and IP-based filtering.

#### Step 4: Associate WAF with AWS Resources

- 1. Associate WAF with CloudFront (Recommended for Global Protection)
  - Go to AWS WAF Console → Web ACLs.
  - Select "MyWebAppWAF" → Click "Add AWS Resources".
  - Select "CloudFront Distribution" → Click "Associate".
- 2. Associate WAF with an Application Load Balancer (ALB)
  - Open AWS EC2 Console → Click "Load Balancers".
  - 2. Select your ALB → Click "Edit Web ACL Association".
  - Choose "MyWebAppWAF" → Click "Associate".
- **Expected Outcome:** WAF is **protecting incoming traffic** to the **web application**.

#### Step 5: Test and Monitor AWS WAF Rules

- 1. Test with a SQL Injection Attack
  - 1. Open a browser → Enter the following URL:

- 2. https://your-web-app.com/?id=1' OR '1'='1'
- 3. If WAF is working, the request should be **blocked**.

#### 2. Test with a DDoS Attack Simulation

- Run an HTTP load testing tool (like Apache Bench or ab command):
- 2. ab -n 1000 -c 50 https://your-web-app.com/
- 3. If **rate-limiting is applied**, requests beyond the **set limit** should be **blocked**.

#### 3. Monitor WAF Logs in CloudWatch

- Open AWS CloudWatch Console → Click "Logs".
- 2. Search for **AWS WAF logs**.
- 3. Check blocked requests and rule matches.

#### Expected Outcome:

You should see **blocked SQL injection attempts, geo-blocked IPs, and rate-limited requests in** WAF logs.

### Step 6: Enable AWS Shield for DDoS Protection (Optional but Recommended)

- 1. Navigate to AWS Shield Console.
- Click "Enable AWS Shield Advanced" (requires AWS Enterprise Support).
- 3. Select **CloudFront or ALB** as protected resources.
- 4. Click "Enable Protection".

\* Expected Outcome: Your web application is now protected against Layer 3/4 DDoS attacks.

#### Step 7: Enable AWS Security Hub for Continuous Monitoring

- Open AWS Security Hub → Click "Enable Security Hub".
- 2. Enable AWS WAF integration.
- Monitor security threats and WAF rule violations.

#### Expected Outcome:

AWS Security Hub provides **detailed insights into web security threats**.

CONCLUSION: SUCCESSFULLY SECURED A WEB APPLICATION USING AWS WAF

By following this guide, you have:

- Created a Web ACL in AWS WAF.
- Added security rules for SQL Injection, XSS, Geo-Blocking, and Rate Limiting.
- Associated WAF with CloudFront or ALB for web application protection.
- Tested WAF by simulating attacks.
- Monitored security events using AWS CloudWatch & Security Hub.

#### FINAL EXERCISE:

1. Create a custom AWS WAF rule to allow traffic only from your corporate IP range.

- 2. Enable bot control rules to block automated traffic.
- 3. Set up a CloudWatch alarm to trigger when WAF blocks more than 100 requests in an hour.

