



**Independent
Skill Development
Mission**



ISDM (INDEPENDENT SKILL DEVELOPMENT MISSION)

MANAGING USERS AND GROUPS IN UNIX

CHAPTER 1: INTRODUCTION TO USER AND GROUP MANAGEMENT

What is User and Group Management in UNIX?

User and group management is a fundamental part of **system administration** in UNIX and Linux. Since UNIX is a **multi-user operating system**, multiple users can access the system simultaneously. To maintain **security, access control, and proper system organization**, users are assigned **unique user accounts** and grouped under **specific groups**.

Proper user and group management helps in:

- **Controlling file access and permissions** – Users can only modify files they have permission for.
- **Enhancing system security** – Prevents unauthorized access to system files.
- **Improving collaboration** – Allows teams to work within shared directories and projects.
- **Managing system resources** – Controls CPU, memory, and disk usage for users.

Each user account is identified by a **User ID (UID)**, and each group has a **Group ID (GID)**. System administrators can **create, modify,**

delete, and manage users and groups using command-line tools like `useradd`, `usermod`, `groupadd`, and `passwd`.

Example

A company with different departments (HR, IT, and Finance) may assign **user groups** to each department. Employees in the IT group can access development files, while Finance employees are restricted from modifying IT-related files.

Exercise

1. List all users on a UNIX system using the `/etc/passwd` file.
2. Create a new user and assign them to an existing group.

Case Study: Implementing Access Control in a Corporate IT Infrastructure

A large company needs to **restrict access to financial records** to only authorized personnel. By implementing **user groups and permissions**, system administrators ensure that only employees in the finance group have access to payroll files, preventing unauthorized changes or data breaches.

CHAPTER 2: CREATING AND MANAGING USERS IN UNIX

Creating a New User

Administrators use the `useradd` command to **create user accounts** in UNIX.

Basic Syntax of `useradd`

```
sudo useradd -m -s /bin/bash username
```

- `-m` → Creates a home directory for the user.

- `-s /bin/bash` → Assigns the Bash shell as the default shell.

Setting a Password for a User

After creating a user, set a password using:

```
sudo passwd username
```

The system will prompt for a new password.

Modifying User Accounts

Use the `usermod` command to **modify an existing user**.

- **Change a user's home directory:**
 - `sudo usermod -d /new/home/directory username`
- **Lock a user account:**
 - `sudo usermod -L username`
- **Unlock a user account:**
 - `sudo usermod -U username`

Deleting a User

To remove a user account, use:

```
sudo userdel -r username
```

- `-r` → Deletes the home directory along with the user account.

Example

A company needs to onboard a new employee, **John Doe**, and assign him a home directory with a default shell:

```
sudo useradd -m -s /bin/bash johndoe
```

```
sudo passwd johndoe
```

Exercise

1. Create a new user and set a password for them.
2. Change the user's default shell to `/bin/zsh` and verify the change.

Case Study: Automating User Account Management in Cloud Servers

Cloud service providers manage thousands of users. Automating **user creation, modification, and deletion** using scripts prevents manual errors and saves time. In this case study, we explore how cloud administrators automate user management using `useradd`, `usermod`, and `userdel`.

CHAPTER 3: MANAGING GROUPS IN UNIX

Creating and Managing Groups

A **group** in UNIX is a collection of users with shared access to files and resources. Groups help in managing **team collaborations, access control, and security policies**.

Creating a Group

To create a new group, use:

```
sudo groupadd developers
```

Adding Users to a Group

A user can be added to a group using `usermod`:

```
sudo usermod -aG developers username
```

- `-aG` → Appends the user to the group without removing them from other groups.

Viewing Group Membership

To check which groups a user belongs to:

`groups username`

Or list all groups in the system:

`cat /etc/group`

Removing a User from a Group

To remove a user from a group:

`sudo deluser username groupname`

Deleting a Group

To remove a group:

`sudo groupdel developers`

Example

A project team working on software development creates a group called `developers` and adds all developers to it:

`sudo groupadd developers`

`sudo usermod -aG developers johndoe`

Exercise

1. Create a group called `admin_team` and add multiple users to it.
2. Remove a user from a group and verify their permissions.

Case Study: Managing Access Control in Multi-User UNIX Environments

A university uses UNIX servers for research and coursework. Professors need access to **grading files**, while students should have access only to **course materials**. By using **user groups**, system administrators enforce role-based access control, improving security and organization.

CHAPTER 4: FILE PERMISSIONS AND OWNERSHIP IN USER AND GROUP MANAGEMENT

Understanding File Ownership

Each file in UNIX has:

1. **Owner** – The user who created the file.
2. **Group** – A group of users who share permissions.
3. **Others** – Any other user on the system.

View file ownership with:

`ls -l filename`

Example output:

```
-rw-r--r-- 1 johndoe developers 2048 Feb 24 12:30 report.txt
```

- **johndoe** → Owner.
- **developers** → Group.
- **Permissions** (rw-r--r--) define who can **read (r)**, **write (w)**, and **execute (x)** the file.

Changing File Ownership and Group

To change a file's owner:

```
sudo chown newuser filename
```

To change a file's group:

```
sudo chgrp newgroup filename
```

Modifying File Permissions

Use chmod to modify permissions:

```
chmod 750 filename
```

- 7 (Owner) = **Read, Write, Execute**
- 5 (Group) = **Read, Execute**
- 0 (Others) = **No access**

Example

A system administrator sets permissions for a confidential document:

```
sudo chown manager report.txt
```

```
sudo chgrp finance report.txt
```

```
chmod 640 report.txt
```

This ensures that **only the manager and finance team can access the file.**

Exercise

1. Create a file and change its owner and group.

2. Modify file permissions to allow only the owner to edit and the group to read.

Case Study: Implementing Secure File Access in a Financial Organization

A financial company stores sensitive payroll data on UNIX servers. Using **strict file permissions (chmod 600)** and **restricted group access**, administrators prevent unauthorized employees from accessing salary records.

CONCLUSION

Understanding **user and group management in UNIX** is essential for **secure and organized system administration**. Key takeaways include:

- ✓ Creating, modifying, and deleting **users and groups**.
- ✓ Assigning **permissions and ownership** to protect files.
- ✓ Enforcing **access control policies** using group-based permissions.

SYSTEM BOOT PROCESS AND RUN LEVELS IN UNIX/LINUX

CHAPTER 1: INTRODUCTION TO THE SYSTEM BOOT PROCESS

What is the System Boot Process?

The **system boot process** is the sequence of steps a UNIX/Linux system follows to start up after being powered on. This process involves initializing hardware components, loading the operating system, and preparing the system for user interaction. A smooth and efficient boot process is essential for **system stability, performance, and security**.

The boot process consists of several key stages:

1. **BIOS/UEFI Initialization** – The system firmware performs hardware checks and loads the bootloader.
2. **Bootloader Execution (GRUB/LILO)** – Loads the kernel into memory and hands over control to the OS.
3. **Kernel Initialization** – The Linux kernel initializes hardware and system processes.
4. **Init/Systemd Process Execution** – The first user-space process that manages system initialization.
5. **Runlevel/Target Execution** – Determines which services and processes should start.

Understanding the boot process is essential for **troubleshooting system startup issues**, optimizing performance, and enhancing security.

Example: Diagnosing a System Boot Failure

If a system fails to boot and shows a **kernel panic error**, it may indicate a corrupted kernel or missing system files. Booting into **recovery mode** and checking logs (dmesg) can help identify the issue.

Exercise

1. Restart your system and enter the **GRUB menu**. Identify the available boot options.
2. View the system logs (dmesg or journalctl -b) to analyze the last boot process.

Case Study: Recovering from a Failed Boot in a Production Server

A company's production server crashes and fails to boot due to a corrupted **GRUB bootloader**. By booting from a **live Linux USB**, system administrators restore the bootloader and recover the system without losing data.

CHAPTER 2: DETAILED STEPS OF THE BOOT PROCESS

Step 1: BIOS/UEFI Initialization

When a system is powered on, the **BIOS (Basic Input/Output System)** or **UEFI (Unified Extensible Firmware Interface)** initializes the hardware.

- **BIOS/UEFI Responsibilities:**
 - Performs **Power-On Self Test (POST)** to check CPU, RAM, and disks.
 - Identifies available **boot devices** (HDD, SSD, USB).
 - Loads the **bootloader** from the primary boot device.

If the BIOS/UEFI cannot find a valid boot device, it displays an error like:

No bootable device found.

Example: Changing the Boot Order in BIOS

To boot from a USB device, enter BIOS (F2, F12, Del, or Esc during startup) and change the **boot order** to prioritize USB.

Exercise

1. Access your system's BIOS/UEFI settings and locate the **boot device priority** menu.
2. Change the boot order and boot the system from a USB drive.

Case Study: Recovering a System by Changing Boot Order

A company's laptop fails to start because the boot order was accidentally changed to **Network Boot** instead of **HDD**. By restoring the correct boot order in BIOS, the issue is resolved quickly.

CHAPTER 3: BOOTLOADER EXECUTION (GRUB/LILO)

What is a Bootloader?

A **bootloader** is a small program responsible for **loading the operating system kernel** into memory. The most common bootloaders in UNIX/Linux are:

- **GRUB (GRand Unified Bootloader)** – Used in most Linux distributions.
- **LILO (Linux Loader)** – An older bootloader, now rarely used.

Functions of a Bootloader:

1. **Loads the Kernel** – Locates and loads the kernel file (vmlinuz).
2. **Provides Boot Options** – Allows users to select different operating systems or kernels.
3. **Passes Boot Parameters** – Sends configuration options to the kernel.

Example: Editing GRUB Boot Parameters

If the system is stuck in a boot loop, you can edit the GRUB menu:

1. Press Esc or Shift during boot to open the GRUB menu.
2. Select a boot entry and press e to edit.
3. Modify the linux line (e.g., add single for single-user mode).
4. Press Ctrl + X to boot with the modified settings.

Exercise

1. Open the GRUB bootloader menu and view available boot entries.
2. Modify a GRUB entry to boot into **single-user mode**.

Case Study: Fixing a Corrupted GRUB Bootloader

A system upgrade corrupts the GRUB bootloader, causing the system to display a **grub rescue prompt**. By using a **Live CD** and running grub-install, administrators restore the bootloader and successfully boot the system.

CHAPTER 4: KERNEL INITIALIZATION

Role of the Kernel in the Boot Process

The **kernel** is the core component of the operating system, responsible for:

- **Initializing hardware drivers** (CPU, RAM, storage).
- **Mounting the root filesystem** (/).
- **Starting system processes** (init/systemd).

If the kernel fails to load, the system may display a **kernel panic** error, requiring manual intervention.

Example: Checking Kernel Messages After Boot

Use the following command to view kernel logs:

```
dmesg | less
```

This helps diagnose **hardware failures or missing drivers**.

Exercise

1. Check the currently running kernel version using:
2. `uname -r`
3. List all available kernels using:
4. `ls /boot/vmlinuz*`

Case Study: Resolving a Kernel Panic on a Server

A new system update causes a **kernel panic**, preventing the server from booting. The administrator boots into an **older working kernel from GRUB**, restores the broken kernel, and ensures system stability.

Chapter 5: Init/Systemd Process Execution

What is Init/Systemd?

Once the kernel loads, it starts the first user-space process:

- **SysV Init (init)** – The traditional init system, used in older UNIX/Linux versions.
- **Systemd (systemctl)** – The modern init system, offering faster boot times and improved process management.

To check the system's init process:

```
ps -p 1
```

Output for Systemd:

```
PID TTY    TIME CMD
```

```
1 ?      00:01:23 systemd
```

Managing System Services with Systemd

Use systemctl to manage services:

- **Check running services:**
 - `systemctl list-units --type=service`
- **Start/Stop a service:**
 - `systemctl start apache2`
 - `systemctl stop apache2`
- **Enable a service at boot:**
 - `systemctl enable ssh`

Example: Restarting a Service After Boot

If the networking service fails to start after boot, restart it manually:

systemctl restart networking

Exercise

1. Check which services are enabled to start at boot using:
2. `systemctl list-unit-files --type=service | grep enabled`
3. Disable an unnecessary service and verify its status.

Case Study: Fixing a Slow Boot Issue Caused by a Failing Service

A company experiences slow boot times. Checking the system logs (`journalctl -b`), administrators find that a **failing service is delaying startup**. By disabling the service, boot time is significantly improved.

CONCLUSION

Understanding the **system boot process and run levels** is essential for:

- **Troubleshooting boot failures** (GRUB errors, kernel panics).
- **Optimizing startup times** by managing services.
- **Enhancing system security** through bootloader configurations.

PACKAGE MANAGEMENT IN UNIX/LINUX (YUM, APT, RPM)

CHAPTER 1: INTRODUCTION TO PACKAGE MANAGEMENT IN UNIX/LINUX

What is Package Management?

Package management is the process of **installing, updating, removing, and managing software packages** on a UNIX/Linux system. Since UNIX/Linux distributions rely heavily on package managers, understanding them is essential for **system administration, software deployment, and maintenance**.

A package manager:

- **Automates software installation and dependency resolution.**
- **Manages package versions and updates securely.**
- **Ensures system stability by preventing broken dependencies.**
- **Verifies package integrity using cryptographic signatures.**

There are two main types of package management systems:

1. **Debian-based (Debian, Ubuntu, etc.)** – Uses dpkg and apt package managers.
2. **Red Hat-based (RHEL, CentOS, Fedora, etc.)** – Uses RPM and YUM/DNF package managers.

These package managers retrieve software from **repositories**, which are official or third-party servers hosting software packages.

Example: Installing Software with a Package Manager

To install the `wget` utility on an Ubuntu system:

```
sudo apt install wget
```

On a Red Hat-based system:

```
sudo yum install wget
```

Exercise

1. Find out which package manager your system uses (`apt`, `yum`, or `dnf`).
2. List all installed packages on your system.

Case Study: Automating System Updates in a Large IT Environment

A company with **hundreds of Linux servers** automates software updates using package managers (`apt` and `yum`). By scheduling regular updates, the company ensures **system security and stability** while reducing manual workload.

CHAPTER 2: RPM PACKAGE MANAGEMENT (RED HAT-BASED SYSTEMS)

What is RPM?

The **RPM Package Manager (RPM)** is the standard package format for Red Hat-based distributions, including:

- **RHEL (Red Hat Enterprise Linux)**
- **CentOS**

- **Fedora**

RPM packages have a .rpm extension and contain **software binaries, dependencies, and metadata**.

Using RPM Commands

RPM does not resolve dependencies automatically. Instead, it installs **only the specified package**, requiring manual installation of dependencies.

1. Installing an RPM Package

```
sudo rpm -ivh package.rpm
```

- -i → Install package.
- -v → Verbose mode.
- -h → Show progress with hash marks.

2. Upgrading an RPM Package

```
sudo rpm -Uvh package.rpm
```

This upgrades an existing package or installs a new one if it's not present.

3. Removing an RPM Package

```
sudo rpm -e package-name
```

4. Querying Installed RPM Packages

- List all installed packages:
- rpm -qa
- Find information about a specific package:
- rpm -qi package-name

Example: Installing and Querying an RPM Package

1. Download an RPM package:
2. `wget`
`http://mirror.centos.org/centos/7/updates/x86_64/Packages/wget-1.14-18.el7.x86_64.rpm`
3. Install the package:
4. `sudo rpm -ivh wget-1.14-18.el7.x86_64.rpm`
5. Verify the installation:
6. `rpm -qa | grep wget`

Exercise

1. Install an RPM package and verify its details using `rpm -qi`.
2. Remove an installed package and confirm its removal.

Case Study: Using RPM for Offline Software Installation in Secure Environments

A high-security government network requires **offline package installation** due to restricted internet access. Administrators use `rpm` to manually install and upgrade software without needing external repositories.

CHAPTER 3: YUM PACKAGE MANAGEMENT (YELLOWDOG UPDATER, MODIFIED)

What is YUM?

YUM (Yellowdog Updater, Modified) is a package manager for Red Hat-based systems that **automatically resolves dependencies**, unlike RPM. It retrieves software from **configured repositories**.

Basic YUM Commands

1. Installing a Package

```
sudo yum install package-name
```

2. Removing a Package

```
sudo yum remove package-name
```

3. Updating All Packages

```
sudo yum update
```

4. Searching for a Package

```
yum search package-name
```

5. Viewing Package Information

```
yum info package-name
```

Example: Installing and Managing a Package with YUM

1. Install the vim package:
2. `sudo yum install vim`
3. Check the installed version:
4. `rpm -q vim`
5. Remove the package:
6. `sudo yum remove vim`

Exercise

1. Search for a package available in YUM repositories and install it.
2. List all installed packages and identify recently installed ones.

Case Study: Automating Patch Management with YUM

A financial institution schedules **automated updates** using yum-cron to ensure all servers receive **security patches and software updates**, reducing vulnerabilities.

CHAPTER 4: APT PACKAGE MANAGEMENT (ADVANCED PACKAGE TOOL)

What is APT?

APT (Advanced Package Tool) is the package manager for **Debian-based systems**, including:

- **Debian**
- **Ubuntu**
- **Linux Mint**

APT simplifies software management by **handling dependencies automatically** and retrieving packages from online repositories.

Basic APT Commands

1. Updating the Package Repository

```
sudo apt update
```

2. Installing a Package

```
sudo apt install package-name
```

3. Removing a Package

```
sudo apt remove package-name
```

4. Upgrading All Installed Packages

```
sudo apt upgrade
```

5. Searching for a Package

```
apt search package-name
```

Example: Managing Packages with APT

1. Update repositories:
2. `sudo apt update`
3. Install the curl package:
4. `sudo apt install curl`
5. Remove the package:
6. `sudo apt remove curl`

Exercise

1. Search for an available package and install it using apt.
2. Upgrade all installed packages and verify the changes.

Case Study: Managing Large-Scale Ubuntu Deployments with APT

A cloud service provider manages **thousands of Ubuntu servers** using **APT automation tools** to deploy, update, and secure software efficiently across the infrastructure.

CHAPTER 5: COMPARISON OF RPM, YUM, AND APT

Feature	RPM	YUM	APT
Distribution	Red Hat-based (RHEL, CentOS, Fedora)	Red Hat-based (RHEL, CentOS, Fedora)	Debian-based (Ubuntu, Debian, Mint)
Dependency Resolution	No (manual)	Yes (automatic)	Yes (automatic)
Installation Command	rpm -ivh package.rpm	yum install package-name	apt install package-name
Package Removal	rpm -e package-name	yum remove package-name	apt remove package-name
Update All Packages	Not applicable	yum update	apt upgrade

Exercise

1. Compare package managers by installing the same software using **RPM, YUM, and APT**.
2. Write a script to **automate system updates using a package manager**.

Case Study: Choosing the Right Package Manager for Enterprise Systems

A tech company decides between **RPM-based (RHEL)** and **APT-based (Ubuntu) distributions** for its servers. After evaluating package management features, they choose Ubuntu with APT for **ease of automation and dependency handling**.

CONCLUSION

This guide covered:

- **Using RPM, YUM, and APT for package management.**
- **Installing, updating, and removing software securely.**
- **Choosing the right package manager for different environments.**

DISK MANAGEMENT AND PARTITIONING IN UNIX/LINUX

CHAPTER 1: INTRODUCTION TO DISK MANAGEMENT AND PARTITIONING

What is Disk Management and Partitioning?

Disk management in UNIX/Linux refers to the process of **configuring, maintaining, and optimizing storage devices**. Proper disk management ensures **efficient use of disk space, data security, and system performance**.

Partitioning is the process of **dividing a physical disk into multiple sections (partitions)**, allowing for better organization and separation of system files, user data, and application files. Partitions enable:

- **Efficient disk space utilization**
- **Isolation of system and user data**
- **Multi-boot system configurations**
- **Improved security and backup strategies**

Types of Disk Partitions

1. **Primary Partition** – A bootable partition that holds the operating system.
2. **Extended Partition** – A special partition that acts as a container for logical partitions.
3. **Logical Partition** – A sub-partition inside an extended partition, used for data storage.

Understanding **disk partitioning schemes** such as **MBR (Master Boot Record)** and **GPT (GUID Partition Table)** is essential for proper disk management.

Example: Viewing Disk Partitions in Linux

Use the following command to check disk partitions:

```
lsblk
```

It displays:

```
NAME MAJ:MIN RM SIZE RO TYPE MOUNTPOINT
```

```
sda 8:0 0 500G 0 disk
```

```
|—sda1 8:1 0 100G 0 part /
```

```
|—sda2 8:2 0 200G 0 part /home
```

```
|—sda3 8:3 0 200G 0 part /var
```

Exercise

1. Use `lsblk` and `fdisk -l` to list partitions on your system.
2. Identify the partition where your Linux OS is installed.

Case Study: Optimizing Disk Partitioning for a Database Server

A financial company runs large databases. By **separating system, database, and log files into different partitions**, administrators optimize **performance and prevent data corruption** in case of a crash.

CHAPTER 2: PARTITIONING A DISK IN LINUX

Understanding Partitioning Tools in UNIX/Linux

To partition a disk in UNIX/Linux, administrators use tools such as:

- **fdisk** – Command-line tool for MBR partitioning.
- **parted** – Used for GPT and MBR partitioning.
- **gparted** – GUI-based tool for partition management.

Using fdisk to Create a New Partition

1. List Available Disks

2. `sudo fdisk -l`

3. Select the Disk to Partition

4. `sudo fdisk /dev/sdb`

5. Create a New Partition

- Press n → Create new partition.
- Choose primary or logical.
- Specify the partition size.
- Press w → Write changes and exit.

6. Format the New Partition

7. `sudo mkfs.ext4 /dev/sdb1`

8. Mount the Partition

9. `sudo mkdir /mnt/newdisk`

10. `sudo mount /dev/sdb1 /mnt/newdisk`

11. Make the Mount Permanent

Edit `/etc/fstab` to automatically mount the partition at boot:

12. `echo "/dev/sdb1 /mnt/newdisk ext4 defaults o o" | sudo tee -a /etc/fstab`

Example: Creating and Formatting a Partition on a New Disk

If a new **500GB disk (/dev/sdc)** is added to a server, you can partition and format it as follows:

```
sudo fdisk /dev/sdc
```

```
# Create new partition → `n`, select partition type → `p`, save → `w`
```

```
sudo mkfs.ext4 /dev/sdc1
```

```
sudo mount /dev/sdc1 /mnt/storage
```

Exercise

1. Create a new partition on an unallocated disk and mount it.
2. Verify partition creation using `lsblk` and `df -h`.

Case Study: Partitioning a Disk for Web Hosting Servers

A web hosting provider **separates system files, website data, and logs** into different partitions. This prevents log files from filling up the root partition, ensuring **server uptime and stability**.

CHAPTER 3: FILESYSTEM MANAGEMENT IN UNIX/LINUX

What is a Filesystem?

A filesystem is the **method used to store and organize data on a partition**. Different filesystems provide varying levels of performance, security, and features.

Common Filesystem Types

Filesystem	Description
ext4	Default Linux filesystem with journaling.
XFS	High-performance filesystem, used for large-scale storage.
NTFS	Windows filesystem, supported in Linux via ntfs-3g.
FAT32	Universal filesystem, compatible with all OS.
Btrfs	Advanced Linux filesystem with snapshots and compression.

Managing Filesystems

- **Check Disk Usage**
- `df -h`
- **Check Filesystem Type**
- `sudo blkid /dev/sdb1`
- **Convert an Existing Filesystem** (Example: Ext4 to XFS)
- `sudo mkfs.xfs /dev/sdb1`

Example: Formatting and Mounting a Filesystem

```
sudo mkfs.ext4 /dev/sdb1
```

```
sudo mkdir /mnt/data
```

```
sudo mount /dev/sdb1 /mnt/data
```

Exercise

1. Check the filesystem type of all partitions using `lsblk -f`.

2. Convert a partition to XFS and mount it.

Case Study: Choosing the Right Filesystem for a Cloud Storage Provider

A cloud storage provider selects **XFS** for large file storage due to its **fast performance and scalability**, ensuring high-speed access to massive datasets.

CHAPTER 4: MANAGING DISK SPACE AND STORAGE OPTIMIZATION

Monitoring Disk Usage

- **Check Free Space:**
- `df -h`
- **Find Large Files:**
- `sudo du -sh /var/log/*`
- **Check Inode Usage:**
- `df -i`

Expanding Storage: Resizing Partitions

Resizing an Ext4 Partition

1. **Unmount the Partition:**
2. `sudo umount /dev/sdb1`
3. **Resize the Partition:**
4. `sudo resize2fs /dev/sdb1`

Example: Increasing Storage for a Full Disk Partition

If /var is full, resize its logical volume:

```
sudo lvextend -l +100%FREE /dev/mapper/centos-var
```

```
sudo resize2fs /dev/mapper/centos-var
```

Exercise

1. Identify which partition is using the most space.
2. Resize an existing partition and verify the change.

Case Study: Preventing Server Downtime Due to Full Disk Issues

A company experiences downtime due to a **full root partition**. By **expanding the partition dynamically**, they prevent system crashes and ensure **continuous uptime**.

CONCLUSION

This guide covered:

- **Creating and managing disk partitions.**
- **Formatting and mounting filesystems.**
- **Monitoring and optimizing disk space.**

MOUNTING FILE SYSTEMS AND NFS IN UNIX/LINUX

CHAPTER 1: INTRODUCTION TO MOUNTING FILE SYSTEMS

What is Mounting in UNIX/Linux?

Mounting is the process of **attaching a storage device or partition to a directory in the UNIX/Linux filesystem** so that users and applications can access it. Unlike Windows, where drives have separate letters (C:, D:), UNIX/Linux integrates all devices into a **single directory structure**.

A mounted filesystem allows:

- **Accessing external storage devices (USB drives, hard disks, network drives).**
- **Expanding system storage by mounting additional partitions.**
- **Sharing resources between multiple users or systems via NFS (Network File System).**

Key Mounting Concepts

1. **Mount Point** – A directory where the filesystem is attached (e.g., /mnt, /media).
2. **Filesystem Type** – Determines the structure and rules of data storage (ext4, XFS, NTFS, NFS).
3. **Mounting and Unmounting** – Attaching and detaching a filesystem using mount and umount commands.

Example: Checking Mounted File Systems

To see all mounted filesystems:

```
mount | column -t
```

Or use:

```
df -h
```

This displays:

Filesystem	Size	Used	Avail	Use%	Mounted on
/dev/sda1	100G	45G	50G	47%	/
/dev/sdb1	500G	120G	350G	25%	/mnt/storage

Exercise

1. List all mounted filesystems on your system using `df -h`.
2. Identify the filesystem type of a mounted partition using `blkid`.

Case Study: Expanding Storage for a Web Server

A company's web server runs out of space. Instead of reinstalling the system, administrators **attach a new hard disk and mount it as `/var/www`**, expanding web storage without downtime.

CHAPTER 2: MANUALLY MOUNTING AND UNMOUNTING FILE SYSTEMS

Mounting a File System

To manually mount a filesystem, use the mount command:

```
sudo mount -t ext4 /dev/sdb1 /mnt/storage
```

- `-t ext4` → Specifies the filesystem type.

- `/dev/sdb1` → The partition being mounted.
- `/mnt/storage` → The mount point.

Making the Mount Permanent

To auto-mount at boot, add an entry to `/etc/fstab`:

```
/dev/sdb1 /mnt/storage ext4 defaults 0 2
```

After modifying `/etc/fstab`, apply changes:

```
sudo mount -a
```

Unmounting a File System

Before removing a device, unmount it:

```
sudo umount /mnt/storage
```

If the device is busy, force unmount:

```
sudo umount -l /mnt/storage
```

Example: Mounting a USB Drive Manually

If a USB drive is detected as `/dev/sdc1`:

```
sudo mkdir /mnt/usb
```

```
sudo mount /dev/sdc1 /mnt/usb
```

```
ls /mnt/usb # Verify access
```

Exercise

1. Mount a newly attached storage device and verify its contents.
2. Unmount the device and check if it's successfully removed.

Case Study: Mounting External Storage for Backup Servers

A backup server needs additional space for archiving files. Administrators **connect an external 1TB disk and mount it as /backup**, ensuring smooth daily backups.

CHAPTER 3: INTRODUCTION TO NETWORK FILE SYSTEM (NFS)

What is NFS?

The **Network File System (NFS)** allows UNIX/Linux systems to **share directories and files over a network**. It enables remote access to files as if they were local, making it useful for:

- **Centralized file storage** – Shared home directories and project files.
- **Efficient data management** – Allowing multiple systems to access shared resources.
- **Cloud and distributed computing** – Mounting remote filesystems in enterprise networks.

How NFS Works

- **NFS Server** – Shares a directory over the network.
- **NFS Client** – Mounts the shared directory and accesses files remotely.
- **NFS Protocol** – Uses TCP/UDP for communication, commonly over port 2049.

Example: Checking NFS Services on a System

Verify if NFS is running:

```
sudo systemctl status nfs-server
```

Check active NFS shares:

```
showmount -e <server-ip>
```

Exercise

1. Find out if your system has nfs-utils or nfs-common installed.
2. Identify available NFS shares on a remote server.

Case Study: Using NFS for Multi-User Data Access in a Research Lab

A university research lab needs a **shared storage solution** for all researchers. By setting up **an NFS server**, all lab computers can access research data in real-time without duplication.

CHAPTER 4: SETTING UP AN NFS SERVER AND CLIENT

Configuring an NFS Server

1. **Install NFS Server Packages**
 2. `sudo apt install nfs-kernel-server` # Debian-based
 3. `sudo yum install nfs-utils` # Red Hat-based
4. **Create a Shared Directory**
 5. `sudo mkdir -p /srv/nfs/shared`
 6. `sudo chown nobody:nogroup /srv/nfs/shared`
7. **Edit the NFS Exports File**

Add the following line to `/etc/exports`:
8. `/srv/nfs/shared 192.168.1.0/24(rw,sync,no_root_squash)`

This allows read/write access for clients in the 192.168.1.x network.

9. Restart NFS Server

10. `sudo systemctl restart nfs-server`

11. `sudo exportfs -rav`

Configuring an NFS Client

1. Install NFS Client Packages

2. `sudo apt install nfs-common # Debian-based`

3. `sudo yum install nfs-utils # Red Hat-based`

4. Mount the NFS Share

5. `sudo mount -t nfs 192.168.1.100:/srv/nfs/shared /mnt/nfs`

6. Make the Mount Persistent

Add to `/etc/fstab`:

7. `192.168.1.100:/srv/nfs/shared /mnt/nfs nfs defaults o o`

Example: Verifying NFS Connectivity

To confirm successful NFS mounting:

`ls /mnt/nfs`

Exercise

1. Set up a temporary NFS share on your local machine and access it from another system.
2. Configure `/etc/fstab` to auto-mount an NFS share at boot.

Case Study: Using NFS for Distributed Software Development

A software development company sets up an **NFS server** to store source code. Developers across multiple locations mount the **same shared directory**, enabling real-time collaboration.

CHAPTER 5: MANAGING AND TROUBLESHOOTING MOUNTING AND NFS ISSUES

Common Mounting Issues and Fixes

1. "Device is busy" error when unmounting

- Find active processes using the mount:
- `sudo lsof /mnt/storage`
- Kill the process and retry:
- `sudo umount -l /mnt/storage`

2. NFS mount hangs or fails

- Check if the NFS server is running:
- `sudo systemctl status nfs-server`
- Verify firewall rules:
- `sudo ufw allow from 192.168.1.0/24 to any port 2049`

Example: Debugging NFS Mount Failures

To check NFS logs for errors:

```
journalctl -xe | grep nfs
```

Exercise

1. Simulate an NFS failure and troubleshoot it.

2. Verify network connectivity between an NFS server and a client.

Case Study: Resolving NFS Access Issues in a Production Environment

A company's shared NFS directory **becomes inaccessible** due to a firewall update. By checking logs and updating firewall rules, administrators restore access without downtime.

CONCLUSION

This guide covered:

- **Mounting local and network file systems.**
- **Setting up and managing NFS shares.**
- **Troubleshooting common mounting and NFS issues.**

PROCESS AND JOB SCHEDULING IN UNIX/LINUX

CHAPTER 1: INTRODUCTION TO PROCESS AND JOB SCHEDULING

What is Process and Job Scheduling?

In UNIX/Linux, **process and job scheduling** refers to the **management and execution of processes and tasks** at different times based on user or system requirements. A **process** is an instance of a running program, while **job scheduling** allows users to execute tasks at a specified time automatically.

Efficient process and job scheduling is crucial for:

- **Managing system performance** by prioritizing tasks.
- **Automating repetitive tasks** like backups, updates, and maintenance.
- **Optimizing CPU usage** by distributing workloads efficiently.
- **Running background processes** without user intervention.

The UNIX/Linux kernel manages processes and job scheduling using process states, job control commands (fg, bg, jobs), and scheduling tools like cron and at.

Example: Listing Running Processes

Use the ps command to list active processes:

```
ps aux
```

This displays:

```
USER  PID %CPU %MEM  COMMAND
```



```
root 1234 0.1 2.5 /usr/bin/apache2
```

```
john 5678 1.2 1.0 /usr/bin/firefox
```

Exercise

1. Use `ps -ef` and `top` to analyze active processes.
2. Identify the highest CPU-consuming process using `top`.

Case Study: Managing System Performance in a High-Traffic Web Server

A company's web server slows down due to **high CPU usage**. By analyzing process scheduling and setting priority levels (`nice`, `renice`), administrators optimize performance by reducing CPU allocation for background processes.

CHAPTER 2: UNDERSTANDING PROCESS MANAGEMENT IN UNIX/LINUX

Process Lifecycle and States

Each process in UNIX/Linux has a **lifecycle** and goes through different states:

1. **Created (New)** – The process is created but not yet running.
2. **Ready** – The process is waiting for CPU time.
3. **Running** – The process is executing instructions.
4. **Waiting (Sleeping)** – The process is waiting for an event (e.g., I/O operation).
5. **Terminated (Zombie)** – The process has completed but still occupies memory.

Use the following command to check process states:

```
ps -eo pid,stat,cmd
```

- R → Running
- S → Sleeping
- Z → Zombie

Example: Killing a Process Using kill

Find and terminate a process with:

```
ps aux | grep firefox
```

```
kill 5678
```

Exercise

1. Create a process using `sleep 100 &` and list it using `ps -ef`.
2. Terminate the process using `kill` and verify its status.

Case Study: Handling Unresponsive Applications in Enterprise Systems

A financial institution runs **real-time stock trading applications**. When a process hangs, administrators use `kill -9` to force termination and ensure **continuous system availability**.

CHAPTER 3: JOB SCHEDULING AND BACKGROUND PROCESSES

Managing Background and Foreground Jobs

A job can run **in the foreground** (requiring user interaction) or **in the background** (running independently).

Running a Process in the Background

Append & to a command:

```
sleep 100 &
```

Check background jobs using:

```
jobs
```

Bringing Jobs to Foreground and Background

Move a job from background to foreground:

```
fg %1
```

Send a job to the background:

```
bg %1
```

Example: Running a Long Task in the Background

If a user is downloading a large file:

```
wget https://example.com/largefile.iso &
```

This allows the user to continue working while the download progresses in the background.

Exercise

1. Start a process in the background and bring it to the foreground using fg.
2. List all running background jobs and terminate one.

Case Study: Running Multiple Computational Jobs on a Research Cluster

A research institute runs **machine learning models** that take hours to complete. By running them in the background using `nohup` and `bg`, scientists can **log out while the processes continue running**.

CHAPTER 4: SCHEDULING TASKS WITH CRON JOBS

What is cron?

The cron daemon automates **scheduled execution of tasks**. It reads the **crontab (cron table)**, which contains scheduled jobs for users and the system.

Creating a Cron Job

1. Open the crontab editor:
2. `crontab -e`
3. Add a cron job to run a backup script daily at midnight:
4. `0 0 * * * /home/user/backup.sh`
5. Save and exit.

Understanding Cron Syntax

`* * * * * command_to_execute`

| | | | |

| | | | | _____ Day of the week (0-7, 0=Sunday)

| | | | _____ Month (1-12)

| | | _____ Day of the month (1-31)

| | _____ Hour (0-23)

_____ Minute (0-59)

Common Cron Job Examples

Task	Cron Syntax
Run every 10 minutes	<code>*/10 * * * * /script.sh</code>
Run every Monday at 5 AM	<code>0 5 * * 1 /script.sh</code>
Run on the 1st of every month	<code>0 0 1 * * /script.sh</code>

Example: Scheduling a Disk Cleanup Task

```
0 3 * * 0 rm -rf /var/log/*.log
```

This deletes log files **every Sunday at 3 AM**.

Exercise

1. Schedule a script to run **every 5 minutes**.
2. Modify a cron job to execute a script **at 10 PM daily**.

Case Study: Automating Server Updates with cron

A cloud provider schedules **automated system updates** at midnight using cron jobs, ensuring that all servers receive **critical security patches without manual intervention**.

CHAPTER 5: ONE-TIME JOB SCHEDULING WITH AT AND BATCH

Using at for One-Time Jobs

The at command schedules a **one-time execution** of a command.

Example: Running a Task at a Specific Time

1. Open the at prompt:

2. at 5:30 PM
3. Enter the command:
4. `echo "System update" >> /var/log/sysupdate.log`
5. Press Ctrl + D to save.

Viewing and Managing Scheduled at Jobs

- List pending jobs:
- `atq`
- Remove a scheduled job:
- `atrm job_id`

Using batch for System Load-Based Scheduling

The batch command runs jobs when system load is low.

```
echo "tar -czf backup.tar.gz /home/user" | batch
```

Example: Scheduling a Database Backup with at

at midnight

```
mysqldump -u root -p database > /backup/db.sql
```

Exercise

1. Schedule a task to **run 30 minutes from now** using at.
2. Use batch to schedule a job when CPU usage is low.

Case Study: Scheduling a One-Time Job for Software Installation

A system administrator uses at to install **security updates outside working hours**, ensuring minimal disruption to services.

CONCLUSION

This guide covered:

- ✓ **Managing processes (ps, kill, jobs).**
- ✓ **Running background tasks (bg, fg, nohup).**
- ✓ **Automating recurring jobs (cron) and one-time tasks (at).**

SYSTEM LOGS AND MONITORING IN UNIX/LINUX

CHAPTER 1: INTRODUCTION TO SYSTEM LOGS AND MONITORING

What are System Logs and Monitoring?

System logs and monitoring in UNIX/Linux play a crucial role in **troubleshooting, security auditing, and performance optimization**. Logs record system activities, user actions, and hardware or software issues, while monitoring tools help track system health in real time.

System logs and monitoring help in:

- **Detecting system errors and failures**
- **Analyzing security incidents**
- **Optimizing performance and resource utilization**
- **Troubleshooting software and hardware issues**

Log files in UNIX/Linux are stored in `/var/log/`, and monitoring tools like `top`, `htop`, `vmstat`, and `dstat` provide **real-time insights** into system performance.

Example: Viewing System Logs

To check system logs, use:

```
sudo cat /var/log/syslog
```

Or, to view logs dynamically:

```
sudo tail -f /var/log/syslog
```

Exercise

1. Check the latest system logs using `tail -f /var/log/syslog`.
2. Identify login attempts using `grep "sshd" /var/log/auth.log`.

Case Study: Investigating System Crashes with Logs

A company's database server crashes unexpectedly. By analyzing `/var/log/messages` and `/var/log/syslog`, administrators identify **high memory usage as the root cause**, preventing future crashes.

CHAPTER 2: UNDERSTANDING SYSTEM LOGS IN UNIX/LINUX

Common System Log Files

Log File	Purpose
<code>/var/log/syslog</code>	General system messages (Debian-based)
<code>/var/log/messages</code>	General system messages (RHEL-based)
<code>/var/log/auth.log</code>	User authentication logs
<code>/var/log/secure</code>	Security-related logs
<code>/var/log/dmesg</code>	Kernel messages at boot
<code>/var/log/cron</code>	Cron job execution logs
<code>/var/log/httpd/access.log</code>	Apache web server access logs
<code>/var/log/httpd/error.log</code>	Apache error logs

Checking System Logs

1. View logs using `cat`, `less`, or `tail`:

2. `sudo less /var/log/syslog`
3. Filter logs for specific events:
4. `grep "error" /var/log/syslog`
5. View authentication logs:
6. `sudo cat /var/log/auth.log`

Example: Checking User Login Attempts

```
sudo grep "Failed password" /var/log/auth.log
```

This helps detect unauthorized login attempts.

Exercise

1. Find the last failed login attempt in `/var/log/auth.log`.
2. Check kernel messages using `dmesg | tail -20`.

Case Study: Investigating Unauthorized Login Attempts

A company notices multiple failed SSH login attempts. By analyzing `/var/log/auth.log`, they identify a brute force attack and secure the system using **fail2ban**.

CHAPTER 3: REAL-TIME SYSTEM MONITORING TOOLS

Monitoring System Performance

Several UNIX/Linux tools help monitor system performance:

- **top** – Displays real-time CPU and memory usage.
- **htop** – Interactive version of top with color-coded output.
- **vmstat** – Monitors CPU, memory, and disk usage.

- **iostat** – Monitors disk I/O statistics.
- **netstat** – Displays network connections.

Using the top Command

The top command provides a **dynamic view** of system resources:

top

- PID – Process ID
- %CPU – CPU usage
- %MEM – Memory usage
- TIME+ – Total CPU time used

To sort processes by memory usage, press M.

Using htop for Interactive Monitoring

Install htop (if not installed):

```
sudo apt install htop # Debian-based
```

```
sudo yum install htop # RHEL-based
```

Run htop for a user-friendly process monitoring interface.

Example: Monitoring CPU and Memory Usage with vmstat

```
vmstat 5
```

This updates CPU, memory, and I/O stats every 5 seconds.

Exercise

1. Run top and identify the process using the most CPU.
2. Use htop to find processes consuming the most memory.

Case Study: Identifying Performance Bottlenecks in a Web Server

A company's web application runs slowly. Using `top` and `htop`, administrators find **Apache processes consuming excessive memory**. Restarting the service restores normal performance.

CHAPTER 4: NETWORK AND DISK MONITORING

Monitoring Network Activity

To check active network connections:

```
netstat -tunlp
```

To monitor live network traffic:

```
sudo iftop
```

Checking Disk Usage with `df` and `du`

- **View free disk space:**
 - `df -h`
- **Find the largest files in a directory:**
 - `sudo du -sh /var/*`

Example: Identifying Disk Space Issues

```
df -h | grep '/dev/sda'
```

This helps detect **full partitions** that need cleanup.

Exercise

1. List the top 10 largest files in `/var/log`.
2. Monitor live network connections using `netstat`.

Case Study: Preventing Disk Overflow in Database Servers

A database server crashes due to a **full /var/log partition**. Using `df -h` and `du -sh`, administrators clear old logs and schedule **automatic log rotation**.

CHAPTER 5: LOG MANAGEMENT AND AUTOMATION

Rotating Logs with Logrotate

Log files grow over time and must be managed efficiently. **Logrotate** automatically rotates, compresses, and removes old logs.

Configuring Log Rotation

Check existing log rotation rules:

```
cat /etc/logrotate.conf
```

To rotate logs daily, edit `/etc/logrotate.d/custom`:

```
/var/log/custom.log {
```

```
    daily
```

```
    rotate 7
```

```
    compress
```

```
    missingok
```

```
    notifempty
```

```
}
```

This rotates logs daily, keeps 7 copies, and compresses old logs.

Automating Monitoring with monit

monit automatically restarts failed services. Install it using:

```
sudo apt install monit # Debian-based
```

```
sudo yum install monit # RHEL-based
```

Configure monit to monitor Apache:

check process apache with pidfile /var/run/apache2.pid

```
start program = "/etc/init.d/apache2 start"
```

```
stop program = "/etc/init.d/apache2 stop"
```

Restart monit:

```
sudo systemctl restart monit
```

Example: Automating Log Cleanup

Schedule log cleanup with cron:

```
0 3 * * * sudo find /var/log -name "*.log" -mtime +7 -delete
```

This deletes logs older than 7 days **every night at 3 AM**.

Exercise

1. Configure log rotation for a custom log file.
2. Set up a cron job to clear logs every week.

Case Study: Automating System Health Monitoring in a Cloud Infrastructure

A cloud provider uses **Logrotate for log management** and monit to **restart failed services automatically**, ensuring **99.99% uptime**.

CONCLUSION

This guide covered:

- ✓ Understanding system logs (syslog, auth.log, dmesg).
- ✓ Real-time monitoring with top, htop, vmstat.
- ✓ Network and disk monitoring (netstat, df, du).
- ✓ Automating log management (logrotate, monit).

ISDM-NxT

ASSIGNMENT SOLUTION: SETTING UP AND MANAGING USERS WITH DIFFERENT ACCESS LEVELS IN UNIX/LINUX

Objective

This assignment provides a step-by-step guide to **create, modify, and manage users with different access levels** in UNIX/Linux. You will:

- ✓ Create users and assign them to different groups.
- ✓ Configure file permissions for access control.
- ✓ Set up **sudo privileges** for administrative users.
- ✓ Manage restricted and privileged access levels.

STEP 1: CREATE USERS WITH DIFFERENT ACCESS LEVELS

1. Create Standard Users

Standard users **have limited access** and cannot execute administrative tasks.

To create users:

```
sudo useradd -m -s /bin/bash user1
```

```
sudo useradd -m -s /bin/bash user2
```

- -m → Creates a home directory.
- -s /bin/bash → Assigns the default shell.

Set passwords for the users:

```
sudo passwd user1
```



```
sudo passwd user2
```

2. Create Administrative Users (With Sudo Access)

Admins can execute **root-level** commands.

```
sudo useradd -m -s /bin/bash adminuser
```

```
sudo passwd adminuser
```

```
sudo usermod -aG sudo adminuser
```

- `-aG sudo` → Adds adminuser to the **sudo group** for administrative privileges.

3. Verify User Access Levels

Check user groups:

```
groups user1
```

```
groups adminuser
```

STEP 2: MANAGING USER GROUPS AND ACCESS CONTROL

1. Create Custom Groups

Create groups for specific access control:

```
sudo groupadd developers
```

```
sudo groupadd hr
```

2. Assign Users to Groups

Add user1 to **developers** and user2 to **hr**:

```
sudo usermod -aG developers user1
```

```
sudo usermod -aG hr user2
```

Verify group assignments:

```
groups user1
```

```
groups user2
```

STEP 3: SET FILE AND DIRECTORY PERMISSIONS

1. Restrict Access to Sensitive Files

Make a directory accessible **only** to developers:

```
sudo mkdir /project
```

```
sudo chown :developers /project
```

```
sudo chmod 770 /project
```

- `chown :developers /project` → Assigns group ownership.
- `chmod 770 /project` → Only **developers** can access it.

2. Protect System Files from Standard Users

```
sudo chmod 640 /etc/shadow
```

- This prevents non-root users from reading **password hashes**.

3. Test Access Restrictions

Log in as user1 and attempt to access /project:

```
cd /project
```

If user1 is in developers, access is granted. Otherwise, permission is denied.

STEP 4: CONFIGURE SUDO PRIVILEGES FOR SPECIFIC USERS

1. Grant Limited Sudo Access

Edit sudoers file using visudo:

```
sudo visudo
```

Add this line to allow adminuser to restart the web server but prevent other admin commands:

```
adminuser ALL=(ALL) NOPASSWD: /bin/systemctl restart apache2
```

- NOPASSWD: → Allows execution without password entry.
- systemctl restart apache2 → Limits commands to web server restarts.

2. Test Sudo Access

Switch to adminuser and try:

```
sudo systemctl restart apache2
```

If configured correctly, the command will execute without a password prompt.

STEP 5: MANAGING USER ACCOUNTS

1. Lock an Inactive User Account

To temporarily disable user2:

```
sudo usermod -L user2
```

To unlock the account:

```
sudo usermod -U user2
```

2. Delete a User and Their Home Directory

If an employee leaves the company:

```
sudo userdel -r user1
```

- -r → Deletes user's home directory.

STEP 6: VERIFY AND MONITOR USER ACCESS

1. Check Active Users

```
who
```

Lists logged-in users.

2. Check Login Attempts

```
sudo cat /var/log/auth.log | grep "user1"
```

Displays login activity for user1.

3. Monitor User Activity

```
sudo lastlog
```

Shows last login times for all users.

CONCLUSION

- ✓ Created and managed users with **different access levels**.
- ✓ Configured **file permissions and group access**.
- ✓ Assigned **sudo privileges** for administrative tasks.
- ✓ Implemented **account security measures**.

ASSIGNMENT SOLUTION: CONFIGURING AND MANAGING SYSTEM LOGS FOR ERROR TRACKING IN UNIX/LINUX

Objective

This assignment provides a step-by-step guide to configuring and managing system logs for **error tracking, monitoring system health, and troubleshooting issues** in UNIX/Linux. You will:

- ✓ Configure **log rotation** to manage system logs.
- ✓ Use **system logging tools** (rsyslog, journalctl) for error tracking.
- ✓ Implement **custom logging** for applications.
- ✓ Automate **log monitoring and alerts**.

STEP 1: UNDERSTANDING SYSTEM LOGS

1. Identify Key Log Files

System logs are stored in `/var/log/`. Some important logs include:

Log File	Purpose
<code>/var/log/syslog</code>	General system events (Debian-based)
<code>/var/log/messages</code>	General system events (RHEL-based)
<code>/var/log/auth.log</code>	User authentication logs
<code>/var/log/secure</code>	Security logs
<code>/var/log/dmesg</code>	Kernel messages at boot
<code>/var/log/cron</code>	Cron job execution logs
<code>/var/log/apache2/error.log</code>	Apache web server errors

2. View System Logs

Check system logs for errors:

```
sudo cat /var/log/syslog | grep "error"
```

Monitor logs in real time:

```
sudo tail -f /var/log/syslog
```

Exercise

1. Identify the last 10 login attempts using:
2. `sudo tail -10 /var/log/auth.log`
3. Search for failed SSH login attempts:
4. `sudo grep "Failed password" /var/log/auth.log`

STEP 2: CONFIGURING LOGGING WITH RSYSLOG

1. Verify and Enable rsyslog

Most modern Linux distributions use rsyslog for logging. Check if rsyslog is running:

```
sudo systemctl status rsyslog
```

If not running, start it:

```
sudo systemctl start rsyslog
```

```
sudo systemctl enable rsyslog
```

2. Configure Custom Log Rules

Edit the configuration file:

```
sudo nano /etc/rsyslog.conf
```

Add a rule to log kernel errors to a custom file:

```
kern.* /var/log/kernel_errors.log
```

Restart rsyslog to apply changes:

```
sudo systemctl restart rsyslog
```

3. Test Custom Logging

Send a test message to the log file:

```
logger -p kern.err "Test kernel error log"
```

Verify the log entry:

```
sudo cat /var/log/kernel_errors.log
```

Exercise

1. Configure rsyslog to log authentication failures to /var/log/auth_failures.log.
2. Test the logging rule using logger.

STEP 3: LOG ROTATION FOR EFFICIENT LOG MANAGEMENT

1. Configure Log Rotation Using Logrotate

Logrotate prevents logs from **consuming excessive disk space** by rotating and compressing logs.

Check existing log rotation settings:

```
cat /etc/logrotate.conf
```

2. Configure Log Rotation for a Custom Log File

Create a new configuration file:

```
sudo nano /etc/logrotate.d/custom_logs
```

Add the following configuration:

```
/var/log/kernel_errors.log {
```

```
weekly  
rotate 4  
compress  
missingok  
notifempty  
}
```

- weekly → Rotates logs weekly.
- rotate 4 → Keeps 4 backups.
- compress → Compresses old logs.

3. Test Log Rotation

Manually run log rotation:

```
sudo logrotate -f /etc/logrotate.d/custom_logs
```

Exercise

1. Set up log rotation for /var/log/auth_failures.log to run daily.
2. Verify the rotated log files in /var/log/.

STEP 4: MONITORING LOGS WITH JOURNALCTL

1. View Recent System Logs

```
journalctl -xe
```

2. Filter Logs by Priority

Show only error logs:

```
journalctl -p err
```


3. Display Logs for a Specific Service

Example: View SSH logs

```
journalctl -u sshd --since "1 hour ago"
```

Exercise

1. Find system errors that occurred in the last 30 minutes.
2. List log entries related to nginx service.

STEP 5: AUTOMATING LOG MONITORING AND ALERTS

1. Set Up Logwatch for Log Analysis

Install logwatch (if not installed):

```
sudo apt install logwatch # Debian-based
```

```
sudo yum install logwatch # RHEL-based
```

Generate a log report:

```
sudo logwatch --detail high --mailto root --range yesterday
```

2. Configure Email Alerts for Log Events

Use logwatch to send email alerts for errors:

Edit `/etc/logwatch/conf/logwatch.conf`:

MailTo = admin@example.com

Range = yesterday

Detail = High

3. Set Up a Cron Job for Daily Log Reports

Open crontab:

```
crontab -e
```

Add the following entry to send reports daily at midnight:

```
0 0 * * * /usr/sbin/logwatch --output mail
```

Exercise

1. Configure logwatch to send error reports to your email.
2. Set up a cron job to automate log analysis.

STEP 6: SECURING AND MANAGING LOGS

1. Protect Log Files from Unauthorized Access

Ensure logs are only accessible to administrators:

```
sudo chmod 640 /var/log/auth.log
```

2. Restrict Remote Log Access

To disable remote logging, edit `/etc/rsyslog.conf`:

```
# Remove or comment the following line
```

```
# $ModLoad imudp
```

```
# $UDPServerRun 514
```

Restart the logging service:

```
sudo systemctl restart rsyslog
```

Exercise

1. Modify permissions for `/var/log/syslog` to restrict access.
2. Disable remote logging in `rsyslog` and restart the service.

STEP 7: TROUBLESHOOTING LOG ISSUES

1. Logs Not Updating? Check Disk Space

If logs are missing or not updating, check available disk space:

```
df -h
```

If /var is full, free space by clearing old logs:

```
sudo rm -rf /var/log/*.gz
```

2. rsyslog Not Running? Restart the Service

Check the status:

```
sudo systemctl status rsyslog
```

Restart it:

```
sudo systemctl restart rsyslog
```

3. Logs Not Rotating? Debug logrotate

Manually force rotation and check errors:

```
sudo logrotate -d /etc/logrotate.conf
```

Exercise

1. Check system logs for errors related to rsyslog.
2. Identify the last failed login attempt and its timestamp.

CONCLUSION

- ✓ Configured system logs for **error tracking**.
- ✓ Implemented **log rotation** using logrotate.
- ✓ Automated **log analysis and email alerts** using logwatch.
- ✓ Secured logs from **unauthorized access**.

ISDM-NxT

ISDM-NxT

ISDM-NxT