



ISDM (INDEPENDENT SKILL DEVELOPMENT MISSION

AWS NETWORKING

AWS VPC (VIRTUAL PRIVATE CLOUD) BASICS – STUDY MATERIAL

INTRODUCTION TO AWS VPC

What is AWS VPC?

AWS **Virtual Private Cloud (VPC)** is a **logically isolated** network in AWS that allows users to launch and manage AWS resources securely. It functions like a **private data center in the cloud**, providing control over networking, security, and internet access.

Key Features of AWS VPC

- ✓ Logical Network Isolation Users can define their own IP address ranges.
- ✓ **Subnetting** Allows creation of **public and private subnets** for different workloads.
- ✓ Internet Gateway (IGW) Enables internet access for resources inside VPC.
- ✓ Security Groups & Network ACLs Control inbound and outbound traffic.
- ✓ Route Tables Define how network traffic is directed.

✓ VPN & Direct Connect – Securely connect on-premises networks to AWS.

CHAPTER 1: AWS VPC ARCHITECTURE & COMPONENTS

1. Key Components of AWS VPC

Component	Description	
VPC (Virtual Private	A logically isolated network in AWS.	
Cloud)		
Subnets	Smaller network divisions within a VPC	
	(Public & Private Subnets).	
Internet Gateway (IGW)	Allows internet access for public subnets.	
NAT Gateway	Enables private subnet instances to access the internet.	
Route Tables	Define how traffic flows between subnets	
	and the internet.	
Security Groups	Controls inbound/outbound traffic for EC2	
	instances.	
Network ACLs	Stateless firewall controlling subnet-level	
(NACLs)	traffic.	
VPC Peering	Connects two VPCs for private	
	communication.	
VPN & Direct	Securely connects AWS VPC with on-	
Connect	premises networks.	

* Example:

A web application may use a public subnet for a web server and a private subnet for a database to ensure security.

CHAPTER 2: CREATING AN AWS VPC STEP-BY-STEP

Step 1: Create a VPC

- 1. Open **AWS Console** → Navigate to **VPC Dashboard**.
- 2. Click "Create VPC".
- 3. **VPC Name:** MyAppVPC.
- 4. **IPv4 CIDR Block:** 10.0.0.0/16 (Allows 65,536 IPs).
- 5. Click "Create VPC".

Expected Outcome:

A **new VPC** is created with a **private IP range** (10.0.0.0/16).

Step 2: Create Subnets

- 1. Open VPC Console → Click "Subnets" → Create Subnet.
- 2. Subnet 1 (Public Subnet):
 - Name: Public-Subnet.
 - o CIDR Block: 10.0.1.0/24.
 - Enable Auto-Assign Public IPv4: Yes.
- 3. Subnet 2 (Private Subnet):
 - Name: Private-Subnet.
 - o CIDR Block: 10.0.2.0/24.

Enable Auto-Assign Public IPv4: No.

Expected Outcome:

Two subnets are created:

- ✓ Public Subnet (10.0.1.0/24) Internet access enabled.
- ✓ Private Subnet (10.0.2.0/24) No direct internet access.

Step 3: Attach an Internet Gateway (IGW)

- Navigate to VPC Console → Internet Gateways.
- Click "Create Internet Gateway" → Name it MyIGW.
- Click "Attach to VPC" → Select MyAppVPC.

Expected Outcome:

The **Internet Gateway is attached**, allowing public subnet resources to access the internet.

Step 4: Configure Route Tables

- Navigate to VPC Console → Route Tables.
- 2. Public Route Table:
 - Select Public Route Table → Click "Edit Routes".
 - Add o.o.o.o/o \rightarrow Target: Internet Gateway (MyIGW).
- 3. Private Route Table:
 - o Used for internal traffic (no internet access).

Expected Outcome:

✓ Public Subnet → Can access the internet via Internet Gateway.

✓ Private Subnet → No direct internet access (for database security).

Step 5: Add a NAT Gateway (For Private Subnet Internet Access)

- Navigate to VPC Console → NAT Gateways → Click "Create NAT Gateway".
- 2. Select Public Subnet (10.0.1.0/24).
- 3. Elastic IP: Allocate a new Elastic IP (EIP).
- 4. Click Create NAT Gateway.
- 5. Update Private Route Table:
 - o Go to Route Tables → Select Private Route Table.
 - Add o.o.o.o/o → Target: NAT Gateway.

Expected Outcome:

✓ Private Subnet instances can access the internet but remain hidden from incoming connections.

CHAPTER 3: CONFIGURING SECURITY GROUPS & NETWORK ACLS

- 1. Creating a Security Group for Web Server
 - 1. Navigate to EC2 Console → Security Groups.
 - 2. Click "Create Security Group".
 - 3. Security Group Name: WebServerSG.
 - 4. Inbound Rules:
 - \circ Allow HTTP (80) and HTTPS (443) from **o.o.o.o/o**.
 - Allow SSH (22) only from your IP.

***** Expected Outcome:

The Web Server can receive HTTP traffic but restricts SSH to specific IPs.

- 2. Creating a Security Group for Database
 - 1. Security Group Name: DatabaseSG.
 - 2. Inbound Rules:
 - Allow MySQL (3306) traffic only from WebServerSG.
- Expected Outcome:

The Database is secured and accessible only from the Web Server.

- 3. Configuring Network ACLs (Optional Extra Security)
 - Navigate to VPC Console → Network ACLs.
 - Click Create Network ACL → Associate it with MyAppVPC.
 - 3. Configure Inbound Rules:
 - Allow HTTP (80) and HTTPS (443).
 - Allow SSH (22) from your IP.
 - 4. Configure Outbound Rules:
 - Allow all traffic (o.o.o.o/o).
- ***** Expected Outcome:

A stateless firewall adds extra security to subnets.

CHAPTER 4: CONNECTING TO EC2 INSTANCE INSIDE THE VPC

1. Launch an EC2 instance in the Public Subnet.

- Assign WebServerSG security group to allow web traffic.
- 3. Connect to EC2 via SSH:
- 4. ssh -i my-key.pem ec2-user@your-public-ip
- 5. Check Internet Connectivity:
- 6. curl https://www.google.com
- * Expected Outcome: The public EC2 instance can access the internet.

CHAPTER 5: AWS VPC BEST PRACTICES

- ✓ Use Public Subnets for Web Servers and Private Subnets for Databases.
- ✓ Restrict SSH Access to Your IP using Security Groups.
- ✓ Use NAT Gateway for Private Subnet Internet Access.
- ✓ Enable VPC Flow Logs for Network Monitoring.
- ✓ Use AWS Direct Connect for Secure On-Premise Connectivity.
- * Example:

A finance company enforces strict security policies using private subnets for databases and public subnets for web applications.

CONCLUSION: MASTERING AWS VPC

By following this guide, you have:

- Created a VPC and Subnets.
- Configured Internet and NAT Gateways.
- Implemented Security Groups and Route Tables.
- Launched an EC2 instance inside the VPC.

FINAL EXERCISE:

- 1. Create a second VPC and set up VPC Peering between two VPCs.
- 2. **Enable AWS VPC Flow Logs** to monitor network traffic.
- 3. Set up a Site-to-Site VPN to connect an on-premises network to AWS.



AWS ROUTE 53 (DOMAIN NAME SYSTEM) – STUDY MATERIAL

INTRODUCTION TO AWS ROUTE 53

What is AWS Route 53?

AWS **Route 53** is a highly scalable and reliable **Domain Name System (DNS) web service** that helps manage domain registration,

DNS routing, and health checking of AWS-hosted applications. It

translates **human-readable domain names** (e.g.,

www.mywebsite.com) into **IP addresses** (e.g., 192.168.1.1) that

computers use to communicate.

Key Features of AWS Route 53

- ✓ **Domain Registration** Buy and manage domain names directly from AWS.
- ✓ **DNS Routing** Directs internet traffic to AWS resources or external servers.
- √ Traffic Routing Policies Supports latency-based, geolocation, failover, and weighted routing.
- ✓ Health Checks & Failover Monitors application health and redirects traffic when failures occur.
- ✓ Highly Available & Scalable AWS Route 53 is a global DNS service that provides 99.99% uptime.
- ✓ Secure & Compliant Integrated with AWS IAM, DNSSEC, and AWS Shield for security.

CHAPTER 1: AWS ROUTE 53 ARCHITECTURE & KEY CONCEPTS

1. AWS Route 53 Components

Component	Description
Hosted Zones	A collection of DNS records for a domain.
DNS Records	Defines how traffic is routed (A, CNAME, MX, TXT, etc.).
Traffic Policies	Controls how DNS queries are routed.
Health Checks	Monitors application availability & failover.
Domain Registration	Allows users to register new domains within AWS.

***** Example:

A website (www.myapp.com) hosted on EC2 uses Route 53 to map its domain name to an Elastic Load Balancer (ALB) IP address.

CHAPTER 2: REGISTERING A DOMAIN WITH AWS ROUTE 53

Step 1: Purchase a Domain Name

- Open AWS Console → Navigate to Route 53.
- 2. Click "Registered Domains" → Click "Register Domain".
- 3. Enter the desired domain name (e.g., myapp.com).
- 4. Select domain extension (e.g., .com, .org, .net).
- 5. Provide contact details (email, phone).
- 6. Click **"Purchase"** and wait for domain registration confirmation.

Expected Outcome:

Your domain is registered and ready to use in Route 53.

Step 2: Create a Hosted Zone

- 1. Open Route 53 Console → Click "Hosted Zones".
- 2. Click "Create Hosted Zone".
- 3. Enter the **Domain Name** (myapp.com).
- 4. Choose **Public Hosted Zone** (for internet-facing websites).
- 5. Click "Create".

Expected Outcome:

A new Hosted Zone is created with default NS (Name Server) and SOA (Start of Authority) records.

Chapter 3: Configuring DNS Records in Route 53

Common DNS Record Types

Record Type	Purpose
A Record (Address	Maps a domain to an IPv4 address.
Record)	
AAAA Record	Maps a domain to an IPv6 address.
CNAME Record	Maps a domain alias to another
(Canonical Name)	domain name.
MX Record (Mail	Routes emails to a mail server.
Exchange)	
TXT Record	Stores text data (used for domain
	verification & security).

* Example:

A business maps www.myapp.com to an EC2 instance using an A Record.

Type: A

Name: www

Value: 192.168.1.1

2. Adding an A Record to Point to an EC2 Instance

- 1. Open Route 53 Console → Select Hosted Zone.
- 2. Click "Create Record".
- 3. **Record Type:** A (Address Record).
- 4. Name: www.myapp.com.
- 5. Value: Enter the EC2 public IP (e.g., 54.32.11.100).
- 6. Click "Create Record".

Expected Outcome:

Users accessing www.myapp.com will be directed to your **EC2** instance.

3. Using a CNAME Record for Load Balancer Routing

- 1. Record Type: CNAME.
- 2. Name: www.myapp.com.
- 3. **Value:** Enter the **ALB DNS Name** (e.g., myapp-alb-123456.us-east-1.elb.amazonaws.com).
- 4. Click "Create Record".

Expected Outcome:

Users accessing www.myapp.com will be routed to the **Load Balancer**, improving scalability.

CHAPTER 4: IMPLEMENTING ROUTE 53 TRAFFIC ROUTING POLICIES

1. Simple Routing (Default Routing)

- Directs all traffic to a single IP or domain.
- Example: www.myapp.com → EC2 instance IP.

2. Weighted Routing (Load Distribution)

- Divides traffic across multiple resources based on assigned weights.
- Example: Send 70% of traffic to AWS ALB and 30% to an onpremises server.

3. Latency-Based Routing (Global Performance Optimization)

- Routes users to the nearest AWS region for better performance.
- Example: Users from US access us-east-1 while users from Europe access eu-west-1.

4. Failover Routing (High Availability)

- Redirects traffic to a secondary resource when the primary is unhealthy.
- Example: If primary EC2 instance fails, Route 53 directs traffic to a backup server.

5. Geolocation Routing (Region-Based Access Control)

- Routes users based on their geographic location.
- Example: Users from India access the India AWS region,
 while US users access the US AWS region.

CHAPTER 5: CONFIGURING HEALTH CHECKS & FAILOVER

- 1. Create a Health Check for an EC2 Instance
 - 1. Open Route 53 Console → Click "Health Checks".
 - 2. Click "Create Health Check".
 - 3. Monitor Endpoint: Enter the EC2 public IP or ALB URL.
 - 4. Protocol: HTTP/HTTPS.
 - 5. **Failure Threshold:** 3 consecutive failures.
 - 6. Create Health Check → Attach it to Failover Routing Policy.

Expected Outcome:

If the **primary server fails**, Route 53 redirects traffic to a **backup** server.

CHAPTER 6: ROUTE 53 SECURITY BEST PRACTICES

- ✓ Enable Route 53 DNSSEC (Domain Name System Security Extensions).
- √ Restrict changes to Route 53 records using IAM policies.
- ✓ Use AWS WAF with Route 53 to block malicious traffic.
- ✓ Monitor DNS query logs using AWS CloudTrail & CloudWatch.
- ✓ Use TTL (Time-to-Live) values efficiently to control DNS caching.

* Example:

A banking website enables DNSSEC and Route 53 logging to detect suspicious DNS requests.

CHAPTER 7: REAL-WORLD USE CASES OF AWS ROUTE 53

1. E-Commerce Website with Global Traffic

- ✓ Uses latency-based routing to send users to the closest AWS region.
- ✓ Implements **failover routing** to keep the website available.
- ✓ Secures DNS records using **AWS WAF & DNSSEC**.
- 2. API Gateway Routing for Microservices
- ✓ Uses subdomains (e.g., api.myapp.com → AWS API Gateway).
- ✓ Implements weighted routing to test new API versions.
- ✓ Monitors API health using Route 53 health checks.
- 3. Multi-Region Disaster Recovery Strategy
- ✓ Uses **failover routing** between AWS US and Europe regions.
- ✓ Implements **geolocation routing** for region-specific applications.

CONCLUSION: MASTERING AWS ROUTE 53

By using **AWS Route 53**, businesses can:

- Manage DNS records efficiently.
- Improve website performance with intelligent traffic routing.
- ✓ Implement failover for high availability.
- Secure DNS queries using AWS security best practices.

FINAL EXERCISE:

- 1. Create a hosted zone for a custom domain in Route 53.
- 2. Set up weighted routing between two AWS regions.
- 3. Enable DNSSEC for enhanced security.

AWS DIRECT CONNECT & VPN – STUDY MATERIAL

INTRODUCTION TO AWS DIRECT CONNECT & VPN

What is AWS Direct Connect?

AWS Direct Connect is a dedicated network connection between an on-premises data center and AWS. It bypasses the public internet, providing a more secure, reliable, and high-speed network link to AWS services.

What is AWS VPN (Virtual Private Network)?

AWS **VPN** is a **secure**, **encrypted tunnel** that connects on-premises networks or remote users to AWS. It provides **site-to-site or client-based** connectivity using **IPsec encryption** over the internet.

CHAPTER 1: AWS DIRECT CONNECT VS. AWS VPN – KEY DIFFERENCES

Feature	AWS Direct Connect	AWS VPN
Connection Type	Dedicated private network	Encrypted internet tunnel
Latency & Speed	Low latency (1Gbps – 100Gbps)	Higher latency (depends on ISP)
Security	Private connection (not over the internet)	Encrypted data transmission

Use Case	High-speed AWS access for enterprises	Secure remote access for small to medium businesses
Setup Time	Weeks (physical link required)	Hours (software-based setup)

***** Example:

A large enterprise with heavy cloud workloads uses Direct
Connect for high-speed AWS access, while a remote development
team uses AWS VPN for secure remote work.

CHAPTER 2: AWS DIRECT CONNECT - HOW IT WORKS

1. AWS Direct Connect Components

Component	Description	
AWS Direct Connect Location	A physical connection point between AWS and on-premises.	
Virtual Interface (VIF)	Logical interface that connects to AWS resources (Private, Public, Transit).	
DX Gateway	Connects Direct Connect to multiple VPCs across AWS Regions.	
Partner Network	AWS Direct Connect can be provisioned via an AWS Partner instead of AWS directly.	

2. Setting Up AWS Direct Connect

Step 1: Request a Direct Connection

- 1. Open **AWS Console** → Navigate to **Direct Connect**.
- 2. Click "Create Connection".
- 3. **Select AWS Direct Connect Location** (e.g., US-East-1, Equinix DC).
- 4. **Specify Connection Speed** (e.g., 1 Gbps).
- 5. Click "Request Connection".

Expected Outcome:

AWS approves the connection request and provides details for cross-connect setup.

Step 2: Configure the Direct Connect Link

- 1. Work with your **ISP or data center provider** to set up a **cross-connect** from your router to AWS's router.
- AWS provides a LOA-CFA (Letter of Authorization -Connecting Facility Assignment) for setup.
- 3. Configure **BGP** (**Border Gateway Protocol**) on your router to establish communication.

Expected Outcome:

A private network connection is now established between your data center and AWS.

Step 3: Create a Virtual Interface (VIF)

- 1. Open AWS Direct Connect Console.
- 2. Click "Create Virtual Interface".

3. Choose **VIF Type**:

- Private VIF Connects to VPCs via VGW.
- Public VIF Connects to AWS services like S₃ & DynamoDB.
- Click "Create VIF".

Expected Outcome:

Your Direct Connect link is ready to send private traffic to AWS.

CHAPTER 3: AWS VPN – How IT WORKS

1. AWS VPN Types

AWS VPN

Description

Type

Site-to-Site

Connects on-premises networks to AWS VPC over

VPN the internet.

Client VPN

Securely connects remote users (laptops, mobile

devices) to AWS.

Example:

A company HQ in New York uses Site-to-Site VPN to connect to AWS, while remote employees use Client VPN to work securely.

Setting Up a Site-to-Site VPN

Step 1: Create a Virtual Private Gateway (VGW)

Open AWS VPC Console → Click "Virtual Private Gateways".

- 2. Click "Create Virtual Private Gateway".
- Enter VGW Name (MyVPNGateway).
- 4. Click "Create" → Attach it to your VPC.

Expected Outcome:

The VGW is ready to handle VPN traffic.

Step 2: Create a Customer Gateway (CGW)

- Open AWS VPC Console → Click "Customer Gateways".
- 2. Click "Create Customer Gateway".
- 3. **IP Address:** Enter the **public IP** of your on-premises router.
- 4. Routing Type: Choose Static (for fixed routes) or Dynamic (for BGP).
- 5. Click "Create".

Expected Outcome:

AWS recognizes your on-premises router as a VPN peer.

Step 3: Create a Site-to-Site VPN Connection

- Open AWS VPC Console → Click "VPN Connections".
- 2. Click "Create VPN Connection".
- 3. Select Virtual Private Gateway (VGW) created earlier.
- 4. Select Customer Gateway (CGW) created earlier.
- 5. Choose Routing Option (Static or Dynamic).

Click "Create VPN Connection".

Expected Outcome:

A **secure VPN tunnel** is established between AWS and your onpremises network.

Step 4: Configure On-Premises Router (IPsec Tunnel)

- Download the VPN Configuration File from AWS.
- 2. Use Cisco, Juniper, or Fortinet routers to establish an IPsec tunnel.
- 3. Configure **IKE Phase 1 & 2 settings** (encryption, authentication).
- 4. Test the VPN tunnel status in AWS Console.

Expected Outcome:

The on-premises router successfully connects to AWS via VPN.

CHAPTER 4: MONITORING AWS DIRECT CONNECT & VPN

- 1. Monitoring Direct Connect with CloudWatch
- ✓ View network utilization, packet loss, and BGP status.
- ✓ Set CloudWatch Alarms for network failures.
- ✓ Use **AWS Direct Connect Resiliency Toolkit** for failover testing.

* Example:

A financial company monitors Direct Connect to ensure low latency for trading applications.

2. Monitoring AWS VPN Connections

- ✓ View VPN connection status in AWS Console.
- ✓ Enable VPC Flow Logs to track VPN traffic.
- ✓ Use **AWS Network Manager** for centralized VPN monitoring.

***** Example:

A global enterprise tracks VPN traffic logs to detect unauthorized access attempts.

CHAPTER 5: AWS DIRECT CONNECT & VPN BEST PRACTICES

- ✓ Use AWS Direct Connect for low-latency, high-bandwidth workloads.
- ✓ Enable VPN as a backup for Direct Connect in case of failure.
- √ Configure BGP routing for dynamic route updates.
- ✓ Use multi-region Direct Connect links for redundancy.
- ✓ Monitor connection health with CloudWatch & VPC Flow Logs.

* Example:

A media streaming company combines Direct Connect for highspeed uploads and VPN as a backup.

CHAPTER 6: REAL-WORLD USE CASES OF AWS DIRECT CONNECT & VPN

1. Hybrid Cloud Deployment

- ✓ **Direct Connect links on-premises data centers** to AWS for high-performance applications.
- ✓ VPN provides secure failover in case of Direct Connect failures.

2. Remote Workforce Security

- ✓ AWS Client VPN allows employees to work remotely with encrypted access.
- ✓ Multi-Factor Authentication (MFA) enhances security.
- 3. Disaster Recovery & Business Continuity
- ✓ Use Direct Connect for real-time database replication.
- ✓ Enable VPN as a backup for AWS access in emergencies.

CONCLUSION: MASTERING AWS DIRECT CONNECT & VPN

By implementing AWS Direct Connect & VPN, businesses can:

- Securely connect on-premises networks to AWS.
- Reduce latency and improve cloud performance.
- Ensure failover using VPN as a backup for Direct Connect.
- Enhance security with encryption and traffic monitoring.

FINAL EXERCISE:

- Create a Site-to-Site VPN Connection in AWS.
- 2. Set up Direct Connect using an AWS Partner.
- 3. Monitor VPN connection health using CloudWatch.

AWS DEVOPS TOOLS

AWS CODECOMMIT, CODEBUILD, AND CODEPIPELINE – STUDY MATERIAL

INTRODUCTION TO AWS CI/CD SERVICES

What are AWS CodeCommit, CodeBuild, and CodePipeline?

AWS provides a **fully managed Continuous Integration & Continuous Deployment (CI/CD) pipeline** through its DevOps services:

- ✓ AWS CodeCommit A private Git repository service for source code management.
- ✓ AWS CodeBuild A serverless build service that compiles, tests, and packages applications.
- ✓ AWS CodePipeline A CI/CD automation service that orchestrates build, test, and deployment workflows.

Key Features of AWS CI/CD Services

- ✓ Fully managed and scalable AWS handles infrastructure and scaling.
- ✓ Integration with AWS IAM, S₃, EC₂, Lambda, and other AWS services.
- ✓ **Automated testing & deployment** Streamlines DevOps processes.
- ✓ Supports multiple programming languages (Java, Python, Node.js, etc.).
- ✓ **Security & Compliance** IAM-based access control and encryption.

CHAPTER 1: AWS CODECOMMIT – SOURCE CODE MANAGEMENT

1. What is AWS CodeCommit?

AWS **CodeCommit** is a **fully managed, secure Git repository** service that enables teams to store and manage their source code in AWS.

2. Benefits of AWS CodeCommit

- ✓ Unlimited private repositories No size restrictions.
- ✓ **High availability** Fully managed, backed by AWS infrastructure.
- ✓ Integrated with AWS IAM Granular permission control for users.
- ✓ **Supports Git CLI** Works with existing Git commands.

3. Setting Up AWS CodeCommit

Step 1: Create a CodeCommit Repository

- Open AWS CodeCommit Console → Click "Create Repository".
- Enter Repository Name (MyAppRepo).
- Add an optional description.
- 4. Click "Create".

Step 2: Connect to the Repository

- 1. Clone using HTTPS (For IAM users with AWS credentials)
- git clone https://git-codecommit.us-east-1.amazonaws.com/v1/repos/MyAppRepo
- cd MyAppRepo
- 4. Clone using SSH (For SSH key authentication)

git clone ssh://git-codecommit.us-east-1.amazonaws.com/v1/repos/MyAppRepo

Step 3: Push Code to CodeCommit

git add.

git commit -m "Initial commit"

git push origin main

***** Expected Outcome:

Your source code is now stored in AWS CodeCommit.

CHAPTER 2: AWS CODEBUILD – AUTOMATED BUILD AND TESTING

1. What is AWS CodeBuild?

AWS **CodeBuild** is a fully managed **build service** that compiles, tests, and packages source code.

- 2. Benefits of AWS CodeBuild
- ✓ Serverless & Scalable No need to manage build servers.
- ✓ Supports multiple environments Works with Java, Python, Node.js, .NET, Go, etc.
- ✓ Integrates with CodeCommit, GitHub, Bitbucket.
- √ Generates real-time logs in Amazon CloudWatch.

3. Setting Up AWS CodeBuild

Step 1: Create a Build Project

- Open AWS CodeBuild Console → Click "Create Build Project".
- 2. Enter Project Name: MyAppBuild.

Source Provider: Select AWS CodeCommit → Choose MyAppRepo.

4. Environment:

- Runtime: Amazon Linux 2
- Build Image: aws/codebuild/standard:5.0
- o Compute Type: **Small**

Step 2: Define Build Commands in buildspec.yml

Create a buildspec.yml file in the repository root:

version: 0.2

phases:

install:

runtime-versions:

nodejs: 14

build:

commands:

- echo "Building Application..."
- npm install
- npm run build

artifacts:

files:

- '**/*'

Step 3: Start the Build

1. Click "Start Build".

2. Monitor build logs in CloudWatch Logs.

* Expected Outcome:

The application code is **compiled and tested successfully**.

CHAPTER 3: AWS CODEPIPELINE - CI/CD AUTOMATION

1. What is AWS CodePipeline?

AWS **CodePipeline** is a CI/CD **orchestration tool** that automates build, test, and deployment workflows.

- 2. Benefits of AWS CodePipeline
- ✓ Automates software release cycles.
- ✓ Integrates with CodeCommit, GitHub, Bitbucket.
- ✓ Supports parallel deployments for multi-region applications.
- ✓ Secure with IAM-based access control.

3. Setting Up AWS CodePipeline

Step 1: Create a New Pipeline

- Open AWS CodePipeline Console → Click "Create Pipeline".
- 2. Pipeline Name: MyAppPipeline.
- Select Service Role: Create new role.
- Click "Next".

Step 2: Configure Source Stage

- Source Provider: Select AWS CodeCommit.
- 2. Choose Repository: MyAppRepo.
- 3. Branch Name: main.

4. Click "Next".

Step 3: Configure Build Stage

- Build Provider: Select AWS CodeBuild.
- 2. Choose **Build Project**: MyAppBuild.
- 3. Click "Next".

Step 4: Configure Deployment Stage (Example: S₃ or Elastic Beanstalk)

- Deployment Provider: Select AWS S₃ or AWS Elastic Beanstalk.
- Choose S3 Bucket/Environment Name for deployment.
- 3. Click "Next" → "Create Pipeline".

Expected Outcome:

Whenever code is **pushed to CodeCommit**, the pipeline will:

- ✓ Automatically build the application using CodeBuild.
- ✓ Deploy the built artifacts to AWS S3 or Elastic Beanstalk.

CHAPTER 4: MONITORING & TROUBLESHOOTING AWS CI/CD SERVICES

- 1. Monitoring CodeBuild & CodePipeline
- ✓ View real-time logs in AWS CloudWatch.
- ✓ Check build status in the CodeBuild console.
- ✓ Set CloudWatch Alarms for failed builds or deployments.
- 2. Troubleshooting Common Issues
- ✓ Authentication Issues Ensure IAM policies allow CodeBuild and CodePipeline to access S₃, CodeCommit, and EC₂.

- ✓ Build Failures Check buildspec.yml for syntax errors.
- ✓ **Pipeline Execution Failure** Verify source repository and build artifacts.

CHAPTER 5: AWS CI/CD BEST PRACTICES

- √ Use CodeCommit for secure source code management.
- ✓ Enable CodeBuild caching for faster builds.
- ✓ Automate deployments using CodePipeline.
- ✓ Use IAM roles to enforce least-privilege access control.
- ✓ Enable CloudWatch monitoring for pipeline visibility.

***** Example:

A **fintech company** automates code deployment using:

- CodeCommit for version control
- CodeBuild for secure builds
- CodePipeline for automated deployment

CHAPTER 6: REAL-WORLD USE CASES OF AWS CODECOMMIT,
CODEBUILD & CODEPIPELINE

- 1. CI/CD for Web Applications
- ✓ Developers push code to CodeCommit.
- ✓ CodeBuild compiles and tests the code.
- ✓ CodePipeline deploys to S₃ or Elastic Beanstalk.
- 2. CI/CD for Serverless Applications (AWS Lambda)
- ✓ Build & test Lambda functions using CodeBuild.
- ✓ Deploy to AWS Lambda automatically using CodePipeline.

3. Automated Security Scanning in CI/CD

✓ Use CodeBuild with security tools (SonarQube, AWS Inspector) for vulnerability scans.

CONCLUSION: MASTERING AWS CODECOMMIT, CODEBUILD & CODEPIPELINE

By using AWS CI/CD services, businesses can:

- Automate software builds, testing, and deployments.
- Securely manage source code using CodeCommit.
- Reduce deployment errors with CodePipeline workflows.

FINAL EXERCISE:

- 1. Create a CodePipeline that deploys a web application to S3.
- 2. Enable CloudWatch Alarms for failed builds in CodeBuild.
- 3. Implement a Lambda function to notify failures via SNS.

CI/CD PIPELINES WITH AWS DEVOPS – STUDY MATERIAL

INTRODUCTION TO CI/CD IN AWS DEVOPS

What is CI/CD?

CI/CD (Continuous Integration/Continuous Deployment) is a DevOps practice that enables frequent and automated software releases by integrating code, running tests, and deploying applications seamlessly.

What is AWS DevOps?

AWS provides a suite of **DevOps tools** to automate **software development, testing, deployment, and monitoring,** allowing teams to efficiently implement CI/CD pipelines.

Key AWS CI/CD Tools

- ✓ AWS CodeCommit Private Git repository for version control.
- ✓ AWS CodeBuild Automated build and test service.
- ✓ AWS CodeDeploy Automated deployment service for EC2, Lambda, and ECS.
- ✓ **AWS CodePipeline** CI/CD workflow automation tool.
- ✓ AWS Elastic Beanstalk Simplified application deployment and management.
- ✓ **AWS** CloudFormation Infrastructure as Code (IaC) for resource provisioning.

CHAPTER 1: UNDERSTANDING CI/CD PIPELINES

1. What is a CI/CD Pipeline?

A **CI/CD pipeline** is an **automated workflow** that integrates, tests, and deploys code changes.

2. Phases of a CI/CD Pipeline

Stage	Description	AWS Service Used
Source	Version control for	AWS CodeCommit, GitHub,
	code changes	Bitbucket
Build	Compile and test	AWS CodeBuild
	code	
Test	Run automated tests	AWS CodeBuild, AWS Lambda,
		Selenium
Deploy	Deploy application to	AWS CodeDeploy, AWS Elastic
	AWS	Beanstalk, AWS Lambda
Monitor	Track performance	Amazon CloudWatch, AWS X-
	and errors	Ray

* Example:

A CI/CD pipeline for a web application would:

Fetch code from CodeCommit → 2. Build and test using
 CodeBuild → 3. Deploy via CodeDeploy to an EC2 instance.

CHAPTER 2: SETTING UP A CI/CD PIPELINE WITH AWS CODEPIPELINE

1. Step-by-Step Guide to Create a CI/CD Pipeline

Step 1: Create a CodeCommit Repository

 Open AWS CodeCommit Console → Click "Create Repository".

- 2. Enter Repository Name: MyAppRepo.
- 3. Click "Create" \rightarrow Clone the repo:
- git clone https://git-codecommit.us-east-1.amazonaws.com/v1/repos/MyAppRepo
- 5. cd MyAppRepo

Step 2: Create a CodeBuild Project

- Open AWS CodeBuild Console → Click "Create Build Project".
- 2. **Source:** Choose **CodeCommit** → Select MyAppRepo.
- 3. Environment:
 - Image: aws/codebuild/standard:5.0
 - Runtime: Amazon Linux 2
- 4. **Buildspec file:** Create a buildspec.yml file in the repo:
- 5. version: 0.2
- 6. phases:
- 7. install:
- 8. runtime-versions:
- 9. nodejs: 14
- 10. build:
- 11. commands:
- 12. echo "Building Application..."
- 13. npm install
- 14. npm run build

15.artifacts:

- 16. files:
- 17. '**/*'
- 18. Click "Create Build Project".

Expected Outcome:

The build service automates code compilation and testing.

Step 3: Create an Application Deployment via CodeDeploy

- Open AWS CodeDeploy Console → Click "Create Application".
- Enter Application Name: MyAppDeploy.
- 3. Choose Compute Platform: EC2/On-Premises.
- 4. Create a **Deployment Group** → Attach an **IAM role**.

Expected Outcome:

AWS CodeDeploy is ready to deploy new versions of the application.

Step 4: Create a CodePipeline for CI/CD Automation

- Open AWS CodePipeline Console → Click "Create Pipeline".
- 2. **Pipeline Name:** MyAppPipeline.
- Source Provider: Select AWS CodeCommit → Choose MyAppRepo.
- Build Provider: Select AWS CodeBuild → Choose MyAppBuild.

- Deployment Provider: Select AWS CodeDeploy → Choose MyAppDeploy.
- 6. Click "Create Pipeline".

Expected Outcome:

Whenever new code is pushed, AWS automatically builds, tests, and deploys it.

CHAPTER 3: AUTOMATING INFRASTRUCTURE DEPLOYMENT WITH AWS CLOUDFORMATION

1. What is AWS CloudFormation?

AWS **CloudFormation** automates the provisioning of AWS resources using **Infrastructure as Code (IaC)**.

- 2. Creating a CloudFormation Template for CI/CD
 - Create a CloudFormation template (cicd.yaml):
 - 2. Resources:
 - 3. MyCodePipeline:
 - 4. Type: AWS::CodePipeline::Pipeline
 - 5. Properties:
 - 6. Name: MyAppPipeline
 - 7. RoleArn:
 arn:aws:iam::123456789012:role/AWSCodePipelineServiceRol
 - 8. ArtifactStore:
 - 9. Type: S₃
 - 10. Location: my-pipeline-artifacts

11.	Stages:
12.	- Name: Source
13.	Actions:
14.	- Name: SourceAction
15.	ActionTypeId:
16.	Category: Source
17.	Owner: AWS
18.	Provider: CodeCommit
19.	Version: 1
20.	Configuration:
21.	RepositoryName: MyAppRepo
22.	BranchName: main
23.	OutputArtifacts:
24.	- Name: SourceArtifact
25.	Deploy CloudFormation stack:
26.	aws cloudformation create-stackstack-name yCICDPipelinetemplate-body file://cicd.yaml

AWS automatically provisions CI/CD pipelines using this script.

CHAPTER 4: MONITORING AND LOGGING IN CI/CD PIPELINES

1. AWS CloudWatch for Build Monitoring

- ✓ View logs from CodeBuild, CodeDeploy, and EC2 instances.
- ✓ **Set up alarms** for failed deployments.
- 2. AWS X-Ray for Debugging CI/CD Pipelines
- √ Tracks application requests and errors.
- ✓ Identifies performance bottlenecks.
- ***** Example:

A CodeDeploy failure triggers a CloudWatch Alarm, alerting developers.

CHAPTER 5: CI/CD BEST PRACTICES WITH AWS DEVOPS

- √ Use Infrastructure as Code (IaC) with CloudFormation.
- ✓ Enable automated testing in CI/CD pipelines.
- ✓ Enforce IAM least-privilege access for CI/CD users.
- ✓ Use CodeDeploy blue/green deployments for zero-downtime updates.
- ✓ Monitor pipelines with CloudWatch and X-Ray.
- * Example:

A fintech company implements CI/CD best practices to automate software releases while maintaining security.

CHAPTER 6: REAL-WORLD USE CASES OF CI/CD WITH AWS DEVOPS

- 1. CI/CD for Microservices (ECS & Lambda)
- ✓ Deploy microservices using AWS Fargate & CodePipeline.
- ✓ Use AWS Lambda for serverless automation.
- 2. CI/CD for Machine Learning Models

- ✓ Retrain models with CodeBuild when new data arrives.
- ✓ Deploy models using AWS SageMaker CI/CD pipelines.
- 3. Enterprise DevOps Pipeline
- ✓ Multi-region deployments using AWS Global Accelerator.
- ✓ Blue/Green deployments for zero downtime.

CONCLUSION: MASTERING CI/CD PIPELINES WITH AWS DEVOPS By using AWS DevOps tools, businesses can:

- Automate code integration, testing, and deployment.
- Improve release efficiency and software quality.
- Scale and secure CI/CD pipelines with AWS best practices.

FINAL EXERCISE:

- Create a CodePipeline for deploying an S3-hosted static website.
- 2. Set up a CodeDeploy blue/green deployment for an EC2 app.
- 3. Enable CloudWatch alarms for failed build/test steps.

INFRASTRUCTURE AS CODE (IAC)

INTRODUCTION TO AWS CLOUDFORMATION – STUDY MATERIAL

What is AWS CloudFormation?

AWS CloudFormation is an Infrastructure as Code (IaC) service that allows users to define and provision AWS resources automatically using templates. It helps automate infrastructure management, ensuring consistency, scalability, and repeatability.

CHAPTER 1: UNDERSTANDING AWS CLOUDFORMATION

1. Why Use AWS CloudFormation?

- ✓ Infrastructure as Code (IaC) Define AWS resources in YAML or JSON templates.
- ✓ Automated Resource Management Deploy, update, and delete AWS infrastructure easily.
- ✓ Scalability & Consistency Deploy the same infrastructure across multiple environments.
- ✓ Supports Most AWS Services EC2, S3, RDS, Lambda, VPC, IAM, and more.
- ✓ **Version Control & Change Management** Track infrastructure changes using templates.

* Example:

A **startup automates AWS resource creation** by defining an EC2 instance, S₃ bucket, and security group in a CloudFormation template.

2. Key CloudFormation Concepts

Concept	Description
Stack	A collection of AWS resources created from a
	CloudFormation template.
Template	A YAML/JSON file defining AWS infrastructure.
StackSet	Allows deployment of stacks across multiple AWS
	accounts/regions.
Parameters	Customizable input values for templates (e.g.,
	instance type).
Mappings	Static key-value pairs used to define configurations.
Conditions	Define rules to include/exclude resources based on
	conditions.
Outputs	Returns values after stack creation (e.g., EC2
	instance ID).
Change	Shows what will change before applying updates to
Sets	stacks.

CHAPTER 2: CREATING A CLOUDFORMATION TEMPLATE

1. Basic CloudFormation Template Structure

A CloudFormation template consists of **five main sections**:

AWSTemplateFormatVersion: "2010-09-09"

Description: "Basic CloudFormation Template Example"

Resources:

MyS3Bucket:

Type: "AWS::S3::Bucket"

Properties:

BucketName: "my-cloudformation-bucket"

- ✓ AWSTemplateFormatVersion Defines the CloudFormation template version.
- ✓ **Description** Provides details about the template.
- ✓ **Resources** Defines AWS resources like EC₂, S₃, RDS, and Lambda.

Expected Outcome:

This template **creates an S3 bucket** named my-cloudformation-bucket.

2. Defining an EC2 Instance in CloudFormation

Resources:

MyEC2Instance:

Type: "AWS::EC2::Instance"

Properties:

lmageld: "ami-oabcdef1234567890"

InstanceType: "t2.micro"

KeyName: "my-key-pair"

SecurityGroups:

- Ref: "MySecurityGroup"

MySecurityGroup:

Type: "AWS::EC2::SecurityGroup"

Properties:

GroupDescription: "Enable SSH Access"

SecurityGroupIngress:

- IpProtocol: "tcp"

FromPort: 22

ToPort: 22

Cidrlp: "o.o.o.o/o"

***** Expected Outcome:

✓ Deploys an **EC2 instance** with a security group allowing SSH access.

CHAPTER 3: DEPLOYING A CLOUD FORMATION STACK

Step 1: Upload Template to CloudFormation

- Open AWS Console → Navigate to CloudFormation.
- Click "Create Stack" → "With New Resources".
- 3. **Choose Template Source:** Upload a **YAML/JSON** template file.
- 4. Click "Next".

Step 2: Configure Stack Details

- 1. **Stack Name:** MyCloudFormationStack.
- 2. Enter Parameters (if any).

3. Click "Next" → Configure Stack Options.

Step 3: Review and Deploy Stack

- 1. Review settings and click "Create Stack".
- 2. Wait for **Stack Status = CREATE_COMPLETE**.

Expected Outcome:

✓ AWS resources (EC2, S3, RDS, etc.) are created automatically.

CHAPTER 4: USING PARAMETERS, MAPPINGS, AND OUTPUTS

1. Using Parameters in CloudFormation

Allows users to **customize stack settings** dynamically.

Parameters:

InstanceType:

Type: String

Default: "t2.micro"

AllowedValues:

- "t2.micro"
- "t2.small"
- "t2.medium"

Resources:

MyEC2Instance:

Type: "AWS::EC2::Instance"

Properties:

InstanceType: !Ref InstanceType

✓ Users **choose an instance type** when launching the stack.

2. Using Mappings for Configurations

Define static key-value pairs.

Mappings:

RegionMap:

us-east-1:

AMI: "ami-oabcdef1234567890"

us-west-1:

AMI: "ami-ofedcbao987654321"

Resources:

MyEC2Instance:

Type: "AWS::EC2::Instance"

Properties:

Imageld: !FindInMap [RegionMap, !Ref "AWS::Region", AMI]

Expected Outcome:

✓ Selects the AMI based on the AWS region.

3. Using Outputs to Retrieve Stack Details

Outputs:

InstanceID:

Description: "The ID of the EC2 instance"

Value: !Ref MyEC2Instance

Expected Outcome:

✓ Outputs the EC2 instance ID after deployment.

CHAPTER 5: UPDATING AND DELETING CLOUDFORMATION STACKS

1. Updating a Stack

- Open CloudFormation Console → Select Stack.
- Click "Update" → Choose "Replace Current Template".
- Modify the YAML/JSON template → Click "Update Stack".

Expected Outcome:

✓ AWS automatically **modifies the existing stack** without downtime.

2. Deleting a Stack

- Open CloudFormation Console → Select Stack.
- Click "Delete Stack".
- 3. Confirm deletion.

Expected Outcome:

✓ AWS removes all resources created by the stack.

Chapter 6: Monitoring and Troubleshooting CloudFormation

1. Monitoring CloudFormation Stacks

- ✓ Use CloudFormation Events to check stack creation status.
- ✓ Enable AWS CloudTrail to log API calls.
- ✓ Use Amazon CloudWatch for performance monitoring.

2. Common CloudFormation Errors & Solutions

Error	Cause	Solution
ROLLBACK_COMPLETE	Deployment	Check logs, update
	failed	stack with correct
		settings
IAM Policy Errors	Insufficient	Assign correct IAM
	permissions	roles
Invalid Parameters	Incorrect values	Validate input values
	in parameters	before deployment

CHAPTER 7: ADVANCED CLOUDFORMATION FEATURES

1. Nested Stacks

- ✓ Use Nested Stacks to manage large deployments.
- ✓ Break a template into smaller, reusable stacks.

***** Example:

A microservices architecture with separate CloudFormation templates for database, backend, and frontend.

2. AWS CloudFormation StackSets

✓ Deploy CloudFormation stacks across multiple AWS accounts and regions.

* Example:

An enterprise manages multi-region VPCs and IAM roles using StackSets.

CHAPTER 8: BEST PRACTICES FOR AWS CLOUDFORMATION

- ✓ Use Version Control Store templates in GitHub, CodeCommit.
- ✓ Implement IAM Least Privilege Assign minimal permissions.
- √ Test in a Dev Environment Validate templates before production deployment.
- ✓ Enable CloudFormation Drift Detection Detect changes made outside CloudFormation.
- ✓ **Use Parameter Store or Secrets Manager** Avoid hardcoding credentials.

CONCLUSION: MASTERING AWS CLOUDFORMATION

By using **AWS CloudFormation**, businesses can:

- Automate infrastructure provisioning.
- Deploy scalable, consistent AWS environments.
- Reduce manual errors with Infrastructure as Code (IaC).
- Easily update and manage AWS resources.

FINAL EXERCISE:

- Create a CloudFormation template to deploy a VPC, EC2 instance, and RDS database.
- 2. Modify an existing stack using Change Sets.
- 3. Deploy a multi-region architecture using StackSets

TERRAFORM FOR AWS – STUDY MATERIAL

INTRODUCTION TO TERRAFORM FOR AWS

What is Terraform?

Terraform is an open-source Infrastructure as Code (IaC) tool that allows you to define, provision, and manage AWS infrastructure using a declarative configuration language (HCL

- HashiCorp Configuration Language).

Why Use Terraform for AWS?

- ✓ Automates AWS resource provisioning No manual setup required.
- ✓ **Declarative syntax** Define the desired state, and Terraform will handle the changes.
- ✓ **Version control** Store infrastructure configurations in Git for tracking.
- ✓ **Reusable infrastructure** Use Terraform modules to deploy standardized environments.
- ✓ **Multi-cloud support** Works with AWS, Azure, Google Cloud, and others.

Chapter 1: Key Terraform Concepts

Concept	Description
Providers	APIs that Terraform interacts with (e.g., AWS, Azure, GCP).
Resources	AWS services created using Terraform (e.g., EC2, S3, VPC).
Variables	Custom inputs used in configurations.

State Files	Tracks deployed infrastructure	
	(terraform.tfstate).	
Modules	Reusable Terraform configurations.	
Outputs	Returns values after Terraform deployment.	
Provisioners	Run scripts on AWS instances during creation.	

***** Example:

A Terraform script can deploy an EC2 instance, create an S3 bucket, and configure security groups automatically.

CHAPTER 2: INSTALLING AND SETTING UP TERRAFORM

1. Install Terraform on Your Machine

For Windows:

- Download Terraform from <u>Terraform Official Website</u>.
- 2. Extract and add it to your **System PATH**.
- 3. Verify installation:
- 4. terraform -version

For Linux/macOS:

- 1. Install using Homebrew (macOS):
- 2. brew install terraform
- 3. Install using APT (Ubuntu):
- 4. sudo apt-get update && sudo apt-get install -y terraform
- 5. Verify installation:
- 6. terraform -version

Expected Outcome: Terraform is successfully installed and ready to use.

CHAPTER 3: WRITING YOUR FIRST TERRAFORM CONFIGURATION FOR AWS

1. Create a Terraform Project Directory

mkdir terraform-aws-demo cd terraform-aws-demo

2. Define AWS Provider and Resources

Create a file named main.tf with the following configuration:

- ***** Explanation:
- ✓ Defines AWS as the provider.
- ✓ Creates an EC2 instance with t2.micro instance type.
- ✓ Assigns the tag "Terraform-EC2" to the instance.

3. Initialize Terraform

Run the following command to initialize Terraform:

terraform init

Expected Outcome: Downloads AWS provider plugins and initializes Terraform.

4. Plan Infrastructure Deployment

Run the following command to preview what Terraform will create:

terraform plan

* Expected Outcome: Displays planned AWS resources without creating them.

5. Apply Terraform Configuration

To create AWS resources, run:

terraform apply

- ✓ Type "yes" when prompted to confirm deployment.
- **Expected Outcome:** AWS EC2 instance is successfully created.

6. Destroy Infrastructure (Cleanup)

To remove all AWS resources created by Terraform:

terraform destroy

- ✓ Type "yes" when prompted.
- **Expected Outcome:** Terraform removes all AWS resources it created.

CHAPTER 4: USING VARIABLES IN TERRAFORM

1. Define Variables

Modify main.tf to use variables instead of hardcoded values:

```
variable "aws_region" {
  default = "us-east-1"
}

variable "instance_type" {
  default = "t2.micro"
}

provider "aws" {
  region = var.aws_region
}
```

```
resource "aws_instance" "my_ec2" {
  ami = "ami-oabcdef1234567890"
  instance_type = var.instance_type

tags = {
  Name = "Terraform-EC2"
  }
}
```

2. Override Variables at Runtime

To specify variables dynamically:

terraform apply -var="instance_type=t3.micro"

Expected Outcome:

Terraform deploys an **EC2** instance with a t3.micro type instead of the default t2.micro.

CHAPTER 5: TERRAFORM STATE MANAGEMENT

1. What is Terraform State?

Terraform maintains the **state of the infrastructure** in a file called terraform.tfstate. This helps track changes over time.

2. Storing State in S3 for Collaboration

Modify main.tf to store state in an S₃ bucket:

```
terraform {
 backend "s3" {
```

```
bucket = "my-terraform-state-bucket"
key = "terraform.tfstate"
region = "us-east-1"
}
```

Terraform will **store the state file in S3**, allowing multiple users to collaborate.

CHAPTER 6: USING TERRAFORM MODULES FOR REUSABILITY

1. What Are Terraform Modules?

Terraform modules allow you to reuse infrastructure configurations.

2. Creating a Simple Terraform Module

```
Create a folder structure:

mkdir modules

mkdir modules/ec2_instance

Inside modules/ec2_instance/main.tf, define an EC2 module:

variable "instance_type" {}

variable "ami" {}

resource "aws_instance" "my_ec2" {

ami = var.ami
```

```
instance_type = var.instance_type
}
Now, call this module in main.tf:
module "ec2_instance" {
  source = "./modules/ec2_instance"
  instance_type = "t2.micro"
  ami = "ami-oabcdef1234567890"
}
```

Terraform reuses the module to deploy multiple instances.

CHAPTER 7: MANAGING TERRAFORM WORKSPACES

1. What Are Workspaces?

Terraform workspaces allow managing multiple environments (e.g., Dev, Staging, Production) with the same codebase.

- 2. Creating Workspaces
- 1. Create a New Workspace
- 2. terraform workspace new dev
- 3. Switch to a Different Workspace
- 4. terraform workspace select production
- 5. List All Workspaces
- 6. terraform workspace list

Terraform manages separate environments using workspaces.

CHAPTER 8: BEST PRACTICES FOR TERRAFORM IN AWS

- ✓ Use Terraform Cloud or S₃ Backend to manage state remotely.
- ✓ Implement Terraform Modules to reuse configurations.
- ✓ Use Terraform Variables & Outputs for flexible deployment.
- ✓ **Use Terraform Workspaces** for multiple environments.
- ✓ Enable IAM Least Privilege to restrict Terraform access.

***** Example:

A DevOps team uses Terraform Cloud to collaborate and manage AWS infrastructure securely.

CHAPTER 9: REAL-WORLD USE CASES OF TERRAFORM FOR AWS

- 1. Multi-Cloud Deployments
- ✓ Deploy AWS & Azure resources from the same Terraform configuration.
- Auto Scaling & Load Balancers
- ✓ Use Terraform to create an AWS ALB with EC₂ Auto Scaling.
- 3. Secure Cloud Environments
- ✓ Automate IAM policies, security groups, and VPC configurations.
- 4. Disaster Recovery & Backups
- ✓ Use Terraform to spin up AWS resources in case of failure.

CONCLUSION: MASTERING TERRAFORM FOR AWS

By using **Terraform for AWS**, businesses can:

- Automate AWS infrastructure provisioning.
- **Ensure consistent deployments** across environments.
- Manage AWS at scale with infrastructure as code (IaC).

FINAL EXERCISE:

- 1. Create a Terraform module to deploy an S₃ bucket with encryption.
- 2. Use Terraform Workspaces to manage Dev, Staging, and Production environments.
- 3. Store Terraform state in S₃ and enable locking with DynamoDB.

ASSIGNMENT

CONFIGURE A SECURE VPC WITH SUBNETS



SOLUTION: CONFIGURE A SECURE VPC WITH SUBNETS (STEP-BY-STEP GUIDE)

This step-by-step guide will walk you through **configuring a secure AWS VPC with subnets,** including **public and private subnets, security groups, network ACLs,** and **internet access controls**.

Step 1: Log in to AWS and Navigate to VPC Console

- 1. Open the AWS Management Console.
- 2. Search for "VPC" and click on VPC Dashboard.
- 3. Click "Create VPC".

Step 2: Create a VPC

- VPC Name: Secure-VPC
- 2. **IPv4 CIDR Block:** 10.0.0.0/16 (Allows 65,536 IP addresses)
- 3. IPv6 CIDR Block: No IPv6 for now (optional).
- 4. Tenancy: Default
- 5. Click "Create VPC".

***** Expected Outcome:

A VPC is created with a private network range of 10.0.0.0/16.

Step 3: Create Public and Private Subnets

Create Public Subnet

Go to VPC Console → Subnets → Click "Create Subnet".

- 2. Subnet Name: Public-Subnet-1.
- 3. **VPC:** Select Secure-VPC.
- 4. **Availability Zone:** Choose us-east-1a (or any AZ).
- 5. **IPv4 CIDR Block:** 10.0.1.0/24 (256 IP addresses).
- 6. Click "Create Subnet".

Create Private Subnet

- 1. Click "Create Subnet" again.
- 2. Subnet Name: Private-Subnet-1.
- 3. VPC: Select Secure-VPC.
- 4. **Availability Zone:** Choose us-east-1b.
- 5. **IPv4 CIDR Block:** 10.0.2.0/24.
- 6. Click "Create Subnet".
- Expected Outcome:
- ✓ A public subnet (10.0.1.0/24) for internet-facing resources.
- ✓ A private subnet (10.0.2.0/24) for internal applications (no internet access).

Step 4: Attach an Internet Gateway (IGW) for Public Subnet

- Go to VPC Console → Internet Gateways → Click "Create Internet Gateway".
- 2. Name: Secure-IGW.
- 3. Click "Create".
- 4. Attach IGW to Secure-VPC:

- Select Secure-IGW → Click "Actions" → "Attach to VPC".
- Choose Secure-VPC and click "Attach".

The **public subnet can connect to the internet** through the **Internet Gateway**.

Step 5: Configure Route Tables

Create a Public Route Table

- Go to VPC Console → Route Tables → Click "Create Route Table".
- 2. Name: Public-RT.
- 3. **VPC:** Select Secure-VPC.
- 4. Click "Create".

Add Route for Internet Access

- Select Public-RT → Click "Edit Routes".
- 2. Click "Add Route".
- Destination: o.o.o.o/o (All Internet Traffic).
- 4. **Target:** Select **Secure-IGW** (Internet Gateway).
- 5. Click "Save Changes".

Associate Public Route Table with Public Subnet

- Click "Subnet Associations".
- Click "Edit Subnet Associations".
- Select Public-Subnet-1 → Click "Save".

✓ The public subnet (10.0.1.0/24) now has internet access via the Internet Gateway.

Step 6: Configure a NAT Gateway for Private Subnet (Optional)

A **NAT Gateway** allows private subnet instances to **access** the **internet** (for updates, etc.) without being exposed to incoming internet traffic.

- Go to VPC Console → NAT Gateways → Click "Create NAT Gateway".
- 2. **Subnet:** Select Public-Subnet-1.
- 3. Elastic IP: Click "Allocate New Elastic IP".
- Click "Create NAT Gateway".

Add Route for NAT Gateway in Private Route Table

- 1. Go to VPC Console → Route Tables → Click "Create Route Table".
- 2. Name: Private-RT.
- 3. VPC: Secure-VPC.
- 4. Click "Create".

Add NAT Gateway Route

- Select Private-RT → Click "Edit Routes".
- Click "Add Route".
- 3. **Destination:** 0.0.0.0/0.
- 4. Target: Select NAT Gateway.

5. Click "Save Changes".

Associate Private Route Table with Private Subnet

- Click "Subnet Associations" → Click "Edit Subnet Associations".
- 2. Select Private-Subnet-1 → Click "Save".

***** Expected Outcome:

✓ Instances in the private subnet can access the internet (for software updates) but cannot receive incoming traffic.

Step 7: Configure Security Groups

- 1. Create a Security Group for Public EC2 Instances
 - Go to VPC Console → Security Groups → Click "Create Security Group".
 - 2. Name: Public-SG.
 - 3. VPC: Secure-VPC.
 - 4. Inbound Rules:
 - Allow HTTP (80) from o.o.o.o/o.
 - Allow HTTPS (443) from 0.0.0.0/0.
 - Allow SSH (22) only from your IP (your-public-ip/32).
 - 5. Click "Create Security Group".

Expected Outcome:

✓ Public EC2 instances can receive web traffic but are secured against unauthorized access.

2. Create a Security Group for Private EC2 Instances

1. Name: Private-SG.

2. **VPC:** Secure-VPC.

3. Inbound Rules:

- Allow MySQL (3306) from Public-SG (for database access).
- Allow internal SSH access (22) only from Public-SG.
- 4. Click "Create Security Group".

Expected Outcome:

✓ Private EC2 instances (e.g., databases) are only accessible from public instances.

Step 8: Configure Network ACLs (Optional, for Additional Security)

- Go to VPC Console → Network ACLs → Click "Create Network ACL".
- 2. Name: Secure-NACL.
- 3. VPC: Secure-VPC.
- 4. Click "Create".

Edit Inbound Rules:

- 1. Allow HTTP (80), HTTPS (443), and SSH (22) only from **trusted** sources.
- 2. Deny all other inbound traffic.

✓ Provides an extra layer of security to prevent unauthorized traffic.

Step 9: Launch and Test an EC2 Instance in Public Subnet

- Go to EC2 Console → Click "Launch Instance".
- 2. Choose Amazon Linux 2 or Ubuntu AMI.
- 3. Select **Instance Type:** t2.micro.
- 4. Choose **Network:** Secure-VPC.
- 5. Choose **Subnet:** Public-Subnet-1.
- 6. Select **Security Group:** Public-SG.
- 7. Click "Launch".

Connect to EC2 via SSH

ssh -i my-key.pem ec2-user@your-public-ip

Expected Outcome:

✓ Successfully connected to the EC2 instance in the public subnet.

CONCLUSION: SUCCESSFULLY CONFIGURED A SECURE VPC WITH SUBNETS

- Created a VPC with public and private subnets.
- Configured an Internet Gateway for public access.
- Set up a NAT Gateway for private subnet internet access.
- Applied Security Groups and Network ACLs for secure networking.

FINAL EXERCISE:

- 1. Launch an EC2 instance in the private subnet and test internet access through NAT Gateway.
- 2. Deploy an RDS instance in the private subnet with security group restrictions.
- 3. Enable AWS VPC Flow Logs for monitoring network traffic.



DEPLOY AN APPLICATION USING AWS CODEPIPELINE



SOLUTION: DEPLOY AN APPLICATION USING AWS CODEPIPELINE (STEP-BY-STEP GUIDE)

This guide will walk you through deploying an application using AWS CodePipeline, integrating CodeCommit (source), CodeBuild (build), and CodeDeploy (deployment) for continuous integration and deployment (CI/CD).

Step 1: Set Up an AWS CodeCommit Repository

AWS **CodeCommit** is a **Git-based repository** used for storing application source code.

1. Create a CodeCommit Repository

- Open the AWS CodeCommit Console.
- 2. Click "Create Repository".
- 3. Repository Name: MyAppRepo.
- 4. Click "Create".

2. Clone the Repository Locally

- 1. Copy the repository clone URL from CodeCommit.
- 2. Open the terminal and run:
- git clone https://git-codecommit.us-east-1.amazonaws.com/v1/repos/MyAppRepo
- 4. cd MyAppRepo

3. Add Sample Application Code

1. Create an index.html file:

- 2. <html>
- 3. <body>
- 4. <h1>Deployed via AWS CodePipeline!</h1>
- 5. </body>
- 6. </html>
- 7. Commit and push the changes:
- 8. git add.
- 9. git commit -m "Initial commit"
- 10. git push origin main

Your source code is now stored in AWS CodeCommit.

Step 2: Create an AWS CodeBuild Project

AWS CodeBuild compiles and packages your application.

Create a buildspec.yml File

Inside the repository (MyAppRepo), create buildspec.yml:

version: 0.2

phases:

install:

runtime-versions:

nodejs: 14

build:

commands:

- echo "Building application..."
- mkdir artifacts
- cp index.html artifacts/

artifacts:

files:

- artifacts/index.html

- 2. Configure CodeBuild in AWS Console
 - Open AWS CodeBuild Console → Click "Create Build Project".
 - 2. Project Name: MyAppBuild.
 - 3. **Source Provider:** Select **AWS CodeCommit,** then select MyAppRepo.
 - 4. Environment:
 - Operating System: Amazon Linux 2
 - Runtime: Standard
 - Image: aws/codebuild/standard:5.0
 - Compute: Small
 - 5. Buildspec File: Choose "Use a buildspec file".
 - 6. Click "Create Build Project".
- ***** Expected Outcome:

CodeBuild automates application build and artifact creation.

Step 3: Set Up AWS CodeDeploy

AWS **CodeDeploy** automates deploying applications to **EC2**, **ECS**, or Lambda.

- 1. Launch an EC2 Instance for Deployment
 - Open AWS EC2 Console → Click "Launch Instance".
 - 2. Choose Amazon Linux 2 or Ubuntu AMI.
 - 3. Select Instance Type: t2.micro.
 - 4. Create or Select Key Pair (to SSH into the instance).
 - 5. Network Settings:
 - Choose VPC and Public Subnet.
 - Enable Auto-Assign Public IP.
 - 6. Add Storage: Default (8GB).
 - 7. Configure Security Group:
 - Allow HTTP (80), HTTPS (443), and SSH (22).
 - 8. Click "Launch".
- 2. Install CodeDeploy Agent on EC2

SSH into the EC2 instance:

ssh -i my-key.pem ec2-user@your-ec2-public-ip

Install the AWS CodeDeploy agent:

sudo yum update -y

sudo yum install ruby -y

sudo yum install wget -y

cd /home/ec2-user

wget https://aws-codedeploy-us-east-1.s3.amazonaws.com/latest/install chmod +x ./install sudo ./install auto

sudo systemctl start codedeploy-agent

sudo systemctl enable codedeploy-agent

3. Create an IAM Role for CodeDeploy

- Open IAM Console → Click "Roles" → "Create Role".
- 2. Select **EC2** as a trusted entity.
- Attach the AWSCodeDeployRole policy.
- 4. Name the role CodeDeployEC2Role → Click "Create".
- 5. Assign this role to the EC2 instance.

Expected Outcome:

EC2 instance is ready to receive deployments from AWS CodeDeploy.

4. Create an AWS CodeDeploy Application

- Open AWS CodeDeploy Console → Click "Create Application".
- 2. **Application Name:** MyAppDeploy.
- 3. **Compute Platform:** EC2/On-Premises.
- 4. Click "Create Application".

5. Create a Deployment Group

1. Open the newly created **MyAppDeploy** application.

- 2. Click "Create Deployment Group".
- 3. **Deployment Group Name:** MyAppDeploymentGroup.
- 4. Service Role: Select AWSCodeDeployRole.
- 5. Environment Configuration:
 - Select EC2 Instances.
 - Choose instances tagged with CodeDeployInstance.
- 6. Click "Create Deployment Group".

CodeDeploy is ready to automate deployments to EC2 instances.

Step 4: Create AWS CodePipeline

AWS CodePipeline automates the CI/CD workflow by integrating CodeCommit, CodeBuild, and CodeDeploy.

- 1. Create a New Pipeline
 - Open AWS CodePipeline Console → Click "Create Pipeline".
 - 2. Pipeline Name: MyAppPipeline.
 - 3. Service Role: Choose "Create New Role".
 - 4. Click "Next".
- 2. Configure Source Stage
 - Source Provider: Select AWS CodeCommit.
 - 2. Choose Repository: MyAppRepo.
 - 3. Branch: main.
 - 4. Click "Next".

3. Configure Build Stage

- Build Provider: Select AWS CodeBuild.
- 2. Choose **Build Project**: MyAppBuild.
- 3. Click "Next".

4. Configure Deploy Stage

- 1. Deploy Provider: Select AWS CodeDeploy.
- 2. Choose **Application Name**: MyAppDeploy.
- Choose Deployment Group: MyAppDeploymentGroup.
- 4. Click "Next" → "Create Pipeline".

Expected Outcome:

CodePipeline automates deployments from CodeCommit to EC2 via CodeDeploy.

Step 5: Trigger the CI/CD Pipeline

1. Make a Code Change and Push to CodeCommit

echo "<h1>Updated Deployment via AWS CodePipeline!</h1>" > index.html

git add.

git commit -m "Updated index.html"

git push origin main

- 2. Monitor Pipeline Execution
 - 1. Open AWS CodePipeline Console.
 - 2. Click on MyAppPipeline.

3. Monitor Source, Build, and Deploy Stages.

* Expected Outcome:

✓ The pipeline automatically detects the code change, builds the application, and deploys it to EC2.

Step 6: Verify Deployment on EC2 Instance

- Open AWS EC2 Console → Select your instance.
- 2. Click "Public IP Address" to open the app in a browser.

Expected Outcome:

Your updated application is deployed via AWS CodePipeline! 🞉



CONCLUSION: SUCCESSFULLY DEPLOYED AN APPLICATION USING AWS CODEPIPELINE

- Created a CodeCommit repository for version control.
- Configured CodeBuild to compile and package the application.
- Deployed the application to EC2 using CodeDeploy.
- Automated the CI/CD pipeline using AWS CodePipeline.

FINAL EXERCISE:

- 1. Modify buildspec.yml to include testing before deployment.
- 2. Use AWS Lambda as a deployment target instead of EC2.
- 3. Integrate Amazon S₃ as a deployment destination for static websites.

WRITE AN AWS CLOUDFORMATION TEMPLATE



SOLUTION: WRITE AN AWS CLOUDFORMATION TEMPLATE (STEP-BY-STEP GUIDE)

This guide will walk you through writing an AWS CloudFormation template to deploy an EC2 instance, an S3 bucket, and a security group.

Step 1: Understand AWS CloudFormation

What is AWS CloudFormation?

AWS CloudFormation allows you to **define infrastructure as code** (IaC) using YAML or JSON templates. It automates **resource** creation, updates, and management.

Step 2: Install and Set Up AWS CLI (Optional)

You can deploy CloudFormation templates using the AWS Console or AWS CLI.

- Install AWS CLI (<u>Download Here</u>).
- 2. Configure AWS CLI:
- 3. aws configure

Enter AWS Access Key, Secret Key, Region (e.g., us-east-1), and output format (json).

Step 3: Write the CloudFormation Template (YAML Format)

1. Create a YAML File for CloudFormation

Create a file named cloudformation-template.yaml.

2. Define the Template Structure

Every CloudFormation template has these key sections:

AWSTemplateFormatVersion: "2010-09-09"

Description: "AWS CloudFormation Template to deploy EC2, S3, and Security Group"

Resources:

MyS₃Bucket:

Type: "AWS::S3::Bucket"

Properties:

BucketName: "my-cloudformation-bucket-12345"

MySecurityGroup:

Type: "AWS::EC2::SecurityGroup"

Properties:

GroupDescription: "Allow SSH and HTTP access"

SecurityGroupIngress:

- IpProtocol: "tcp"

FromPort: 22

ToPort: 22

Cidrlp: "o.o.o.o/o"

- IpProtocol: "tcp"

FromPort: 80

ToPort: 80

Cidrlp: "o.o.o.o/o"

MyEC2Instance:

Type: "AWS::EC2::Instance"

Properties:

InstanceType: "t2.micro"

ImageId: "ami-oabcdef1234567890" # Replace with a valid AMI

ID

SecurityGroups:

- Ref: "MySecurityGroup"

KeyName: "my-key-pair" # Replace with an existing key pair

Outputs:

EC2InstancePublicIP:

Description: "Public IP of the EC2 Instance"

Value: !GetAtt MyEC2Instance.PublicIp

 $S_3 Bucket Name: \\$

Description: "Name of the created S₃ bucket"

Value: !Ref MyS3Bucket

Step 4: Validate the Template

Run the following command to validate the YAML syntax:

aws cloudformation validate-template --template-body file://cloudformation-template.yaml

Expected Outcome:

If the template is correct, AWS will return a success message.

Step 5: Deploy the CloudFormation Stack

1. Using AWS Console

- Open AWS Console → Navigate to CloudFormation.
- Click "Create Stack" → Select "With new resources".
- 3. Choose "Upload a template file", then select cloudformation-template.yaml.
- 4. Click "Next" and enter:
 - Stack Name: MyCloudFormationStack.
- 5. Click "Next" → "Create Stack".

2. Using AWS CLI

Run the following command:

aws cloudformation create-stack --stack-name MyCloudFormationStack --template-body file://cloudformationtemplate.yaml

Expected Outcome:

AWS will provision the EC2 instance, S3 bucket, and security group.

Step 6: Verify the Stack Deployment

1. Using AWS Console

- Open AWS CloudFormation Console.
- 2. Click on MyCloudFormationStack.
- 3. Check the "Resources" tab to see deployed resources.
- 4. Go to EC2 Console \rightarrow Find the Instance ID and Public IP.
- 5. Go to S3 Console → Find the created S3 bucket.

2. Using AWS CLI

To check the stack status:

aws cloudformation describe-stacks --stack-name
MyCloudFormationStack

To find the EC2 public IP:

aws ec2 describe-instances --query

'Reservations[*].Instances[*].PublicIpAddress' --output text

***** Expected Outcome:

✓ An EC2 instance is running with SSH & HTTP access.

✓ An S₃ bucket is created.

Step 7: Update the CloudFormation Stack

To modify resources, edit cloudformation-template.yaml, then run:

aws cloudformation update-stack --stack-name MyCloudFormationStack --template-body file://cloudformationtemplate.yaml

Example Update:

Change InstanceType from t2.micro to t3.micro.

Step 8: Delete the CloudFormation Stack

To **clean up resources**, run:

aws cloudformation delete-stack -- stack-name MyCloudFormationStack

Expected Outcome:

✓ AWS removes all resources created by the stack.

CONCLUSION: SUCCESSFULLY CREATED AN AWS

CLOUDFORMATION TEMPLATE

- Defined a YAML CloudFormation template.
- Deployed an EC2 instance, S3 bucket, and Security Group.
- Validated and updated the CloudFormation stack.
- Deleted the stack to clean up resources.

FINAL EXERCISE:

- 1. Modify the template to create an RDS database in a private subnet.
- 2. Use CloudFormation Outputs to return the instance's private IP.
- 3. Enable automatic stack rollback on failure.