



ISDM (INDEPENDENT SKILL DEVELOPMENT MISSION

BIGQUERY – ANALYZING LARGE DATASETS

CHAPTER 1: INTRODUCTION TO BIGQUERY

1.1 What is BigQuery?

BigQuery is a **fully managed**, **serverless**, **and highly scalable data warehouse** provided by **Google Cloud Platform (GCP)** for **analyzing large datasets** in real-time. It allows users to execute **SQL-like queries** on petabyte-scale datasets efficiently.

1.2 Key Features of BigQuery

- ✓ Serverless No infrastructure management required.
- ✓ Fast SQL Queries Uses Dremel technology for quick query execution.
- ✓ Built-in Machine Learning (BigQuery ML) Enables ML model training using SQL.
- √ Streaming Data Ingestion Supports real-time analytics.
- ✓ **Security & Compliance** Supports **IAM-based access control** and encryption.

***** Example:

An **e-commerce company** uses **BigQuery** to analyze **customer purchasing behavior** and improve product recommendations.

CHAPTER 2: BIGQUERY ARCHITECTURE

2.1 How BigQuery Works?

BigQuery follows a **distributed architecture** with the following components:

Component	Description	
Storage Engine	Stores structured and semi-structured data	
	efficiently.	
Query Engine	Uses Dremel execution for fast SQL queries .	
BigQuery ML	Enables machine learning model training	
	within BigQuery.	
BigQuery BI	Optimizes dashboard performance in Looker,	
Engine	Data Studio, etc.	
Data Transfer	Automates data movement from Google Ads,	
Service	YouTube, and other sources.	

📌 Example:

A marketing team imports Google Analytics data into BigQuery for real-time customer insights.

CHAPTER 3: SETTING UP BIGQUERY

3.1 Prerequisites

- √ Google Cloud Project with billing enabled.
- ✓ Enable the **BigQuery API**.
- ✓ Install Google Cloud SDK (for CLI access).

3.2 Creating a BigQuery Dataset

Step 1: Open BigQuery Console

- Navigate to Google Cloud Console → BigQuery.
- 2. Click **Create Dataset** \rightarrow Name it my_dataset.

Step 2: Upload Data to BigQuery

- 1. Open the **BigQuery Console**.
- 2. Click **Create Table** → Choose **Upload File**.
- 3. Select CSV, JSON, or Avro file → Click Create Table.

***** Example:

A finance company uploads transaction logs to Big Query for fraud detection analysis.

CHAPTER 4: QUERYING DATA IN BIGQUERY

4.1 Running Basic SQL Queries

BigQuery supports **SQL-like syntax** to retrieve data.

Example: Selecting Data

SELECT customer_id, total_spent

FROM 'my_project.my_dataset.sales_data'

WHER<mark>E total_spe</mark>nt > 1000

ORDER BY total_spent DESC

LIMIT 10;

Example: Aggregating Data

SELECT category, COUNT(*) AS num_orders

FROM 'my_project.my_dataset.orders'

GROUP BY category

ORDER BY num_orders DESC;



* Example:

An online retailer queries BigQuery to identify the top-selling product categories.

CHAPTER 5: OPTIMIZING QUERY PERFORMANCE

5.1 Best Practices for Faster Queries

- ✓ Partition Tables Store data in time-based partitions.
- ✓ **Use Clustering** Group frequently queried fields together.
- ✓ Avoid SELECT * Query only required columns.
- ✓ **Use Query Caching** Reduce cost and execution time.

Example: Using Partitioning in BigQuery

CREATE TABLE my_project.my_dataset.sales_data

PARTITION BY DATE(transaction_date)

AS

SELECT * FROM `source_table`;



Example:

A logistics company improves shipment tracking performance using partitioned tables.

CHAPTER 6: BIGQUERY ML – MACHINE LEARNING IN SQL

6.1 What is BigQuery ML?

BigQuery ML enables users to train machine learning models using SQL queries.

6.2 Creating an ML Model in BigQuery

CREATE MODEL my_project.my_dataset.customer_churn_model OPTIONS(model_type='logistic_reg') AS

SELECT * FROM my_project.my_dataset.customer_data;

***** Example:

A telecom company predicts customer churn using BigQuery ML.

CHAPTER 7: SECURITY & ACCESS CONTROL IN BIGQUERY
7.1 Implementing IAM Policies

✓ Role-Based Access Control (RBAC) – Assign roles such as Viewer, Editor, Owner.

- ✓ Row-Level Security Restrict access at row-level.
- ✓ Data Encryption Supports Customer-Managed Encryption Keys (CMEK).

***** Example:

A healthcare organization ensures HIPAA compliance by restricting patient data access.

CHAPTER 8: EXERCISE & REVIEW QUESTIONS

Exercise:

- Upload a dataset to BigQuery and create a table.
- 2. **Run SQL queries** to filter and aggregate data.
- 3. **Optimize performance** using partitions and clustering.
- 4. **Train an ML model** using BigQuery ML.

Review Questions:

- 1. What are the key components of **BigQuery Architecture**?
- 2. How does **BigQuery ML** enable machine learning?
- 3. What is the difference between **partitioning and clustering**?
- 4. What are the best practices for optimizing BigQuery queries?
- 5. How can **IAM roles** enhance security in BigQuery?

CONCLUSION: UNLOCKING THE POWER OF BIGQUERY

- ✓ **BigQuery simplifies large-scale data analytics w**ith SQL-based querying.
- ✓ Optimizations like partitioning & clustering improve query performance.
- ✓ Built-in ML capabilities enable predictive analytics.
- Mastering BigQuery helps businesses make data-driven decisions efficiently!

GOOGLE CLOUD DATAFLOW – STREAM & BATCH PROCESSING PIPELINES

CHAPTER 1: INTRODUCTION TO GOOGLE CLOUD DATAFLOW

1.1 What is Google Cloud Dataflow?

Google Cloud Dataflow is a **fully managed service** for **stream and batch data processing**, built on **Apache Beam**. It allows developers
to **process, analyze, and transform large datasets** in **real-time**(stream processing) or as **scheduled jobs** (batch processing).

1.2 Key Features of Dataflow

- ✓ Serverless Execution No need to manage infrastructure.
- ✓ Unified Model Supports both batch and stream processing.
- ✓ Auto-Scaling Dynamically scales resources based on demand.
- ✓ Pay-for-Use Pricing Charges only for the compute resources used.
- ✓ Integration with BigQuery, Pub/Sub, Cloud Storage, and AI tools.

* Example:

An e-commerce company uses Dataflow with Pub/Sub to process real-time sales transactions and generate real-time analytics dashboards.

CHAPTER 2: UNDERSTANDING BATCH & STREAM PROCESSING

2.1 What is Batch Processing?

Batch processing handles large volumes of data at once, often at scheduled intervals.

√ Use Cases:

- Processing log files from a website.
- Performing ETL (Extract, Transform, Load) jobs.
- Generating daily reports from stored data.

√ Example:

A finance company processes daily stock market transactions using Dataflow batch processing to calculate end-of-day summaries.

2.2 What is Stream Processing?

Stream processing processes data in real time as it arrives.

√ Use Cases:

- Real-time fraud detection for banking transactions.
- Monitoring IoT sensor data.
- Processing live customer interactions (e.g., chatbots, recommendations).

√ Example:

A media streaming service uses Dataflow to analyze live video viewership trends and adjust recommendations dynamically.

CHAPTER 3: ARCHITECTURE OF GOOGLE CLOUD DATAFLOW

3.1 Core Components of Dataflow

Component	Description
Pipeline	The flow of data from input to output.

Transform	Operations applied to data (filtering, aggregating,
	etc.).
PCollection	A distributed dataset in Dataflow.
Source	The input data source (e.g., Pub/Sub, Cloud Storage,
	BigQuery).
Sink	The output destination (e.g., BigQuery, Cloud
	Storage, Datastore).

***** Example:

A retail company uses Dataflow to process customer transaction logs from Cloud Storage and load the results into BigQuery.

3.2 Dataflow vs Other Google Cloud Services

Feature	Dataflow	BigQuery	Dataproc
Processing	Batch &	Query-based	Managed
Model	Stream	(SQL)	Hadoop/Spark
Scalability	Auto-scales	Auto-scales	Cluster-based
Use Case	Real-time &	Data	Big data analytics
	batch	warehousing	
	processing		
Best For	ETL, real-time	SQL analytics	Custom ML & big
	analytics		data jobs

***** Example:

A telecom provider uses Dataflow to process call records in realtime, while storing historical data in **BigQuery**.

CHAPTER 4: SETTING UP A DATAFLOW PIPELINE

4.1 Prerequisites

- ✓ Enable the **Dataflow API** in **Google Cloud Console**.
- ✓ Install Google Cloud SDK (gcloud init).
- √ Install Apache Beam SDK:

pip install apache-beam[gcp]

4.2 Creating a Dataflow Pipeline Using Python (Apache Beam)

Step 1: Define the Pipeline

import apache_beam as beam

from apache_beam.options.pipeline_options import
PipelineOptions

pipeline_options = PipelineOptions()

p = beam.Pipeline(options=pipeline_options)

Read data from Cloud Storage

input_<mark>d</mark>ata = 'gs://my-bucket/input-data.csv'

lines = p | "Read from Cloud Storage" >> beam.io.ReadFromText(input_data)

Apply transformation (e.g., filtering data)

filtered_lines = lines | "Filter" >> beam.Filter(lambda line: "error" in line.lower())

Write output to Cloud Storage

output_data = 'gs://my-bucket/output-data.txt'

filtered_lines | "Write to Cloud Storage" >> beam.io.WriteToText(output_data)

p.run()

Step 2: Deploy the Pipeline on Dataflow

python my_pipeline.py --runner DataflowRunner --project my-gcpproject --temp_location gs://my-bucket/temp/

***** Example:

A log management system filters error logs from Cloud Storage and writes filtered logs back to Cloud Storage.

CHAPTER 5: INTEGRATING DATAFLOW WITH OTHER GOOGLE CLOUD SERVICES

5.1 Dataflow with Pub/Sub for Real-Time Processing

- 1. Pub/Sub receives streaming data (e.g., IoT, transactions).
- 2. Dataflow processes the stream (e.g., filtering, aggregating).
- 3. Data is written to BigQuery for analysis.

√ Use Case:

- Fraud detection in financial transactions.
- Monitoring real-time customer orders.

Sample Code: Streaming Data from Pub/Sub to BigQuery

from apache_beam.io.gcp.pubsub import ReadFromPubSub from apache_beam.io.gcp.bigquery import WriteToBigQuery

```
p = beam.Pipeline()

messages = (
    p
    | "Read from Pub/Sub" >>
ReadFromPubSub(subscription="projects/my-
project/subscriptions/my-subscription")
    | "Transform Data" >> beam.Map(lambda msg: {"message":
    msg.decode("utf-8")})
    | "Write to BigQuery" >> WriteToBigQuery("my-
project:dataset.table")
)
```

Example:

A cybersecurity firm uses Dataflow with Pub/Sub to detect unauthorized access attempts in real time.

CHAPTER 6: BEST PRACTICES FOR DATAFLOW PIPELINES

✓ **Use Windowing for Streaming Data** – Helps group real-time data into time-based windows.

✓ Optimize Resource Allocation – Use auto-scaling to control

costs.

- ✓ Use Dead-Letter Queues Redirect failed records to Cloud Storage or Pub/Sub.
- ✓ Enable Logging & Monitoring Use Cloud Logging & Stackdriver for debugging.
- ✓ Secure Dataflow Pipelines Implement IAM roles to restrict access.

***** Example:

A healthcare provider ensures data security by using IAM roles to restrict access to Dataflow pipelines processing patient records.

CHAPTER 7: EXERCISE & REVIEW QUESTIONS

Exercise:

- Deploy a batch pipeline that processes CSV files from Cloud Storage.
- 2. **Create a streaming pipeline** that ingests data from Pub/Sub and writes to BigQuery.
- Optimize a Dataflow job by setting auto-scaling and monitoring logs.

Review Questions:

- 1. What is the difference between batch processing and stream processing?
- 2. How does **Apache Beam** enable portability in Dataflow pipelines?
- 3. What are **PCollections** in Dataflow?
- 4. How can **windowing** improve real-time data processing?

5. What is the role of Pub/Sub in streaming pipelines?

CONCLUSION: LEVERAGING DATAFLOW FOR SCALABLE DATA PROCESSING

- ✓ Google Cloud Dataflow enables real-time and batch data processing at scale.
- ✓ Integration with Pub/Sub, BigQuery, and Cloud Storage makes it ideal for ETL workflows.
- ✓ Serverless execution, auto-scaling, and monitoring features optimize performance.
- Mastering Dataflow helps businesses process, analyze, and transform data efficiently in Google Cloud!

CLOUD DATAPROC – RUNNING SPARK & HADOOP ON GCP

CHAPTER 1: INTRODUCTION TO CLOUD DATAPROC

1.1 What is Cloud Dataproc?

Cloud Dataproc is a **fully managed, cloud-native service** provided by Google Cloud for running **Apache Hadoop, Apache Spark, Apache Flink, Presto, and other big data frameworks**. It enables organizations to process large-scale data workloads **cost-effectively and efficiently**.

1.2 Why Use Cloud Dataproc?

- ✓ Fully Managed No need to set up or manage infrastructure.
- ✓ Fast Cluster Deployment Clusters can be created in under 90 seconds.
- ✓ Scalable Auto-scaling allows dynamic resource allocation.
- ✓ Cost-Effective Pay only for what you use with per-second billing.
- ✓ Integrated with GCP Services Works with BigQuery, Cloud Storage, Cloud Dataflow, and AI/ML services.

1.3 Key Use Cases

- ✓ **Data Transformation & ETL Pipelines** Process structured and unstructured data.
- ✓ Machine Learning & Al Pipelines Train models using big data frameworks.
- ✓ Real-time & Batch Data Processing Process streaming or batch workloads.

√ Ad-hoc Querying & Data Exploration – Run Spark SQL queries
on large datasets.

***** Example:

A retail company uses **Cloud Dataproc to analyze customer purchase trends** by processing large datasets stored in **Google Cloud Storage** using **Apache Spark**.

CHAPTER 2: SETTING UP CLOUD DATAPROC CLUSTER

- 2.1 Prerequisites
- ✓ Google Cloud Project with billing enabled.
- ✓ Enable Cloud Dataproc API:

gcloud services enable dataproc.googleapis.com

- ✓ **Set Up IAM Permissions** for managing Dataproc clusters.
- ✓ Install Google Cloud SDK and authenticate using:

gcloud auth login

2.2 Creating a Dataproc Cluster

Method 1: Using Google Cloud Console

- Open Google Cloud Console → Cloud Dataproc.
- 2. Click Create Cluster.
- 3. Set the **Cluster Name** (e.g., my-dataproc-cluster).
- 4. Choose the **Region** and **Zone**.
- 5. Select Cluster Mode:
 - Standard Multi-node cluster (Default).
 - Single Node Best for testing.

- **High Availability** Multiple master nodes for reliability.
- 6. Configure Machine Types & Nodes.
- 7. Click **Create** to launch the cluster.

Method 2: Using Google Cloud CLI

gcloud dataproc clusters create my-dataproc-cluster \

- --region=us-central1\
- --zone=us-central1-a\
- --master-machine-type=n1-standard-4\
- --worker-machine-type=n1-standard-2 \
- --num-workers=2

* Example:

A finance company sets up a Dataproc cluster with two worker nodes to process large financial transactions in real time.

CHAPTER 3: RUNNING SPARK & HADOOP JOBS ON DATAPROC

3.1 Running an Apache Spark Job

Submit a Spark Job Using the Cloud Console

- Open Cloud Dataproc → Jobs.
- 2. Click **Submit Job** \rightarrow Select **Spark** as the Job Type.
- Specify the main application file (e.g., a JAR file stored in Cloud Storage).
- 4. Provide necessary arguments & parameters.
- 5. Click Submit.

Submit a Spark Job Using CLI

gcloud dataproc jobs submit spark \

- --cluster=my-dataproc-cluster\
- --region=us-central1\
- --class=org.apache.spark.examples.SparkPi\
- --jars=file:///usr/lib/spark/examples/jars/spark-examples.jar\
- -- 1000

* Example:

A **social media company** runs a **Spark job** on Dataproc to analyze trending hashtags from millions of tweets.

3.2 Running a Hadoop Job

Submit a Hadoop Job Using CLI

gcloud dataproc jobs submit hadoop \

- --cluster=my-dataproc-cluster \
- --region=us-central1\
- --class=org.apache.hadoop.examples.WordCount \
- --jars=file:///usr/lib/hadoop-mapreduce/hadoop-mapreduceexamples.jar \
 - -- input.txt output-dir

📌 Example:

A **news website** runs a **Hadoop MapReduce job** to count the number of words in an article dataset stored in Cloud Storage.

CHAPTER 4: DATA PROCESSING WITH CLOUD DATAPROC

4.1 Integrating Dataproc with Cloud Storage

Dataproc uses **Cloud Storage (GCS) as a storage backend** instead of HDFS.

gcloud storage buckets create my-dataproc-bucket --region=uscentral1

✓ Upload input data to Cloud Storage:

gcloud storage cp local-data.csv gs://my-dataproc-bucket/

✓ Read data in Spark using Python (PySpark):

df = spark.read.csv("gs://my-dataproc-bucket/local-data.csv", header=True, inferSchema=True)

df.show()

🖈 Example:

An **e-commerce company** stores sales data in **GCS** and processes it using **Apache Spark on Dataproc**.

CHAPTER 5: OPTIMIZING & SCALING CLOUD DATAPROC

5.1 Autoscaling Dataproc Clusters

Enable auto-scaling for cost efficiency:

gcloud dataproc clusters update my-dataproc-cluster \

--region=us-central1\

--autoscaling-policy=default

5.2 Optimizing Spark Performance on Dataproc

- ✓ Use cluster preemptible VMs to reduce costs.
- ✓ Enable dynamic resource allocation in Spark.
- ✓ Optimize shuffle performance using SSD disks.

***** Example:

A **biotech company** reduces cloud computing costs by enabling **auto-scaling and using preemptible nodes**.

CHAPTER 6: SECURITY & MONITORING IN CLOUD DATAPROC

- 6.1 Securing Dataproc Clusters
- ✓ Use IAM roles to control access.
- ✓ Enable **VPC Service Controls** for network security.
- ✓ Use Cloud Storage ACLs to manage data permissions.
- 6.2 Monitoring Dataproc Jobs
- ✓ View logs in Cloud Logging:

gcloud dataproc jobs list --region=us-central1

✓ Enable Cloud Monitoring & Alerting.

* Example:

A banking organization enables role-based IAM access to restrict sensitive data processing.

CHAPTER 7: EXERCISE & REVIEW QUESTIONS

Exercise:

- 1. Create a Dataproc Cluster with two worker nodes.
- 2. **Run a Spark job** using PySpark.

- 3. **Process a dataset** stored in Cloud Storage.
- 4. **Enable Autoscaling** for cost optimization.

Review Questions:

- 1. What are the benefits of using Cloud Dataproc over onpremise Hadoop clusters?
- 2. How can Cloud Storage be used instead of HDFS in Dataproc?
- 3. What are the differences between running Spark and Hadoop jobs on Dataproc?
- 4. How does auto-scaling improve cost efficiency in Cloud Dataproc?
- 5. What are best practices for securing Dataproc clusters?

CONCLUSION: CLOUD DATAPROC FOR SCALABLE DATA PROCESSING

- ✓ Google Cloud Dataproc provides an efficient way to run big data frameworks like Spark & Hadoop.
- ✓ Integrates seamlessly with GCS, BigQuery, and AI/ML services.
- ✓ Cost-effective with auto-scaling & per-second billing.
- ✓ Secure, managed, and scalable for modern data processing workflows.
- Mastering Cloud Dataproc helps organizations accelerate big data analytics and machine learning in the cloud!

CLOUD PUB/SUB – REAL-TIME MESSAGING & EVENT PROCESSING

CHAPTER 1: INTRODUCTION TO CLOUD PUB/SUB

1.1 WHAT IS CLOUD PUB/SUB?

Cloud Pub/Sub is a **real-time messaging service** provided by Google Cloud that enables **asynchronous communication** between independent applications. It follows the **publish-subscribe messaging model**, ensuring **scalable**, **reliable**, **and event-driven** communication.

- ✓ **Real-time event processing** Enables real-time data streaming.
- ✓ Asynchronous messaging Decouples producers (publishers)
 and consumers (subscribers).
- ✓ Scalable & fault-tolerant Handles millions of messages per second.
- ✓ **Push & Pull delivery** Supports multiple message consumption models.

1.2 Use Cases of Cloud Pub/Sub

- ✓ **Data Streaming** Process IoT device data in real-time.
- ✓ Event-driven Systems Notify services when new events occur.
- ✓ Log Aggregation Collect and analyze logs from multiple sources.
- ✓ **Decoupling Microservices** Enable communication between microservices.
- ✓ Real-time Analytics Process and analyze financial transactions, stock trades, etc.

* Example:

A financial services company uses Cloud Pub/Sub to process stock market data and update dashboards in real-time.

CHAPTER 2: CLOUD PUB/SUB ARCHITECTURE

2.1 Key Components of Cloud Pub/Sub

Component	Description	
Publisher	Produces messages and sends them to a	
	topic.	
Topic	A named resource where messages are	
	sent.	
Subscription	A connection point that allows consumers	
	to receive messages from a topic.	
Subscriber	Applications or services that process	
	messages from a subscription.	
Message	Subscribers acknowledge messages to	
Acknowledgment	remove them from the queue.	

* Example:

A ride-sharing app uses Cloud Pub/Sub to update ride statuses, ensuring drivers and riders get real-time updates.

CHAPTER 3: SETTING UP CLOUD PUB/SUB

3.1 Prerequisites

- ✓ A Google Cloud account with billing enabled.
- ✓ Install Google Cloud SDK and authenticate using gcloud auth

login.

✓ Enable the Pub/Sub API:

gcloud services enable pubsub.googleapis.com

3.2 Creating a Pub/Sub Topic

- Open Google Cloud Console → Navigate to Pub/Sub.
- 2. Click **Create Topic** \rightarrow Enter my-topic.
- Choose No message retention (default) or enable message retention.
- 4. Click Create.

Example:

A weather monitoring system creates a topic called weatherupdates where IoT sensors publish temperature readings.

CHAPTER 4: PUBLISHING & SUBSCRIBING TO MESSAGES

4.1 Creating a Subscription

- Open Google Cloud Console → Go to Pub/Sub → Subscriptions.
- 2. Click Create Subscription → Select my-topic.
- 3. Choose Delivery Type:
 - Pull Subscribers fetch messages manually.
 - Push Messages are automatically pushed to an HTTP endpoint.
- 4. Click Create.

4.2 Publishing a Message

Using gcloud CLI:

gcloud pubsub topics publish my-topic --message "Hello, Pub/Sub!"

4.3 Pulling Messages from a Subscription

gcloud pubsub subscriptions pull my-subscription --auto-ack

* Example:

An e-commerce platform publishes an event when a new order is placed, and a warehouse system subscribes to process orders in real-time.

CHAPTER 5: CLOUD PUB/SUB MESSAGE DELIVERY MODELS

5.1 Push vs Pull Subscriptions

Feature	Push Subscription	Pull Subscription
Delivery	Google Cloud pushes messages to a subscriber's HTTP endpoint	Subscriber pulls messages manually
	Titti chaponit	
Use	Real-time processing (e.g.,	Batch processing (e.g.,
Case	webhooks)	analytics, log
		aggregation)
Latency	Lower latency	Higher latency

5.2 Dead Letter Topics (DLT)

- ✓ Helps store undeliverable messages for debugging.
- ✓ Messages exceeding the **maximum delivery attempts** are sent to a **dead-letter topic**.

***** Example:

A chat application uses a pull subscription to process user messages asynchronously.

CHAPTER 6: SECURING & MONITORING PUB/SUB

6.1 Security Best Practices

- ✓ Use IAM roles (roles/pubsub.publisher & roles/pubsub.subscriber).
- ✓ Enable **encryption** for sensitive messages.
- ✓ Implement access control policies for topics and subscriptions.
- ✓ Use Google Cloud Audit Logs to track message activity.

6.2 Monitoring with Cloud Logging & Metrics

- ✓ View message delivery statistics in Cloud Monitoring.
- ✓ Use **Cloud Logging** to track failed deliveries.
- ✓ Set up **alerts** for failed subscriptions.

***** Example:

A healthcare system encrypts Pub/Sub messages to secure patient records shared across hospitals.

CHAPTER 7: INTEGRATING CLOUD PUB/SUB WITH OTHER SERVICES

Google Cloud	Integration Purpose
Service	
Cloud Functions	Trigger functions on new Pub/Sub
	messages.
Cloud Run	Process messages with serverless
	applications.
BigQuery	Ingest and analyze streaming data.
Cloud Dataflow	Transform and process messages at scale.

* Example:

A smart home automation system integrates Pub/Sub with Cloud **Functions** to **trigger alarms** when motion is detected.

CHAPTER 8: EXERCISE & REVIEW QUESTIONS

Exercise:

- 1. Create a Pub/Sub topic and publish a message.
- 2. **Set up a push subscription** and receive messages via an HTTP endpoint.
- Monitor Pub/Sub message delivery using Cloud Logging.

Review Questions:

- What are the key components of Cloud Pub/Sub?
- 2. How does **push vs pull delivery** work?
- 3. What is the purpose of dead letter topics (DLT)?
- 4. How can IAM roles secure Pub/Sub access?
- 5. What are real-world applications of Cloud Pub/Sub?

CONCLUSION: REAL-TIME EVENT PROCESSING WITH CLOUD PUB/SUB

- √ Cloud Pub/Sub enables scalable, real-time event-driven architecture.
- √ Supports push & pull subscriptions for flexible message consumption.
- ✓ Seamlessly integrates with Google Cloud services for big data processing & microservices.

Mastering Cloud Pub/Sub helps build highly responsive, scalable cloud applications!



CLOUD COMPOSER – MANAGED WORKFLOW AUTOMATION

CHAPTER 1: INTRODUCTION TO CLOUD COMPOSER

1.1 What is Cloud Composer?

Cloud Composer is a **fully managed workflow orchestration service** built on **Apache Airflow**. It allows users to **create, schedule, and monitor data pipelines** in the cloud.

1.2 Key Features of Cloud Composer

- ✓ Managed Apache Airflow No need to manage infrastructure.
- ✓ Scalability Automatically scales based on workflow demand.
- ✓ **Hybrid & Multi-Cloud Support** Connects to on-premises and external cloud services.
- ✓ Built-in Security Uses IAM and VPC Service Controls for security.
- ✓ Monitoring & Logging Integrated with Cloud Logging and Cloud Monitoring.

* Example:

A financial services company uses Cloud Composer to automate data pipeline workflows for fraud detection.

CHAPTER 2: CLOUD COMPOSER ARCHITECTURE

2.1 Key Components

Component	Description

DAG (Directed Acyclic	Defines the sequence of tasks in a
Graph)	workflow.
Scheduler	Determines when and how workflows run.
Workers	Execute the tasks defined in DAGs.
Airflow Metadata	Stores workflow history and task
Database	execution logs.
Cloud Storage Bucket	Stores DAG files and workflow logs.

***** Example:

A retail company uses DAGs in Cloud Composer to automate daily sales report generation.

CHAPTER 3: SETTING UP CLOUD COMPOSER

3.1 Prerequisites

- ✓ Google Cloud Project with billing enabled.
- ✓ Enable the Cloud Composer API.
- ✓ Assign IAM roles: Composer Admin, Viewer, and Service Account
 User.

3.2 Deploying a Cloud Composer Environment

Step 1: Open Cloud Composer in Google Cloud Console

- Navigate to Cloud Composer → Environments.
- 2. Click Create Environment.
- 3. Choose a **region** and specify **environment name**.

Step 2: Configure Composer Environment

- ✓ Machine Type Choose CPU & Memory allocation.
- ✓ Networking Select VPC and Subnet.
- ✓ Python Packages Install additional Python libraries (if needed).

📌 Example:

A healthcare company deploys Cloud Composer to automate patient data synchronization across hospitals.

CHAPTER 4: CREATING WORKFLOWS IN CLOUD COMPOSER

4.1 Understanding DAGs in Apache Airflow

A DAG (Directed Acyclic Graph) is a collection of tasks defining workflow execution order.

- 4.2 Creating a Simple DAG in Cloud Composer
 - 1. Upload DAG to Cloud Storage
 - 2. Edit DAG in Python
 - 3. Deploy DAG to Cloud Composer

Example DAG: Hello World Workflow

from airflow import DAG

from airflow.operators.dummy_operator import DummyOperator from datetime import datetime

```
default_args = {
  'owner': 'airflow',
  'start_date': datetime(2024, 1, 1),
  'retries': 1,
```

}

```
dag = DAG('hello_world',
    default_args=default_args,
    description='A simple Hello World DAG',
    schedule_interval='@daily')
```

start = DummyOperator(task_id='start', dag=dag)
end = DummyOperator(task_id='end', dag=dag)

start >> end # Defines the sequence

***** Example:

A media company automates video processing workflows using Cloud Composer DAGs.

CHAPTER 5: INTEGRATING CLOUD COMPOSER WITH OTHER SERVICES 5.1 Common Cloud Composer Integrations

Service	Use Case
BigQuery	Automate data transformation jobs.
Cloud Storage	Move files between cloud and on-premises.
Cloud Pub/Sub	Trigger workflows on real-time events.
Dataproc	Manage Spark and Hadoop jobs.

Vertex AI	Schedule ML model training workflows.

Example: Automating BigQuery Data Pipeline

from airflow.providers.google.cloud.operators.bigquery import BigQueryOperator

```
bq_query = BigQueryOperator(
   task_id='bq_query_task',
   sql='SELECT * FROM my_project.my_dataset.my_table',
   use_legacy_sql=False,
   dag=dag,
)
```

Example:

A banking institution automates customer transaction analysis using Cloud Composer and BigQuery.

CHAPTER 6: MONITORING & DEBUGGING WORKFLOWS

6.1 Using Cloud Composer Logs & Monitoring

- ✓ Airflow UI View DAG execution status.
- ✓ Cloud Logging Track task failures and warnings.
- ✓ Cloud Monitoring Set up alerts for workflow failures.

6.2 Debugging Failed Workflows

- ✓ Check **Airflow UI logs** for task failures.
- ✓ Use retry mechanisms in DAGs.
- ✓ Validate **permissions and IAM roles** for external services.

* Example:

A social media company monitors ad performance workflows in Cloud Composer.

CHAPTER 7: BEST PRACTICES FOR CLOUD COMPOSER

- ✓ Use Environment Variables Store API keys and credentials securely.
- ✓ Optimize DAG Execution Break large workflows into modular tasks.
- ✓ **Use Airflow Sensors** Wait for external events before triggering tasks.
- ✓ **Leverage Airflow XComs** Share data between tasks dynamically.

Example: Using XCom to Pass Data Between Tasks

from airflow.operators.python_operator import PythonOperator

```
def process_data(**kwargs):
    return 'processed_data'

task1 = PythonOperator(
    task_id='task1',
    python_callable=process_data,
    provide_context=True,
    dag=dag
)
```

* Example:

A gaming company automates daily user engagement reports using optimized DAG execution.

CHAPTER 8: EXERCISE & REVIEW QUESTIONS

Exercise:

- Deploy a Cloud Composer Environment in Google Cloud.
- 2. Create a DAG to automate a file transfer workflow.
- Integrate Cloud Composer with BigQuery for data analysis.
- Set up logging & monitoring to track DAG performance.

Review Questions:

- 1. What is the role of **DAGs in Cloud Composer**?
- 2. How does Cloud Composer integrate with BigQuery?
- 3. What are best practices for debugging failed workflows?
- 4. How does IAM improve security in Cloud Composer?
- 5. Why is Airflow XCom used in DAGs?

CONCLUSION: AUTOMATING WORKFLOWS WITH CLOUD COMPOSER

- √ Cloud Composer automates complex workflows using Apache Airflow
- √ Seamless integrations with BigQuery, Cloud Storage, and Dataproc.
- ✓ Monitoring tools like Cloud Logging and Monitoring ensure workflow reliability.

Mastering Cloud Composer helps organizations build scalable, automated data pipelines efficiently!



GOOGLE DATA STUDIO – DATA VISUALIZATION & REPORTING

CHAPTER 1: INTRODUCTION TO GOOGLE DATA STUDIO

1.1 What is Google Data Studio?

Google Data Studio is a **free**, **cloud-based data visualization tool** that helps users create **interactive reports and dashboards**. It enables businesses to **connect**, **visualize**, **and share data from multiple sources** with an easy-to-use interface.

1.2 Key Features of Google Data Studio

- ✓ **Drag-and-Drop Interface** No coding required.
- ✓ Connects to Multiple Data Sources Supports Google Analytics, BigQuery, Sheets, SQL databases, and more.
- ✓ Customizable Reports & Dashboards Offers filters, charts, and interactive elements.
- ✓ **Real-Time Data Updates** Ensures up-to-date reporting.
- ✓ Collaboration & Sharing Similar to Google Docs, allowing multiple users to edit reports.

Example:

A marketing team uses Google Data Studio to create real-time dashboards for tracking website traffic and conversion rates from Google Analytics.

CHAPTER 2: CONNECTING DATA SOURCES IN GOOGLE DATA STUDIO

2.1 Supported Data Sources

Google Data Studio supports both **Google and third-party data** sources.

Google Data Sources	Third-Party Data Sources
Google Analytics	MySQL, PostgreSQL, SQL Server
Google Ads	Facebook Ads
Google BigQuery	Salesforce
Google Sheets	Shopify
Google Search Console	Stripe, Mailchimp

2.2 How to Connect a Data Source

- Open Google Data Studio (https://datastudio.google.com).
- 2. Click Create → Data Source.
- 3. Select a data connector (e.g., Google Sheets, BigQuery).
- 4. Grant permissions and authorize access.
- 5. Configure data fields and metrics.
- 6. Click Add to Report.

***** Example:

A finance team connects Google Sheets to Data Studio to generate monthly financial reports dynamically.

CHAPTER 3: CREATING DATA VISUALIZATIONS

3.1 Common Chart Types in Data Studio

Chart	Usage	
Туре		
/ -		

Bar Chart	Compare categorical data (e.g., sales by region).
Line Chart	Show trends over time (e.g., website traffic).
Pie Chart	Display proportions (e.g., revenue by product category).
Geo Map	Visualize location-based data (e.g., user traffic by country).
Tables	Present structured data (e.g., product sales).

3.2 Steps to Create a Chart in Google Data Studio

- 1. Open a report in Google Data Studio.
- 2. Click + Add a chart.
- 3. Select a chart type (e.g., Bar, Line, Pie).
- 4. Choose the data source.
- 5. Configure dimensions (categories) and metrics (values).
- 6. Customize the colors, labels, and styles.
- 7. Click **Apply** to add it to the dashboard.

***** Example:

A sales team creates a line chart in Data Studio to track monthly revenue growth trends.

CHAPTER 4: CUSTOMIZING REPORTS & DASHBOARDS

4.1 Adding Filters & Controls

Filters allow users to interact with reports dynamically.

- ✓ Date Filters Enable users to view data for different time periods.
- ✓ Dropdown Filters Allow filtering by product, region, or

department.

✓ **Search Filters** – Let users search specific data points.

4.2 Steps to Add Filters

- 1. Click + Add a Control in the report.
- 2. Choose Date Range, Drop-down List, or Search Box.
- 3. Select the **data source** and assign relevant fields.
- 4. Adjust styling and apply changes.

***** Example:

A retail company adds filters to a sales dashboard to let users select different product categories.

4.3 Using Calculated Fields for Custom Metrics

Calculated fields allow users to create **custom formulas** for advanced analysis.

√ Example Formulas:

- **Profit Margin:** (Revenue Cost) / Revenue
- Click-Through Rate (CTR): (Clicks / Impressions) * 100
- Customer Retention Rate: (Returning Customers / Total Customers) * 100

Steps to Create a Calculated Field

- 1. Go to the **Data Source** settings.
- 2. Click + Add a Field.
- 3. Enter a Formula (e.g., Revenue Expenses).
- 4. Click **Apply** and use it in visualizations.

A **digital marketing agency** creates a **CTR metric** using calculated fields to track **advertising campaign performance**.

CHAPTER 5: SHARING & COLLABORATING ON REPORTS

5.1 How to Share a Google Data Studio Report

- 1. Click **Share** in the top-right corner.
- 2. Add email addresses of collaborators.
- 3. Set permissions:
 - Viewer (Read-only)
 - Editor (Can edit and modify reports)
- 4. Click Send Invitation.

* Example:

A **CEO shares a performance dashboard** with department heads to track **monthly KPIs**.

5.2 Embedding & Exporting Reports

- ✓ Embed in Websites Copy iframe code to embed reports in a website.
- ✓ Export as PDF Download reports for offline analysis.
- ✓ Schedule Email Reports Send automated reports to stakeholders.

Example:

A news website embeds Google Data Studio charts to show live election results.

CHAPTER 6: BEST PRACTICES FOR DATA VISUALIZATION IN DATA STUDIO

- ✓ **Keep Reports Simple & Clear** Avoid overloading dashboards with too much data.
- ✓ **Use Consistent Colors & Formatting** Maintain uniform styling for better readability.
- ✓ **Highlight Key Metrics with Scorecards** Show important KPIs at the top.
- ✓ **Use Filters to Enhance Interactivity** Allow users to explore data dynamically.
- ✓ Optimize Performance Reduce load time by using preaggregated data sources.

***** Example:

A marketing manager designs a minimalistic dashboard with key performance indicators (KPIs) at the top.

CHAPTER 7: EXERCISE & REVIEW QUESTIONS

Exercise:

- Create a sales report in Google Data Studio using Google Sheets.
- 2. Add a pie chart showing sales distribution by region.
- 3. **Implement filters** for date range and product categories.
- 4. **Share the report with your team** and schedule automated email updates.

Review Questions:

1. What is the main purpose of Google Data Studio?

- 2. How can you connect Google Analytics to Google Data Studio?
- 3. What are calculated fields, and how can they be used?
- 4. What is the difference between a bar chart and a pie chart?
- 5. How do you **share a report with collaborators** in Google Data Studio?

CONCLUSION: MASTERING DATA VISUALIZATION WITH GOOGLE
DATA STUDIO

- ✓ Google Data Studio simplifies business intelligence with realtime, interactive dashboards.
- ✓ Connect multiple data sources, create custom metrics, and build engaging visual reports.
- ✓ Collaboration and automation make reporting efficient for business teams and analysts.
- Mastering Google Data Studio enables professionals to transform raw data into actionable insights!

ASSIGNMENT

PROCESS STREAMING DATA USING CLOUD DATAFLOW



SOLUTION: PROCESS STREAMING DATA USING CLOUD DATAFLOW

Overview

Google Cloud Dataflow is a fully managed stream and batch data processing service built on Apache Beam. It enables real-time analytics, ETL (Extract, Transform, Load) workflows, and AI/ML model streaming by processing live data from sources like Pub/Sub, Kafka, and Cloud Storage.

Step 1: Set Up the Google Cloud Environment

1.1 Enable Required Services

Ensure the following services are **enabled** in Google Cloud:

gcloud services enable dataflow.googleapis.com pubsub.googleapis.com storage.googleapis.com bigquery.googleapis.com

1.2 Set Up IAM Permissions

Assign the **necessary roles** to the user or service account:

- Dataflow Admin
- Pub/Sub Editor
- BigQuery Data Editor
- Cloud Storage Admin

gcloud projects add-iam-policy-binding [PROJECT_ID] \

--member=user:[YOUR_EMAIL] --role=roles/dataflow.admin



A **log monitoring system** assigns IAM roles to **Dataflow processing agents** to handle real-time logs.

Step 2: Create a Streaming Data Source (Pub/Sub)

2.1 Create a Pub/Sub Topic & Subscription

Create a Pub/Sub topic to receive streaming data:

gcloud pubsub topics create streaming-topic

Create a **subscription** to allow Dataflow to pull messages:

gcloud pubsub subscriptions create streaming-sub \

--topic=streaming-topic



Example:

A retail company sets up a Pub/Sub topic to receive real-time sales transactions.

Step 3: Develop a Streaming Pipeline Using Apache Beam

3.1 Install Apache Beam SDK

Ensure you have Apache Beam installed:

pip install apache-beam[gcp]

3.2 Write a Python Pipeline to Process Streaming Data

Create a script dataflow_streaming.py:

import apache_beam as beam

from apache_beam.options.pipeline_options import PipelineOptions, StandardOptions

```
# Define pipeline options
pipeline_options = PipelineOptions(
 streaming=True,
 project='[PROJECT_ID]',
 region='us-central1',
 job_name='streaming-dataflow-job',
 temp_location='qs://[BUCKET_NAME]/temp',
 staging_location='gs://[BUCKET_NAME]/staging'
)
# Define a processing function
def transform_data(element):
 import json
 record = json.loads(element)
                          record['user_id'],
                                                'purchase_amount':
            {'user_id':
  return
record['amount']}
# Build the Dataflow pipeline
with beam. Pipeline (options = pipeline _ options) as pipeline:
 (
   pipeline
```

```
'Read
                             from
                                           Pub/Sub'
beam.io.ReadFromPubSub(topic='projects/[PROJECT_ID]/topics/str
eaming-topic')
   | 'Decode UTF-8' >> beam.Map(lambda msg: msg.decode('utf-8'))
   | 'Transform Data' >> beam.Map(transform_data)
   | 'Write to BigQuery' >> beam.io.WriteToBigQuery(
     '[PROJECT_ID]:dataset_name.table_name',
     schema='user_id:STRING, purchase_amount:FLOAT',
write_disposition=beam.io.BigQueryDisposition.WRITE_APPEND
   )
                                                       Example:
A banking application processes real-time transactions and writes
the output to BigQuery for fraud detection.
```

Step 4: Deploy the Streaming Pipeline to Cloud Dataflow

4.1 Up<mark>load Pyth</mark>on Script to Cloud Storage

gsutil cp dataflow_streaming.py gs://[BUCKET_NAME]/dataflow/

4.2 Run Dataflow Job Using Cloud CLI

 $python - m \ apache_beam.runners.dataflow.dataflow_runner \setminus$

- --project=[PROJECT_ID] \
- --region=us-central1\
- --staging_location=gs://[BUCKET_NAME]/staging/\

- --temp_location=gs://[BUCKET_NAME]/temp/\
- --runner=DataflowRunner



A **stock trading company** deploys a Dataflow job to **track real-time stock price changes**.

Step 5: Monitor and Optimize Streaming Jobs

5.1 Monitor Dataflow Pipeline

- Open Google Cloud Console → Dataflow
- View job status, logs, and data throughput
- Use Cloud Logging to troubleshoot failures

gcloud dataflow jobs list

5.2 Optimize Streaming Jobs

- ✓ Autoscaling: Use flexible worker count to optimize resources
- ✓ **Use Dataflow Shuffle**: Reduces memory footprint for large datasets
- ✓ Dead-letter queue: Handle failed messages by routing them to Pub/Sub or Cloud Storage



Example:

A **smart home automation** company optimizes streaming jobs **to scale up processing during peak hours**.

Step 6: Exercise & Review Questions

Exercise:

1. Create a Pub/Sub topic and publish test messages.

- 2. **Deploy a streaming Dataflow pipeline** that writes to BigQuery.
- 3. Optimize Dataflow performance using autoscaling.

Review Questions:

- 1. What are the **benefits of Cloud Dataflow** for streaming data?
- 2. How does Pub/Sub work as a data ingestion source for Dataflow?
- 3. What are the best practices for scaling Dataflow jobs?
- 4. How can **Cloud Monitoring help optimize** streaming pipelines?
- 5. How does **BigQuery integration enhance real-time analytics**?

CONCLUSION: CLOUD DATAFLOW FOR REAL-TIME PROCESSING

- ✓ Fully managed Apache Beam execution for streaming & batch workloads
- ✓ Integrates with Pub/Sub, BigQuery, Cloud Storage, and Alservices
- ✓ Autoscaling & fault-tolerant streaming for enterprise-grade applications
- ✓ Secure, cost-effective, and highly scalable solution for big data analytics
- Mastering Cloud Dataflow enables real-time insights and datadriven decision-making!

ANALYZE LARGE DATASETS USING BIGQUERY



SOLUTION: ANALYZE LARGE DATASETS USING BIGQUERY

Step 1: Understanding BigQuery

1.1 What is BigQuery?

Google BigQuery is a **serverless**, **highly scalable**, **and cost-effective data warehouse** designed for **fast SQL-based analytics** on large datasets.

- ✓ **Serverless** No infrastructure management required.
- ✓ Massive Scalability Handles petabyte-scale data analysis.
- ✓ Real-time Analytics Supports streaming data ingestion.
- ✓ Built-in AI & ML Capabilities Runs machine learning models directly in SQL.

*

Example:

A **retail company** uses BigQuery to **analyze customer purchase trends** from millions of transactions.

Step 2: Setting Up BigQuery

2.1 Prerequisites

- ✓ A Google Cloud account with billing enabled.
- ✓ Enable BigQuery API in the Google Cloud Console:

gcloud services enable bigquery.googleapis.com

✓ Install and authenticate the Google Cloud SDK:

gcloud auth login

2.2 Creating a BigQuery Dataset

- Open Google Cloud Console → Navigate to BigQuery.
- 2. Click **Create Dataset** → Enter sales_dataset.
- Choose a location (e.g., US).
- 4. Click Create.



A **finance company** creates a dataset named transactions_dataset to store stock market transaction records.

Step 3: Importing Data into BigQuery

3.1 Loading Data from Google Cloud Storage

- Go to BigQuery → Select your dataset (sales_dataset).
- 2. Click Create Table → Choose Source: Cloud Storage.
- 3. Enter the Cloud Storage URI:
- 4. gs://my-bucket/sales_data.csv
- 5. Select File Format: CSV.
- 6. Click Create Table.

3.2 Loading Data from a Local File

Using gcloud CLI, upload a CSV file:

bq load --source_format=CSV sales_dataset.sales_table ./sales_data.csv

3.3 Loading Data from Google Sheets

- 1. In **BigQuery Console**, click **Create Table**.
- 2. Select **Source: Google Drive** → Enter the Google Sheet URL.

- Choose File Format: Google Sheets.
- 4. Click Create Table.



An **IoT company** loads **sensor data** into BigQuery for **real-time temperature monitoring**.

Step 4: Querying Large Datasets in BigQuery

4.1 Running a Simple Query

Example: Get total sales per region

SELECT region, SUM(sales) AS total_sales

FROM 'sales_dataset.sales_table'

GROUP BY region

ORDER BY total_sales DESC;

4.2 Filtering Data Efficiently

Example: Find orders over \$500 in the last month

SELECT order_id, customer_name, total_amount

FROM 'sales_dataset.sales_table'

WHERE total_amount > 500

AND order_date >= DATE_SUB(CURRENT_DATE(), INTERVAL 30 DAY);

4.3 Using Partitioning for Faster Queries

CREATE TABLE sales_dataset.sales_partitioned

PARTITION BY DATE(order_date) AS

SELECT * FROM 'sales_dataset.sales_table';



Example:

A healthcare company runs queries on partitioned patient records to retrieve last month's appointments.

Step 5: Optimizing Queries for Performance

5.1 Best Practices for Query Optimization

- ✓ Use Partitioned Tables Speeds up queries on large datasets.
- ✓ Avoid SELECT * Query only the required columns.
- ✓ Use Filter Conditions Reduces scanned data.
- ✓ Leverage Clustering Groups data to improve indexing.



Example:

A media company uses partitioned logs to quickly analyze user behavior on their streaming platform.

Step 6: Visualizing Data in Google Data Studio

6.1 Connecting BigQuery to Data Studio

- Open Google Data Studio → Click Create Report.
- Click Add Data → Select BigQuery.
- 3. Choose sales_dataset.sales_table.
- 4. Click Add to Report.
- 5. Create **charts & dashboards** to visualize **sales trends**.



Example:

A marketing team visualizes advertising performance using a BigQuery-powered dashboard.

Step 7: Automating Data Pipelines with BigQuery

7.1 Scheduling Queries in BigQuery

- In BigQuery Console, go to Scheduled Queries.
- 2. Click Create a New Scheduled Query.
- 3. Enter SQL Query:
- 4. SELECT customer_id, COUNT(order_id) AS order_count
- FROM `sales_dataset.sales_table`
- 6. WHERE order_date >= CURRENT_DATE()
- 7. GROUP BY customer_id;
- 8. Set Run Frequency (e.g., Daily at Midnight).
- 9. Click **Create**.

Example:

An **HR team** sets up **automated reporting** for **employee attendance** data.

Step 8: Using BigQuery ML for Predictive Analysis

8.1 Training a Machine Learning Model in BigQuery

Predict customer churn using a logistic regression model:

CREATE OR REPLACE MODEL sales_dataset.churn_model

OPTIONS(model_type='logistic_reg') AS

SELECT age, total_spent, visit_frequency, churn_label

FROM 'sales_dataset.sales_table';

8.2 Running Predictions

SELECT customer_id, predicted_churn

FROM ML.PREDICT(MODEL sales_dataset.churn_model,

(SELECT age, total_spent, visit_frequency FROM `sales_dataset.sales_table`));



Example:

A **telecom company** builds a **churn prediction model** in **Big**Query ML to identify customers likely to switch providers.

Step 9: Exercise & Review Questions

Exercise:

- 1. Create a BigQuery dataset and import a sample CSV file.
- 2. **Run SQL queries** to analyze sales performance.
- 3. Set up a scheduled query to automate reporting.
- 4. Create a dashboard in Google Data Studio using BigQuery data.
- 5. **Train a machine learning model** using BigQuery ML.

Review Questions:

- 1. What are the advantages of **serverless BigQuery** over traditional databases?
- 2. How does **partitioning and clustering** improve query performance?
- 3. How can **BigQuery ML** be used for predictive analytics?
- 4. What are the different ways to import data into BigQuery?

5. How does **Google Data Studio** enhance BigQuery data visualization?

CONCLUSION: SCALING ANALYTICS WITH BIGQUERY

- √ BigQuery enables fast and scalable data analysis.
- ✓ Supports structured, semi-structured, and streaming data.
- ✓ Built-in ML capabilities allow predictive analytics with SQL.
- ✓ Integrates seamlessly with Google Data Studio for visualization.