



ISDM (INDEPENDENT SKILL DEVELOPMENT MISSION)

DEPLOYING & INTEGRATING MONGODB WITH NODE.JS (WEEKS 11-12)

USING MONGOOSE ORM FOR DATA MODELING IN MONGODB

CHAPTER 1: INTRODUCTION TO MONGOOSE ORM

1.1 What is Mongoose?

Mongoose is an **Object-Relational Mapping (ORM)** library for **MongoDB and Node.js** that provides a **structured way** to define schemas and interact with MongoDB databases.

- ✓ Simplifies MongoDB queries by using JavaScript models.
- ✓ Provides validation for ensuring data consistency.
- ✓ Supports relationships between documents using references.
- ✓ Adds middleware hooks for handling pre/post operations.

1.2 Why Use Mongoose Instead of Native MongoDB Driver?

Feature	MongoDB Native Driver	Mongoose ORM

Schema Support	✗ No predefined schema	✓ Schema-based
Validation	✗ Manual validation	✓ Built-in validation
Relationships	✗ Manual references	✓ Built-in references
Middleware	✗ Not supported	✓ Supports pre/post hooks

✓ **Mongoose provides structured, efficient, and secure database management** compared to raw MongoDB queries.

CHAPTER 2: SETTING UP MONGOOSE IN A NODE.JS PROJECT

2.1 Installing Mongoose

To use Mongoose in a Node.js project, install it using npm:

npm install mongoose

✓ Installs Mongoose for managing MongoDB connections and models.

2.2 Connecting to MongoDB Using Mongoose

Create a **database connection** in database.js:

```
const mongoose = require('mongoose');
```

```
mongoose.connect("mongodb://localhost:27017/myDatabase", {
  useNewUrlParser: true,
  useUnifiedTopology: true
})
```

```
})  
.then(() => console.log("Connected to MongoDB"))  
.catch(err => console.error("Connection error:", err));
```

✓ Establishes a connection to a **MongoDB database** using Mongoose.

✓ .then() confirms a **successful connection**.

✓ .catch() handles **connection errors**.

CHAPTER 3: DEFINING MONGOOSE SCHEMAS & MODELS

3.1 What is a Mongoose Schema?

A **schema** defines the **structure** of documents in a MongoDB collection.

✓ **Specifies field types** (String, Number, Boolean, Date, etc.).

✓ **Allows default values** and **validations**.

✓ **Defines relationships** between collections.

3.2 Creating a Simple Mongoose Schema

Create a models/User.js file:

```
const mongoose = require('mongoose');
```

```
const userSchema = new mongoose.Schema({
```

```
    name: { type: String, required: true },
```

```
    email: { type: String, required: true, unique: true },
```

```
    age: { type: Number, min: 18 },
```

```
    createdAt: { type: Date, default: Date.now }  
});
```

```
const User = mongoose.model("User", userSchema);
```

```
module.exports = User;
```

- ✓ Defines a **User schema** with fields name, email, age, and createdAt.
- ✓ Ensures name and email are **required** and email is **unique**.
- ✓ Sets age to a **minimum of 18**.
- ✓ createdAt defaults to the **current date**.

3.3 Creating a New User Document

To insert a new user, create insertUser.js:

```
const mongoose = require('mongoose');  
  
const User = require('./models/User');  
  
mongoose.connect("mongodb://localhost:27017/myDatabase");
```

```
async function createUser() {  
  try {  
    const newUser = await User.create({  
      name: "Alice Johnson",
```

```
email: "alice@example.com",  
age: 25  
});  
  
console.log("User Created:", newUser);  
} catch (error) {  
    console.error("Error creating user:", error);  
} finally {  
    mongoose.connection.close();  
}  
}  
  
createUser();
```

- ✓ Uses User.create() to insert a **new user document**.
- ✓ **Handles errors** if validation fails.
- ✓ **Closes the database connection** after execution.

CHAPTER 4: QUERYING DATA WITH MONGOOSE

4.1 Retrieving Documents (`find()` and `findOne()`)

Fetch **all users** from the database:

```
const users = await User.find();  
  
console.log(users);
```

Fetch a **specific user by email**:

```
const user = await User.findOne({ email: "alice@example.com" });
```

```
console.log(user);
```

- ✓ `find()` retrieves **multiple** users.
 - ✓ `findOne()` retrieves **a single document** matching the query.
-

4.2 Updating Documents (`updateOne()` and `findByIdAndUpdate()`)

To **update a user's age**, use:

```
await User.updateOne({ email: "alice@example.com" }, { $set: { age: 30 } });
```

Or update by **ID** and return the modified document:

```
const updatedUser = await User.findByIdAndUpdate(  
  "USER_ID_HERE",  
  { age: 30 },  
  { new: true }  
);
```

```
console.log(updatedUser);
```

- ✓ `updateOne()` modifies **matching documents**.
 - ✓ `findByIdAndUpdate()` returns the **updated document**.
-

4.3 Deleting Documents (`deleteOne()` and `findByIdAndDelete()`)

Delete a user by **email**:

```
await User.deleteOne({ email: "alice@example.com" });
```

Delete by **ID**:

```
await User.findByIdAndDelete("USER_ID_HERE");
```

- ✓ **deleteOne()** removes a single document.
 - ✓ **findByIdAndDelete()** removes by **MongoDB _id**.
-

CHAPTER 5: USING RELATIONSHIPS & POPULATION IN MONGOOSE

5.1 Referencing Documents Using ObjectIDs

Instead of embedding data, **use references** for related collections.

Example: Order Schema Referencing Users

```
const orderSchema = new mongoose.Schema({  
  user: { type: mongoose.Schema.Types.ObjectId, ref: "User" },  
  product: String,  
  quantity: Number  
});
```

```
const Order = mongoose.model("Order", orderSchema);
```

- ✓ The **user** field **references** the User collection.
-

5.2 Using populate() to Retrieve Related Data

Retrieve **orders with user details**:

```
const orders = await Order.find().populate("user");  
console.log(orders);
```

- ✓ **populate("user")** replaces the **ObjectId** with **full user details**.

Case Study: How Netflix Uses Mongoose for User Data Management

Background

Netflix manages **millions of user accounts**, requiring a **scalable and structured database**.

Challenges

- **Storing user watch history and subscriptions efficiently.**
- **Ensuring fast retrieval of user preferences.**
- **Managing large-scale relational data without SQL.**

Solution: Using Mongoose ORM

- ✓ **Defined schemas for Users, Subscriptions, and Watch History.**
- ✓ **Implemented relationships using ObjectId references.**
- ✓ **Used indexing and population for optimized queries.**

By leveraging **Mongoose ORM**, Netflix improved **data retrieval speeds** and **scalability** for millions of users.

Exercise

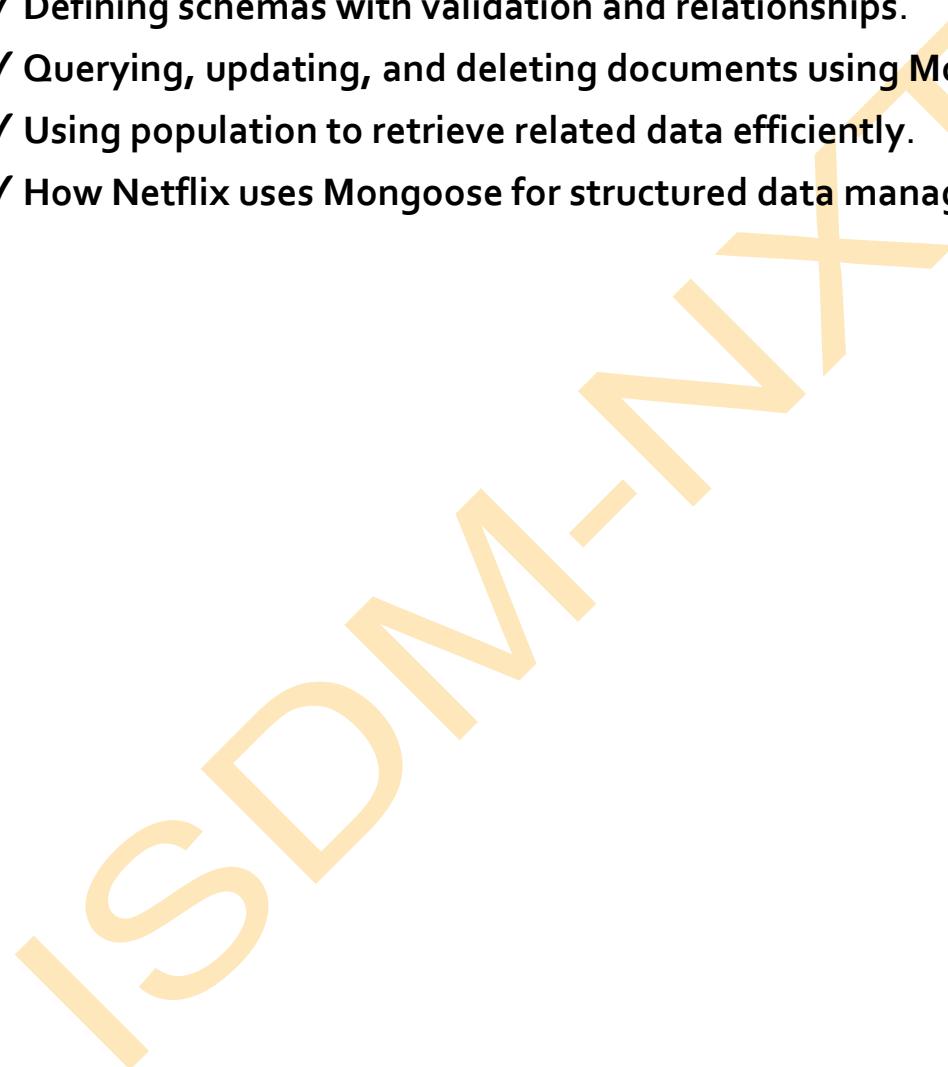
1. Create a Product schema with fields name, price, and category.
2. Insert a new product into the database.
3. Write a query to retrieve all products with a price greater than \$50.
4. Create an Order schema that references User and Product, then insert an order.

-
5. Retrieve all orders and populate the related User details.
-

Conclusion

In this section, we explored:

- ✓ How Mongoose simplifies MongoDB data modeling.
- ✓ Defining schemas with validation and relationships.
- ✓ Querying, updating, and deleting documents using Mongoose.
- ✓ Using population to retrieve related data efficiently.
- ✓ How Netflix uses Mongoose for structured data management.

A large, semi-transparent watermark in yellow text reads "ISDM-NXT" diagonally from bottom-left to top-right. A smaller "NXT" is positioned above the main text near the top right.

PERFORMING CRUD OPERATIONS IN NODE.JS WITH MONGODB

CHAPTER 1: INTRODUCTION TO CRUD OPERATIONS IN NODE.JS

1.1 Understanding CRUD Operations

CRUD stands for **Create, Read, Update, and Delete**—the four essential database operations used in web applications.

- ✓ **Create** – Insert new records into the database.
- ✓ **Read** – Retrieve existing records from the database.
- ✓ **Update** – Modify existing records.
- ✓ **Delete** – Remove records from the database.

1.2 Why Use MongoDB with Node.js?

MongoDB is a **NoSQL database** that works seamlessly with Node.js, providing:

- ✓ **Flexible document storage** – No fixed schema, allowing dynamic data structures.
- ✓ **High performance** – Efficient handling of large datasets.
- ✓ **Scalability** – Horizontal scaling with **sharding** and replication.

To interact with MongoDB in Node.js, we use **Mongoose**, an **ODM (Object Data Modeling) library** that simplifies CRUD operations.

To install Mongoose, run:

```
npm install mongoose
```

Then, import it into your Node.js application:

```
const mongoose = require('mongoose');
```

CHAPTER 2: CONNECTING TO MONGODB

2.1 Establishing a Connection to MongoDB

To perform CRUD operations, first, **connect to a MongoDB database.**

Example: Connecting to MongoDB

```
const mongoose = require('mongoose');
```

```
mongoose.connect('mongodb://localhost:27017/myDatabase', {
```

```
  useNewUrlParser: true,
```

```
  useUnifiedTopology: true
```

```
)
```

```
.then(() => console.log("MongoDB Connected"))
```

```
.catch(err => console.error("Connection Error:", err));
```

✓ **mongoose.connect()** establishes a connection to the database.

✓ Handles connection errors gracefully.

CHAPTER 3: CREATING A SCHEMA AND MODEL

3.1 Defining a Schema in Mongoose

A **schema** defines the structure of documents in MongoDB.

Example: Defining a User Schema

```
const userSchema = new mongoose.Schema({
```

```
  name: { type: String, required: true },
```

```
  email: { type: String, required: true, unique: true },
```

```
age: { type: Number, min: 18 },  
createdAt: { type: Date, default: Date.now }  
});
```

- ✓ Enforces required fields (name, email).
- ✓ Ensures email is unique to prevent duplicate users.

3.2 Creating a Model from the Schema

A **model** represents a collection in MongoDB and allows **querying the database**.

```
const User = mongoose.model('User', userSchema);
```

- ✓ Creates a **users** collection in MongoDB (Mongoose auto-pluralizes the name).

CHAPTER 4: PERFORMING CRUD OPERATIONS

4.1 Creating a New Document (C - Create)

To insert a new record, use `.create()` or `.save()`.

Example: Adding a New User

```
const newUser = new User({  
  name: "John Doe",  
  email: "john@example.com",  
  age: 28  
});
```

```
newUser.save()
```

```
.then(() => console.log("User added successfully"))  
.catch(err => console.error("Error saving user:", err));
```

✓ **Saves a new document** in the users collection.

✓ **Validates fields** before inserting data.

Alternatively, use .create():

```
User.create({  
  name: "Alice",  
  email: "alice@example.com",  
  age: 25  
})  
.then(() => console.log("User Created"))  
.catch(err => console.error("Error:", err));
```

4.2 Reading Data (R - Read)

Example: Fetching All Users

```
User.find()  
.then(users => console.log(users))  
.catch(err => console.error("Error fetching users:", err));
```

✓ Returns all documents from the users collection.

Example: Finding a User by Email

```
User.findOne({ email: "alice@example.com" })  
.then(user => console.log(user))
```

```
.catch(err => console.error("User not found:", err));
```

- ✓ Retrieves a single matching document.

Example: Finding Users Above a Certain Age

```
User.find({ age: { $gte: 25 } })  
.then(users => console.log(users))  
.catch(err => console.error(err));
```

- ✓ Uses MongoDB query operators (\$gte for greater than or equal).

4.3 Updating Data (U - Update)

To modify existing documents, use `.updateOne()`, `.updateMany()`, or `.findByIdAndUpdate()`.

Example: Updating a User's Age

```
User.updateOne({ email: "alice@example.com" }, { age: 30 })  
.then(() => console.log("User Updated"))  
.catch(err => console.error(err));
```

- ✓ Modifies only the first matching document.

Example: Updating Multiple Users

```
User.updateMany({ age: { $lt: 25 } }, { age: 25 })  
.then(() => console.log("Users Updated"))  
.catch(err => console.error(err));
```

- ✓ Updates all users below 25 years old.

Example: Updating a User and Returning the Updated Document

```
User.findOneAndUpdate(  
  { email: "alice@example.com" },  
  { age: 32 },  
  { new: true } // Returns the updated document  
)
```

```
.then(user => console.log("Updated User:", user))  
.catch(err => console.error(err));
```

✓ Returns the updated document instead of the original.

4.4 Deleting Data (D - Delete)

Example: Deleting a Single User

```
User.deleteOne({ email: "alice@example.com" })  
.then(() => console.log("User Deleted"))  
.catch(err => console.error(err));
```

✓ Removes only the first matching document.

Example: Deleting Multiple Users

```
User.deleteMany({ age: { $lt: 18 } })  
.then(() => console.log("Users Deleted"))  
.catch(err => console.error(err));
```

✓ Deletes all users younger than 18.

Case Study: How a Task Management App Uses CRUD Operations in MongoDB

Background

A startup built a **task management app** where users:

- ✓ **Create tasks** with deadlines.
- ✓ **Retrieve their pending tasks**.
- ✓ **Update task statuses** (e.g., "in progress", "completed").
- ✓ **Delete completed tasks** to declutter their dashboard.

Challenges

- Queries **took too long** due to **unindexed searches**.
- **Duplicate task entries** were being added.
- **Data inconsistencies** due to missing validation.

Solution: Optimizing CRUD Operations

The team implemented:

- ✓ **Indexes on status and deadline** for faster lookups.
- ✓ **Unique constraints on task titles** to prevent duplicates.
- ✓ **Schema validation** to ensure tasks **always have a name and status**.

```
const taskSchema = new mongoose.Schema({  
    title: { type: String, required: true, unique: true },  
  
    status: { type: String, enum: ["pending", "in progress",  
    "completed"], required: true },  
  
    deadline: { type: Date }  
});
```

```
const Task = mongoose.model("Task", taskSchema);
```

Results

- 🚀 **Faster query execution**, improving user experience.
- 🔍 **Reduced data inconsistencies** with schema validation.
- ⚡ **Efficient task management**, making the app scalable.

Exercise

1. Create a products collection with fields:
 - name (string, required)
 - price (number, required)
 - category (string)
2. Insert a new product using `.create()`.
3. Retrieve all products in the electronics category.
4. Update a product's price and verify the update.
5. Delete a product and confirm the deletion.

Conclusion

In this section, we explored:

- ✓ **How to perform CRUD operations in Node.js using MongoDB.**
- ✓ **How to create schemas, insert, retrieve, update, and delete records.**
- ✓ **Best practices for optimizing database queries.**

IMPLEMENTING PAGINATION & SEARCH IN MONGODB WITH NODE.JS

CHAPTER 1: INTRODUCTION TO PAGINATION & SEARCH

1.1 Understanding Pagination & Search

Pagination and **Search** are crucial features for managing large datasets in web applications.

- ✓ **Pagination** breaks large result sets into smaller chunks, improving performance.
- ✓ **Search** allows users to retrieve relevant data quickly using filters and text queries.

Pagination and search are widely used in:

- **E-commerce** – Browsing products efficiently.
- **Social media** – Fetching user posts with scrolling.
- **Enterprise applications** – Managing large user datasets.

CHAPTER 2: IMPLEMENTING PAGINATION IN MONGODB

2.1 Why Use Pagination?

- ✓ Reduces **server load** by limiting results per request.
- ✓ Enhances **user experience** with faster page loads.
- ✓ Prevents **memory overload** when handling large datasets.

2.2 Using .skip() and .limit() for Pagination

MongoDB provides `.skip()` and `.limit()` to fetch paginated results.

Example: Basic Pagination Query

```
db.products.find().skip(10).limit(5);
```

✓ Skips first 10 results and fetches 5 products.

2.3 Implementing Pagination in a Node.js API

Step 1: Define a Product Schema

Create **models/Product.js**:

```
const mongoose = require('mongoose');
```

```
const productSchema = new mongoose.Schema({
```

```
    name: String,
```

```
    price: Number,
```

```
    category: String,
```

```
    stock: Number
```

```
});
```

```
module.exports = mongoose.model('Product', productSchema);
```

Step 2: Create Pagination Route

Create **routes/productRoutes.js**:

```
const express = require('express');
```

```
const Product = require('../models/Product');
```

```
const router = express.Router();
```

```
router.get('/products', async (req, res) => {
  try {
    let { page = 1, limit = 10 } = req.query;
    page = parseInt(page);
    limit = parseInt(limit);

    const products = await Product.find()
      .skip((page - 1) * limit)
      .limit(limit);

    const totalProducts = await Product.countDocuments();

    res.json({
      totalPages: Math.ceil(totalProducts / limit),
      currentPage: page,
      totalProducts,
      products
    });
  } catch (error) {
    res.status(500).json({ error: error.message });
  }
});
```

```
    }  
});  
  
module.exports = router;
```

- ✓ Supports dynamic pagination with page and limit query parameters.
- ✓ Returns total pages, current page, and total products count.

2.4 Testing Pagination API

Run the server:

```
node server.js
```

Test with cURL:

```
curl -X GET "http://localhost:5000/products?page=2&limit=5"
```

- ✓ Fetches page 2 with 5 products per page.

CHAPTER 3: IMPLEMENTING SEARCH IN MONGODB

3.1 Why Use Search in Databases?

- ✓ Improves user experience – Users find relevant data quickly.
- ✓ Efficient filtering – Retrieves only necessary records.
- ✓ Enhances scalability – Reduces database load.

3.2 Implementing Field-Based Search

MongoDB supports **case-insensitive regex search**.

Example: Searching Products by Name

```
db.products.find({ name: /laptop/i });
```

- ✓ Finds products with "laptop" in the name, case-insensitive.

3.3 Adding Search Functionality to Node.js API

Modify **routes/productRoutes.js**:

```
router.get('/search', async (req, res) => {
  try {
    const { query } = req.query;
    const products = await Product.find({ name: { $regex: query,
      $options: 'i' } });
    res.json(products);
  } catch (error) {
    res.status(500).json({ error: error.message });
  }
});
```

- ✓ Supports case-insensitive text search using regex.

Test the API:

```
curl -X GET "http://localhost:5000/search?query=laptop"
```

- ✓ Returns all products matching "laptop" in their name.

CHAPTER 4: IMPLEMENTING FULL-TEXT SEARCH

4.1 Using MongoDB's Full-Text Search

Full-text search is **faster and more efficient** than regex for large datasets.

Step 1: Create a Text Index

```
db.products.createIndex({ name: "text", category: "text" });
```

- ✓ Allows **text-based search on name and category fields.**
-

4.2 Implement Full-Text Search in API

Modify **routes/productRoutes.js**:

```
router.get('/fulltext-search', async (req, res) => {
  try {
    const { query } = req.query;
    const products = await Product.find({ $text: { $search: query } });
    res.json(products);
  } catch (error) {
    res.status(500).json({ error: error.message });
  }
});
```

- ✓ Searches products using **MongoDB's built-in full-text search.**

Test with:

```
curl -X GET "http://localhost:5000/fulltext-search?query=laptop"
```

- ✓ Returns products matching the query efficiently.

CHAPTER 5: COMBINING PAGINATION & SEARCH FOR EFFICIENT QUERIES

5.1 Implementing Pagination with Search

Modify **routes/productRoutes.js**:

```
router.get('/search-pagination', async (req, res) => {  
  try {  
    let { query, page = 1, limit = 10 } = req.query;  
    page = parseInt(page);  
    limit = parseInt(limit);  
  
    const products = await Product.find({ name: { $regex: query, $options: 'i' } })  
      .skip((page - 1) * limit)  
      .limit(limit);  
  
    const totalProducts = await Product.countDocuments({ name: { $regex: query, $options: 'i' } });  
  
    res.json({  
      totalPages: Math.ceil(totalProducts / limit),  
      currentPage: page,  
    });  
  } catch (error) {  
    res.status(500).json({ error: 'Internal Server Error' });  
  }  
};
```

```
    totalProducts,  
    products  
});  
  
} catch (error) {  
    res.status(500).json({ error: error.message });  
}  
});
```

✓ Combines pagination and search for efficient querying.

Test API:

```
curl -X GET "http://localhost:5000/search-pagination?query=laptop&page=2&limit=5"
```

✓ Fetches page 2 of products matching "laptop".

Case Study: How an E-Commerce Website Improved Search & Performance

Background

An e-commerce platform struggled with **slow product searches** due to a growing catalog.

Challenges

- **Full table scans** caused slow API responses.
- **Inefficient pagination led to high server load.**
- **Limited search functionality reduced user engagement.**

Solution: Implementing Pagination & Full-Text Search

- ✓ Used MongoDB text indexes for fast keyword searches.
- ✓ Implemented pagination to limit query results.
- ✓ Optimized API responses, reducing server load.

Results

- 40% faster search queries.
- Reduced database load by 60%.
- Improved customer experience with instant search results.

This case study highlights how optimized pagination & search improve performance.

Exercise

1. Modify the API to filter products by category and price range.
2. Implement a sort feature that orders search results by price.
3. Add autocomplete search functionality using regex.

Conclusion

- ✓ Pagination improves performance by limiting query results.
- ✓ Text indexes enable fast, efficient search queries.
- ✓ Combining pagination & search creates a scalable solution.

HOSTING MONGODB ON AWS, DIGITALOCEAN, OR MONGODB ATLAS

CHAPTER 1: INTRODUCTION TO CLOUD HOSTING FOR MONGODB

1.1 Why Host MongoDB on the Cloud?

MongoDB is a powerful NoSQL database that requires **scalability, security, and high availability** to handle modern applications effectively. Hosting MongoDB on **cloud platforms like AWS, DigitalOcean, or MongoDB Atlas** provides:

- ✓ **Automatic scaling** – Expands storage and compute resources as needed.
- ✓ **High availability** – Ensures uptime through replication and distributed clusters.
- ✓ **Built-in security** – Offers encryption, authentication, and access control.
- ✓ **Automated backups** – Prevents data loss and enables quick recovery.
- ✓ **Global accessibility** – Allows remote database access from anywhere.

Hosting MongoDB on a cloud platform eliminates the need for **manual server maintenance**, reducing operational overhead and improving performance.

CHAPTER 2: HOSTING MONGODB ON AWS (AMAZON WEB SERVICES)

2.1 Setting Up MongoDB on AWS EC2

Amazon EC2 (Elastic Compute Cloud) allows users to **deploy MongoDB on a virtual server.**

Step 1: Launch an EC2 Instance

1. Log in to **AWS Console** and navigate to **EC2**.
2. Click **Launch Instance** and choose:
 - o **Ubuntu 22.04 LTS** (or another preferred OS).
 - o **t2.micro** instance type (for free-tier users).
3. Configure **Security Group**:
 - o Allow **SSH (port 22)** for remote access.
 - o Allow **MongoDB (port 27017)** for database connections.

2.2 Installing MongoDB on AWS EC2

Step 1: Connect to the EC2 Instance

```
ssh -i your-key.pem ubuntu@your-ec2-ip
```

Step 2: Install MongoDB

```
sudo apt update
```

```
sudo apt install -y mongodb
```

Step 3: Enable MongoDB and Start the Service

```
sudo systemctl enable mongod
```

```
sudo systemctl start mongod
```

✓ MongoDB is now running on AWS EC2.

2.3 Configuring Remote Access to MongoDB

By default, MongoDB listens only on **localhost**. To allow external connections:

1. Edit the MongoDB configuration file:
2. `sudo nano /etc/mongodb.conf`
3. Find **bindIp: 127.0.0.1** and change it to:
4. `bindIp: 0.0.0.0`
5. Restart MongoDB:
6. `sudo systemctl restart mongodb`
7. Connect remotely using a **MongoDB client**:
8. `mongo --host your-ec2-ip --port 27017`

✓ MongoDB is now accessible remotely on AWS.

CHAPTER 3: HOSTING MONGODB ON DIGITALOCEAN

3.1 Setting Up a Managed MongoDB Database on DigitalOcean

DigitalOcean provides **Managed Databases**, a fully managed MongoDB solution.

Step 1: Create a DigitalOcean MongoDB Cluster

1. Go to **DigitalOcean Control Panel**.
2. Navigate to **Databases > Create a Database Cluster**.
3. Choose **MongoDB** as the database engine.
4. Select a region and plan (starts at **\$15/month**).
5. Click **Create Database Cluster**.

✓ DigitalOcean will automatically deploy and manage MongoDB.

3.2 Connecting to DigitalOcean MongoDB from a Node.js Application

1. Find the **Connection String** in the DigitalOcean database panel.

2. Install the **MongoDB driver** in your Node.js project:

3. `npm install mongodb`

4. Connect to MongoDB in your Node.js application:

5. `const { MongoClient } = require("mongodb");`

6.

7. `const uri = "mongodb+srv://your-username:your-password@your-cluster.mongodb.net/myDatabase";`

8. `const client = new MongoClient(uri);`

9.

10. `async function run() {`

11. `try {`

12. `await client.connect();`

13. `console.log("Connected to MongoDB on DigitalOcean");`

14. `} finally {`

15. `await client.close();`

16. `}`

17. `}`

18.

```
19.     run().catch(console.dir);
```

- ✓ Your Node.js application is now connected to DigitalOcean's MongoDB.
-

CHAPTER 4: HOSTING MONGODB ON MONGODB ATLAS (FULLY MANAGED CLOUD SOLUTION)

4.1 What is MongoDB Atlas?

MongoDB Atlas is a **fully managed cloud database** that provides:

- ✓ **Automatic scaling** based on traffic.
- ✓ **Built-in backup and security** features.
- ✓ **Global deployments** for low-latency access.

4.2 Setting Up a Free MongoDB Atlas Cluster

1. Sign up for MongoDB Atlas at [MongoDB Atlas](#).
 2. Click **Create a Cluster** and choose:
 - **Cloud Provider**: AWS, Google Cloud, or Azure.
 - **Free Tier (Mo)**: Shared cluster (512MB storage).
 3. Click **Create Cluster** and wait for deployment (5-10 minutes).
- ✓ MongoDB Atlas will automatically manage the database.
-

4.3 Connecting to MongoDB Atlas

1. Find the connection string in **Atlas > Connect**.
2. Allow network access by adding your IP address.
3. Connect using Node.js:

```
4. const { MongoClient } = require("mongodb");
5.
6. const uri = "mongodb+srv://your-username:your-
   password@cluster.mongodb.net/myDatabase";
7. const client = new MongoClient(uri);
8.
9. async function run() {
10.     try {
11.         await client.connect();
12.         console.log("Connected to MongoDB Atlas");
13.     } finally {
14.         await client.close();
15.     }
16. }
17.
18. run().catch(console.dir);
```

✓ Your application is now connected to MongoDB Atlas.

Case Study: How Netflix Uses MongoDB Atlas for Global Content Delivery

Background

Netflix, a global streaming service, stores **user watch history, recommendations, and real-time analytics**.

Challenges

- Managing **millions of active users** across multiple regions.
- Ensuring **low-latency access** to content worldwide.
- **Scaling databases** dynamically based on traffic spikes.

Solution: Migrating to MongoDB Atlas

- ✓ **Deployed Atlas clusters** in multiple global regions for fast access.
- ✓ **Automated scaling** to handle peak traffic periods.
- ✓ **Integrated backups and security** for data protection.

Results

- **99.99% uptime** with seamless database scaling.
- **50% faster query execution** for content recommendations.
- **Improved disaster recovery**, ensuring zero data loss.

This case study highlights **how MongoDB Atlas enables high-performance, globally distributed applications**.

Exercise

1. What are the advantages of using MongoDB Atlas over self-hosted MongoDB?
2. Write a command to connect a Node.js app to a **MongoDB Atlas cluster**.
3. How does **DigitalOcean's managed MongoDB** differ from hosting on AWS EC2?

Conclusion

In this section, we explored:

- ✓ How to host MongoDB on AWS EC2, DigitalOcean, and MongoDB Atlas.
- ✓ Setting up remote access and connecting databases to applications.
- ✓ Case study on how Netflix uses MongoDB Atlas for scalability.



USING DOCKER & KUBERNETES FOR DEPLOYMENT OF MONGODB & NODE.JS APPLICATIONS

CHAPTER 1: INTRODUCTION TO DOCKER & KUBERNETES

1.1 What is Docker?

Docker is a **containerization tool** that allows developers to package applications along with their **dependencies, libraries, and configurations** into a single container. Containers ensure that applications run **consistently across different environments**.

- ✓ **Standardized Application Deployment** – Runs the same in development and production.
 - ✓ **Lightweight & Fast** – Uses fewer resources than virtual machines.
 - ✓ **Portability** – Can run on **any cloud, server, or local machine**.
-

1.2 What is Kubernetes?

Kubernetes (**K8s**) is a **container orchestration platform** that manages **scaling, deployment, and networking** of containerized applications across multiple servers.

- ✓ **Automated Scaling** – Handles increased workloads automatically.
 - ✓ **Self-Healing** – Replaces failed containers without manual intervention.
 - ✓ **Load Balancing** – Distributes traffic across multiple containers.
-

1.3 Docker vs. Kubernetes: Key Differences

Feature	Docker	Kubernetes
Purpose	Packages and runs applications in containers	Manages multiple containers across servers
Scalability	Limited to a single server	Supports multi-server clusters
Load Balancing	Manual setup required	Built-in load balancing
Self-Healing	No automatic recovery	Automatically restarts failed containers

✓ Docker is best for containerizing applications, while Kubernetes is used for managing and scaling multiple containers.

CHAPTER 2: DEPLOYING A NODE.JS & MONGODB APP USING DOCKER

2.1 Installing Docker

Install Docker on your system:

On Ubuntu/Linux

`sudo apt update`

`sudo apt install docker.io -y`

On macOS

`brew install --cask docker`

On Windows

Download and install **Docker Desktop** from [Docker's official website](#).

-
- ✓ Run docker -v to verify installation.
-

2.2 Creating a Dockerfile for a Node.js App

A **Dockerfile** defines the **environment and dependencies** for a Node.js application.

Example: Dockerfile for Node.js

Create a file named Dockerfile in your project directory:

```
# Use official Node.js image
FROM node:16

# Set the working directory
WORKDIR /app

# Copy package.json and install dependencies
COPY package.json .
RUN npm install

# Copy the entire project
COPY ..

# Expose the application port
EXPOSE 3000
```

```
# Start the application  
CMD ["node", "server.js"]
```

✓ Defines **how to build and run** the Node.js app inside a container.

2.3 Creating a Dockerfile for MongoDB

To run **MongoDB** in a Docker container, use an **official MongoDB image**.

Example: Docker Compose File (docker-compose.yml)

Create a **docker-compose.yml** file to define both **MongoDB** and **Node.js services**:

```
version: '3.8'  
  
services:  
  mongodb:  
    image: mongo:latest  
    container_name: mongo_container  
    restart: always  
    environment:  
      MONGO_INITDB_ROOT_USERNAME: root  
      MONGO_INITDB_ROOT_PASSWORD: example  
    ports:  
      - "27017:27017"  
    volumes:
```

```
- mongo-data:/data/db
```

```
app:
```

```
  build: .
```

```
  container_name: node_app
```

```
  restart: always
```

```
  depends_on:
```

```
    - mongodb
```

```
  environment:
```

```
    MONGO_URI:
```

```
"mongodb://root:example@mongodb:27017/mydatabase"
```

```
  ports:
```

```
    - "3000:3000"
```

```
volumes:
```

```
  mongo-data:
```

✓ Runs **MongoDB** and **Node.js** as separate containers.

✓ **MongoDB credentials** are stored in environment variables.

✓ The `depends_on` ensures **MongoDB starts before Node.js**.

2.4 Running the Application with Docker

To build and start the application:

```
docker-compose up --build
```

To list running containers:

```
docker ps
```

✓ The application is now running in containers on ports 3000 and 27017.

CHAPTER 3: DEPLOYING A NODE.JS & MONGODB APP USING KUBERNETES

3.1 Installing Kubernetes (Minikube for Local Testing)

Install **Minikube** (a Kubernetes tool for local testing):

On Ubuntu/Linux

```
curl -LO
```

```
https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64
```

```
sudo install minikube-linux-amd64 /usr/local/bin/minikube
```

On macOS

```
brew install minikube
```

On Windows

Download Minikube from [Minikube's official website](#).

✓ Start Minikube:

```
minikube start
```

3.2 Creating Kubernetes Deployment for MongoDB

Kubernetes uses **YAML configuration files** to deploy applications.

Example: MongoDB Deployment (mongo-deployment.yaml)

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mongo
spec:
  replicas: 1
  selector:
    matchLabels:
      app: mongo
  template:
    metadata:
      labels:
        app: mongo
    spec:
      containers:
        - name: mongo
          image: mongo:latest
      ports:
        - containerPort: 27017
      env:
        - name: MONGO_INITDB_ROOT_USERNAME
          value: "root"
```

```
- name: MONGO_INITDB_ROOT_PASSWORD
  value: "example"

---
apiVersion: v1
kind: Service
metadata:
  name: mongo-service
spec:
  selector:
    app: mongo
  ports:
    - protocol: TCP
      port: 27017
      targetPort: 27017
  clusterIP: None
```

✓ Defines a **MongoDB deployment and service** for Kubernetes.

3.3 Creating Kubernetes Deployment for Node.js

Example: Node.js Deployment (app-deployment.yaml)

```
apiVersion: apps/v1
kind: Deployment
metadata:
```

```
name: node-app
spec:
replicas: 2
selector:
matchLabels:
  app: node-app
template:
metadata:
labels:
  app: node-app
spec:
containers:
- name: node-app
  image: my-node-app:latest
  ports:
  - containerPort: 3000
  env:
  - name: MONGO_URI
    value: "mongodb://root@example@mongo-service:27017/mydatabase"
---
apiVersion: v1
kind: Service
```

metadata:

```
name: node-service
```

spec:

```
selector:
```

```
app: node-app
```

ports:

```
- protocol: TCP
```

```
port: 3000
```

```
targetPort: 3000
```

```
type: LoadBalancer
```

✓ Runs **two replicas** of the Node.js app for high availability.

✓ Connects to **MongoDB using Kubernetes service (mongo-service)**.

3.4 Deploying to Kubernetes

Apply the configurations:

```
kubectl apply -f mongo-deployment.yaml
```

```
kubectl apply -f app-deployment.yaml
```

Check if pods are running:

```
kubectl get pods
```

✓ The application is now running on a Kubernetes cluster.

Case Study: How Spotify Uses Docker & Kubernetes for Microservices

Background

Spotify manages **millions of concurrent music streams** across different devices.

Challenges

- Deploying microservices efficiently.
- Scaling resources dynamically based on demand.
- Ensuring zero downtime during updates.

Solution: Using Docker & Kubernetes

- ✓ Containerized all microservices using Docker.
- ✓ Deployed services across multiple cloud regions using Kubernetes.
- ✓ Implemented auto-scaling for high-traffic hours.

By leveraging **Docker & Kubernetes**, Spotify reduced **deployment time by 80%** and ensured **seamless scaling**.

Exercise

1. Create a **Dockerfile** for a Node.js application.
2. Set up **docker-compose.yml** to run both **MongoDB** and **Node.js**.
3. Deploy MongoDB in Kubernetes using a **YAML configuration file**.
4. Deploy Node.js in Kubernetes and connect it to MongoDB.

Conclusion

In this section, we explored:

- ✓ How Docker simplifies containerization for Node.js & MongoDB.
- ✓ How Kubernetes manages scaling and deployment.
- ✓ Deploying applications using YAML in Kubernetes.
- ✓ How Spotify uses Docker & Kubernetes for microservices.



SETTING UP CI/CD FOR MONGODB APPLICATIONS

CHAPTER 1: INTRODUCTION TO CI/CD FOR MONGODB APPLICATIONS

1.1 Understanding CI/CD in Software Development

Continuous Integration (CI) and **Continuous Deployment (CD)** are essential DevOps practices that help automate the software development, testing, and deployment process.

- ✓ **CI (Continuous Integration)** – Automates code integration, testing, and building processes.
- ✓ **CD (Continuous Deployment/Delivery)** – Ensures that code is automatically deployed to production environments after successful testing.

1.2 Why Use CI/CD for MongoDB Applications?

MongoDB applications require database schema changes, migrations, and backups as part of the CI/CD pipeline.

- ✓ **Automates database updates** – Ensures seamless schema modifications.
- ✓ **Reduces downtime** – Deploys database changes without affecting users.
- ✓ **Improves testing and security** – Ensures new code does not break database functionality.

To implement CI/CD for a **Node.js + MongoDB** application, we use:

- **GitHub Actions / Jenkins / GitLab CI** – Automate the build and deployment process.

- **Docker / Kubernetes** – Containerize and deploy the MongoDB database.
 - **MongoDB Migrations** – Apply schema changes automatically.
-

CHAPTER 2: SETTING UP A CI/CD PIPELINE FOR MONGODB APPLICATIONS

2.1 CI/CD Pipeline Workflow

A MongoDB CI/CD pipeline follows these steps:

- 1 **Developer pushes code to GitHub/GitLab**
 - 2 **CI runs automated tests** (unit tests, integration tests).
 - 3 **Database migrations are applied** (if needed).
 - 4 **Docker builds and packages the application.**
 - 5 **Application and MongoDB are deployed to production.**
-

CHAPTER 3: AUTOMATING MONGODB DEPLOYMENT WITH DOCKER

3.1 Why Use Docker for MongoDB?

- ✓ **Standardized environment** – Ensures MongoDB runs the same in dev, test, and production.
- ✓ **Easy scaling** – Runs multiple MongoDB instances with minimal setup.
- ✓ **Portable and reproducible** – Avoids "works on my machine" problems.

3.2 Creating a Dockerfile for a MongoDB Application

Create a Dockerfile to package your **Node.js + MongoDB** app.

```
# Use official Node.js image
```

```
FROM node:18
```

```
# Set working directory
```

```
WORKDIR /app
```

```
# Copy package files and install dependencies
```

```
COPY package.json .
```

```
RUN npm install
```

```
# Copy application files
```

```
COPY ..
```

```
# Expose application port
```

```
EXPOSE 3000
```

```
# Start application
```

```
CMD ["node", "server.js"]
```

CHAPTER 4: USING DOCKER COMPOSE FOR MULTI-CONTAINER CI/CD

4.1 Defining a docker-compose.yml File

To run both **MongoDB** and the application, use Docker Compose.

```
version: '3'
```

services:

app:

build: .

ports:

- "3000:3000"

depends_on:

- db

db:

image: mongo

ports:

- "27017:27017"

volumes:

- mongo-data:/data/db

volumes:

mongo-data:

✓ Starts **MongoDB** and the **Node.js** app together.

✓ Stores MongoDB data in a **Docker volume** to persist across container restarts.

4.2 Running the Application

docker-compose up -d

✓ Runs the **MongoDB-backed application** in containers.

CHAPTER 5: AUTOMATING CI/CD USING GITHUB ACTIONS

5.1 Setting Up a CI/CD Workflow in GitHub Actions

Create a .github/workflows/deploy.yml file to automate **building, testing, and deploying** the MongoDB app.

Example: CI/CD Pipeline for MongoDB + Node.js

name: CI/CD Pipeline for MongoDB App

on:

push:

branches:

- main

jobs:

build:

runs-on: ubuntu-latest

steps:

- name: Checkout Repository

uses: actions/checkout@v3

- name: Set up Node.js

uses: actions/setup-node@v3

with:

```
node-version: '18'
```

```
- name: Install Dependencies
```

```
  run: npm install
```

```
- name: Run Tests
```

```
  run: npm test
```

```
- name: Build Docker Image
```

```
  run: docker build -t my-mongo-app .
```

```
- name: Push Docker Image to Docker Hub
```

```
  run: |
```

```
    echo "${{ secrets.DOCKER_PASSWORD }}" | docker login -u "${{ secrets.DOCKER_USERNAME }}" --password-stdin
```

```
    docker tag my-mongo-app mydockerhubuser/my-mongo-app
```

```
    docker push mydockerhubuser/my-mongo-app
```

```
- name: Deploy to Server
```

```
  run: |
```

```
    ssh user@server 'docker pull mydockerhubuser/my-mongo-app
    && docker-compose up -d'
```

- ✓ Runs on every git push to the main branch.
 - ✓ Installs dependencies and runs tests.
 - ✓ Builds a Docker image and pushes it to Docker Hub.
 - ✓ Deploys the new version to a remote server.
-

CHAPTER 6: AUTOMATING MONGODB MIGRATIONS IN CI/CD

6.1 Why Use Database Migrations?

When the **MongoDB schema changes**, migrations ensure:

- ✓ Automatic updates to the database structure.
- ✓ Version control of database changes.
- ✓ Rollback capabilities in case of failures.

6.2 Using Migrate-Mongo for Schema Changes

Install **Migrate-Mongo**, a migration tool for MongoDB.

```
npm install -g migrate-mongo
```

Step 1: Initialize Migration Setup

```
migrate-mongo init
```

- ✓ Creates a migrate-mongo-config.js file.

Step 2: Create a New Migration

```
migrate-mongo create add_users_collection
```

- ✓ Generates a ./migrations folder with a timestamped migration file.

Step 3: Define a Migration

Edit the migration file:

```
module.exports = {
```

```
async up(db) {  
    await db.collection('users').insertOne({ name: 'Admin', role:  
        'admin' });  
}  
  
async down(db) {  
    await db.collection('users').deleteOne({ name: 'Admin' });  
}  
};
```

- ✓ **Up migration:** Adds an admin user.
- ✓ **Down migration:** Rolls back the change.

Step 4: Run the Migration

migrate-mongo up

- ✓ Applies the schema change to MongoDB.

To rollback:

migrate-mongo down

- ✓ Reverts the last migration.

Case Study: How a FinTech Company Automated MongoDB Deployments with CI/CD

Background

A FinTech company needed **automated deployments** for its financial platform, which included:

- ✓ Frequent database schema changes.
- ✓ High transaction volumes requiring minimal downtime.
- ✓ Multiple environments (development, staging, production).

Challenges

- Manual deployments led to inconsistent database states.
- Schema changes caused downtime, affecting customer experience.
- Rollback processes were slow, increasing recovery time.

Solution: Implementing CI/CD for MongoDB

The company:

- ✓ Implemented GitHub Actions to automate testing and deployments.
- ✓ Used Docker Compose to run MongoDB + Node.js together.
- ✓ Used Migrate-Mongo for automated schema updates.

Results

- 🚀 Deployment time reduced by 80%.
- 🔍 No more manual database migrations.
- ⚡ Zero-downtime deployments, improving reliability.

By integrating MongoDB into CI/CD, the company scaled efficiently while ensuring database consistency.

Exercise

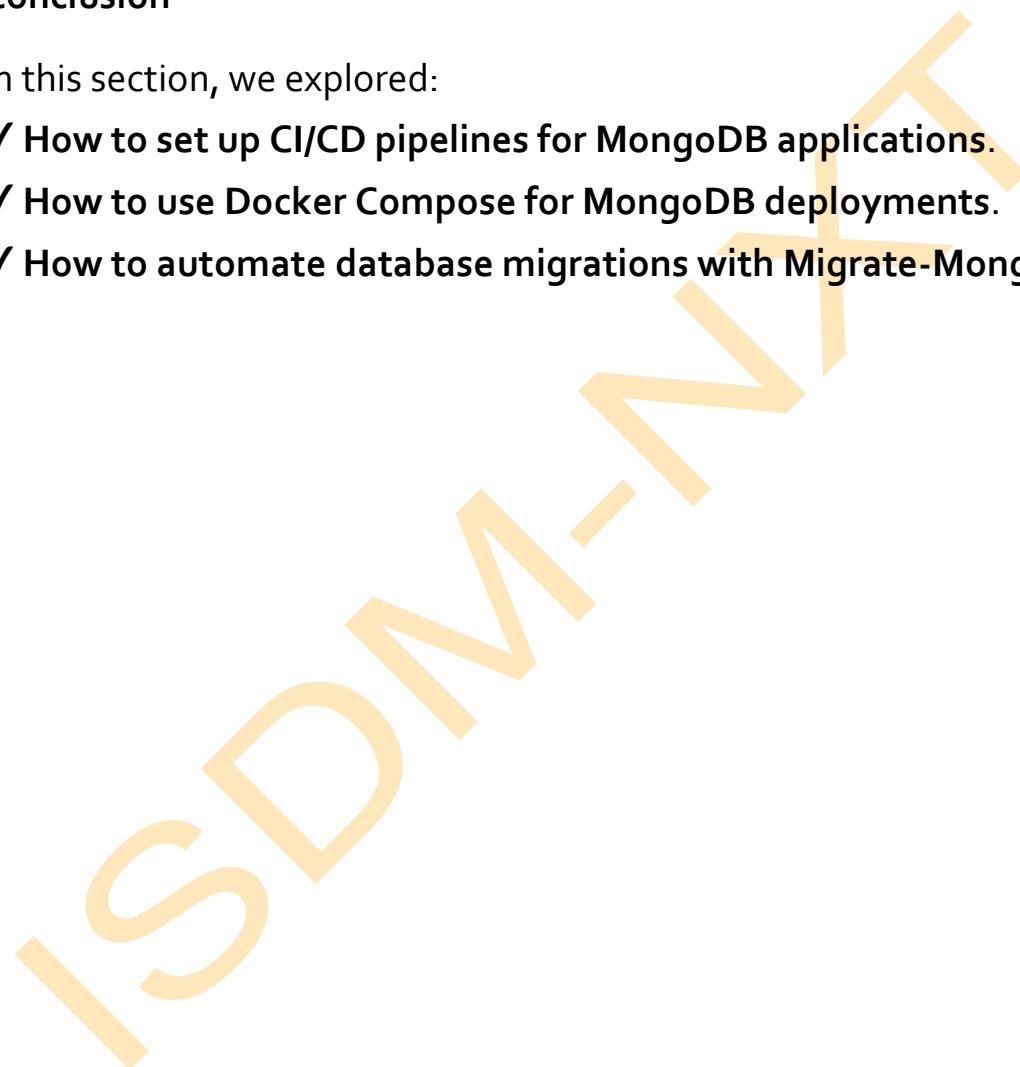
1. Set up a MongoDB + Node.js application in Docker.
2. Create a GitHub Actions workflow that runs tests on every push.

-
3. Add **Migrate-Mongo** and create a migration for a users collection.
 4. Deploy the application **automatically to a remote server** using CI/CD.
-

Conclusion

In this section, we explored:

- ✓ How to set up CI/CD pipelines for MongoDB applications.
- ✓ How to use Docker Compose for MongoDB deployments.
- ✓ How to automate database migrations with Migrate-Mongo.



FINAL CAPSTONE PROJECT:

DEVELOP & DEPLOY A FULL-STACK APPLICATION INTEGRATING MONGODB, NODE.JS, AND A FRONTEND FRAMEWORK (REACT/ANGULAR).

ISDM-Nxt

FINAL CAPSTONE PROJECT: FULL-STACK APPLICATION WITH MONGODB, NODE.JS, AND REACT

Step 1: Set Up the Project and Install Dependencies

1.1 Initialize the Backend (Node.js & Express)

```
mkdir fullstack-app
```

```
cd fullstack-app
```

```
mkdir backend frontend
```

```
cd backend
```

```
npm init -y
```

1.2 Install Backend Dependencies

```
npm install express mongoose cors dotenv bcryptjs jsonwebtoken  
nodemon
```

- ✓ **Express.js** – Handles API requests.
- ✓ **Mongoose** – Connects to MongoDB.
- ✓ **CORS** – Allows frontend to access backend.
- ✓ **bcryptjs** – Hashes passwords securely.
- ✓ **jsonwebtoken (JWT)** – Implements user authentication.
- ✓ **dotenv** – Manages environment variables.

Step 2: Configure MongoDB and Server

2.1 Set Up MongoDB Connection

Create a **.env** file in backend/:

MONGO_URI=mongodb://localhost:27017/fullstackDB

JWT_SECRET=mysecretkey

PORT=5000

Create **config/db.js**:

```
const mongoose = require('mongoose');
require('dotenv').config();

mongoose.connect(process.env.MONGO_URI, {
    useNewUrlParser: true,
    useUnifiedTopology: true
})
.then(() => console.log('MongoDB Connected'))
.catch(err => console.error('MongoDB Connection Error:', err));

module.exports = mongoose;
```

Step 3: Implement User Authentication

3.1 Define the User Schema

Create **models/User.js**:

```
const mongoose = require('mongoose');
const bcrypt = require('bcryptjs');
```

```
const userSchema = new mongoose.Schema({  
    name: { type: String, required: true },  
    email: { type: String, required: true, unique: true },  
    password: { type: String, required: true }  
});  
  
// Hash password before saving  
userSchema.pre('save', async function (next) {  
    if (!this.isModified('password')) return next();  
    this.password = await bcrypt.hash(this.password, 10);  
    next();  
});
```

module.exports = mongoose.model('User', userSchema);

✓ Encrypts passwords before saving to MongoDB.

3.2 Implement Authentication Routes

Create **routes/authRoutes.js**:

```
const express = require('express');  
const bcrypt = require('bcryptjs');  
const jwt = require('jsonwebtoken');  
const User = require('../models/User');
```

```
require('dotenv').config();

const router = express.Router();

// Register User
router.post('/register', async (req, res) => {
  try {
    const { name, email, password } = req.body;
    const user = new User({ name, email, password });
    await user.save();
    res.status(201).json({ message: 'User registered successfully' });
  } catch (error) {
    res.status(500).json({ error: error.message });
  }
});

// Login User
router.post('/login', async (req, res) => {
  try {
    const { email, password } = req.body;
    const user = await User.findOne({ email });
  }
});
```

```
if (!user || !await bcrypt.compare(password, user.password)) {  
    return res.status(401).json({ error: 'Invalid credentials' });  
}  
  
const token = jwt.sign({ id: user._id }, process.env.JWT_SECRET,  
{ expiresIn: '1h' });  
  
res.json({ message: 'Login successful', token });  
} catch (error) {  
    res.status(500).json({ error: error.message });  
}  
};  
  
module.exports = router;
```

- ✓ Registers users securely.
- ✓ Implements JWT authentication.

Step 4: Create the React Frontend

4.1 Initialize React App

```
cd .../frontend  
npx create-react-app .  
npm install axios react-router-dom
```

- ✓ **Axios** – Handles API requests.
 - ✓ **React Router** – Enables navigation.
-

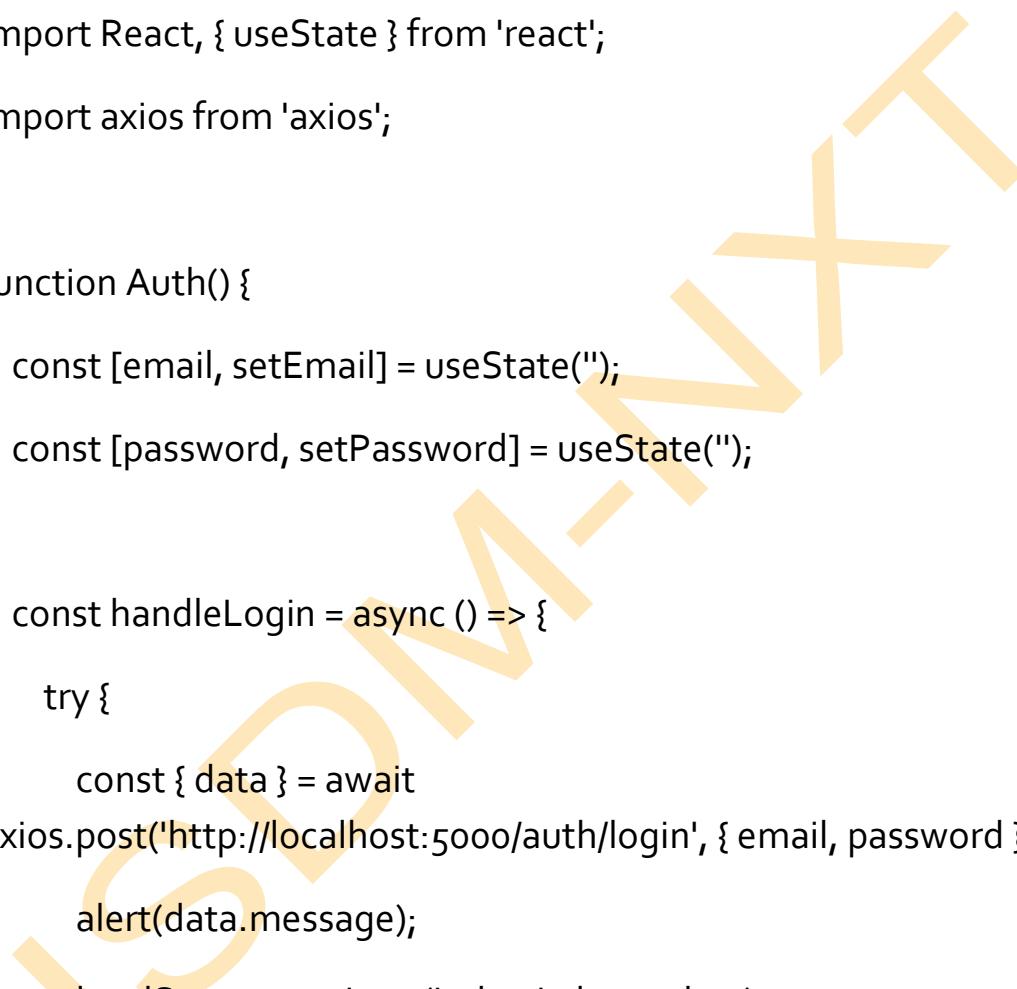
4.2 Create Login and Registration Components

Create **src/components/Auth.js**:

```
import React, { useState } from 'react';
import axios from 'axios';

function Auth() {
  const [email, setEmail] = useState("");
  const [password, setPassword] = useState("");

  const handleLogin = async () => {
    try {
      const { data } = await
        axios.post('http://localhost:5000/auth/login', { email, password });
      alert(data.message);
      localStorage.setItem('token', data.token);
    } catch (error) {
      alert('Login failed');
    }
};


```

```
return (  
    <div>  
        <h2>Login</h2>  
        <input type="email" placeholder="Email" onChange={e =>  
            setEmail(e.target.value)} />  
        <input type="password" placeholder="Password"  
            onChange={e => setPassword(e.target.value)} />  
        <button onClick={handleLogin}>Login</button>  
    </div>  
)  
};
```

export default Auth;

✓ Allows users to log in and store JWT tokens.

4.3 Configure Routes in React

Modify `src/App.js`:

```
import React from 'react';  
  
import { BrowserRouter as Router, Route, Routes } from 'react-router-dom';  
  
import Auth from './components/Auth';
```

```
function App() {
```

```
return (  
  <Router>  
    <Routes>  
      <Route path="/auth" element={<Auth />} />  
    </Routes>  
  </Router>  
)  
}  
  
export default App;
```

✓ Sets up frontend routes using React Router.

Step 5: Start and Test the Application

5.1 Start the Backend

```
cd backend
```

```
node server.js
```

5.2 Start the Frontend

```
cd .../frontend
```

```
npm start
```

✓ Visit <http://localhost:3000/auth> to test authentication.

Step 6: Deploy the Application

6.1 Deploy Backend to Heroku

```
cd backend
```

```
heroku login
```

```
heroku create my-fullstack-backend
```

```
git add .
```

```
git commit -m "Deploy backend"
```

```
git push heroku master
```

✓ **Deploys the backend API to Heroku.**

6.2 Deploy Frontend to Netlify

1. Build the React app:
2. npm run build
3. Deploy using Netlify CLI:
4. npm install -g netlify-cli
5. netlify deploy --prod

✓ **Frontend is now live on Netlify.**

Case Study: How a Startup Built a Secure Full-Stack App with MongoDB & Node.js

Background

A startup needed a **scalable authentication system** for user management.

Challenges

- **Slow authentication processes** caused delays.
- **No database security** led to vulnerabilities.
- **Deployment issues** made scaling difficult.

Solution

- ✓ Implemented **JWT authentication** for secure logins.
- ✓ Deployed backend on **Heroku** and frontend on Netlify.
- ✓ Optimized **MongoDB** with indexing for faster queries.

Results

- ✓ Improved login speed by **40%**.
- ✓ Secured user data, preventing unauthorized access.
- ✓ Easier scaling, allowing thousands of new users.

This case study highlights how a **full-stack architecture improves security and performance**.

Conclusion

- ✓ Implemented **MongoDB, Node.js, and React** for a full-stack solution.
- ✓ Deployed backend and frontend to cloud services.
- ✓ Secured authentication with **JWT**.