



## ISDM (INDEPENDENT SKILL DEVELOPMENT MISSION)



# INTRODUCTION TO TIME SERIES ANALYSIS

- ◆ CHAPTER 1: WHAT IS TIME SERIES ANALYSIS?

### 1.1 Definition

Time series analysis is a statistical technique used to analyze **sequential data points** collected over time. It helps in identifying trends, seasonal patterns, and making predictions based on historical data.

### 1.2 Characteristics of Time Series Data

- **Time-Dependent:** Observations are collected at different time intervals.
- **Ordered Data:** Unlike regular datasets, time series data maintains a chronological order.
- **Patterns and Trends:** Data may show upward/downward trends or seasonal effects.

### 1.3 Examples of Time Series Data

- **Finance:** Stock prices, exchange rates, inflation rates.
- **Weather:** Temperature, rainfall, wind speed.

- **Economics:** GDP growth, unemployment rates.
  - **Healthcare:** Patient counts, disease outbreak trends.
  - **Retail:** Sales data, website traffic.
- 

## ◆ CHAPTER 2: COMPONENTS OF TIME SERIES DATA

### 2.1 Trend

A long-term increase or decrease in the data over time. ✓ **Example:** The increase in global temperatures over decades.

### 2.2 Seasonality

Recurring patterns at fixed intervals (daily, weekly, monthly, yearly). ✓ **Example:** Higher ice cream sales in summer.

### 2.3 Cyclical Patterns

Fluctuations that occur over long time periods due to economic cycles. ✓ **Example:** Business cycles affecting market trends.

### 2.4 Irregular (Residual) Variations

Unpredictable fluctuations due to unforeseen events (e.g., natural disasters, pandemics).

---

## ◆ CHAPTER 3: TYPES OF TIME SERIES MODELS

### 3.1 Deterministic vs. Stochastic Models

- **Deterministic Models:** Predictable changes (e.g., linear trends).
- **Stochastic Models:** Include random variations (e.g., stock prices).

## 3.2 Additive vs. Multiplicative Models

- **Additive Model:**

$$Y(t) = \text{Trend} + \text{Seasonality} + \text{Noise}$$

✓ Used when variations are constant over time.

- **Multiplicative Model:**

$$Y(t) = \text{Trend} \times \text{Seasonality} \times \text{Noise}$$

✓ Used when variations grow over time.

## ◆ CHAPTER 4: DATA PREPROCESSING FOR TIME SERIES ANALYSIS

### 4.1 Handling Missing Data

Missing values can disrupt time series models. Techniques include: ✓

**Forward Fill:** Use the previous value.

**✓ Interpolation:** Estimate missing values.

**✓ Dropping Missing Values:** If few values are missing.

### 4.2 Resampling Data

Changing the time frequency (e.g., daily to weekly) for better analysis.

```
df.resample('M').mean() # Resample daily data into monthly averages
```

### 4.3 Smoothing Techniques

Used to reduce noise: ✓ **Moving Averages**

**✓ Exponential Smoothing**

## ◆ CHAPTER 5: TIME SERIES VISUALIZATION

### 5.1 Line Plot

A basic method to visualize trends.

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
df = pd.read_csv("timeseries_data.csv", parse_dates=["Date"],  
index_col="Date")
```

```
df.plot(figsize=(10, 5), title="Time Series Data")
```

```
plt.show()
```

### 5.2 Rolling Mean (Moving Average)

Used to smooth out short-term fluctuations.

```
df["Rolling_Mean"] = df["Value"].rolling(window=7).mean()
```

```
df[["Value", "Rolling_Mean"]].plot(figsize=(10, 5))
```

```
plt.show()
```

### 5.3 Seasonal Decomposition

Breaks down a time series into trend, seasonality, and residuals.

```
from statsmodels.tsa.seasonal import seasonal_decompose
```

```
result = seasonal_decompose(df["Value"], model="additive")
```

```
result.plot()
```

```
plt.show()
```

## ◆ CHAPTER 6: FORECASTING TECHNIQUES

### 6.1 Naïve Forecasting

Uses the last observed value as the next prediction. ✓ Example:  
Tomorrow's stock price is assumed to be today's stock price.

### 6.2 Moving Averages

Predicts using the average of past observations.

```
df["Forecast"] = df["Value"].rolling(window=5).mean()  
df.plot(figsize=(10,5))  
plt.show()
```

### 6.3 Exponential Smoothing

Gives more weight to recent observations.

```
from statsmodels.tsa.holtwinters import SimpleExpSmoothing
```

```
model = SimpleExpSmoothing(df["Value"]).fit()  
df["Forecast"] = model.fittedvalues  
df.plot(figsize=(10, 5))  
plt.show()
```

### 6.4 ARIMA (AutoRegressive Integrated Moving Average)

A powerful forecasting model.

```
from statsmodels.tsa.arima.model import ARIMA
```

```
model = ARIMA(df["Value"], order=(2,1,2)) # ARIMA(p,d,q)  
model_fit = model.fit()
```

```
df["Forecast"] = model_fit.fittedvalues  
df.plot(figsize=(10,5))  
plt.show()
```

## ◆ CHAPTER 7: PERFORMANCE EVALUATION METRICS

### 7.1 Mean Absolute Error (MAE)

Measures the average absolute difference between actual and predicted values.

$$MAE = \frac{1}{n} \sum |y_i - \hat{y}_i|$$

```
from sklearn.metrics import mean_absolute_error  
  
mae = mean_absolute_error(y_actual, y_predicted)  
  
print(f"MAE: {mae}")
```

### 7.2 Mean Squared Error (MSE)

Squares the errors before averaging.

$$MSE = \frac{1}{n} \sum (y_i - \hat{y}_i)^2$$

```
from sklearn.metrics import mean_squared_error  
  
mse = mean_squared_error(y_actual, y_predicted)  
  
print(f"MSE: {mse}")
```

### 7.3 Root Mean Squared Error (RMSE)

$$RMSE = \sqrt{MSE}$$

```
rmse = np.sqrt(mse)  
print(f"RMSE: {rmse}")
```

## 7.4 Mean Absolute Percentage Error (MAPE)

$$MAPE = \frac{1}{n} \sum \left| \frac{y_i - \hat{y}_i}{y_i} \right| \times 100$$

```
mape = np.mean(np.abs((y_actual - y_predicted) / y_actual)) * 100  
print(f"MAPE: {mape}%")
```

### ◆ CHAPTER 8: REAL-WORLD APPLICATIONS OF TIME SERIES ANALYSIS

#### 8.1 Stock Market Prediction

- ✓ Predicting future stock prices based on historical trends.

#### 8.2 Sales Forecasting

- ✓ Predicting product demand for inventory management.

#### 8.3 Weather Forecasting

- ✓ Predicting temperatures, rainfall, and storms.

#### 8.4 Energy Consumption

- ✓ Forecasting electricity demand to optimize power grids.

#### 8.5 Healthcare

- ✓ Monitoring patient data trends for disease prediction.

◆ CHAPTER 9: SUMMARY

- ✓ **Time Series Analysis** is crucial for understanding sequential data.
- ✓ **Key components** include trend, seasonality, and cyclic patterns.
- ✓ **Visualization techniques** (line plots, decomposition) help understand data.
- ✓ **Forecasting methods** range from simple moving averages to advanced ARIMA models.
- ✓ **Performance metrics** (MAE, MSE, RMSE) evaluate model accuracy.

◆ CHAPTER 10: NEXT STEPS

- Try advanced models like **LSTMs** for time series forecasting.
- Use real-world datasets like stock prices or weather data.
- Experiment with **hyperparameter tuning** for ARIMA models.

# ☑ FORECASTING MODELS: ARIMA, SARIMA & EXPONENTIAL SMOOTHING

## 📌 CHAPTER 1: INTRODUCTION TO FORECASTING MODELS

### 1.1 What is Forecasting?

Forecasting is the process of predicting **future values** based on historical data. It is widely used in **finance, supply chain, stock markets, sales, weather prediction, and economic analysis**.

### 1.2 Why is Forecasting Important?

- ✓ Helps businesses make **data-driven decisions**.
- ✓ Reduces **uncertainty** in planning and resource allocation.
- ✓ Improves **inventory management** and demand prediction.

### 1.3 Common Types of Forecasting Models

1. **Time Series Models** - ARIMA, SARIMA, Exponential Smoothing
2. **Machine Learning Models** - Random Forest, XGBoost, LSTMs
3. **Deep Learning Models** - RNNs, Transformer-based models

This study material focuses on **three traditional time series models**:

1. **ARIMA (AutoRegressive Integrated Moving Average)**
2. **SARIMA (Seasonal ARIMA)**
3. **Exponential Smoothing (ETS)**

## 📌 CHAPTER 2: ARIMA (AUTOREGRESSIVE INTEGRATED MOVING AVERAGE)

### 2.1 What is ARIMA?

**ARIMA (AutoRegressive Integrated Moving Average)** is a powerful forecasting model used for **univariate time series** data.

### 2.2 How ARIMA Works?

ARIMA is defined by three components:

- **AutoRegressive (AR)** - Uses past values to predict future values.
- **Integrated (I)** - Differencing is applied to make data stationary.
- **Moving Average (MA)** - Uses past forecast errors for predictions.

ARIMA is denoted as **ARIMA(p, d, q)**:

- **p (AR order)** - Number of past observations used.
- **d (Differencing order)** - Number of times differencing is applied.
- **q (MA order)** - Number of past forecast errors used.

### 2.3 When to Use ARIMA?

- ✓ When the time series is **non-seasonal**.
- ✓ When data follows a **trend but no clear seasonal pattern**.
- ✓ When forecasting **short-term future values**.

### 2.4 Advantages of ARIMA

- ✓ Works well for **short-term forecasting**.
- ✓ Captures **trends and relationships** effectively.
- ✓ Can be fine-tuned using **p, d, q parameters**.

## 2.5 Disadvantages of ARIMA

- ✗ Does not handle **seasonality** well (use SARIMA instead).
- ✗ Requires **manual parameter tuning**.
- ✗ Sensitive to **outliers and missing values**.

## 2.6 Implementing ARIMA in Python

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.arima.model import ARIMA

# Load dataset (Example: Monthly sales data)
df = pd.read_csv("sales_data.csv", index_col='Date',
parse_dates=True)

# Fit ARIMA model (p=2, d=1, q=2)
model = ARIMA(df['Sales'], order=(2,1,2))
model_fit = model.fit()

# Forecast next 12 periods
forecast = model_fit.forecast(steps=12)
```

```
# Plot results  
plt.figure(figsize=(10,5))  
plt.plot(df['Sales'], label="Actual Sales")  
plt.plot(forecast, label="Forecast", linestyle="dashed")  
plt.legend()  
plt.title("ARIMA Forecast")  
plt.show()
```

## 📌 CHAPTER 3: SARIMA (SEASONAL ARIMA)

### 3.1 What is SARIMA?

**SARIMA (Seasonal ARIMA)** extends ARIMA to handle **seasonality** in time series data.

### 3.2 How SARIMA Works?

SARIMA is denoted as **SARIMA(p, d, q) × (P, D, Q, m)**:

- **p, d, q** - Standard ARIMA components.
- **P, D, Q** - Seasonal AR, Differencing, and MA components.
- **m** - Seasonal period (e.g., 12 for monthly data, 7 for weekly data).

### 3.3 When to Use SARIMA?

- ✓ When the time series exhibits **seasonal patterns** (e.g., sales, temperature, demand).
- ✓ When ARIMA fails due to **seasonal fluctuations**.

### 3.4 Advantages of SARIMA

- ✓ Handles both **trends and seasonality**.
- ✓ Works well for **time series forecasting with seasonal fluctuations**.

### 3.5 Disadvantages of SARIMA

- ✗ Computationally **expensive** for large datasets.
- ✗ Parameter tuning requires **trial and error**.

### 3.6 Implementing SARIMA in Python

```
from statsmodels.tsa.statespace.sarimax import SARIMAX

# Fit SARIMA model (p=2, d=1, q=2) x (P=1, D=1, Q=1, m=12)
sarima_model = SARIMAX(df['Sales'], order=(2,1,2),
seasonal_order=(1,1,1,12))

sarima_fit = sarima_model.fit()

# Forecast next 12 periods
sarima_forecast = sarima_fit.forecast(steps=12)

# Plot results
plt.figure(figsize=(10,5))

plt.plot(df['Sales'], label="Actual Sales")

plt.plot(sarima_forecast, label="SARIMA Forecast",
linestyle="dashed")

plt.legend()

plt.title("SARIMA Forecast")
```

```
plt.show()
```

## 📌 CHAPTER 4: EXPONENTIAL SMOOTHING

### 4.1 What is Exponential Smoothing?

**Exponential Smoothing (ETS)** is a simple but effective method for forecasting by assigning **exponentially decreasing weights** to past observations.

### 4.2 Types of Exponential Smoothing

1. **Simple Exponential Smoothing (SES)** - For data with **no trend or seasonality**.
2. **Holt's Exponential Smoothing** - Handles **trends** in data.
3. **Holt-Winters (Triple) Exponential Smoothing** - Handles both **trends and seasonality**.

### 4.3 When to Use Exponential Smoothing?

- When the time series **lacks strong patterns** (for SES).
- When data exhibits **trend and seasonality** (Holt-Winters).
- When a **simple and fast method** is required.

### 4.4 Advantages of Exponential Smoothing

- ✓ **Easy to implement and computationally efficient.**
- ✓ Works well for **short-term forecasting**.
- ✓ Can handle **trend and seasonal components** (Holt-Winters).

### 4.5 Disadvantages of Exponential Smoothing

- ✗ Does not handle **complex relationships** like ARIMA/SARIMA.
- ✗ Requires **manual tuning of parameters**.

## 4.6 Implementing Exponential Smoothing in Python

```
from statsmodels.tsa.holtwinters import ExponentialSmoothing

# Fit Holt-Winters Exponential Smoothing

hw_model = ExponentialSmoothing(df['Sales'], trend='add',
seasonal='add', seasonal_periods=12)

hw_fit = hw_model.fit()

# Forecast next 12 periods

hw_forecast = hw_fit.forecast(steps=12)

# Plot results

plt.figure(figsize=(10,5))

plt.plot(df['Sales'], label="Actual Sales")

plt.plot(hw_forecast, label="Holt-Winters Forecast",
linestyle="dashed")

plt.legend()

plt.title("Exponential Smoothing Forecast")

plt.show()
```

## 📌 CHAPTER 5: COMPARISON OF FORECASTING MODELS

Feature	ARIMA	SARIMA	Exponential Smoothing
<b>Best For</b>	Trend-based data	Seasonal data	Simple patterns
<b>Handles Seasonality?</b>	✗ No	✓ Yes	✓ Yes (Holt-Winters)
<b>Computational Cost</b>	Medium	High	Low
<b>Short-Term Forecasting</b>	✓ Yes	✓ Yes	✓ Yes
<b>Long-Term Forecasting</b>	✗ No	✓ Yes	✗ No

## 📌 CHAPTER 6: SUMMARY & CONCLUSION

### 6.1 Key Takeaways

- ✓ ARIMA is best for **trend-based, non-seasonal** data.
- ✓ SARIMA is best for **seasonal time series**.
- ✓ Exponential Smoothing is best for **simple short-term forecasting**.

### 6.2 Next Steps

- Try **auto-tuning parameters** using **Auto-ARIMA**.
- Apply **deep learning-based forecasting (LSTMs, Transformer models)**.

- Experiment with **real-world datasets** (e.g., stock prices, sales data).

ISDM-NxT



# LSTM FOR TIME SERIES PREDICTION

## 📌 CHAPTER 1: INTRODUCTION TO TIME SERIES FORECASTING

### 1.1 What is Time Series Prediction?

Time series prediction involves forecasting future values based on past observations. A **time series** is a sequence of data points recorded at regular time intervals (e.g., daily stock prices, temperature readings, sales data).

### 1.2 Why Use LSTMs for Time Series?

Traditional machine learning models (e.g., ARIMA, Linear Regression) fail to capture **long-term dependencies**. **Long Short-Term Memory (LSTM) networks** are designed to handle **sequential data**, making them **ideal for time series forecasting**.

### 1.3 Applications of Time Series Forecasting

- Stock Market Prediction
- Weather Forecasting
- Sales & Demand Forecasting
- Energy Consumption Prediction
- Healthcare Monitoring (ECG, BP Trends)



## CHAPTER 2: UNDERSTANDING LSTMs

### 2.1 What is an LSTM?

**LSTM (Long Short-Term Memory)** is a special type of **Recurrent Neural Network (RNN)** that solves the **vanishing gradient** problem, allowing it to learn long-term dependencies.

## 2.2 LSTM Cell Structure

An LSTM cell consists of:

1. **Forget Gate**: Decides which past information to forget.
2. **Input Gate**: Updates the cell state with new information.
3. **Cell State**: Stores long-term memory.
4. **Output Gate**: Produces the final output.

## 2.3 Why LSTMs for Time Series?

- Can learn patterns from sequential data.
- Handles **long-term dependencies** better than RNNs.
- Suitable for **multivariate forecasting** (multiple features).

## CHAPTER 3: PREPARING DATA FOR LSTM

### 3.1 Steps in Time Series Forecasting

1. **Load Dataset** 
2. **Preprocess Data** (scaling, handling missing values)
3. **Convert Data into Sequences** (LSTMs need sequential inputs)
4. **Train the LSTM Model** 
5. **Make Predictions & Evaluate Performance** 

## CHAPTER 4: IMPLEMENTING LSTM FOR TIME SERIES PREDICTION

We will use **Python, TensorFlow, and Keras** to build an LSTM model for predicting stock prices.

#### 4.1 Install Required Libraries

```
!pip install numpy pandas matplotlib scikit-learn tensorflow
```

#### 4.2 Import Libraries

```
import numpy as np  
  
import pandas as pd  
  
import matplotlib.pyplot as plt  
  
from sklearn.preprocessing import MinMaxScaler  
  
from tensorflow.keras.models import Sequential  
  
from tensorflow.keras.layers import LSTM, Dense
```

#### 4.3 Load & Visualize Dataset

```
# Load dataset (Example: Apple stock prices)  
  
df = pd.read_csv('AAPL.csv', usecols=['Date', 'Close'])  
  
df['Date'] = pd.to_datetime(df['Date'])  
  
df.set_index('Date', inplace=True)  
  
# Plot closing prices  
  
plt.figure(figsize=(10,5))  
  
plt.plot(df['Close'], label="Closing Price")  
  
plt.xlabel("Date")  
  
plt.ylabel("Price")  
  
plt.title("Stock Price Over Time")
```

```
plt.legend()  
plt.show()
```

## 📌 CHAPTER 5: PREPROCESSING FOR LSTM

### 5.1 Scale the Data

LSTMs perform better with **scaled data**.

```
scaler = MinMaxScaler(feature_range=(0,1))
```

```
df_scaled = scaler.fit_transform(df[['Close']])
```

### 5.2 Convert Data into Sequences

LSTMs require **input sequences (X)** and **target values (Y)**.

```
def create_sequences(data, time_steps=50):  
    X, Y = [], []  
  
    for i in range(len(data) - time_steps):  
        X.append(data[i:i+time_steps])  
        Y.append(data[i+time_steps])  
  
    return np.array(X), np.array(Y)
```

```
# Define time steps
```

```
time_steps = 50
```

```
X, Y = create_sequences(df_scaled, time_steps)
```

```
# Split into training & testing sets
```

```
split = int(0.8 * len(X))

X_train, Y_train = X[:split], Y[:split]

X_test, Y_test = X[split:], Y[split:]

print("Training Data Shape:", X_train.shape, Y_train.shape)

print("Testing Data Shape:", X_test.shape, Y_test.shape)
```

## 📌 CHAPTER 6: BUILD & TRAIN LSTM MODEL

### 6.1 Define the LSTM Model

```
model = Sequential([
    LSTM(50, activation='relu', return_sequences=True,
         input_shape=(time_steps, 1)),
    LSTM(50, activation='relu', return_sequences=False),
    Dense(25),
    Dense(1)
])

# Compile the model
model.compile(optimizer='adam', loss='mean_squared_error')

# Display model summary
model.summary()
```

### 6.2 Train the Model

```
history = model.fit(X_train, Y_train, epochs=20, batch_size=32,  
validation_data=(X_test, Y_test))
```

## 📌 CHAPTER 7: EVALUATE & MAKE PREDICTIONS

### 7.1 Plot Training Loss

```
plt.plot(history.history['loss'], label='Training Loss')  
plt.plot(history.history['val_loss'], label='Validation Loss')  
plt.legend()  
plt.title("Model Loss Over Epochs")  
plt.show()
```

### 7.2 Make Predictions

```
Y_pred = model.predict(X_test)  
Y_pred = scaler.inverse_transform(Y_pred) # Convert back to  
original scale  
Y_test_actual = scaler.inverse_transform(Y_test.reshape(-1,1))  
  
# Plot actual vs predicted values  
plt.figure(figsize=(10,5))  
plt.plot(Y_test_actual, label="Actual Price")  
plt.plot(Y_pred, label="Predicted Price")  
plt.xlabel("Time")  
plt.ylabel("Price")  
plt.title("Stock Price Prediction")
```

```
plt.legend()
```

```
plt.show()
```

---

## 📌 CHAPTER 8: ADVANCED LSTM TECHNIQUES

### 8.1 Improving Performance

- **Use More Data** – Train on a longer time period.
- **Hyperparameter Tuning** – Adjust LSTM layers, neurons, and learning rate.
- **Bidirectional LSTM** – Captures both past and future trends.

### 8.2 Using LSTM for Multivariate Forecasting

Instead of using **only the closing price**, add **features like volume, moving averages**.

```
df['MA_50'] = df['Close'].rolling(50).mean() # 50-day Moving Average
```

```
df.dropna(inplace=True) # Drop NaN values
```

Modify `create_sequences()` to include multiple features.

---

## 📌 CHAPTER 9: SUMMARY & NEXT STEPS

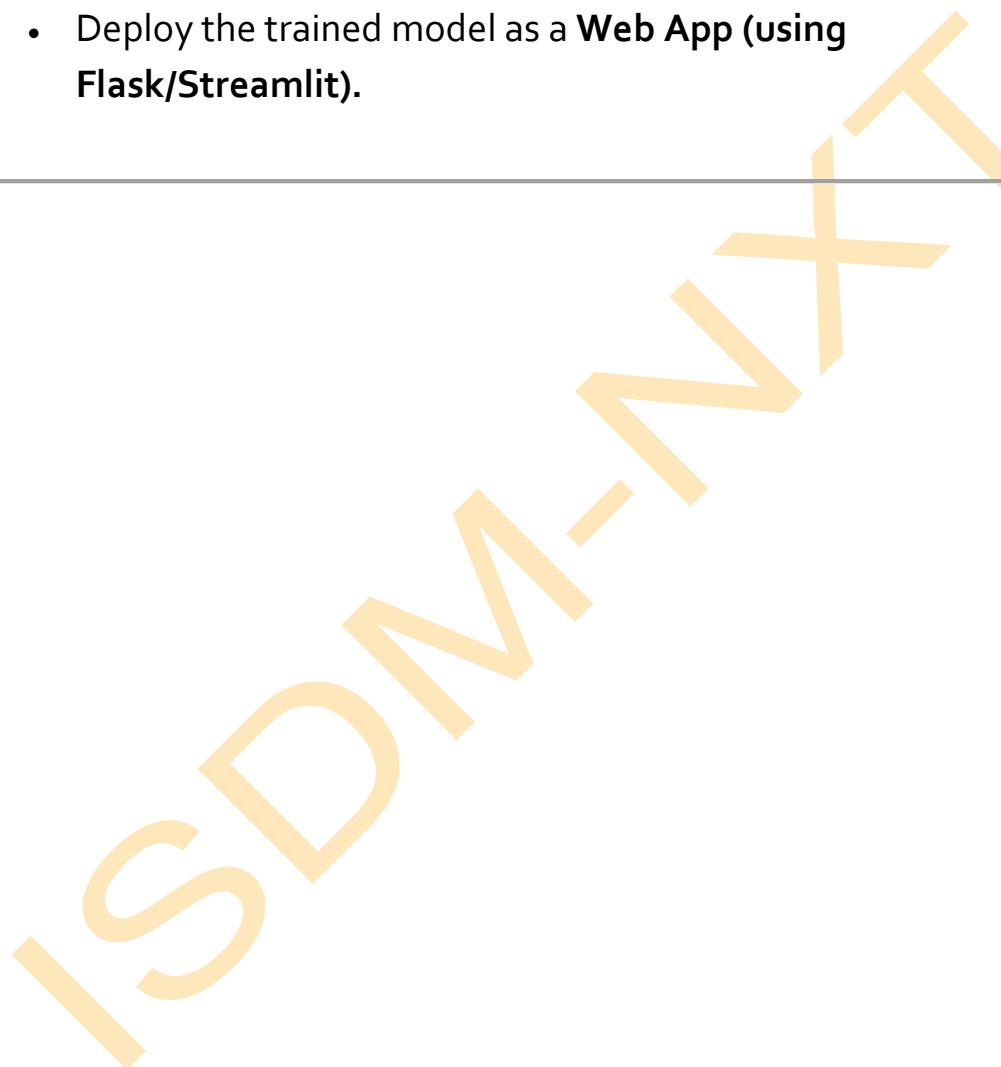
### 9.1 Key Takeaways

- ✓ **LSTMs excel in time series prediction** by capturing long-term dependencies.
- ✓ **Preprocessing is crucial** (scaling, sequence creation).
- ✓ **Hyperparameter tuning** improves model performance.

- ✓ Advanced LSTMs (Bidirectional, Attention-based) enhance accuracy.

## 9.2 Next Steps

- Try **Bidirectional LSTM** for better context learning.
  - Use **Transformer models (like BERT for time series)**.
  - Deploy the trained model as a **Web App (using Flask/Streamlit)**.
- 

A large, faint watermark reading "ISDM-Nxt" diagonally across the page. The letters are yellow and have a blocky, stylized font. The "Nxt" part is smaller than the "ISDM".

---

# STOCK PRICE PREDICTION USING MACHINE LEARNING

---

## CHAPTER 1: INTRODUCTION TO STOCK PRICE PREDICTION

### 1.1 What is Stock Price Prediction?

Stock price prediction involves using **machine learning (ML) models** to analyze historical stock data and forecast future stock prices. It combines **data science, financial analysis, and predictive modeling**.

### 1.2 Why Use Machine Learning for Stock Prediction?

-  Identifies **patterns and trends** in stock prices.
-  Reduces **human biases** in decision-making.
-  Helps in **risk management** and portfolio optimization.
-  Improves **short-term and long-term trading strategies**.

### 1.3 Challenges in Stock Price Prediction

- **Stock prices are highly volatile** (affected by news, market trends, etc.).
- **Data quality matters** (historical prices, indicators, news sentiment).
- **Risk of overfitting** if models are too complex.

---

## CHAPTER 2: TYPES OF MACHINE LEARNING MODELS FOR STOCK PREDICTION

### 2.1 Traditional vs. Machine Learning Approaches

Approach	Description	Examples
<b>Traditional</b>	Uses statistical methods	Moving Averages, ARIMA
<b>ML-Based</b>	Uses AI to learn patterns	Linear Regression, LSTMs

## 2.2 Common ML Models Used

1. **Linear Regression** – Predicts future prices based on historical trends.
2. **Decision Trees & Random Forests** – Handles non-linearity in data.
3. **Support Vector Machines (SVMs)** – Finds optimal trends in stock price data.
4. **Recurrent Neural Networks (RNNs) / LSTMs** – Captures time-series dependencies.
5. **XGBoost** – Powerful for structured financial data.

## CHAPTER 3: STOCK PRICE PREDICTION USING PYTHON

### 3.1 Install Required Libraries

```
!pip install pandas numpy scikit-learn matplotlib yfinance
```

### 3.2 Import Libraries

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import yfinance as yf
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.linear_model import LinearRegression
```

```
from sklearn.metrics import mean_squared_error
```

## 📌 CHAPTER 4: LOAD AND PREPARE STOCK DATA

### 4.1 Download Stock Data using Yahoo Finance

```
# Download Apple (AAPL) stock data from Yahoo Finance
df = yf.download('AAPL', start='2015-01-01', end='2023-01-01')

# Display first few rows
print(df.head())

4.2 Select Features

# Selecting only 'Close' price for prediction
df = df[['Close']]

# Create a new column 'Target' (Shifted Close Price)
df['Target'] = df['Close'].shift(-1)

# Drop NaN values
df.dropna(inplace=True)

# Display dataset
print(df.head())
```

## CHAPTER 5: SPLITTING DATA AND TRAINING A MACHINE LEARNING MODEL

### 5.1 Splitting the Data

```
# Splitting into features (X) and target (y)
```

```
X = df[['Close']]
```

```
y = df['Target']
```

```
# Splitting into training and test sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
random_state=42)
```

### 5.2 Train a Linear Regression Model

```
# Train the model
```

```
model = LinearRegression()
```

```
model.fit(X_train, y_train)
```

```
# Make predictions
```

```
y_pred = model.predict(X_test)
```

### 5.3 Evaluate the Model

```
# Calculate Mean Squared Error
```

```
mse = mean_squared_error(y_test, y_pred)
```

```
print(f"Mean Squared Error: {mse:.2f}")
```

## 📌 CHAPTER 6: VISUALIZING THE PREDICTIONS

```
plt.figure(figsize=(10,5))

plt.plot(y_test.values, label='Actual Prices', color='blue')

plt.plot(y_pred, label='Predicted Prices', color='red',
         linestyle='dashed')

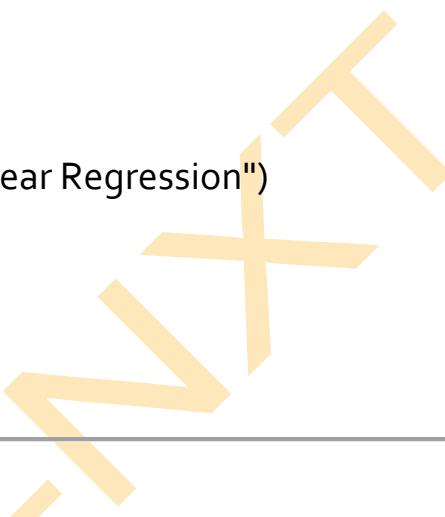
plt.xlabel("Time")

plt.ylabel("Stock Price")

plt.title("Stock Price Prediction - Linear Regression")

plt.legend()

plt.show()
```



## 📌 CHAPTER 7: ADVANCED TECHNIQUES

### 7.1 Using Moving Averages for Smoother Predictions

```
df['SMA_50'] = df['Close'].rolling(window=50).mean()

df['SMA_200'] = df['Close'].rolling(window=200).mean()

plt.figure(figsize=(10,5))

plt.plot(df['Close'], label='Stock Price', color='blue')

plt.plot(df['SMA_50'], label='50-Day SMA', color='green')

plt.plot(df['SMA_200'], label='200-Day SMA', color='red')

plt.title("Stock Price with Moving Averages")

plt.legend()

plt.show()
```

## 7.2 Using LSTM for Deep Learning-Based Prediction

LSTMs (Long Short-Term Memory networks) are used for time-series forecasting.

```
from tensorflow.keras.models import Sequential
```

```
from tensorflow.keras.layers import LSTM, Dense
```

```
# Reshape data for LSTM
```

```
X_train_reshaped = np.reshape(X_train.values, (X_train.shape[0],  
X_train.shape[1], 1))
```

```
X_test_reshaped = np.reshape(X_test.values, (X_test.shape[0],  
X_test.shape[1], 1))
```

```
# Define LSTM model
```

```
model = Sequential([
```

```
    LSTM(units=50, return_sequences=True,  
input_shape=(X_train_reshaped.shape[1], 1)),
```

```
    LSTM(units=50, return_sequences=False),
```

```
    Dense(units=25),
```

```
    Dense(units=1)
```

```
])
```

```
# Compile model
```

```
model.compile(optimizer='adam', loss='mean_squared_error')
```

```
# Train the model
```

```
model.fit(X_train_reshaped, y_train, epochs=10, batch_size=32)
```

---

## 📌 CHAPTER 8: SUMMARY

Model	Best For	Advantages	Disadvantages
Linear Regression	Basic trend prediction	Simple & interpretable	Can't capture complex patterns
Decision Trees	Handling non-linear data	Works with missing data	Prone to overfitting
Random Forest	Robust stock forecasting	Reduces overfitting	Computationally expensive
LSTMs (Deep Learning)	Time-series forecasting	Captures dependencies in stock data	Needs large datasets

---

## 📌 CHAPTER 9: CONCLUSION & NEXT STEPS

### 9.1 Key Takeaways

- ✓ Stock prices can be predicted using **ML techniques like Linear Regression, Decision Trees, and LSTMs.**
- ✓ **Feature engineering** (like Moving Averages) improves model accuracy.
- ✓ **LSTMs are powerful** but require **more data and computing power.**

### 9.2 Next Steps

- 🚀 Try different ML models (SVM, XGBoost, Deep Learning).
- 🚀 Use fundamental/technical indicators (RSI, MACD).
- 🚀 Apply NLP to analyze stock-related news sentiment.

ISDM-NxT

# AI FOR FRAUD DETECTION & RISK MANAGEMENT

## CHAPTER 1: INTRODUCTION TO AI IN FRAUD DETECTION & RISK MANAGEMENT

### 1.1 What is Fraud Detection & Risk Management?

Fraud detection and risk management refer to the processes used to **identify, prevent, and mitigate fraudulent activities** and financial risks in various industries. These activities include **credit card fraud, insurance fraud, money laundering, identity theft, and cyber fraud.**

### 1.2 Why Use AI for Fraud Detection & Risk Management?

Traditional fraud detection methods rely on **rule-based systems**,

which are:  **Slow and reactive**

 **Prone to human errors**

 **Inefficient at detecting new fraud patterns**

 **AI-based fraud detection uses Machine Learning (ML) and Deep Learning to:**  **Detect fraud in real-time.**

 **Identify hidden patterns** in large datasets.

 **Adapt to new fraud techniques** without human intervention.

 **Reduce false positives** in fraud detection.

### 1.3 Key Applications of AI in Fraud & Risk Management

Industry	Application
<b>Banking &amp; Finance</b>	Credit card fraud, transaction monitoring
<b>E-commerce</b>	Fake accounts, fraudulent transactions

<b>Healthcare</b>	Insurance fraud, medical claims fraud
<b>Cybersecurity</b>	Identity theft, phishing attacks
<b>Gaming &amp; Gambling</b>	Account takeovers, bonus abuse fraud

## 📌 CHAPTER 2: AI TECHNIQUES USED IN FRAUD DETECTION

### 2.1 Machine Learning (ML) for Fraud Detection

ML algorithms can analyze **large datasets** and **learn fraud patterns**. They classify transactions as **fraudulent (1)** or **legitimate (0)** based on historical data.

### 2.2 Supervised vs. Unsupervised Learning

Approach	Description	Example Algorithm
<b>Supervised Learning</b>	Uses labeled fraud data to train models	Logistic Regression, Decision Trees, Random Forest
<b>Unsupervised Learning</b>	Detects anomalies in <b>unlabeled</b> data	K-Means Clustering, Autoencoders, Isolation Forest
<b>Reinforcement Learning</b>	Learns from past fraud cases and improves decisions	Deep Q-Learning

### 2.3 Deep Learning for Fraud Detection

- ✓ **Recurrent Neural Networks (RNNs)** – Used for detecting fraud in **sequential transactions** (e.g., credit card fraud).
- ✓ **Long Short-Term Memory (LSTM)** – Detects **time-series fraud patterns**.

- ✓ **Autoencoders** – Used for **anomaly detection** in cybersecurity threats.

## 2.4 Natural Language Processing (NLP)

- ✓ **Text analysis for fraud detection** in emails, chat messages, and online reviews.
- ✓ **Detects phishing attempts** by analyzing patterns in scam messages.

## CHAPTER 3: FRAUD DETECTION WORKFLOW USING AI

### 3.1 Data Collection

- **Transaction logs** (banking, e-commerce, cryptocurrency)
- **Customer behavior data** (login locations, time of transactions)
- **Financial records** (credit history, loan applications)

### 3.2 Data Preprocessing

- **Handling missing values**
- **Feature engineering** (e.g., calculating transaction frequency)
- **Converting categorical data** (e.g., transaction type) into numerical values

### 3.3 Training AI Models

- **Choose an algorithm** (e.g., Random Forest, Neural Networks)
- **Split data into training & test sets** (80%-20% rule)
- **Train model to classify transactions as 'fraud' or 'legitimate'**

### 3.4 Real-Time Fraud Detection

- AI model **monitors transactions in real-time**

- If fraud probability > **threshold**, trigger **alert/block transaction**
  - If fraud probability **low**, transaction proceeds normally
- 



## CHAPTER 4: IMPLEMENTING AI FOR FRAUD DETECTION

### 4.1 Install Dependencies

```
!pip install numpy pandas scikit-learn matplotlib seaborn tensorflow  
keras
```

### 4.2 Load and Explore Fraud Dataset

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
# Load dataset (example: Credit Card Fraud dataset)
```

```
df = pd.read_csv("creditcard.csv")
```

```
# Display first few rows
```

```
print(df.head())
```

```
# Check for missing values
```

```
print(df.isnull().sum())
```

### 4.3 Data Preprocessing

```
from sklearn.model_selection import train_test_split  
from sklearn.preprocessing import StandardScaler
```

```
# Features and target variable
```

```
X = df.drop(columns=['Class']) # Features  
y = df['Class'] # Target: 0 (legit), 1 (fraud)
```

```
# Scale the data
```

```
scaler = StandardScaler()  
X_scaled = scaler.fit_transform(X)
```

```
# Split data into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,  
test_size=0.2, random_state=42)
```

```
print("Training set size:", X_train.shape)
```

```
print("Test set size:", X_test.shape)
```

### 4.4 Train a Machine Learning Model for Fraud Detection

```
from sklearn.ensemble import RandomForestClassifier  
from sklearn.metrics import accuracy_score, classification_report
```

```
# Train a Random Forest model  
  
model = RandomForestClassifier(n_estimators=100,  
random_state=42)  
  
model.fit(X_train, y_train)
```

```
# Make predictions  
  
y_pred = model.predict(X_test)  
  
# Evaluate model performance  
  
print("Accuracy:", accuracy_score(y_test, y_pred))  
print(classification_report(y_test, y_pred))
```

---

#### 4.5 Using Deep Learning for Fraud Detection (LSTM)

```
import tensorflow as tf  
  
from tensorflow.keras.models import Sequential  
  
from tensorflow.keras.layers import LSTM, Dense, Dropout  
  
# Reshape data for LSTM  
  
X_train_lstm = np.reshape(X_train, (X_train.shape[0],  
X_train.shape[1], 1))  
  
X_test_lstm = np.reshape(X_test, (X_test.shape[0], X_test.shape[1],  
1))
```

```
# Build LSTM model  
  
lstm_model = Sequential([  
  
    LSTM(64, return_sequences=True, input_shape=(X_train.shape[1],  
    1)),  
  
    Dropout(0.2),  
  
    LSTM(32, return_sequences=False),  
  
    Dense(1, activation='sigmoid')  
])  
  
# Compile model  
  
lstm_model.compile(loss='binary_crossentropy', optimizer='adam',  
metrics=['accuracy'])  
  
# Train model  
  
lstm_model.fit(X_train_lstm, y_train, epochs=5, batch_size=32,  
validation_data=(X_test_lstm, y_test))
```

## 📌 CHAPTER 5: REAL-WORLD AI FRAUD DETECTION USE CASES

### 5.1 AI in Banking

- Credit card fraud detection** using ML and deep learning models.
- Loan risk assessment** – AI evaluates customers' credit scores for loan approvals.

### 5.2 AI in Cybersecurity

- AI-powered anomaly detection** to stop hacking attempts.
- AI detects phishing attacks** based on email patterns.

### 5.3 AI in Insurance

- Automated claims processing** using AI-powered fraud detection.
- AI flagging suspicious claims** based on customer history.

---

## 📌 CHAPTER 6: CHALLENGES & FUTURE OF AI IN FRAUD DETECTION

### 6.1 Challenges

- Imbalanced Datasets** – Fraud cases are rare, making training difficult.
- False Positives** – Overly aggressive fraud detection can block real customers.
- Evolving Fraud Techniques** – Fraudsters constantly develop new methods.

### 6.2 Future Trends

- AI-Powered Blockchain for Security**
- Explainable AI (XAI) for better decision-making**
- Deep Reinforcement Learning for adaptive fraud detection**

---

## 📌 CHAPTER 7: SUMMARY

- AI fraud detection reduces risks in banking, e-commerce, and cybersecurity.**
- ML and DL techniques detect fraud patterns in large datasets.**
- Real-time fraud detection is essential for stopping cyber**

threats.

- LSTM models work well for sequential fraud analysis.**
- 

📌 **CHAPTER 8: NEXT STEPS**

- 🚀 Try implementing **AI-powered fraud detection on real-world datasets.**
- 🚀 Use **Autoencoders for Anomaly Detection** in cybersecurity.
- 🚀 Deploy a **Fraud Detection API** for banking applications.

ISDM-NXT

---

📌 ⚡ **ASSIGNMENT 1:**  
🎯 **TRAIN AN ARIMA MODEL FOR STOCK  
PRICE PREDICTION.**

ISDM-NXT



# ASSIGNMENT SOLUTION 1: TRAIN AN ARIMA MODEL FOR STOCK PRICE PREDICTION

## 🎯 Objective

The goal of this assignment is to **train an ARIMA model** to predict stock prices based on historical stock market data. ARIMA (AutoRegressive Integrated Moving Average) is a powerful statistical method for **time series forecasting**.

We will:

- Load historical stock price data
- Perform **data preprocessing** and visualization
- Check **stationarity** and apply transformations if needed
- Train an **ARIMA model** and evaluate its performance
- Make **future predictions**

## 🛠 Step 1: Install and Import Required Libraries

We will use **pandas**, **matplotlib**, **statsmodels**, and **yfinance** to train our ARIMA model.

### 1.1 Install Dependencies (if not installed)

```
!pip install pandas numpy matplotlib statsmodels yfinance
```

### 1.2 Import Required Libraries

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt  
import statsmodels.api as sm  
import yfinance as yf  
from statsmodels.tsa.stattools import adfuller  
from statsmodels.tsa.arima.model import ARIMA  
from sklearn.metrics import mean_absolute_error,  
mean_squared_error
```

## Step 2: Load Stock Price Data

We will use **Yahoo Finance (yfinance)** to fetch historical stock price data.

### 2.1 Load Stock Data for a Specific Company

```
# Define the stock symbol and time period  
stock_symbol = "AAPL" # Apple Inc.  
start_date = "2018-01-01"  
end_date = "2023-12-31"  
  
# Fetch stock data  
df = yf.download(stock_symbol, start=start_date, end=end_date)  
  
# Display first few rows  
print(df.head())
```

### 2.2 Select the Closing Price

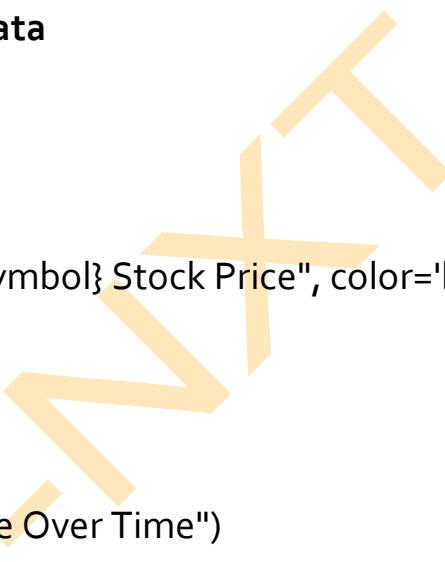
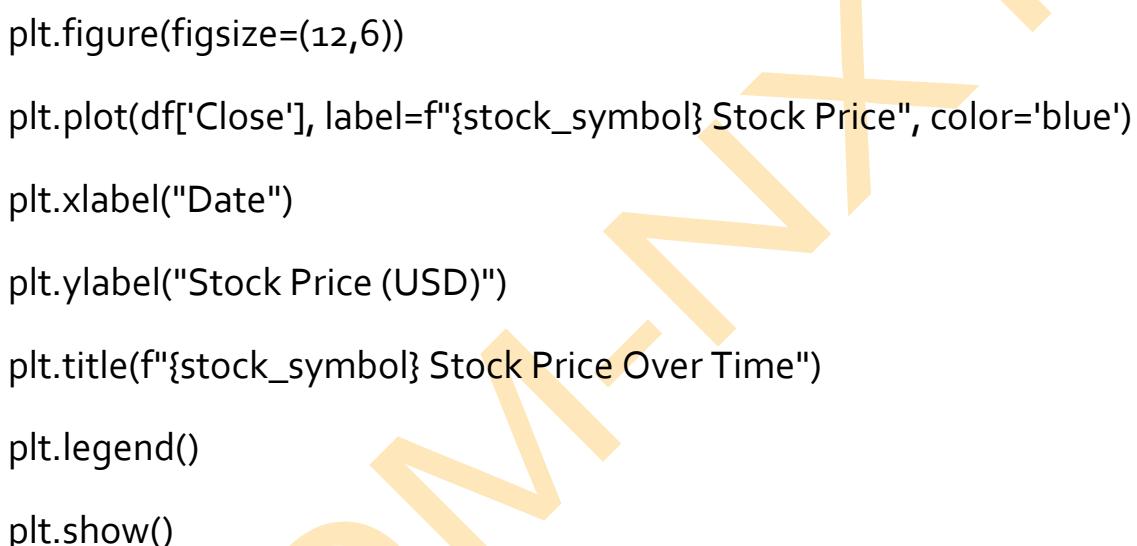
```
df = df[['Close']] # We only need the closing price  
df = df.asfreq('B') # Convert to business days (removing non-trading  
days)  
df = df.fillna(method='ffill') # Fill missing values by forward filling
```

---

### Step 3: Visualize Stock Price Data

#### 3.1 Plot the Stock Price Over Time

```
plt.figure(figsize=(12,6))  
plt.plot(df['Close'], label=f'{stock_symbol} Stock Price', color='blue')  
plt.xlabel("Date")  
plt.ylabel("Stock Price (USD)")  
plt.title(f'{stock_symbol} Stock Price Over Time')  
plt.legend()  
plt.show()
```



### Step 4: Check for Stationarity

A stationary time series has **constant mean and variance** over time, which is essential for ARIMA modeling.

#### 4.1 Perform the Augmented Dickey-Fuller (ADF) Test

```
def adf_test(series):  
  
    result = adfuller(series)  
  
    print("ADF Statistic:", result[0])  
  
    print("p-value:", result[1])
```

```
print("Critical Values:", result[4])  
if result[1] < 0.05:  
    print(" ✅ The data is stationary (p < 0.05).")  
else:  
    print(" ❌ The data is NOT stationary (p > 0.05). Differencing is required.")
```

```
# Run ADF test  
adf_test(df['Close'])
```

## ⌚ Step 5: Make Data Stationary (If Needed)

If the ADF test shows **non-stationarity**, we apply **differencing**.

### 5.1 Apply First Order Differencing

```
df['Close_diff'] = df['Close'].diff().dropna()
```

```
# Re-run ADF test  
adf_test(df['Close_diff'].dropna())
```

```
# Plot differenced data  
plt.figure(figsize=(12,6))  
plt.plot(df['Close_diff'], label="Differenced Stock Price",  
color='green')  
plt.xlabel("Date")
```

```
plt.ylabel("Differenced Closing Price")  
plt.title(f"{{stock_symbol}} Stock Price After Differencing")  
plt.legend()  
plt.show()
```

---

## 📌 Step 6: Determine ARIMA Model Parameters (p, d, q)

ARIMA has three components:

- **p (AutoRegressive Order):** Number of lag observations
- **d (Differencing Order):** Number of times differencing is applied
- **q (Moving Average Order):** Number of lagged forecast errors

### 6.1 Plot AutoCorrelation Function (ACF) and Partial AutoCorrelation Function (PACF)

```
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
```

```
plt.figure(figsize=(12,6))  
plt.subplot(121)  
plot_acf(df['Close_diff'].dropna(), lags=20, ax=plt.gca()) # ACF plot  
plt.subplot(122)  
plot_pacf(df['Close_diff'].dropna(), lags=20, ax=plt.gca()) # PACF plot  
plt.show()
```

### 6.2 Choose ARIMA Order Based on Plots

- **p (AR)**: Determined from the PACF plot (significant lags).
- **q (MA)**: Determined from the ACF plot.
- **d**: Already determined (first differencing = 1).

From the plots, let's assume:

$p = 2$  # Number of significant lags in PACF

$d = 1$  # Differencing order from ADF test

$q = 2$  # Number of significant lags in ACF

## Step 7: Train the ARIMA Model

### 7.1 Fit ARIMA Model

```
model = ARIMA(df['Close'], order=(p, d, q))
```

```
model_fit = model.fit()
```

```
# Display model summary
```

```
print(model_fit.summary())
```

## Step 8: Forecast Future Stock Prices

### 8.1 Make Predictions on Test Data

```
df['Forecast'] = model_fit.fittedvalues
```

```
plt.figure(figsize=(12,6))
```

```
plt.plot(df['Close'], label="Actual Stock Price", color='blue')
```

```
plt.plot(df['Forecast'], label="ARIMA Forecast", color='red')
plt.xlabel("Date")
plt.ylabel("Stock Price (USD)")
plt.title(f"ARIMA Model - {stock_symbol} Stock Price Prediction")
plt.legend()
plt.show()
```

## 8.2 Evaluate Model Performance

```
actual = df['Close'].dropna()
predicted = df['Forecast'].dropna()

mae = mean_absolute_error(actual, predicted)
mse = mean_squared_error(actual, predicted)
rmse = np.sqrt(mse)

print(f"Mean Absolute Error (MAE): {mae:.2f}")
print(f"Mean Squared Error (MSE): {mse:.2f}")
print(f"Root Mean Squared Error (RMSE): {rmse:.2f}")
```

## ➡ Step 9: Make Future Predictions

### 9.1 Forecast Stock Prices for the Next 30 Days

```
forecast_steps = 30 # Number of days to forecast
future_dates = pd.date_range(df.index[-1],
periods=forecast_steps+1, freq='B')[1:]
```

```
forecast = model_fit.forecast(steps=forecast_steps)

forecast_series = pd.Series(forecast, index=future_dates)

# Plot the forecast

plt.figure(figsize=(12,6))

plt.plot(df['Close'], label="Historical Stock Price", color='blue')

plt.plot(forecast_series, label="Future Prediction", color='orange',
         linestyle="dashed")

plt.xlabel("Date")

plt.ylabel("Stock Price (USD)")

plt.title(f"Future Stock Price Prediction for {stock_symbol}")

plt.legend()

plt.show()
```

#### FINAL SUMMARY

Step	Description
<b>Step 1</b>	Install and import necessary libraries
<b>Step 2</b>	Load stock price data using Yahoo Finance
<b>Step 3</b>	Visualize stock price trends
<b>Step 4</b>	Check stationarity using the <b>ADF test</b>
<b>Step 5</b>	Apply differencing if data is non-stationary

<b>Step 6</b>	Determine <b>ARIMA(p,d,q)</b> using ACF & PACF plots
<b>Step 7</b>	Train an <b>ARIMA model</b>
<b>Step 8</b>	Evaluate model performance & visualize predictions
<b>Step 9</b>	Forecast future stock prices

## CONCLUSION

- ARIMA is a powerful tool for **stock price forecasting**.
- **Stationarity is crucial** for accurate predictions.
- **Hyperparameter tuning (p,d,q) is needed** to improve accuracy.
- The model can be extended to **real-time stock analysis**.

ISDM

---

📌 ⚡ ASSIGNMENT 2:  
🎯 IMPLEMENT LSTM FOR FORECASTING  
ELECTRICITY CONSUMPTION TRENDS.

ISDM-NxT

---



## ASSIGNMENT SOLUTION 2: IMPLEMENT LSTM FOR FORECASTING ELECTRICITY CONSUMPTION TRENDS

---

### 🎯 Objective

The goal of this assignment is to **implement a Long Short-Term Memory (LSTM) model** using **TensorFlow/Keras** to **forecast electricity consumption trends** based on historical data. LSTMs are powerful for time series forecasting as they can learn **long-term dependencies** in data.

### ❖ Step 1: Install and Import Necessary Libraries

We will use:

- **Pandas** for data handling.
- **NumPy** for numerical computations.
- **Matplotlib** for visualization.
- **Scikit-learn** for data preprocessing.
- **TensorFlow/Keras** for building the LSTM model.
- ◆ **Install Dependencies (if not installed)**

```
!pip install tensorflow pandas numpy scikit-learn matplotlib
```

#### ◆ Import Required Libraries

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt  
from sklearn.preprocessing import MinMaxScaler  
from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import LSTM, Dense, Dropout  
from tensorflow.keras.optimizers import Adam
```

---

## Step 2: Load and Explore the Dataset

For this assignment, we assume the dataset contains electricity consumption data in **hourly or daily intervals**.

### ◆ Load the Electricity Consumption Data

```
# Load dataset (Assume CSV file with 'Date' and 'Consumption'  
columns)  
  
df = pd.read_csv("electricity_consumption.csv",  
parse_dates=['Date'], index_col='Date')
```

```
# Display first few rows
```

```
print(df.head())
```

```
# Check for missing values
```

```
print(df.isnull().sum())
```

### ◆ Visualize the Data

```
plt.figure(figsize=(12,5))  
plt.plot(df, label="Electricity Consumption")
```

```
plt.xlabel("Date")
plt.ylabel("Consumption (kWh)")
plt.title("Electricity Consumption Over Time")
plt.legend()
plt.show()
```

### Step 3: Data Preprocessing

LSTM models require **normalized data** and **input sequences** for training.

#### ◆ Normalize the Data

```
scaler = MinMaxScaler(feature_range=(0,1)) # Scale between 0 and 1
df_scaled = scaler.fit_transform(df)
# Convert back to DataFrame
df_scaled = pd.DataFrame(df_scaled, index=df.index, columns=['Consumption'])
```

#### ◆ Create Sequences for LSTM

We need to prepare the data as **input-output sequences** where:

- **X (Input)** contains past n timestamps.
- **y (Output)** is the next consumption value.

```
def create_sequences(data, time_steps=24):
```

```
    X, y = [], []
```

```
for i in range(len(data) - time_steps):  
    X.append(data[i:i+time_steps]) # Last 'time_steps' values  
    y.append(data[i+time_steps]) # Next time-step value  
return np.array(X), np.array(y)
```

```
# Define time steps (e.g., 24 hours for daily forecasting)
```

```
time_steps = 24
```

```
# Create sequences
```

```
X, y = create_sequences(df_scaled.values, time_steps)
```

```
# Display shape
```

```
print(f"Input Shape: {X.shape}, Output Shape: {y.shape}")
```

---

#### Step 4: Train-Test Split

Split data into **training and testing sets** (80% train, 20% test).

```
train_size = int(len(X) * 0.8)
```

```
# Train and test sets
```

```
X_train, y_train = X[:train_size], y[:train_size]
```

```
X_test, y_test = X[train_size:], y[train_size:]
```

```
# Display shapes  
print(f"Training Data: {X_train.shape}, Testing Data: {X_test.shape}")
```

---

## ■ Step 5: Build the LSTM Model

Now, let's create a **stacked LSTM model** with:

- **3 LSTM layers** with units=50
- **Dropout layers** to prevent overfitting
- **Dense output layer** for predicting electricity consumption

```
# Define LSTM model
```

```
model = Sequential([  
    LSTM(units=50, return_sequences=True,  
        input_shape=(X_train.shape[1], 1)),  
    Dropout(0.2),  
    LSTM(units=50, return_sequences=True),  
    Dropout(0.2),  
    LSTM(units=50),  
    Dropout(0.2),  
    Dense(units=1) # Output layer  
])
```

```
# Compile the model
```

```
model.compile(optimizer=Adam(learning_rate=0.001), loss='mse')
```

---

```
# Display model summary
```

```
model.summary()
```

---

### ➡ Step 6: Train the Model

```
# Train the model
```

```
history = model.fit(X_train, y_train, epochs=20, batch_size=32,  
validation_data=(X_test, y_test))
```

---

### 📈 Step 7: Evaluate Model Performance

- ◆ Plot Training Loss

```
plt.figure(figsize=(10,5))  
  
plt.plot(history.history['loss'], label="Training Loss")  
  
plt.plot(history.history['val_loss'], label="Validation Loss")  
  
plt.title("LSTM Training Loss")  
  
plt.xlabel("Epochs")  
  
plt.ylabel("Loss")  
  
plt.legend()  
  
plt.show()
```

- ◆ Make Predictions

```
# Make predictions  
y_pred = model.predict(X_test)  
  
# Rescale predictions back to original scale  
y_pred_rescaled = scaler.inverse_transform(y_pred)  
y_test_rescaled = scaler.inverse_transform(y_test.reshape(-1,1))  
  
# Plot actual vs. predicted  
plt.figure(figsize=(12,5))  
plt.plot(y_test_rescaled, label="Actual Consumption",  
         linestyle="solid")  
plt.plot(y_pred_rescaled, label="Predicted Consumption",  
         linestyle="dashed")  
plt.title("Electricity Consumption Forecasting with LSTM")  
plt.xlabel("Time")  
plt.ylabel("Consumption (kWh)")  
plt.legend()  
plt.show()
```

## ➡ Step 8: Forecast Future Electricity Consumption

To forecast **future consumption**, we use the last time\_steps values.

```
# Get the last known values from the dataset
```

```
last_known_values = df_scaled.values[-time_steps:]
```

```
# Reshape input for LSTM  
  
input_seq = np.array([last_known_values])  
  
print(f"Input Sequence Shape: {input_seq.shape}")  
  
  
# Predict next consumption value  
  
future_prediction = model.predict(input_seq)  
  
  
# Convert back to original scale  
  
future_value = scaler.inverse_transform(future_prediction.reshape(-1,1))  
  
  
print(f"Forecasted Electricity Consumption: {future_value[0][0]:.2f} kWh")
```

#### FINAL SUMMARY

Step	Description
Step 1	Install and import necessary libraries
Step 2	Load and explore the electricity consumption dataset
Step 3	Normalize the data and create LSTM sequences
Step 4	Split data into training and testing sets
Step 5	Build the LSTM model using TensorFlow/Keras

<b>Step 6</b>	Train the model on the dataset
<b>Step 7</b>	Evaluate model performance and visualize results
<b>Step 8</b>	Forecast future electricity consumption

## CONCLUSION

- **LSTMs effectively capture temporal dependencies** in electricity consumption data.
- The trained model can predict future electricity trends using historical data.
- The accuracy can be improved using hyperparameter tuning and more data.
- This model can be extended to real-world energy forecasting applications.

## NEXT STEPS

- Tune LSTM hyperparameters (e.g., units, epochs, batch\_size).
- Compare LSTM with other models (ARIMA, SARIMA, XGBoost).
- Deploy the model using Flask or FastAPI for real-time predictions.
- Use multivariate forecasting (e.g., include temperature, weather conditions, and holidays).