



ISDM (INDEPENDENT SKILL DEVELOPMENT MISSION)

◊ TEXT PREPROCESSING: TOKENIZATION, STEMMING, LEMMATIZATION

📌 CHAPTER 1: INTRODUCTION TO TEXT PREPROCESSING

◆ 1.1 What is Text Preprocessing?

Text preprocessing is the process of preparing and cleaning text data before applying **Natural Language Processing (NLP)** techniques. Raw text contains **noise, variations, and redundancies**, making it difficult for machine learning models to interpret. Preprocessing converts text into a **structured, meaningful format**.

◆ 1.2 Why is Text Preprocessing Important?

✓ **Improves Model Accuracy** – Cleaner data leads to better predictions in NLP models.

✓ **Reduces Dimensionality** – Removes unnecessary words, reducing computational cost.

✓ **Enhances Text Normalization** – Unifies words with similar meanings.

✓ **Prepares Text for Tokenization, Sentiment Analysis, and Topic Modeling.**

◆ 1.3 Common Steps in Text Preprocessing

- 1 **Tokenization** – Splitting text into words or sentences.
- 2 **Lowercasing** – Standardizing text by converting it to lowercase.
- 3 **Removing Punctuation & Special Characters** – Eliminating unnecessary symbols.
- 4 **Stopword Removal** – Filtering out common words like "the", "is", "and".
- 5 **Stemming & Lemmatization** – Converting words to their root form.

📌 Example:

Raw Text:

→ "The runners were running in the park happily."

Processed Text:

→ "runner run park happy"

💡 Conclusion:

Text preprocessing helps transform raw text into **clean and structured data** for NLP applications.

📌 CHAPTER 2: TOKENIZATION

◆ 2.1 What is Tokenization?

Tokenization is the process of splitting text into **smaller units** (tokens), such as words or sentences. These tokens are then analyzed for patterns and relationships.

✓ **Word Tokenization** – Splitting text into individual words.

✓ **Sentence Tokenization** – Splitting text into sentences.

◆ 2.2 Implementing Tokenization in Python

Install Required Libraries

```
pip install nltk spacy
```

Word Tokenization using NLTK

```
import nltk  
nltk.download('punkt')  
from nltk.tokenize import word_tokenize  
  
text = "Natural Language Processing (NLP) is exciting!"  
tokens = word_tokenize(text)  
  
print("Word Tokens:", tokens)
```

📌 Output:

```
['Natural', 'Language', 'Processing', '(', 'NLP', ')', 'is', 'exciting', '!']
```

Sentence Tokenization using NLTK

```
from nltk.tokenize import sent_tokenize
```

```
text = "Hello! How are you? This is an NLP tutorial."
```

```
sentences = sent_tokenize(text)
```

```
print("Sentence Tokens:", sentences)
```

❖ Output:

```
['Hello!', 'How are you?', 'This is an NLP tutorial.']}
```

◆ 2.3 Tokenization using SpaCy

```
import spacy  
nlp = spacy.load("en_core_web_sm")
```

```
text = "Text preprocessing is a crucial step in NLP."
```

```
doc = nlp(text)
```

```
tokens = [token.text for token in doc]
```

```
print("Word Tokens:", tokens)
```

❖ Output:

```
['Text', 'preprocessing', 'is', 'a', 'crucial', 'step', 'in', 'NLP', '.']
```

💡 Conclusion:

Tokenization **splits text into words or sentences**, making it **easier to analyze**.

❖ CHAPTER 3: STEMMING

◆ 3.1 What is Stemming?

Stemming reduces words to their **root form** by **removing suffixes**. However, it may not always produce real words.

- ✓ Reduces **inflected words** to their base form.
 - ✓ Useful for **search engines and text mining**.
 - ✓ Doesn't always create meaningful words.
-

◆ 3.2 Implementing Stemming in Python

Using Porter Stemmer (NLTK)

```
from nltk.stem import PorterStemmer
```

```
stemmer = PorterStemmer()
```

```
words = ["running", "flies", "easily", "happiness"]
```

```
stemmed_words = [stemmer.stem(word) for word in words]
```

```
print("Stemmed Words:", stemmed_words)
```

◆ Output:

```
['run', 'fli', 'easili', 'happi']
```

✓ "running" → "run"

✓ "flies" → "fli" (not a real word)

Using Lancaster Stemmer (More Aggressive)

```
from nltk.stem import LancasterStemmer
```

```
lancaster = LancasterStemmer()
```

```
words = ["running", "flies", "easily", "happiness"]  
stemmed_words = [lancaster.stem(word) for word in words]  
  
print("Stemmed Words:", stemmed_words)
```

📌 **Output:**

```
['run', 'fli', 'easy', 'happy']
```

💡 **Conclusion:**

Stemming reduces words to their root but may not retain meaning. It is useful for text classification and search engines.

📌 **CHAPTER 4: LEMMATIZATION**

◆ **4.1 What is Lemmatization?**

Lemmatization converts words to their **dictionary base form (lemma)**. Unlike stemming, lemmatization produces **real words**.

- ✓ Uses a **vocabulary-based approach**.
- ✓ More **accurate than stemming**.
- ✓ Requires **Part-of-Speech (POS) tagging**.

◆ **4.2 Implementing Lemmatization in Python**

Using WordNetLemmatizer (NLTK)

```
from nltk.stem import WordNetLemmatizer  
from nltk.corpus import wordnet
```

```
nltk.download('wordnet')
```

```
lemmatizer = WordNetLemmatizer()
```

```
words = ["running", "flies", "easily", "happiness"]
```

```
lemmas = [lemmatizer.lemmatize(word, pos="v") for word in words]  
# 'v' = verb
```

```
print("Lemmatized Words:", lemmas)
```

📌 **Output:**

```
['run', 'fly', 'ease', 'happiness']
```

- ✓ "running" → "run"
- ✓ "flies" → "fly"
- ✓ "easily" → "ease"

Using Spacy for Lemmatization

```
import spacy
```

```
nlp = spacy.load("en_core_web_sm")
```

```
text = "The runners are running faster than usual."
```

```
doc = nlp(text)
```

```
lemmas = [token.lemma_ for token in doc]
```

```
print("Lemmatized Text:", lemmas)
```

📌 **Output:**

```
['the', 'runner', 'be', 'run', 'fast', 'than', 'usual', '.']
```

💡 **Conclusion:**

Lemmatization **preserves the correct meaning** of words and is **better for NLP tasks** than stemming.

📌 **CHAPTER 5: STEMMING VS. LEMMATIZATION**

Feature	Stemming	Lemmatization
Definition	Removes suffixes	Converts to dictionary form
Speed	Faster	Slower
Accuracy	Less accurate	More accurate
Use Case	Search engines, quick NLP tasks	Sentiment analysis, AI chatbots

📌 **Example:**

- ✓ "Caring" → Stemmed: "car", Lemmatized: "care"
- ✓ "Better" → Stemmed: "better", Lemmatized: "good"

💡 **Conclusion:**

- ✓ Use **stemming** for **quick NLP tasks**.
- ✓ Use **lemmatization** for **AI chatbots, sentiment analysis, and deep learning**.

📌 SUMMARY & NEXT STEPS

✓ Key Takeaways:

- ✓ **Tokenization** splits text into **words/sentences** for NLP.
- ✓ **Stemming** reduces words to **root form** (may not be real words).
- ✓ **Lemmatization** converts words to their **dictionary base form** (more accurate).

📌 Next Steps:

- ◆ Use NLP libraries (**NLTK, SpaCy**) for text preprocessing.
- ◆ Apply stemming and lemmatization on real-world datasets (tweets, news articles).
- ◆ Build a text classification model using preprocessed text. 🚀

ISDM-NXT

◊ SENTIMENT ANALYSIS & NAMED ENTITY RECOGNITION (NER)

📌 CHAPTER 1: INTRODUCTION TO SENTIMENT ANALYSIS & NAMED ENTITY RECOGNITION (NER)

◆ 1.1 What is Sentiment Analysis?

Sentiment Analysis (also called **opinion mining**) is a **Natural Language Processing (NLP) technique** that determines the **emotional tone of a piece of text**. It is widely used in social media monitoring, product reviews, customer feedback analysis, and more.

Types of Sentiment Analysis:

- ✓ **Binary Sentiment Analysis** – Classifies text as **Positive or Negative**.
- ✓ **Multiclass Sentiment Analysis** – Includes **Neutral**, along with **Positive and Negative**.
- ✓ **Fine-grained Sentiment Analysis** – Uses a scale (e.g., **very positive, positive, neutral, negative, very negative**).
- ✓ **Aspect-based Sentiment Analysis** – Analyzes **sentiment towards specific aspects** of a product/service.

📌 Example:

- ✓ "I love this phone! The camera is amazing!" → **Positive**
- ✓ "The battery life is terrible." → **Negative**
- ✓ "The movie was okay, but too long." → **Neutral**

💡 Conclusion:

Sentiment Analysis helps businesses **understand customer emotions, track brand reputation, and improve user experiences**.

◆ 1.2 What is Named Entity Recognition (NER)?

Named Entity Recognition (NER) is an NLP technique that identifies and classifies entities in text into predefined categories, such as:

- ✓ **Person Names** → "Elon Musk", "Barack Obama"
- ✓ **Organizations** → "Google", "NASA", "Apple Inc."
- ✓ **Locations** → "New York", "Mount Everest"
- ✓ **Dates & Time** → "January 1st, 2023", "5 PM"
- ✓ **Monetary Values** → "\$100", "₹5000"
- ✓ **Product Names** → "iPhone 13", "Tesla Model S"

📌 Example:

✓ **Sentence:** "Amazon was founded by Jeff Bezos in 1994."

✓ **NER Output:**

- **Organization:** Amazon
- **Person:** Jeff Bezos
- **Date:** 1994

💡 Conclusion:

NER is crucial for automated information extraction, chatbot development, and knowledge graph creation.

📌 CHAPTER 2: SENTIMENT ANALYSIS USING NLP

◆ 2.1 Preprocessing Text for Sentiment Analysis

Before applying Sentiment Analysis, we need to **clean** and **preprocess** the text.

Key Preprocessing Steps:

- ✓ **Lowercasing** – Convert text to lowercase.
- ✓ **Removing Punctuation & Special Characters** – Remove , . ? ! etc.
- ✓ **Tokenization** – Splitting text into words.
- ✓ **Stopword Removal** – Removing common words (e.g., "the", "is", "and").
- ✓ **Lemmatization** – Convert words to base form ("running" → "run").

Example of Preprocessing:

```
import re
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem import WordNetLemmatizer
nltk.download('stopwords')
nltk.download('punkt')
nltk.download('wordnet')

# Define preprocessing function
def preprocess_text(text):
    text = text.lower() # Convert to lowercase
    text = re.sub(r'^\w\s', '', text) # Remove punctuation
    tokens = word_tokenize(text) # Tokenization
```

```
tokens = [word for word in tokens if word not in  
stopwords.words('english')] # Remove stopwords  
  
lemmatizer = WordNetLemmatizer()  
  
tokens = [lemmatizer.lemmatize(word) for word in tokens] #  
Lemmatization  
  
return " ".join(tokens)
```

Example usage

```
sample_text = "The movie was absolutely amazing, I loved it!"  
print(preprocess_text(sample_text))
```

📌 **Output:** "movie absolutely amazing loved"

◆ 2.2 Sentiment Analysis using TextBlob (Simple Approach)

```
from textblob import TextBlob
```

Function to predict sentiment

```
def get_sentiment(text):  
    score = TextBlob(text).sentiment.polarity  
  
    if score > 0:  
  
        return "Positive"  
  
    elif score < 0:  
  
        return "Negative"  
  
    else:
```

```
return "Neutral"

# Test cases

print(get_sentiment("The product is fantastic!"))

print(get_sentiment("I hate the slow customer service."))

print(get_sentiment("It's an okay experience."))
```

❖ **Output:**

- ✓ "Positive"
- ✓ "Negative"
- ✓ "Neutral"

◆ **2.3 Sentiment Analysis using Machine Learning (TF-IDF + Logistic Regression)**

```
import pandas as pd

from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression

from sklearn.metrics import accuracy_score
```

Sample dataset

```
data = {'text': ["I love this phone!", "This is the worst movie ever.",
"The food was okay.",

"Amazing experience!", "Horrible customer service."],

'label': [1, 0, 2, 1, 0]} # 1 = Positive, 0 = Negative, 2 = Neutral
```

```
df = pd.DataFrame(data)

# Convert text into numerical features
vectorizer = TfidfVectorizer()
X = vectorizer.fit_transform(df['text'])
y = df['label']

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

# Train model
model = LogisticRegression()
model.fit(X_train, y_train)

# Make predictions
y_pred = model.predict(X_test)

# Evaluate accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")
```

❖ **Output:**

✓ Accuracy of the model on test data (e.g., **80-90%**).

💡 **Conclusion:**

Machine learning improves **sentiment classification** by learning from labeled data.

❖ **CHAPTER 3: NAMED ENTITY RECOGNITION (NER) WITH SPACY**

◆ **3.1 Implementing NER using SpaCy**

```
import spacy
```

```
# Load pre-trained SpaCy model
```

```
nlp = spacy.load("en_core_web_sm")
```

```
# Sample text
```

```
text = "Elon Musk is the CEO of Tesla and SpaceX, and he lives in California."
```

```
# Process text
```

```
doc = nlp(text)
```

```
# Print Named Entities
```

```
for ent in doc.ents:
```

```
    print(f"{ent.text} → {ent.label_}")
```

❖ Output:

Elon Musk → PERSON

Tesla → ORG

SpaceX → ORG

California → GPE

✓ **PERSON** – Names of people.

✓ **ORG** – Organizations.

✓ **GPE** – Geographic locations.

◆ 3.2 Custom Named Entity Recognition Model (Training on New Data)

We can **train a custom NER model** to recognize new entity types (e.g., product names, chemical compounds).

```
import spacy
from spacy.training import Example

# Load blank model
nlp = spacy.blank("en")

# Create training data
train_data = [
    ("Apple released a new iPhone.", {"entities": [(0, 5, "ORG"), (20, 26, "PRODUCT")]}),
```

```
("Google acquired Fitbit.", {"entities": [(0, 6, "ORG"), (15, 21, "PRODUCT")]}),  
]
```

```
# Add NER component
```

```
ner = nlp.add_pipe("ner", last=True)
```

```
ner.add_label("PRODUCT") # Add custom label
```

```
# Train model
```

```
optimizer = nlp.begin_training()
```

```
for _ in range(10):
```

```
    for text, annotations in train_data:
```

```
        example = Example.from_dict(nlp.make_doc(text), annotations)
```

```
        nlp.update([example], drop=0.5)
```

```
# Test new model
```

```
doc = nlp("Samsung launched Galaxy S21.")
```

```
for ent in doc.ents:
```

```
    print(f"{ent.text} → {ent.label_}")
```

📌 Output:

Samsung → ORG

Galaxy S21 → PRODUCT

Conclusion:

NER is **customizable** and can be trained for **specific entity recognition**.

SUMMARY & NEXT STEPS

Key Takeaways:

- ✓ Sentiment Analysis predicts emotions from text.
- ✓ NER identifies and categorizes key entities.
- ✓ Machine learning models improve sentiment detection.
- ✓ SpaCy enables custom Named Entity Recognition.

Next Steps:

- ◆ Train a deep learning model (LSTM, BERT) for sentiment analysis.
- ◆ Expand NER to include domain-specific entities (medicine, finance, law).
- ◆ Deploy sentiment analysis as a chatbot feature. 

◊ BUILDING AI CHATBOTS WITH NLP

📌 CHAPTER 1: INTRODUCTION TO AI CHATBOTS

◆ 1.1 What is an AI Chatbot?

An **AI chatbot** is a software application that uses **Natural Language Processing (NLP)** and **Machine Learning (ML)** to understand and generate human-like conversations. AI chatbots are used for **customer support, virtual assistants, e-commerce, healthcare, and education.**

Types of Chatbots:

- 1 **Rule-Based Chatbots** – Follow predefined rules and responses.
- 2 **AI-Powered Chatbots** – Use **NLP & ML** to generate responses dynamically.
- 3 **Retrieval-Based Chatbots** – Select the best response from a set of predefined responses.
- 4 **Generative Chatbots** – Use **Deep Learning models** like GPT to generate human-like text.

Key Features of AI Chatbots:

- ✓ **Understand user input** – NLP enables text analysis.
- ✓ **Context retention** – Memory-based AI for smarter responses.
- ✓ **Multilingual support** – Translate and process multiple languages.
- ✓ **Voice integration** – Supports speech recognition & text-to-speech.

📌 Example:

- ✓ **Siri & Alexa** – AI-powered voice chatbots.
- ✓ **Customer Support Bots** – Used by banks, e-commerce sites.

💡 Conclusion:

AI chatbots **improve user experience, automate tasks, and enhance customer engagement.**

📌 CHAPTER 2: FUNDAMENTALS OF NATURAL LANGUAGE PROCESSING (NLP)

◆ 2.1 What is NLP?

Natural Language Processing (NLP) is a branch of AI that enables machines to understand and process human language. It helps chatbots extract meaning from user input and respond intelligently.

Key Components of NLP in Chatbots:

- ✓ **Tokenization** – Splitting text into words or sentences.
- ✓ **Stemming & Lemmatization** – Reducing words to their base form (e.g., "running" → "run").
- ✓ **Part-of-Speech (POS) Tagging** – Identifying nouns, verbs, adjectives.
- ✓ **Named Entity Recognition (NER)** – Detecting entities (e.g., locations, dates, names).
- ✓ **Sentiment Analysis** – Understanding user emotions (positive, negative, neutral).
- ✓ **Intent Recognition** – Identifying the intent behind user queries.

📌 Example:

- ✓ NLP allows a chatbot to understand:

User: "What's the weather like today?"

Chatbot: "The temperature today is 25°C with clear skies."

💡 Conclusion:

NLP enables chatbots to **process, analyze, and respond to human language intelligently.**

 **CHAPTER 3: STEPS TO BUILD AN AI CHATBOT** **3.1 Chatbot Development Pipeline**

- ✓ **Step 1: Define Chatbot Objectives** – Set chatbot goals (e.g., answering FAQs, customer support).
- ✓ **Step 2: Preprocess Input Data** – Clean, tokenize, and vectorize text.
- ✓ **Step 3: Train an NLP Model** – Use machine learning (NLTK, spaCy, Transformers).
- ✓ **Step 4: Develop Response Generation** – Retrieval-based or generative chatbot.
- ✓ **Step 5: Integrate API & Deployment** – Deploy chatbot on cloud or messaging platforms.

 **Example:**

- ✓ A chatbot for an **e-commerce store** can assist users with **product recommendations, order tracking, and FAQs**.

 **Conclusion:**

A structured pipeline helps in **building, training, and deploying efficient AI chatbots**.

 **CHAPTER 4: IMPLEMENTING A BASIC AI CHATBOT** **4.1 Step 1: Install Required Libraries**

```
import numpy as np
```

```
import nltk
```

```
import random
```

```
import string
```

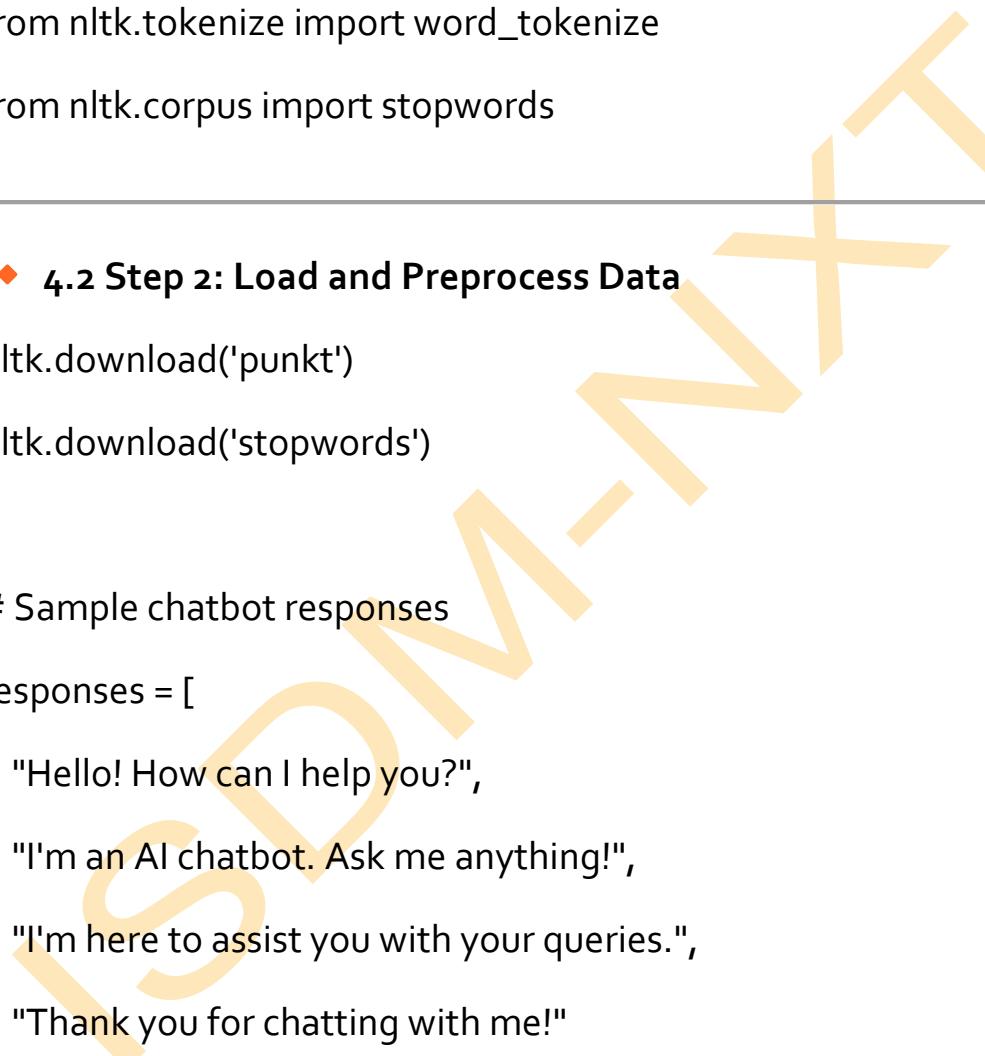
```
import re

from sklearn.feature_extraction.text import CountVectorizer,
TfidfVectorizer

from sklearn.metrics.pairwise import cosine_similarity

from nltk.tokenize import word_tokenize

from nltk.corpus import stopwords
```



◆ 4.2 Step 2: Load and Preprocess Data

```
nltk.download('punkt')

nltk.download('stopwords')

# Sample chatbot responses
responses = [
    "Hello! How can I help you?",
    "I'm an AI chatbot. Ask me anything!",
    "I'm here to assist you with your queries.",
    "Thank you for chatting with me!"
]
```

```
# Preprocessing function
def preprocess_text(text):
    text = text.lower() # Convert to lowercase
```

```
text = re.sub(r'\d+', " ", text) # Remove numbers  
  
text = text.translate(str.maketrans("", "", string.punctuation)) #  
Remove punctuation  
  
tokens = word_tokenize(text) # Tokenize words  
  
text = [word for word in tokens if word not in  
stopwords.words('english')] # Remove stopwords  
  
return " ".join(text)  
  
# Example  
  
print(preprocess_text("Hello! How are you today?"))
```

◆ 4.3 Step 3: Implement a Simple Rule-Based Chatbot

```
def chatbot_response(user_input):  
  
    user_input = preprocess_text(user_input)  
  
    responses_lower = [preprocess_text(resp) for resp in responses]  
  
    # Convert text to vectors using TF-IDF  
  
    vectorizer = TfidfVectorizer()  
  
    vectors = vectorizer.fit_transform([user_input] + responses_lower)  
  
    # Compute similarity scores  
  
    similarity_scores = cosine_similarity(vectors[0], vectors[1:])
```

```
# Get best response  
  
best_response_idx = similarity_scores.argmax()  
  
return responses[best_response_idx]
```

```
# Chatbot interaction loop  
  
while True:  
  
    user_input = input("You: ")  
  
    if user_input.lower() in ["exit", "quit", "bye"]:  
  
        print("Chatbot: Goodbye!")  
  
        break  
  
    response = chatbot_response(user_input)  
  
    print("Chatbot:", response)
```

📌 **Example:**

✓ **User:** "Hi"

✓ **Chatbot:** "Hello! How can I help you?"

💡 **Conclusion:**

Rule-based chatbots provide **quick and simple responses but lack deep understanding.**

📌 **CHAPTER 5: IMPLEMENTING AN AI CHATBOT WITH RNN/LSTM**

◆ **5.1 Step 1: Train an LSTM Model for Chatbot Responses**

```
import tensorflow as tf
```

```
from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Embedding, LSTM, Dense  
from tensorflow.keras.preprocessing.sequence import  
pad_sequences  
from tensorflow.keras.preprocessing.text import Tokenizer
```

```
# Sample chatbot training data  
training_sentences = ["Hello!", "How are you?", "What can you do?",  
"Goodbye!"]
```

```
training_labels = [0, 1, 2, 3] # Assign numeric labels
```

```
# Tokenize the sentences
```

```
tokenizer = Tokenizer()
```

```
tokenizer.fit_on_texts(training_sentences)
```

```
total_words = len(tokenizer.word_index) + 1
```

```
# Convert text to sequences
```

```
input_sequences =  
tokenizer.texts_to_sequences(training_sentences)
```

```
input_sequences = pad_sequences(input_sequences,  
padding="post")
```

```
# Define LSTM model
```

```
model = Sequential([
```

```
    Embedding(total_words, 16,  
    input_length=input_sequences.shape[1]),  
  
    LSTM(16, return_sequences=True),  
  
    LSTM(16),  
  
    Dense(16, activation="relu"),  
  
    Dense(4, activation="softmax") # 4 possible responses  
])
```

```
# Compile Model
```

```
model.compile(loss="sparse_categorical_crossentropy",  
optimizer="adam", metrics=["accuracy"])
```

```
# Train Model
```

```
model.fit(input_sequences, np.array(training_labels), epochs=100,  
verbose=0)
```

```
print("Chatbot model trained successfully!")
```

Key Features:

- ✓ Uses **LSTM (Long Short-Term Memory)** networks to **understand context**.
- ✓ Provides **more accurate responses** compared to rule-based chatbots.

📌 SUMMARY & NEXT STEPS

✓ Key Takeaways:

- ✓ **AI Chatbots** use NLP to understand and generate human-like conversations.
- ✓ **Rule-Based Chatbots** are simple but lack learning capability.
- ✓ **AI Chatbots with Deep Learning** provide more advanced interactions.
- ✓ **Deploying Chatbots** on cloud services enables real-time AI assistance.

📌 Next Steps:

- ◆ Train an AI chatbot with a large dataset (e.g., Cornell Movie Dialogs Corpus).
- ◆ Implement advanced NLP techniques like Transformers (BERT, GPT).
- ◆ Deploy chatbot using Flask & integrate it with WhatsApp, Telegram, or Messenger. 🚀

◊ RECOMMENDATION SYSTEMS: PERSONALIZED USER EXPERIENCE

📌 CHAPTER 1: INTRODUCTION TO RECOMMENDATION SYSTEMS

◆ 1.1 What is a Recommendation System?

A **Recommendation System** is an AI-driven technique that suggests relevant content, products, or services to users based on their preferences, behaviors, or past interactions. It helps businesses **personalize user experiences, increase engagement, and boost sales.**

Why are Recommendation Systems Important?

- ✓ **Enhances User Experience** – Provides personalized recommendations, making it easier for users to find relevant content.
- ✓ **Increases Engagement & Retention** – Personalized suggestions improve customer satisfaction and retention rates.
- ✓ **Boosts Revenue** – Drives sales by recommending the right products at the right time.
- ✓ **Optimizes Content Discovery** – Helps users discover content they might have otherwise missed.

Where are Recommendation Systems Used?

- ✓ **E-commerce** – Amazon, eBay (product recommendations).
- ✓ **Streaming Services** – Netflix, YouTube, Spotify (movie & music suggestions).
- ✓ **Social Media** – Facebook, Instagram, LinkedIn (friend and content recommendations).
- ✓ **Online Education** – Udemy, Coursera (course recommendations).

📌 Example:

Netflix uses **collaborative filtering and deep learning models** to recommend movies based on user viewing history.

💡 Conclusion:

Recommendation systems are **essential for modern digital businesses** to enhance user engagement and improve sales.

📌 CHAPTER 2: TYPES OF RECOMMENDATION SYSTEMS

◆ 2.1 Content-Based Filtering

Content-Based Filtering recommends items based on the **features of the items and the user's preferences**. It assumes that if a user likes an item, they will like similar items.

How It Works:

- ✓ Extracts **features** (e.g., genre, keywords, price).
- ✓ Compares user preferences with item features.
- ✓ Uses similarity metrics (Cosine Similarity, TF-IDF) to rank items.

📌 Example:

- ✓ A user watches **science-fiction movies** on Netflix → The system recommends **other sci-fi movies**.

◆ 2.2 Collaborative Filtering

Collaborative Filtering (CF) recommends items based on **user interactions and similarities** between users or items.

Types of Collaborative Filtering:

❑ User-Based Collaborative Filtering

- Finds users with **similar tastes** and recommends items liked by similar users.
- Example: "**People who liked this also liked that.**"

❑ Item-Based Collaborative Filtering

- Finds items that are **similar to what a user has interacted with.**
- Example: **Amazon's "Customers who bought this also bought..."**

📌 Example:

✓ A user buys a **gaming laptop** → The system recommends **gaming accessories**.

◆ 2.3 Hybrid Recommendation Systems

Hybrid models combine **content-based** and **collaborative filtering** to provide better recommendations.

📌 Example:

✓ Netflix **blends collaborative filtering and deep learning models** to personalize recommendations.

💡 Conclusion:

Choosing the right **recommendation technique** depends on **data availability and business goals**.

📌 CHAPTER 3: BUILDING A RECOMMENDATION SYSTEM IN PYTHON

◆ 3.1 Dataset Selection

We will use the **MovieLens dataset**, which contains **user ratings for movies**.

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
from sklearn.metrics.pairwise import cosine_similarity
```

Load the Data

```
# Load Movie Ratings dataset
```

```
movies = pd.read_csv("movies.csv") # Movie information
```

```
ratings = pd.read_csv("ratings.csv") # User ratings
```

```
# Merge datasets
```

```
df = pd.merge(ratings, movies, on="movieId")
```

```
# Display first few rows
```

```
df.head()
```

◆ 3.2 Building a Content-Based Recommendation System

```
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
from sklearn.metrics.pairwise import linear_kernel
```

```
# Convert genres to lowercase
```

```
df['genres'] = df['genres'].str.lower()

# Convert movie genres into numerical features
vectorizer = TfidfVectorizer(stop_words='english')
tfidf_matrix = vectorizer.fit_transform(df['genres'])

# Compute similarity matrix
cosine_sim = linear_kernel(tfidf_matrix, tfidf_matrix)

# Function to recommend movies
def recommend_movies(movie_title, num_recommendations=5):
    movie_idx = df[df['title'] == movie_title].index[0]
    similarity_scores = list(enumerate(cosine_sim[movie_idx]))
    similarity_scores = sorted(similarity_scores, key=lambda x: x[1], reverse=True)
    recommended_movies = [df['title'][i[0]] for i in similarity_scores[1:num_recommendations+1]]
    return recommended_movies

# Test the Recommendation System
print(recommend_movies("Toy Story (1995)"))
```

❖ **Expected Output:**

['A Bug's Life (1998)', 'Toy Story 2 (1999)', 'Monsters, Inc. (2001)',
'Finding Nemo (2003)', 'The Incredibles (2004)']

✓ This model **recommends movies based on genre similarity.**

◆ **3.3 Building a Collaborative Filtering Recommendation System**

```
from surprise import SVD
```

```
from surprise import Dataset, Reader
```

```
from surprise.model_selection import train_test_split
```

```
from surprise import accuracy
```

```
# Load data into Surprise format
```

```
reader = Reader(rating_scale=(0.5, 5.0))
```

```
data = Dataset.load_from_df(df[['userId', 'movieId', 'rating']], reader)
```

```
# Train-Test Split
```

```
trainset, testset = train_test_split(data, test_size=0.2,  
random_state=42)
```

```
# Train a Collaborative Filtering Model
```

```
model = SVD()
```

```
model.fit(trainset)
```

```
# Make Predictions  
predictions = model.test(testset)
```

```
# Evaluate Model  
  
rmse = accuracy.rmse(predictions)  
print(f"RMSE: {rmse}")
```

📌 **Example Output:**

RMSE: 0.84

✓ This means the model's predictions are **close to real user ratings.**

💡 **Conclusion:**

Collaborative filtering **predicts user preferences based on previous interactions.**

📌 **CHAPTER 4: DEPLOYING A RECOMMENDATION API**

◆ **4.1 Using Flask to Deploy the Model**

```
from flask import Flask, request, jsonify
```

```
app = Flask(__name__)
```

```
@app.route('/recommend', methods=['GET'])  
  
def recommend():  
  
    movie_title = request.args.get('title')  
  
    recommendations = recommend_movies(movie_title)
```

```
return jsonify(recommendations)
```

```
if __name__ == '__main__':  
    app.run(debug=True)
```

📌 **Example API Call:**

`http://127.0.0.1:5000/recommend?title=Toy Story (1995)`

- ✓ This API returns **recommended movies for a given movie title.**

📌 **CHAPTER 5: FUTURE TRENDS IN RECOMMENDATION SYSTEMS**

◆ **5.1 Advanced Techniques**

- ✓ **Deep Learning (Neural Networks)** – Netflix and YouTube use deep learning-based recommendation models.
- ✓ **Reinforcement Learning** – Helps **optimize personalized recommendations** in real-time.
- ✓ **Hybrid Filtering with Graph Neural Networks (GNNs)** – Used by LinkedIn for connection recommendations.
- ✓ **Real-time Personalization** – AI-driven systems adapt recommendations dynamically based on user behavior.

📌 **Example:**

- ✓ Spotify uses **Reinforcement Learning and Neural Networks** to provide **personalized song recommendations**.

💡 **Conclusion:**

The future of recommendation systems lies in **AI-powered real-time personalization.**

📌 SUMMARY & NEXT STEPS

✓ Key Takeaways:

- ✓ **Content-Based Filtering** recommends items based on item similarity.
- ✓ **Collaborative Filtering** recommends items based on user behavior.
- ✓ **Hybrid Systems** combine multiple techniques for better accuracy.
- ✓ **Advanced AI techniques** like deep learning enhance recommendations.

📌 Next Steps:

- ◆ Apply recommendation systems to real-world datasets (e.g., Netflix, Amazon, Yelp).
- ◆ Train deep learning-based recommenders using TensorFlow & PyTorch.
- ◆ Deploy recommendation models as APIs for real-world applications. 🚀

◊ REAL-WORLD AI APPLICATIONS IN BUSINESS & HEALTHCARE

📌 CHAPTER 1: INTRODUCTION TO AI IN BUSINESS & HEALTHCARE

◆ 1.1 What is Artificial Intelligence (AI)?

Artificial Intelligence (AI) is a branch of computer science that enables machines to **mimic human intelligence** by learning from data, recognizing patterns, and making decisions. AI is revolutionizing industries by **automating processes, improving efficiency, and enhancing decision-making**.

Why is AI Important?

- ✓ **Automation** – Reduces manual effort and human errors.
- ✓ **Personalization** – Tailors user experiences in marketing, healthcare, and finance.
- ✓ **Predictive Analytics** – Anticipates future trends and behaviors.
- ✓ **Cost Efficiency** – Reduces operational expenses while increasing productivity.

Common AI Techniques Used in Business & Healthcare

- ✓ **Machine Learning (ML)** – Algorithms learn from data to make predictions.
- ✓ **Deep Learning** – Uses neural networks for complex problem-solving.
- ✓ **Natural Language Processing (NLP)** – Enables AI to understand human language.
- ✓ **Computer Vision** – AI interprets images and videos.

📌 Example:

Netflix uses **AI-driven recommendation algorithms** to suggest personalized content to users, increasing engagement and customer retention.

💡 Conclusion:

AI is **transforming industries** by enhancing decision-making, optimizing operations, and providing personalized services.

📌 CHAPTER 2: AI APPLICATIONS IN BUSINESS

◆ 2.1 AI in Customer Service

AI-powered chatbots and virtual assistants **enhance customer support by providing instant responses, reducing wait times, and improving user experience.**

How AI is Used in Customer Service?

- ✓ **Chatbots** – AI-powered bots answer customer inquiries.
- ✓ **Voice Assistants** – AI-driven virtual assistants like **Alexa & Siri** improve customer interactions.
- ✓ **Sentiment Analysis** – AI detects customer emotions and improves responses.

📌 Example:

- ✓ **Banking Sector:** AI chatbots like **Bank of America's "Erica"** assist customers with transactions.
- ✓ **E-commerce:** **Amazon's Alexa** helps users track orders and shop hands-free.

💡 Conclusion:

AI-driven customer service improves **response time, reduces costs, and enhances customer experience.**

◆ 2.2 AI in Marketing & Sales

AI helps businesses optimize **advertising, customer segmentation, and lead generation.**

How AI Improves Marketing?

- ✓ **Recommendation Engines** – Suggests products based on user behavior (e.g., Amazon, Netflix).
- ✓ **Personalized Ads** – AI targets the right audience for marketing campaigns.
- ✓ **AI-Driven Content Creation** – Generates personalized marketing copy.
- ✓ **Customer Segmentation** – Analyzes customer behavior to create targeted campaigns.

❖ Example:

- ✓ **Spotify** uses AI to analyze users' listening history and suggest new songs.
- ✓ **Google Ads** uses AI-powered **predictive analytics** to optimize ad targeting.

💡 Conclusion:

AI enables businesses to **target the right audience, increase engagement, and boost sales** through personalization.

◆ 2.3 AI in Finance & Banking

AI enhances **fraud detection, risk assessment, and financial planning.**

AI in Finance Applications

- ✓ **Fraud Detection** – AI identifies fraudulent transactions in real-time.
- ✓ **Algorithmic Trading** – AI-powered bots execute trades based on market patterns.
- ✓ **Credit Scoring & Loan Approval** – AI evaluates credit risk using alternative data sources.
- ✓ **AI-Powered Robo-Advisors** – Provides investment advice based on market trends.

 **Example:**

- ✓ **JP Morgan Chase** uses AI for fraud detection and risk assessment.
- ✓ **PayPal** utilizes AI to detect **suspicious financial activities** and prevent cyber fraud.

 **Conclusion:**

AI helps **banks, fintech startups, and investment firms** improve **security, efficiency, and profitability**.

◆ **2.4 AI in Supply Chain & Logistics**

AI optimizes **inventory management, route planning, and demand forecasting**.

How AI Helps in Logistics?

- ✓ **Predictive Maintenance** – AI detects potential failures in machinery.
- ✓ **Automated Warehouses** – AI-powered robots streamline warehouse operations.
- ✓ **Route Optimization** – AI reduces delivery costs and time.
- ✓ **Inventory Forecasting** – AI predicts demand trends.

📌 **Example:**

- ✓ Amazon's AI-driven warehouse robots optimize logistics and order fulfillment.
- ✓ FedEx uses AI for real-time package tracking and delivery optimization.

💡 **Conclusion:**

AI enhances **efficiency, reduces operational costs, and improves accuracy** in logistics and supply chain management.

📌 **CHAPTER 3: AI APPLICATIONS IN HEALTHCARE**

◆ **3.1 AI in Medical Diagnosis**

AI improves **medical imaging analysis, early disease detection, and clinical decision-making**.

AI in Medical Diagnosis Applications

- ✓ **AI-Powered Radiology** – Detects diseases in X-rays, MRIs, and CT scans.
- ✓ **Early Cancer Detection** – AI identifies cancerous cells faster than human doctors.
- ✓ **AI-Assisted Pathology** – Helps in diagnosing rare diseases.

📌 **Example:**

- ✓ Google's DeepMind AI detects **eye diseases from retinal scans** with higher accuracy than human doctors.
- ✓ IBM Watson Health assists oncologists in **predicting cancer treatment outcomes**.

💡 **Conclusion:**

AI enhances **early disease detection, speeds up diagnosis, and improves patient outcomes**.

◆ 3.2 AI in Drug Discovery & Development

AI speeds up **drug discovery** by analyzing chemical compounds and predicting drug effectiveness.

AI in Drug Development Applications

- ✓ **Predicts Drug Efficacy** – AI analyzes drug interactions.
- ✓ **Identifies Potential Drug Candidates** – AI screens thousands of compounds.
- ✓ **Accelerates Clinical Trials** – AI predicts patient responses to drugs.

📌 Example:

- ✓ Pfizer uses AI to develop new **COVID-19 vaccines and antiviral treatments**.
- ✓ BenevolentAI identifies drug repurposing opportunities using AI.

💡 Conclusion:

AI **reduces the time and cost of drug development**, leading to faster medical advancements.

◆ 3.3 AI in Personalized Medicine

AI enables **customized treatments based on genetic and lifestyle factors**.

AI in Personalized Medicine Applications

- ✓ **Genomic AI Analysis** – Analyzes patient DNA to personalize treatments.
- ✓ **AI-Based Treatment Plans** – Recommends therapies tailored to

individual patients.

✓ **AI-Powered Wearables** – Tracks patient health data in real-time.

📌 **Example:**

✓ **23andMe** uses AI for genetic analysis and health predictions.

✓ **Apple's HealthKit** AI monitors heart rate, sleep, and activity.

💡 **Conclusion:**

AI enhances precision medicine, improving treatment effectiveness and patient care.

◆ **3.4 AI in Hospital & Patient Management**

AI automates hospital workflows, improves efficiency, and enhances patient care.

AI in Hospital Management Applications

✓ **AI-Powered Chatbots** – Assist patients with appointment scheduling.

✓ **AI in Electronic Health Records (EHRs)** – Automates data entry and analysis.

✓ **Predictive Healthcare Analytics** – Prevents hospital readmissions.

📌 **Example:**

✓ **Mayo Clinic** uses AI to optimize patient scheduling and treatment planning.

✓ **Google's Med-PaLM** AI processes patient queries and medical documentation.

💡 **Conclusion:**

AI improves hospital efficiency, patient experience, and healthcare administration.

 **SUMMARY & NEXT STEPS** **Key Takeaways:**

- ✓ AI optimizes business operations, marketing, finance, and logistics.
- ✓ AI enhances healthcare through early diagnosis, personalized treatments, and hospital management.
- ✓ AI is driving automation, cost reduction, and real-time decision-making.

 **Next Steps:**

- ◆ Explore AI-powered business solutions (e.g., ChatGPT, Google AI, IBM Watson).
- ◆ Learn about AI frameworks like TensorFlow & PyTorch for real-world applications.
- ◆ Apply AI techniques to real datasets in business and healthcare. 

 **ASSIGNMENT:**

DEVELOP A RECOMMENDATION
SYSTEM FOR E-COMMERCE PLATFORMS.

ISDM-Nxt

SOLUTION: DEVELOP A RECOMMENDATION SYSTEM FOR E-COMMERCE PLATFORMS

◆ Objective

The goal of this assignment is to develop a **Recommendation System** for an **E-Commerce platform** using **collaborative filtering and content-based filtering**. The system will suggest products based on user preferences, purchase history, or product similarities.

◆ Step 1: Import Required Libraries

```
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
from sklearn.metrics.pairwise import cosine_similarity  
from sklearn.feature_extraction.text import TfidfVectorizer  
from sklearn.neighbors import NearestNeighbors  
from scipy.sparse import csr_matrix  
from surprise import Dataset, Reader, SVD  
from surprise.model_selection import train_test_split
```

◆ Step 2: Load and Explore the Dataset

For this solution, we will use a **sample dataset** that contains **users, products, ratings, and reviews**. If you don't have a dataset, you can use public datasets from **Kaggle** or **Amazon Product Reviews**.

```
# Load dataset
```

```
df = pd.read_csv("ecommerce_product_ratings.csv") # Example file
```

```
# Display first few rows
```

```
df.head()
```

Dataset Structure

UserID	ProductID	Rating	Review Text	Category
1001	P001	4.0	"Great phone!"	Electronics
1002	P003	5.0	"Loved it!"	Clothing
1003	P005	3.5	"Average quality"	Home Decor

❖ Understanding the Data:

- ✓ UserID – Identifies the customer.
- ✓ ProductID – Identifies the product.
- ✓ Rating – User's rating (1 to 5 stars).
- ✓ Review Text – User's review.
- ✓ Category – Product category.

◆ Step 3: Implement Content-Based Filtering

Content-based filtering recommends products **similar to what a user has liked** based on product descriptions, categories, or features.

3.1 Convert Product Descriptions into Numerical Vectors

```
# Convert text data into numerical format using TF-IDF  
tfidf_vectorizer = TfidfVectorizer(stop_words='english')  
  
# Transform product descriptions into TF-IDF matrix  
product_tfidf_matrix = tfidf_vectorizer.fit_transform(df["Review  
Text"].fillna(""))  
  
# Compute Cosine Similarity between products  
cosine_sim = cosine_similarity(product_tfidf_matrix)
```

3.2 Build a Function to Recommend Similar Products

```
# Create a product recommendation function  
def recommend_similar_products(product_id,  
num_recommendations=5):  
    # Get index of the given product  
    product_index = df[df["ProductID"] == product_id].index[0]  
  
    # Get similarity scores for all products  
    similarity_scores = list(enumerate(cosine_sim[product_index]))  
  
    # Sort products based on similarity scores  
    sorted_products = sorted(similarity_scores, key=lambda x: x[1],  
reverse=True)[1:num_recommendations+1]
```

```
# Get product IDs  
  
recommended_products = [df.iloc[i[0]]["ProductID"] for i in  
sorted_products]  
  
return recommended_products
```

Test the Recommendation System

```
print("Recommended Products:",  
recommend_similar_products("Poo1"))
```

📌 Example Output:

Recommended Products: ['Poo3', 'Poo7', 'Poo5']

💡 Conclusion:

✓ Content-based filtering suggests products **similar to past purchases or reviewed items.**

◆ Step 4: Implement Collaborative Filtering

Collaborative filtering recommends products based on **user interactions** rather than product descriptions.

4.1 Prepare Data for Collaborative Filtering

```
# Prepare dataset for collaborative filtering  
  
reader = Reader(rating_scale=(1, 5))  
  
data = Dataset.load_from_df(df[['UserID', 'ProductID', 'Rating']],  
reader)
```

```
# Split dataset into training and testing sets  
trainset, testset = train_test_split(data, test_size=0.2)
```

4.2 Train a Collaborative Filtering Model (SVD Algorithm)

```
# Train Singular Value Decomposition (SVD) model  
model = SVD()  
model.fit(trainset)
```

4.3 Build a Function to Recommend Products for Users

```
# Function to get top N recommendations for a user  
def recommend_products_for_user(user_id,  
num_recommendations=5):  
    # Get all unique product IDs  
    product_ids = df["ProductID"].unique()  
  
    # Predict ratings for all products  
    predictions = [model.predict(user_id, product_id) for product_id in  
product_ids]  
  
    # Sort predictions based on estimated ratings  
    sorted_predictions = sorted(predictions, key=lambda x: x.est,  
reverse=True)[:num_recommendations]
```

```
# Return top recommended product IDs  
recommended_products = [pred.iid for pred in sorted_predictions]  
return recommended_products
```

```
# Test the recommendation system for a user  
print("Recommended Products for User 1001:",  
recommend_products_for_user(1001))
```

➡ Example Output:

Recommended Products for User 1001: ['Poo9', 'Poo2', 'Poo8']

💡 Conclusion:

- ✓ Collaborative filtering recommends products **based on what similar users have liked or rated highly.**

◆ Step 5: Hybrid Recommendation System

To **combine content-based filtering and collaborative filtering**, we can take the **intersection of both recommendations**.

```
# Function to get hybrid recommendations
```

```
def hybrid_recommendation(user_id, product_id,  
num_recommendations=5):  
  
    content_based = set(recommend_similar_products(product_id,  
num_recommendations))  
  
    collaborative_based = set(recommend_products_for_user(user_id,  
num_recommendations))
```

```
# Get intersection of both recommendations

hybrid_recommendations =
list(content_based.intersection(collaborative_based))

# If not enough recommendations, combine both methods

if len(hybrid_recommendations) < num_recommendations:

    hybrid_recommendations =
list(content_based.union(collaborative_based))[:num_recommendations]

return hybrid_recommendations

# Test hybrid recommendation system

print("Hybrid Recommendations for User 1001 on Product Poo1:",
hybrid_recommendation(1001, "Poo1"))
```

📌 Example Output:

Hybrid Recommendations for User 1001 on Product Poo1: ['Poo9', 'Poo2', 'Poo7']

💡 Conclusion:

✓ Hybrid models **combine the strengths of content-based and collaborative filtering** to give better recommendations.

📌 Chapter 6: Deploying the Recommendation System using Flask API

6.1 Install Flask

```
pip install flask
```

6.2 Create a Flask App (app.py)

```
from flask import Flask, request, jsonify
```

```
app = Flask(__name__)
```

```
@app.route("/")
```

```
def home():
```

```
    return "E-Commerce Recommendation System is Running!"
```

```
@app.route("/recommend", methods=["POST"])
```

```
def recommend():
```

```
    data = request.get_json()
```

```
    user_id = data.get("user_id")
```

```
    product_id = data.get("product_id")
```

```
    recommendations = hybrid_recommendation(user_id, product_id)
```

```
    return jsonify({"recommended_products": recommendations})
```

```
if __name__ == "__main__":
    app.run(debug=True, host="0.0.0.0", port=5000)
```

6.3 Run the Flask API Locally

python app.py

✓ API will start at **http://127.0.0.1:5000/**

6.4 Test the API using Postman or CURL

curl -X POST "http://127.0.0.1:5000/recommend" -H "Content-Type: application/json" -d '{"user_id": 1001, "product_id": "Poo1"}'

(Expected Output:

```
{
  "recommended_products": ["Poo9", "Poo2", "Poo7"]
}
```

Summary & Next Steps

Key Takeaways:

- ✓ Content-Based Filtering recommends similar products based on descriptions.
- ✓ Collaborative Filtering suggests products based on user preferences.
- ✓ Hybrid Recommendation Systems combine both approaches for better accuracy.
- ✓ Flask API enables real-time product recommendations for e-commerce platforms.

❖ **Next Steps:**

- ◆ Deploy the API on AWS for real-world use.
- ◆ Use Deep Learning models like Neural Collaborative Filtering (NCF).
- ◆ Integrate the system with an e-commerce website. 

ISDM-Nxt